

This file contains the exercises, hints, and solutions for Chapter 1 of the book "Introduction to the Design and Analysis of Algorithms," 2nd edition, by A. Levitin. The problems that might be challenging for at least some students are marked by \triangleright ; those that might be difficult for a majority of students are marked by \blacktriangleright .

Exercises 1.1

1. Do some research on al-Khorezmi (also al-Khwarizmi), the man from whose name the word "algorithm" is derived. In particular, you should learn what the origins of the words "algorithm" and "algebra" have in common.
2. Given that the official purpose of the U.S. patent system is the promotion of the "useful arts," do you think algorithms are patentable in this country? Should they be?
3. a. Write down driving directions for going from your school to your home with the precision required by an algorithm.
 b. Write down a recipe for cooking your favorite dish with the precision required by an algorithm.
4. Design an algorithm for computing $\lfloor \sqrt{n} \rfloor$ for any positive integer n . Besides assignment and comparison, your algorithm may only use the four basic arithmetical operations.
5. a. Find $\text{gcd}(31415, 14142)$ by applying Euclid's algorithm.
 b. Estimate how many times faster it will be to find $\text{gcd}(31415, 14142)$ by Euclid's algorithm compared with the algorithm based on checking consecutive integers from $\min\{m, n\}$ down to $\text{gcd}(m, n)$.
6. \triangleright Prove the equality $\text{gcd}(m, n) = \text{gcd}(n, m \bmod n)$ for every pair of positive integers m and n .
7. What does Euclid's algorithm do for a pair of numbers in which the first number is smaller than the second one? What is the largest number of times this can happen during the algorithm's execution on such an input?
8. a. What is the smallest number of divisions made by Euclid's algorithm among all inputs $1 \leq m, n \leq 10$?
 b. What is the largest number of divisions made by Euclid's algorithm among all inputs $1 \leq m, n \leq 10$?
9. a. Euclid's algorithm, as presented in Euclid's treatise, uses subtractions rather than integer divisions. Write a pseudocode for this version of Euclid's algorithm.

- b.► *Euclid's game* (see [Bog]) starts with two unequal positive numbers on the board. Two players move in turn. On each move, a player has to write on the board a positive number equal to the difference of two numbers already on the board; this number must be new, i.e., different from all the numbers already on the board. The player who cannot move loses the game. Should you choose to move first or second in this game?
10. The *extended Euclid's algorithm* determines not only the greatest common divisor d of two positive integers m and n but also integers (not necessarily positive) x and y , such that $mx + ny = d$.
- a. Look up a description of the extended Euclid's algorithm (see, e.g., [KnuI], p. 13) and implement it in the language of your choice.
- b. Modify your program for finding integer solutions to the Diophantine equation $ax + by = c$ with any set of integer coefficients a , b , and c .
11. *Locker doors* There are n lockers in a hallway numbered sequentially from 1 to n . Initially, all the locker doors are closed. You make n passes by the lockers, each time starting with locker #1. On the i th pass, $i = 1, 2, \dots, n$, you toggle the door of every i th locker: if the door is closed, you open it, if it is open, you close it. For example, after the first pass every door is open; on the second pass you only toggle the even-numbered lockers (#2, #4, ...) so that after the second pass the even doors are closed and the odd ones are opened; the third time through you close the door of locker #3 (opened from the first pass), open the door of locker #6 (closed from the second pass), and so on. After the last pass, which locker doors are open and which are closed? How many of them are open?

Hints to Exercises 1.1

1. It is probably faster to do this by searching the Web, but your library should be able to help too.
2. One can find arguments supporting either view. There is a well established principle pertinent to the matter though: scientific facts or mathematical expressions of them are not patentable. (Why do you think it is the case?) But should this preclude granting patents for all algorithms?
3. You may assume that you are writing your algorithms for a human rather than a machine. Still, make sure that your descriptions do not contain obvious ambiguities. Knuth [KnuI], p.6 provides an interesting comparison between cooking recipes and algorithms.
4. There is a quite straightforward algorithm for this problem based on the definition of $\lfloor \sqrt{n} \rfloor$.
5. a. Just follow Euclid's algorithm as described in the text.
b. Compare the number of divisions made by the two algorithms.
6. Prove that if d divides both m and n (i.e., $m = sd$ and $n = td$ for some positive integers s and t), then it also divides both n and $r = m \bmod n$ and vice versa. Use the formula $m = qn + r$ ($0 \leq r < n$) and the fact that if d divides two integers u and v , it also divides $u + v$ and $u - v$. (Why?)
7. Perform one iteration of the algorithm for two arbitrarily chosen integers $m < n$.
8. The answer to part (a) can be given immediately; the answer to part (b) can be given by checking the algorithm's performance on all pairs $1 < m < n \leq 10$.
9. a. Use the equality

$$\gcd(m, n) = \gcd(m - n, n) \text{ for } m \geq n > 0.$$

- b. The key is to figure out the total number of distinct integers that can be written on the board, starting with an initial pair m, n where $m > n \geq 1$. You should exploit a connection of this question to the question of part (a). Considering small examples, especially those with $n = 1$ and $n = 2$, should help, too.
10. Of course, for some coefficients, the equation will have no solutions.
11. Tracing the algorithm by hand for, say, $n = 10$ and studying its outcome should help answering both questions.

Solutions to Exercises 1.1

1. Al-Khwarizmi (9th century C.E.) was a great Arabic scholar, most famous for his algebra textbook. In fact, the word “algebra” is derived from the Arabic title of this book while the word “algorithm” is derived from a translation of Al-Khwarizmi’s last name (see, e.g., [KnuI], pp. 1-2, [Knu96], pp. 88-92, 114).
2. This legal issue has yet to be settled. The current legal state of affairs distinguishes mathematical algorithms, which are not patentable, from other algorithms, which may be patentable if implemented as computer programs (e.g., [Cha00]).
3. n/a
4. A straightforward algorithm that does not rely on the availability of an approximate value of \sqrt{n} can check the squares of consecutive positive integers until the first square exceeding n is encountered. The answer will be the number’s immediate predecessor. Note: A much faster algorithm for solving this problem can be obtained by using Newton’s method (see Sections 11.4 and 12.4).
5. a. $\gcd(31415, 14142) = \gcd(14142, 3131) = \gcd(3131, 1618) = \gcd(1618, 1513) = \gcd(1513, 105) = \gcd(1513, 105) = \gcd(105, 43) = \gcd(43, 19) = \gcd(19, 5) = \gcd(5, 4) = \gcd(4, 1) = \gcd(1, 0) = 1$.

b. To answer the question, we need to compare the number of divisions the algorithms make on the input given. The number of divisions made by Euclid’s algorithm is 11 (see part a). The number of divisions made by the consecutive integer checking algorithm on each of its 14142 iterations is either 1 and 2; hence the total number of multiplications is between $1 \cdot 14142$ and $2 \cdot 14142$. Therefore, Euclid’s algorithm will be between $1 \cdot 14142 / 11 \approx 1300$ and $2 \cdot 14142 / 11 \approx 2600$ times faster.
6. Let us first prove that if d divides two integers u and v , it also divides both $u + v$ and $u - v$. By definition of division, there exist integers s and t such that $u = sd$ and $v = td$. Therefore

$$u \pm v = sd \pm td = (s \pm t)d,$$

i.e., d divides both $u + v$ and $u - v$.

Also note that if d divides u , it also divides any integer multiple ku of u . Indeed, since d divides u , $u = sd$. Hence

$$ku = k(sd) = (ks)d,$$

i.e., d divides ku .

Now we can prove the assertion in question. For any pair of positive integers m and n , if d divides both m and n , it also divides both n and $r = m \bmod n = m - qn$. Similarly, if d divides both n and $r = m \bmod n = m - qn$, it also divides both $m = r + qn$ and n . Thus, the two pairs (m, n) and (n, r) have the same finite nonempty set of common divisors, including the largest element in the set, i.e., $\gcd(m, n) = \gcd(n, r)$.

7. For any input pair m, n such that $0 \leq m < n$, Euclid's algorithm simply swaps the numbers on the first iteration:

$$\gcd(m, n) = \gcd(n, m)$$

because $m \bmod n = m$ if $m < n$. Such a swap can happen only once since $\gcd(m, n) = \gcd(n, m \bmod n)$ implies that the first number of the new pair (n) will be greater than its second number $(m \bmod n)$ after every iteration of the algorithm.

8. a. For any input pair $m \geq n \geq 1$, in which m is a multiple of n , Euclid's algorithm makes exactly one division; it is the smallest number possible for two positive numbers.

b. The answer is 5 divisions, which is made by Euclid's algorithm in computing $\gcd(5, 8)$. It is not too time consuming to get this answer by examining the number of divisions made by the algorithm on all input pairs $1 < m < n \leq 10$.

Note: A pertinent general result (see [KnuII], p. 360) is that for any input pair m, n where $0 \leq n < N$, the number of divisions required by Euclid's algorithm to compute $\gcd(m, n)$ is at most $\lfloor \log_\phi(3 - \phi)N \rfloor$ where $\phi = (1 + \sqrt{5})/2$.

9. a. Here is a nonrecursive version:

Algorithm *Euclid2*(m, n)

//Computes $\gcd(m, n)$ by Euclid's algorithm based on subtractions

//Input: Two nonnegative integers m and n not both equal to 0

//Output: The greatest common divisor of m and n

while $n \neq 0$ **do**

if $m < n$ *swap*(m, n)

$m \leftarrow m - n$

return m

b. It is not too difficult to prove that the integers that can be written on the board are the integers generated by the subtraction version of Euclid's algorithm and only them. Although the order in which they appear on the board may vary, their total number always stays the same: It is equal to $m/\gcd(m, n)$, where m is the maximum of the initial numbers, which includes two integers of the initial pair. Hence, the total number of possible moves is $m/\gcd(m, n) - 2$. Consequently, if $m/\gcd(m, n)$ is odd, one should choose to go first; if it is even, one should choose to go second.

10. n/a

11. Since all the doors are initially closed, a door will be open after the last pass if and only if it is toggled an odd number of times. Door i ($1 \leq i \leq n$) is toggled on pass j ($1 \leq j \leq n$) if and only if j divides i . Hence, the total number of times door i is toggled is equal to the number of its divisors. Note that if j divides i , i.e. $i = jk$, then k divides i too. Hence all the divisors of i can be paired (e.g., for $i = 12$, such pairs are 1 and 12, 2 and 6, 3 and 4) unless i is a perfect square (e.g., for $i = 16$, 4 does not have another divisor to be matched with). This implies that i has an odd number of divisors if and only if it is a perfect square, i.e., $i = j^2$. Hence doors that are in the positions that are perfect squares and only such doors will be open after the last pass. The total number of such positions not exceeding n is equal to $\lfloor \sqrt{n} \rfloor$: these numbers are the squares of the positive integers between 1 and $\lfloor \sqrt{n} \rfloor$ inclusively.

Exercises 1.2

1. *Old World puzzle* A peasant finds himself on a riverbank with a wolf, a goat, and a head of cabbage. He needs to transport all three to the other side of the river in his boat. However, the boat has room for only the peasant himself and one other item (either the wolf, the goat, or the cabbage). In his absence, the wolf would eat the goat, and the goat would eat the cabbage. Solve this problem for the peasant or prove it has no solution. (Note: The peasant is a vegetarian but does not like cabbage and hence can eat neither the goat nor the cabbage to help him solve the problem. And it goes without saying that the wolf is a protected species.)
2. *New World puzzle* There are four people who want to cross a bridge; they all begin on the same side. You have 17 minutes to get them all across to the other side. It is night, and they have one flashlight. A maximum of two people can cross the bridge at one time. Any party that crosses, either one or two people, must have the flashlight with them. The flashlight must be walked back and forth; it cannot be thrown, for example. Person 1 takes 1 minute to cross the bridge, person 2 takes 2 minutes, person 3 takes 5 minutes, and person 4 takes 10 minutes. A pair must walk together at the rate of the slower person's pace. For example, if person 1 and person 4 walk across first, 10 minutes have elapsed when they get to the other side of the bridge. If person 4 returns the flashlight, a total of 20 minutes have passed and you have failed the mission. (Note: According to a rumor on the Internet, interviewers at a well-known software company located near Seattle have given this problem to interviewees.)
3. Which of the following formulas can be considered an algorithm for computing the area of a triangle whose side lengths are given positive numbers a , b , and c ?
 - a. $S = \sqrt{p(p-a)(p-b)(p-c)}$, where $p = (a+b+c)/2$
 - b. $S = \frac{1}{2}bc \sin A$, where A is the angle between sides b and c
 - c. $S = \frac{1}{2}ah_a$, where h_a is the height to base a
4. Write a pseudocode for an algorithm for finding real roots of equation $ax^2 + bx + c = 0$ for arbitrary real coefficients a , b , and c . (You may assume the availability of the square root function $\text{sqrt}(x)$.)
5. Describe the standard algorithm for finding the binary representation of a positive decimal integer
 - a. in English.
 - b. in a pseudocode.

6. Describe the algorithm used by your favorite ATM machine in dispensing cash. (You may give your description in either English or a pseudocode, whichever you find more convenient.)
7. a. Can the problem of computing the number π be solved exactly?
 b. How many instances does this problem have?
 c. Look up an algorithm for this problem on the World Wide Web.
8. Give an example of a problem other than computing the greatest common divisor for which you know more than one algorithm. Which of them is simpler? Which is more efficient?
9. Consider the following algorithm for finding the distance between the two closest elements in an array of numbers.

Algorithm *MinDistance*($A[0..n-1]$)
 //Input: Array $A[0..n-1]$ of numbers
 //Output: Minimum distance between two of its elements
 $dmin \leftarrow \infty$
for $i \leftarrow 0$ **to** $n-1$ **do**
 for $j \leftarrow 0$ **to** $n-1$ **do**
 if $i \neq j$ **and** $|A[i] - A[j]| < dmin$
 $dmin \leftarrow |A[i] - A[j]|$
return $dmin$

Make as many improvements as you can in this algorithmic solution to the problem. (If you need to, you may change the algorithm altogether; if not, improve the implementation given.)

10. One of the most influential books on problem solving, titled *How To Solve It* [Pol57], was written by the Hungarian-American mathematician George Polya (1887–1985). Polya summarized his ideas in a four-point summary. Find this summary on the Web or, better yet, in his book, and compare it with the plan outlined in Section 1.2. What do they have in common? How are they different?

Hints to Exercises 1.2

1. The peasant would have to make several trips across the river, starting with the only one possible.
2. Unlike the Old World puzzle of Problem 1, the first move solving this puzzle is not obvious.
3. The principal issue here is a possible ambiguity.
4. Your algorithm should work correctly for all possible values of the coefficients, including zeros.
5. You almost certainly learned this algorithm in one of your introductory programming courses. If this assumption is not true, you have a choice between designing such an algorithm on your own or looking it up.
6. You may need to make a field trip to refresh your memory.
7. Question (a) is difficult though the answer to it—discovered in 1760s by the German mathematician Johann Lambert—is well-known. By comparison, question (b) is incomparably simpler.
8. You probably know two or more different algorithms for sorting an array of numbers.
9. You can: decrease the number of times the innermost loop is executed, make that loop run faster (at least for some inputs), or, more significantly, design a faster algorithm from scratch.
10. n/a

Solutions to Exercises 1.2

1. Let P, w, g, and c stand for the peasant, wolf, goat, and cabbage head, respectively. The following is one of the two principal sequences that solve the problem:

	P g	g	Pwg	w	Pw c	w c	Pwgc
Pwgc	w c	Pw c	c	P gc	g	P g	

Note: This problem is revisited later in the book (see Section 6.6).

2. Let 1, 2, 5, 10 be labels representing the men of the problem, f represent the flashlight's location, and the number in the parenthesis be the total amount of time elapsed. The following sequence of moves solves the problem:

	$f,1,2$	2	$f,2,5,10$	5,10	$f,1,2,5,10$
(0)	(2)	(3)	(13)	(15)	(17)
$f,1,2,5,10$	5,10	$f,1,5,10$	1	$f,1,2$	

3. a. The formula can be considered an algorithm if we assume that we know how to compute the square root of an arbitrary positive number.

b. The difficulty here lies in computing $\sin A$. Since the formula says nothing about how it has to be computed, it should not be considered an algorithm. This is true even if we assume, as we did for the square root function, that we know how to compute the sine of a given angle. (There are several algorithms for doing this but only approximately, of course.) The problem is that the formula says nothing about how to compute angle A either.

c. The formula says nothing about how to compute h_a .

4. **Algorithm** *Quadratic*(a, b, c)
//The algorithm finds real roots of equation $ax^2 + bx + c = 0$
//Input: Real coefficients a, b, c
//Output: The real roots of the equation or a message about their absence
if $a \neq 0$
 $D \leftarrow b * b - 4 * a * c$
 if $D > 0$
 $temp \leftarrow 2 * a$
 $x1 \leftarrow (-b + \text{sqrt}(D))/temp$
 $x2 \leftarrow (-b - \text{sqrt}(D))/temp$

```

    return  $x_1, x_2$ 
else if  $D = 0$  return  $-b/(2 * a)$ 
else return 'no real roots'
else //  $a = 0$ 
    if  $b \neq 0$  return  $-c/b$ 
    else //  $a = b = 0$ 
        if  $c = 0$  return 'all real numbers'
        else return 'no real roots'

```

Note: See a more realistic algorithm for this problem in Section 11.4.

5. a. Divide the given number n by 2: the remainder r_n (0 or 1) will be the next (from right to left) digit of the binary representation in question. Replace n by the quotient of the last division and repeat this operation until n becomes 0.

b. **Algorithm** *Binary*(n)

//The algorithm implements the standard method for finding

//the binary expansion of a positive decimal integer

//Input: A positive decimal integer n

//Output: The list $b_k b_{k-1} \dots b_1 b_0$ of n 's binary digits

$k \leftarrow 0$

while $n \neq 0$

$b_k \leftarrow n \bmod 2$

$n \leftarrow \lfloor n/2 \rfloor$

$k \leftarrow k + 1$

6. n/a

7. a. π , as an irrational number, can be computed only approximately.

b. It is natural to consider, as an instance of this problem, computing π 's value with a given level of accuracy, say, with n correct decimal digits. With this interpretation, the problem has infinitely many instances.

8. n/a

9. The following improved version considers the same pair of elements only once and avoids recomputing the same expression in the innermost loop:

Algorithm *MinDistance2*($A[0..n-1]$)

//Input: An array $A[0..n-1]$ of numbers

//Output: The minimum distance d between two of its elements

```

 $dmin \leftarrow \infty$ 
for  $i \leftarrow 0$  to  $n - 2$  do
    for  $j \leftarrow i + 1$  to  $n - 1$  do
         $temp \leftarrow |A[i] - A[j]|$ 
        if  $temp < dmin$ 
             $dmin \leftarrow temp$ 
return  $dmin$ 

```

A faster algorithm is based on the idea of presorting (see Section 6.1).


10. Polya's general four-point approach is:
 1. Understand the problem
 2. Devise a plan
 3. Implement the plan
 4. Look back/check

Exercises 1.3

1. Consider the algorithm for the sorting problem that sorts an array by counting, for each of its elements, the number of smaller elements and then uses this information to put the element in its appropriate position in the sorted array:


Algorithm *ComparisonCountingSort*($A[0..n-1]$, $S[0..n-1]$)
//Sorts an array by comparison counting
//Input: Array $A[0..n-1]$ of orderable values
//Output: Array $S[0..n-1]$ of A 's elements sorted in nondecreasing order
for $i \leftarrow 0$ **to** $n-1$ **do**
 $Count[i] \leftarrow 0$
for $i \leftarrow 0$ **to** $n-2$ **do**
 for $j \leftarrow i+1$ **to** $n-1$ **do**
 if $A[i] < A[j]$
 $Count[j] \leftarrow Count[j] + 1$
 else $Count[i] \leftarrow Count[i] + 1$
for $i \leftarrow 0$ **to** $n-1$ **do**
 $S[Count[i]] \leftarrow A[i]$

- a. Apply this algorithm to sorting the list 60, 35, 81, 98, 14, 47.
 - b. Is this algorithm stable?
 - c. Is it in place?
2. Name the algorithms for the searching problem that you already know. Give a good succinct description of each algorithm in English. (If you know no such algorithms, use this opportunity to design one.)
 3. Design a simple algorithm for the string-matching problem.
 4. *Königsberg bridges* The Königsberg bridge puzzle is universally accepted as the problem that gave birth to graph theory. It was solved by the great Swiss-born mathematician Leonhard Euler (1707–1783). The problem asked whether one could, in a single stroll, cross all seven bridges of the city of Königsberg exactly once and return to a starting point. Following is a sketch of the river with its two islands and seven bridges:




- a. State the problem as a graph problem.

- b. Does this problem have a solution? If you believe it does, draw such a stroll; if you believe it does not, explain why and indicate the smallest number of new bridges that would be required to make such a stroll possible.
5. *Icosian Game* A century after Euler's discovery (see Problem 4), another famous puzzle—this one invented by the renown Irish mathematician Sir William Hamilton (1805-1865)—was presented to the world under the name of the Icosian Game. The game was played on a circular wooden board on which the following graph was carved:



- Find a ***Hamiltonian circuit***—a path that visits all the graph's vertices exactly once before returning to the starting vertex—for this graph.
6. Consider the following problem: Design an algorithm to determine the best route for a subway passenger to take from one designated station to another in a well-developed subway system similar to those in such cities as Washington, D.C., and London, UK.
- The problem's statement is somewhat vague, which is typical of real-life problems. In particular, what reasonable criterion can be used for defining the "best" route?
 - How would you model this problem by a graph?
7. a. Rephrase the traveling salesman problem in combinatorial object terms.
- b. Rephrase the graph-coloring problem in combinatorial object terms.
8. Consider the following map:



- a. Explain how we can use the graph-coloring problem to color the map so that no two neighboring regions are colored the same.
 - b. Use your answer to part (a) to color the map with the smallest number of colors.
9. Design an algorithm for the following problem: Given a set of n points in the Cartesian plane, determine whether all of them lie on the same circumference.
 10. Write a program that reads as its inputs the (x, y) coordinates of the endpoints of two line segments P_1Q_1 and P_2Q_2 and determines whether the segments have a common point.

Hints to Exercises 1.3

1. Trace the algorithm on the input given. Use the definitions of stability and being in place that were introduced in the section.
2. If you do not recall any searching algorithms, you should design a simple searching algorithm (without succumbing to the temptation to find one in the latter chapters of the book).
3. This algorithm is introduced later in the book but you should have no trouble to design it on your own.
4. If you have not encountered this problem in your previous courses, you may look up the answers on the Web or in a discrete structures textbook. The answers are, in fact, surprisingly simple.
5. No efficient algorithm for solving this problem for an arbitrary graph is known. This particular graph does have Hamiltonian circuits which are not difficult to find. (You need to find just one of them.)
6. a. Put yourself (mentally) in a passenger's place and ask yourself what criterion for the "best" route you would use. Then think of people that may have different needs.

b. The representation of the problem by a graph is straightforward. Give some thoughts though to stations where trains can be changed.
7. a. What are tours in the traveling salesman problem?

b. It would be natural to consider vertices colored the same color as elements of the same subset.
8. Create a graph whose vertices represent the map's regions. You will have to decide on the edges on your own.
9. Assume that the circumference in question exists and find its center first. Also, do not forget to give a special answer for $n \leq 2$.
10. Be careful not to miss some special cases of the problem.


Solutions to Exercises 1.3

1. a. Sorting 60, 35, 81, 98, 14, 47 by comparison counting will work as follows:

Array $A[0..5]$		60	35	81	98	14	47
Initially	$Count[]$	0	0	0	0	0	0
After pass $i = 0$	$Count[]$	3	0	1	1	0	0
After pass $i = 1$	$Count[]$		1	2	2	0	1
After pass $i = 2$	$Count[]$			4	3	0	1
After pass $i = 3$	$Count[]$				5	0	1
After pass $i = 4$	$Count[]$					0	2
Final state	$Count[]$	3	1	4	5	0	2

Array $S[0..5]$		14	35	47	60	81	98
-----------------	--	----	----	----	----	----	----


- b. The algorithm is not stable. Consider, as a counterexample, the result of its application to $1', 1''$.
- c. The algorithm is not in place because it uses two extra arrays of size n : $Count$ and S .
2. Answers may vary but most students should be familiar with sequential search, binary search, binary tree search and, possibly, hashing from their introductory programming courses.
3. Align the pattern with the beginning of the text. Compare the corresponding characters of the pattern and the text left-to right until either all the pattern characters are matched (then stop—the search is successful) or the algorithm runs out of the text's characters (then stop—the search is unsuccessful) or a mismatching pair of characters is encountered. In the latter case, shift the pattern one position to the right and resume the comparisons.
4. a. If we represent each of the river's banks and each of the two islands by vertices and the bridges by edges, we will get the following graph:



(This is, in fact, a multigraph, not a graph, because it has more than one edge between the same pair of vertices. But this doesn't matter for the issue at hand.) The question is whether there exists a path (i.e., a sequence of adjacent vertices) in this multigraph that traverses all the edges exactly once and returns to a starting vertex. Such paths are called *Eulerian circuits*; if a path traverses all the edges exactly once but does not return to its starting vertex, it is called an *Eulerian path*.


b. Euler proved that an Eulerian circuit exists in a connected (multi)graph if and only if all its vertices have even degrees, where the degree of a vertex is defined as the number of edges for which it is an endpoint. Also, an Eulerian path exists in a connected (multi)graph if and only if it has exactly two vertices of odd degrees; such a path must start at one of those two vertices and end at the other. Hence, for the multigraph of the puzzle, there exists neither an Eulerian circuit nor an Eulerian path because all its four vertices have odd degrees.

If we are to be satisfied with an Eulerian path, two of the multigraph's vertices must be made even. This can be accomplished by adding one new bridge connecting the same places as the existing bridges. For example, a new bridge between the two islands would make possible, among others, the walk $a - b - c - a - b - d - c - b - d$




If we want a walk that returns to its starting point, all the vertices in the

corresponding multigraph must be even. Since a new bridge/edge changes the parity of two vertices, at least two new bridges/edges will be needed. For example, here is one such “enhancement”:




This would make possible $a - b - c - a - b - d - c - b - d - a$, among several other such walks.

5. A Hamiltonian circuit is marked on the graph below:



6. a. At least three “reasonable” criteria come to mind: the fastest trip, a trip with the smallest number of train stops, and a trip that requires the smallest number of train changes. Note that the first criterion requires the information about the expected traveling time between stations and the time needed for train changes whereas the other two criteria do not require such information.
- b. A natural approach is to mimic subway plans by representing stations by vertices of a graph, with two vertices connected by an edge if there is a train line between the corresponding stations. If the time spent on changing a train is to be taken into account (e.g., because the station in question is on more than one line), the station should be represented by more than one vertex.

7. a. Find a permutation of n given cities for which the sum of the distances between consecutive cities in the permutation plus the distance between its last and first city is as small as possible.
- b. Partition all the graph's vertices into the smallest number of disjoint subsets so that there is no edge connecting vertices from the same subset.
8. a. Create a graph whose vertices represent the map's regions and the edges connect two vertices if and only if the corresponding regions have a common border (and therefore cannot be colored the same color). Here is the graph for the map given:



Solving the graph coloring problem for this graph yields the map's coloring with the smallest number of colors possible.

- b. Without loss of generality, we can assign colors 1 and 2 to vertices c and a , respectively. This forces the following color assignment to the remaining vertices: 3 to b , 2 to d , 3 to f , 4 to e . Thus, the smallest number of colors needed for this map is four.

Note: It's a well-known fact that *any* map can be colored in four colors or less. This problem—known as the Four-Color Problem—has remained unresolved for more than a century until 1976 when it was finally solved by the American mathematicians K. Appel and W. Haken by a combination of mathematical arguments and extensive computer use.

9. If $n = 2$, the answer is always “yes”; so, we may assume that $n \geq 3$. Select three points P_1 , P_2 , and P_3 from the set given. Write an equation of the perpendicular bisector l_1 of the line segment with the endpoints at P_1 and P_2 , which is the locus of points equidistant from P_1 and P_2 . Write an equation of the perpendicular bisector l_2 of the line segment with the endpoints at P_2 and P_3 , which is the locus of points equidistant from P_2 and P_3 . Find the coordinates (x, y) of the intersection point P of the lines l_1 and l_2 by solving the system of two equations in two unknowns x and y . (If the system has no solutions, return “no”: such a circumference does not exist.) Compute the distances (or much better yet the distance squares!) from P to each of the points P_i , $i = 3, 4, \dots, n$ and check whether all of them are the same: if they are, return “yes,” otherwise, return “no”.

Exercises 1.4

1. Describe how one can implement each of the following operations on an array so that the time it takes does not depend on the array's size n .
 - a. Delete the i th element of an array ($1 \leq i \leq n$).
 - b. Delete the i th element of a sorted array (the remaining array has to stay sorted, of course).
2. If you have to solve the searching problem for a list of n numbers, how can you take advantage of the fact that the list is known to be sorted? Give separate answers for
 - a. lists represented as arrays.
 - b. lists represented as linked lists.

3. a. Show the stack after each operation of the following sequence that starts with the empty stack:

push(a), push(b), pop, push(c), push(d), pop

- b. Show the queue after each operation of the following sequence that starts with the empty queue:

enqueue(a), enqueue(b), dequeue, enqueue(c), enqueue(d), dequeue

4. a. Let A be the adjacency matrix of an undirected graph. Explain what property of the matrix indicates that
 - i. the graph is complete.
 - ii. the graph has a loop, i.e., an edge connecting a vertex to itself.
 - iii. the graph has an isolated vertex, i.e., a vertex with no edges incident to it.
- b. Answer the same questions for the adjacency list representation.
5. Give a detailed description of an algorithm for transforming a free tree into a tree rooted at a given vertex of the free tree.
6. Prove the inequalities that bracket the height of a binary tree with n vertices:

$$\lfloor \log_2 n \rfloor \leq h \leq n - 1.$$

7. Indicate how the ADT priority queue can be implemented as
 - a. an (unsorted) array.

- b. a sorted array.
 - c. a binary search tree.
8. How would you implement a dictionary of a reasonably small size n if you knew that all its elements are distinct (e.g., names of 50 states of the United States)? Specify an implementation of each dictionary operation.
 9. For each of the following applications, indicate the most appropriate data structure:
 - a. answering telephone calls in the order of their known priorities.
 - b. sending backlog orders to customers in the order they have been received.
 - c. implementing a calculator for computing simple arithmetical expressions.
 10. *Anagram checking* Design an algorithm for checking whether two given words are anagrams, i.e., whether one word can be obtained by permuting the letters of the other. (For example, the words *tea* and *eat* are anagrams.)

Hints to Exercises 1.4

1. a. Take advantage of the fact that the array is not sorted.

b. We used this trick in implementing one of the algorithms in Section 1.1.
2. a. For a sorted array, there is a spectacularly efficient algorithm you almost certainly have heard about.

b. Unsuccessful searches can be made faster.
3. a. *Push*(x) puts x on the top of the stack, *pop* deletes the item from the top of the stack.

b. *Enqueue*(x) adds x to the rear of the queue, *dequeue* deletes the item from the front of the queue.
4. Just use the definitions of the graph properties in question and data structures involved.
5. There are two well-known algorithms that can solve this problem. The first uses a stack, the second uses a queue. Although these algorithms are discussed later in the book, do not miss this chance to discover them by yourself!
6. The inequality $h \leq n - 1$ follows immediately from the height's definition. The lower bound inequality follows from inequality $2^{h+1} - 1 \geq n$, which can be proved by considering the largest number of vertices a binary tree of height h can have.
7. You need to indicate how each of the three operations of the priority queue will be implemented.
8. Because of insertions and deletions, using an array of the dictionary's elements (sorted or unsorted) is not the best implementation possible.
9. You need to know about the postfix notation in order to answer one of these questions. (If you are not familiar with it, find the information on the Internet.)
10. There are several algorithms for this problem. Keep in mind that the words may contain multiple occurrences of the same letter.

Solutions to Exercises 1.4

1. a. Replace the i th element with the last element and decrease the array size by 1.

b. Replace the i th element with a special symbol that cannot be a value of the array's element (e.g., 0 for an array of positive numbers) to mark the i th position as empty. (This method is sometimes called the “lazy deletion”.)
2. a. Use binary search (see Section 4.3 if you are not familiar with this algorithm).

b. When searching in a sorted linked list, stop as soon as an element greater than or equal to the search key is encountered.

3. a.

$push(a)$		$push(b)$	b	pop		$push(c)$	c	$push(d)$	d	c	pop	c
	a		a		a		a		a		a	a

- b.

$enqueue(a)$	$enqueue(b)$	$dequeue$	$enqueue(c)$	$enqueue(d)$	$dequeue$
a	ab	b	bc	bcd	cd

4. a. For the adjacency matrix representation:

- i. A graph is complete if and only if all the elements of its adjacency matrix except those on the main diagonal are equal to 1, i.e., $A[i, j] = 1$ for every $1 \leq i, j \leq n, i \neq j$.

- ii. A graph has a loop if and only if its adjacency matrix has an element equal to 1 on its main diagonal, i.e., $A[i, i] = 1$ for some $1 \leq i \leq n$.

- iii. An (undirected, without loops) graph has an isolated vertex if and only if its adjacency matrix has an all-zero row.

- b. For the adjacency list representation:

- i. A graph is complete if and only if each of its linked lists contains all the other vertices of the graph.

- ii. A graph has a loop if and only if one of its adjacency lists contains the

vertex defining the list.

iii. An (undirected, without loops) graph has an isolated vertex if and only if one of its adjacency lists is empty.

5. The first algorithm works as follows. Mark a vertex to serve as the root of the tree, make it the root of the tree to be constructed, and initialize a stack with this vertex. Repeat the following operation until the stack becomes empty: If there is an unmarked vertex adjacent to the vertex on the top to the stack, mark the former vertex, attach it as a child of the top's vertex in the tree, and push it onto the stack; otherwise, pop the vertex off the top of the stack.

The second algorithm works as follows. Mark a vertex to serve as the root of the tree, make it the root of the tree to be constructed, and initialize a queue with this vertex. Repeat the following operations until the queue becomes empty: If there are unmarked vertices adjacent to the vertex at the front of the queue, mark all of them, attach them as children to the front vertex in the tree, and add them to the queue; then dequeue the queue.

6. Since the height is defined as the length of the longest simple path from the tree's root to its leaf, such a pass will include no more than n vertices, which is the total number of vertices in the tree. Hence, $h \leq n - 1$.

The binary tree of height h with the largest number of vertices is the full tree that has all its $h + 1$ levels filled with the largest number of vertices possible. The total number of vertices in such a tree is $\sum_{l=0}^h 2^l = 2^{h+1} - 1$. Hence, for any binary tree with n vertices and height h

$$2^{h+1} - 1 \geq n.$$

This implies that

$$2^{h+1} \geq n + 1$$

or, after taking binary logarithms of both hand sides and taking into account that $h + 1$ is an integer,

$$h + 1 \geq \lceil \log_2(n + 1) \rceil.$$

Since $\lceil \log_2(n + 1) \rceil = \lfloor \log_2 n \rfloor + 1$ (see Appendix A), we finally obtain

$$h + 1 \geq \lfloor \log_2 n \rfloor + 1 \text{ or } h \geq \lfloor \log_2 n \rfloor.$$

7. a. Insertion can be implemented by adding the new item after the array's last element. Finding the largest element requires a standard scan

through the array to find its largest element. Deleting the largest element $A[i]$ can be implemented by exchanging it with the last element and decreasing the array's size by 1.

b. We will assume that the array $A[0..n-1]$ representing the priority queue is sorted in ascending order. Inserting a new item of value v can be done by scanning the sorted array, say, left to right until an element $A[j] \geq v$ or the end of the array is reached. (A faster algorithm for finding a place for inserting a new element is binary search discussed in Section 4.3.) In the former case, the new item is inserted before $A[j]$ by first moving $A[n-1], \dots, A[j]$ one position to the right; in the latter case, the new item is simply appended after the last element of the array. Finding the largest element is done by simply returning the value of the last element of the sorted array. Deletion of the largest element is done by decreasing the array's size by one.

c. Insertion of a new element is done by using the standard algorithm for inserting a new element in a binary search tree: recursively, the new key is inserted in the left or right subtree depending on whether it is smaller or larger than the root's key. Finding the largest element will require finding the rightmost element in the binary tree by starting at the root and following the chain of the right children until a vertex with no right subtree is reached. The key of that vertex will be the largest element in question. Deleting it can be done by making the right pointer of its parent to point to the left child of the vertex being deleted; if the rightmost vertex has no left child, this pointer is made "null". Finally, if the rightmost vertex has no parent, i.e., if it happens to be the root of the tree, its left child becomes the new root; if there is no left child, the tree becomes empty.

8. Use a bit vector, i.e., an array on n bits in which the i th bit is 1 if the i th element of the underlying set is currently in the dictionary and 0 otherwise. The search, insertion, and deletion operations will require checking or changing a single bit of this array.
9. Use: (a) a priority queue; (b) a queue; (c) a stack (and *reverse Polish notation*—a clever way of representing arithmetical expressions without parentheses, which is usually studied in a data structures course).
10. The most straightforward solution is to search for each successive letter of the first word in the second one. If the search is successful, delete the first occurrence of the letter in the second word, stop otherwise.

Another solution is to sort the letters of each word and then compare

them in a simple parallel scan.

We can also generate and compare “letter vectors” of the given words: $V_w[i]$ = the number of occurrences of the alphabet’s i th letter in the word w . Such a vector can be generated by initializing all its components to 0 and then scanning the word and incrementing appropriate letter counts in the vector.

This file contains the exercises, hints, and solutions for Chapter 2 of the book "Introduction to the Design and Analysis of Algorithms," 2nd edition, by A. Levitin. The problems that might be challenging for at least some students are marked by \triangleright ; those that might be difficult for a majority of students are marked by \blacktriangleright .

Exercises 2.1

1. For each of the following algorithms, indicate (i) a natural size metric for its inputs; (ii) its basic operation; (iii) whether the basic operation count can be different for inputs of the same size:
 - a. computing the sum of n numbers
 - b. computing $n!$
 - c. finding the largest element in a list of n numbers
 - d. Euclid's algorithm
 - e. sieve of Eratosthenes
 - f. pen-and-pencil algorithm for multiplying two n -digit decimal integers
2.
 - a. Consider the definition-based algorithm for adding two n -by- n matrices. What is its basic operation? How many times is it performed as a function of the matrix order n ? As a function of the total number of elements in the input matrices?
 - b. Answer the same questions for the definition-based algorithm for matrix multiplication.
3. Consider a variation of sequential search that scans a list to return the number of occurrences of a given search key in the list. Will its efficiency differ from the efficiency of classic sequential search?
4.
 - a. *Glove selection* There are 22 gloves in a drawer: 5 pairs of red gloves, 4 pairs of yellow, and 2 pairs of green. You select the gloves in the dark and can check them only after a selection has been made. What is the smallest number of gloves you need to select to have at least one matching pair in the best case? in the worst case? (after [Mos01], #18)
 - b. *Missing socks* Imagine that after washing 5 distinct pairs of socks, you discover that two socks are missing. Of course, you would like to have the largest number of complete pairs remaining. Thus, you are left with 4 complete pairs in the best-case scenario and with 3 complete pairs in the worst case. Assuming that the probability of disappearance for each

of the 10 socks is the same, find the probability of the best-case scenario; the probability of the worst-case scenario; the number of pairs you should expect in the average case. (after [Mos01], #48)

5. a. \triangleright Prove formula (2.1) for the number of bits in the binary representation of a positive integer.
- b. What would be the analogous formula for the number of decimal digits?
- c. Explain why, within the accepted analysis framework, it does not matter whether we use binary or decimal digits in measuring n 's size.
6. Suggest how any sorting algorithm can be augmented in a way to make the best-case count of its key comparisons equal to just $n - 1$ (n is a list's size, of course). Do you think it would be a worthwhile addition to any sorting algorithm?
7. Gaussian elimination, the classic algorithm for solving systems of n linear equations in n unknowns, requires about $\frac{1}{3}n^3$ multiplications, which is the algorithm's basic operation.
 - a. How much longer should you expect Gaussian elimination to work on a system of 1000 equations versus a system of 500 equations?
 - b. You are considering buying a computer that is 1000 times faster than the one you currently have. By what factor will the faster computer increase the sizes of systems solvable in the same amount of time as on the old computer?
8. For each of the following functions, indicate how much the function's value will change if its argument is increased fourfold.

a. $\log_2 n$ b. \sqrt{n} c. n d. n^2 e. n^3 f. 2^n

9. Indicate whether the first function of each of the following pairs has a smaller, same, or larger order of growth (to within a constant multiple) than the second function.

a. $n(n + 1)$ and $2000n^2$	b. $100n^2$ and $0.01n^3$
c. $\log_2 n$ and $\ln n$	d. $\log_2^2 n$ and $\log_2 n^2$
e. 2^{n-1} and 2^n	f. $(n - 1)!$ and $n!$
10. *Invention of chess* According to a well-known legend, the game of chess was invented many centuries ago in northwestern India by a sage named Shashi. When he took his invention to his king, the king liked the game

so much that he offered the inventor any reward he wanted. Sashi asked for some grain to be obtained as follows: just a single grain of wheat was to be placed on the first square of the chess board, two on the second, four on the third, eight on the fourth, and so on, until all 64 squares had been filled. What would the ultimate result of this algorithm have been?

Hints to Exercises 2.1

1. The questions are indeed as straightforward as they appear, though some of them may have alternative answers. Also, keep in mind the caveat about measuring an integer's size.
2. a. The sum of two matrices is defined as the matrix whose elements are the sums of the corresponding elements of the matrices given.

b. Matrix multiplication requires two operations: multiplication and addition. Which of the two would you consider basic and why?
3. Will the algorithm's efficiency vary on different inputs of the same size?
4. a. Gloves are not socks: they can be right-handed and left-handed.

b. You have only two qualitatively different outcomes possible. Count the number of ways to get each of the two.
5. a. Prove first that if a positive decimal integer n has b digits in its binary representation then

$$2^{b-1} \leq n < 2^b.$$

Then take logarithms to base 2 of the terms in this inequality.

- b. The formula will be the same, with just one small adjustment to account for the different radix.
- c. How can we switch from one logarithm base to another?
6. Insert a verification of whether the problem is already solved.
7. A similar question was investigated in the section.
8. Use either the difference between or the ratio of $f(4n)$ and $f(n)$, whichever is more convenient for getting a compact answer. If it is possible, try to get an answer that does not depend on n .
9. If necessary, simplify the functions in question to single out terms defining their orders of growth to within a constant multiple. (We will discuss formal methods for answering such questions in the next section; however, these questions can be answered without knowledge of such methods.)
10. Use the formula $\sum_{i=0}^n 2^i = 2^{n+1} - 1$.

Solutions to Exercises 2.1

1. The answers are as follows.

a. (i) n ; (ii) addition of two numbers; (iii) no

b. (i) the magnitude of n , i.e., the number of bits in its binary representation; (ii) multiplication of two integers; (iii) no

c. (i) n ; (ii) comparison of two numbers; (iii) no (for the standard list scanning algorithm)

d. (i) either the magnitude of the larger of two input numbers, or the magnitude of the smaller of two input numbers, or the sum of the magnitudes of two input numbers; (ii) modulo division; (iii) yes

e. (i) the magnitude of n , i.e., the number of bits in its binary representation; (ii) elimination of a number from the list of remaining candidates to be prime; (iii) no

f. (i) n ; (ii) multiplication of two digits; (iii) no

2. a. Addition of two numbers. It's performed n^2 times (once for each of n^2 elements in the matrix being computed). Since the total number of elements in two given matrices is $N = 2n^2$, the total number of additions can also be expressed as $n^2 = N/2$.

b. Since on most computers multiplication takes longer than addition, multiplication is a better choice for being considered the basic operation of the standard algorithm for matrix multiplication. Each of n^2 elements of the product of two n -by- n matrices is computed as the scalar (dot) product of two vectors of size n , which requires n multiplications. The total number of multiplications is $n \cdot n^2 = n^3 = (N/2)^{3/2}$.

3. This algorithm will always make n key comparisons on every input of size n , whereas this number may vary between n and 1 for the classic version of sequential search.

4. a. The best-case number is, obviously, two. The worst-case number is twelve: one more than the number of gloves of one handedness.

b. There are just two possible outcomes here: the two missing socks make a pair (the best case) and the two missing socks do not make a pair (the worst case). The total number of different outcomes (the ways

to choose the missing socks) is $\binom{10}{2} = 45$. The number of best-case ones is 5; hence its probability is $\frac{5}{45} = \frac{1}{9}$. The number of worst-case ones is $45 - 5 = 40$; hence its probability is $\frac{40}{45} = \frac{8}{9}$. On average, you should expect $4 \cdot \frac{1}{9} + 3 \cdot \frac{8}{9} = \frac{28}{9} = 3\frac{1}{9}$ matching pairs.

5. a. The smallest positive integer that has b binary digits in its binary expansion is $\underbrace{10\dots 0}_{b-1}$, which is 2^{b-1} ; the largest positive integer that has b binary digits in its binary expansion is $\underbrace{11\dots 1}_{b-1}$, which is $2^{b-1} + 2^{b-2} + \dots + 1 = 2^b - 1$. Thus,

$$2^{b-1} \leq n < 2^b.$$

Hence

$$\log_2 2^{b-1} \leq \log_2 n < \log_2 2^b$$

or

$$b - 1 \leq \log_2 n < b.$$

These inequalities imply that $b - 1$ is the largest integer not exceeding $\log_2 n$. In other words, using the definition of the floor function, we conclude that

$$b - 1 = \lfloor \log_2 n \rfloor \text{ or } b = \lfloor \log_2 n \rfloor + 1.$$

- b. $B = \lfloor \log_{10} n \rfloor + 1$.

c. $b = \lfloor \log_2 n \rfloor + 1 \approx \log_2 n = \log_2 10 \log_{10} n \approx (\log_2 10)B$, where $B = \lfloor \log_{10} n \rfloor + 1$. That is, the two size metrics are about equal to within a constant multiple for large values of n .

6. Before applying a sorting algorithm, compare the adjacent elements of its input: if $a_i \leq a_{i+1}$ for every $i = 0, \dots, n - 2$, stop. Generally, it is not a worthwhile addition because it slows down the algorithm on all but very special inputs. Note that some sorting algorithms (notably bubble sort and insertion sort, which are discussed in Sections 3.1 and 5.1, respectively) intrinsically incorporate this test in the body of the algorithm.

7. a. $\frac{T(2n)}{T(n)} \approx \frac{c_M \frac{1}{3}(2n)^3}{c_M \frac{1}{3}n^3} = 8$, where c_M is the time of one multiplication.

b. We can estimate the running time for solving systems of order n on the old computer and that of order N on the new computer as $T_{old}(n) \approx c_M \frac{1}{3}n^3$ and $T_{new}(N) \approx 10^{-3}c_M \frac{1}{3}N^3$, respectively, where c_M is the time of one multiplication on the old computer. Replacing $T_{old}(n)$ and $T_{new}(N)$

by these estimates in the equation $T_{old}(n) = T_{new}(N)$ yields $c_M \frac{1}{3} n^3 \approx 10^{-3} c_M \frac{1}{3} N^3$ or $\frac{N}{n} \approx 10$.

8. a. $\log_2 4n - \log_2 n = (\log_2 4 + \log_2 n) - \log_2 n = 2$.

b. $\frac{\sqrt{4n}}{\sqrt{n}} = 2$.

c. $\frac{4n}{n} = 4$.

d. $\frac{(4n)^2}{n^2} = 4^2$.

e. $\frac{(4n)^3}{n^3} = 4^3$.

f. $\frac{2^{4n}}{2^n} = 2^{3n} = (2^n)^3$.

9. a. $n(n+1) \approx n^2$ has the same order of growth (quadratic) as $2000n^2$ to within a constant multiple.

b. $100n^2$ (quadratic) has a lower order of growth than $0.01n^3$ (cubic).

c. Since changing a logarithm's base can be done by the formula

$$\log_a n = \log_a b \log_b n,$$

all logarithmic functions have the same order of growth to within a constant multiple.

d. $\log_2^2 n = \log_2 n \log_2 n$ and $\log_2 n^2 = 2 \log_2 n$. Hence $\log_2^2 n$ has a higher order of growth than $\log_2 n^2$.

e. $2^{n-1} = \frac{1}{2} 2^n$ has the same order of growth as 2^n to within a constant multiple.

f. $(n-1)!$ has a lower order of growth than $n! = (n-1)!n$.

10. An unimaginable ruin: $\sum_{i=1}^{64} 2^{i-1} = \sum_{j=0}^{63} 2^j = 2^{64} - 1 \approx 1.8 \cdot 10^{19}$. (You may want to estimate the amount of space all this grain would have occupied.)

Exercises 2.2

- Use the most appropriate notation among O , Θ , and Ω to indicate the time efficiency class of sequential search (see Section 2.1)
 - in the worst case.
 - in the best case.
 - in the average case.
- Use the informal definitions of O , Θ , and Ω to determine whether the following assertions are true or false.
 - $n(n+1)/2 \in O(n^3)$
 - $n(n+1)/2 \in O(n^2)$
 - $n(n+1)/2 \in \Theta(n^3)$
 - $n(n+1)/2 \in \Omega(n)$
- For each of the following functions, indicate the class $\Theta(g(n))$ the function belongs to. (Use the simplest $g(n)$ possible in your answers.) Prove your assertions.
 - $(n^2 + 1)^{10}$
 - $\sqrt{10n^2 + 7n + 3}$
 - $2n \lg(n+2)^2 + (n+2)^2 \lg \frac{n}{2}$
 - $2^{n+1} + 3^{n-1}$
 - $\lfloor \log_2 n \rfloor$
- Table 2.1 contains values of several functions that often arise in analysis of algorithms. These values certainly suggest that the functions

$$\log n, \quad n, \quad n \log n, \quad n^2, \quad n^3, \quad 2^n, \quad n!$$

are listed in increasing order of their order of growth. Do these values prove this fact with mathematical certainty?

- Prove that the functions are indeed listed in increasing order of their order of growth.
- Order the following functions according to their order of growth (from the lowest to the highest):

$$(n-2)!, \quad 5 \lg(n+100)^{10}, \quad 2^{2n}, \quad 0.001n^4 + 3n^3 + 1, \quad \ln^2 n, \quad \sqrt[3]{n}, \quad 3^n.$$

- Prove that every polynomial of degree k , $p(n) = a_k n^k + a_{k-1} n^{k-1} + \dots + a_0$, with $a_k > 0$ belongs to $\Theta(n^k)$.
 - Prove that exponential functions a^n have different orders of growth for different values of base $a > 0$.

7. Prove (by using the definitions of the notations involved) or disprove (by giving a specific counterexample) the following assertions.
 - a. If $t(n) \in O(g(n))$, then $g(n) \in \Omega(t(n))$.
 - b. $\Theta(\alpha g(n)) = \Theta(g(n))$, where $\alpha > 0$.
 - c. $\Theta(g(n)) = O(g(n)) \cap \Omega(g(n))$.

d.▷ For any two nonnegative functions $t(n)$ and $g(n)$ defined on the set of nonnegative integers, either $t(n) \in O(g(n))$, or $t(n) \in \Omega(g(n))$, or both.
8. ▷ Prove the section's theorem for
 - a. Ω notation.
 - b. Θ notation.
9. We mentioned in this section that one can check whether all elements of an array are distinct by a two-part algorithm based on the array's presorting.
 - a. If the presorting is done by an algorithm with the time efficiency in $\Theta(n \log n)$, what will be the time efficiency class of the entire algorithm?
 - b. If the sorting algorithm used for presorting needs an extra array of size n , what will be the space efficiency class of the entire algorithm?
10. ▷ *Door in a wall* You are facing a wall that stretches infinitely in both directions. There is a door in the wall, but you know neither how far away nor in which direction. You can see the door only when you are right next to it. Design an algorithm that enables you to reach the door by walking at most $O(n)$ steps where n is the (unknown to you) number of steps between your initial position and the door. [Par95], #652

Hints to Exercises 2.2

1. Use the corresponding counts of the algorithm's basic operation (see Section 2.1) and the definitions of O , Θ , and Ω .
2. Establish the order of growth of $n(n+1)/2$ first and then use the informal definitions of O , Θ , and Ω . (Similar examples were given in the section.)
3. Simplify the functions given to single out the terms defining their orders of growth.
4. a. Check carefully the pertinent definitions.

b. Compute the ratio limits of every pair of consecutive functions on the list.
5. First simplify some of the functions. Then use the list of functions in Table 2.2 to "anchor" each of the functions given. Prove their final placement by computing appropriate limits.
6. a. You can prove this assertion either by computing an appropriate limit or by applying mathematical induction.

b. Compute $\lim_{n \rightarrow \infty} a_1^n / a_2^n$.
7. Prove the correctness of (a), (b), and (c) by using the appropriate definitions; construct a counterexample for (d) (e.g., by constructing two functions behaving differently for odd and even values of their arguments).
8. The proof of part (a) is similar to the one given for the theorem's assertion in Section 2.2. Of course, different inequalities need to be used to bound the sum from below.
9. Follow the analysis plan used in the text when the algorithm was mentioned for the first time.
10. You should walk intermittently left and right from your initial position until the door is reached.

Solutions to Exercises 2.2

1. a. Since $C_{worst}(n) = n$, $C_{worst}(n) \in \Theta(n)$.

b. Since $C_{best}(n) = 1$, $C_{best}(1) \in \Theta(1)$.

c. Since $C_{avg}(n) = \frac{p(n+1)}{2} + n(1-p) = (1-\frac{p}{2})n + \frac{p}{2}$ where $0 \leq p \leq 1$, $C_{avg}(n) \in \Theta(n)$.

2. $n(n+1)/2 \approx n^2/2$ is quadratic. Therefore

a. $n(n+1)/2 \in O(n^3)$ is true. b. $n(n+1)/2 \in O(n^2)$ is true.

c. $n(n+1)/2 \in \Theta(n^3)$ is false. d. $n(n+1)/2 \in \Omega(n)$ is true.

3. a. Informally, $(n^2+1)^{10} \approx (n^2)^{10} = n^{20} \in \Theta(n^{20})$ Formally,

$$\lim_{n \rightarrow \infty} \frac{(n^2+1)^{10}}{n^{20}} = \lim_{n \rightarrow \infty} \frac{(n^2+1)^{10}}{(n^2)^{10}} = \lim_{n \rightarrow \infty} \left(\frac{n^2+1}{n^2} \right)^{10} = \lim_{n \rightarrow \infty} \left(1 + \frac{1}{n^2} \right)^{10} = 1.$$

Hence $(n^2+1)^{10} \in \Theta(n^{20})$.

Note: An alternative proof can be based on the binomial formula and the assertion of Exercise 6a.

b. Informally, $\sqrt{10n^2+7n+3} \approx \sqrt{10n^2} = \sqrt{10}n \in \Theta(n)$. Formally,

$$\lim_{n \rightarrow \infty} \frac{\sqrt{10n^2+7n+3}}{n} = \lim_{n \rightarrow \infty} \sqrt{\frac{10n^2+7n+3}{n^2}} = \lim_{n \rightarrow \infty} \sqrt{10 + \frac{7}{n} + \frac{3}{n^2}} = \sqrt{10}.$$

Hence $\sqrt{10n^2+7n+3} \in \Theta(n)$.

c. $2n \lg(n+2)^2 + (n+2)^2 \lg \frac{n}{2} = 2n2 \lg(n+2) + (n+2)^2(\lg n - 1) \in \Theta(n \lg n) + \Theta(n^2 \lg n) = \Theta(n^2 \lg n)$.

d. $2^{n+1} + 3^{n-1} = 2^n 2 + 3^n \frac{1}{3} \in \Theta(2^n) + \Theta(3^n) = \Theta(3^n)$.

e. Informally, $\lfloor \log_2 n \rfloor \approx \log_2 n \in \Theta(\log n)$. Formally, by using the inequalities $x-1 < \lfloor x \rfloor \leq x$ (see Appendix A), we obtain an upper bound

$$\lfloor \log_2 n \rfloor \leq \log_2 n$$

and a lower bound

$$\lfloor \log_2 n \rfloor > \log_2 n - 1 \geq \log_2 n - \frac{1}{2} \log_2 n \text{ (for every } n \geq 4) = \frac{1}{2} \log_2 n.$$

Hence $\lfloor \log_2 n \rfloor \in \Theta(\log_2 n) = \Theta(\log n)$.

4. a. The order of growth and the related notations O , Ω , and Θ deal with the asymptotic behavior of functions as n goes to infinity. Therefore no specific values of functions within a finite range of n 's values, suggestive as they might be, can establish their orders of growth with mathematical certainty.

$$\text{b. } \lim_{n \rightarrow \infty} \frac{\log_2 n}{n} = \lim_{n \rightarrow \infty} \frac{(\log_2 n)'}{(n)'} = \lim_{n \rightarrow \infty} \frac{\frac{1}{n} \log_2 e}{1} = \log_2 e \lim_{n \rightarrow \infty} \frac{1}{n} = 0.$$

$$\lim_{n \rightarrow \infty} \frac{n}{n \log_2 n} = \lim_{n \rightarrow \infty} \frac{1}{\log_2 n} = 0.$$

$$\lim_{n \rightarrow \infty} \frac{n \log_2 n}{n^2} = \lim_{n \rightarrow \infty} \frac{\log_2 n}{n} = (\text{see the first limit of this exercise}) = 0.$$

$$\lim_{n \rightarrow \infty} \frac{n^2}{n^3} = \lim_{n \rightarrow \infty} \frac{1}{n} = 0.$$

$$\begin{aligned} \lim_{n \rightarrow \infty} \frac{n^3}{2^n} &= \lim_{n \rightarrow \infty} \frac{(n^3)'}{(2^n)'} = \lim_{n \rightarrow \infty} \frac{3n^2}{2^n \ln 2} = \frac{3}{\ln 2} \lim_{n \rightarrow \infty} \frac{n^2}{2^n} = \frac{3}{\ln 2} \lim_{n \rightarrow \infty} \frac{(n^2)'}{(2^n)'} \\ &= \frac{3}{\ln 2} \lim_{n \rightarrow \infty} \frac{2n}{2^n \ln 2} = \frac{6}{\ln^2 2} \lim_{n \rightarrow \infty} \frac{n}{2^n} = \frac{6}{\ln^2 2} \lim_{n \rightarrow \infty} \frac{(n)'}{(2^n)'} \\ &= \frac{6}{\ln^2 2} \lim_{n \rightarrow \infty} \frac{1}{2^n \ln 2} = \frac{6}{\ln^3 2} \lim_{n \rightarrow \infty} \frac{1}{2^n} = 0. \end{aligned}$$

$$\lim_{n \rightarrow \infty} \frac{2^n}{n!} = (\text{see Example 3 in the section}) 0.$$

5. $(n-2)! \in \Theta((n-2)!)$, $5 \lg(n+100)^{10} = 50 \lg(n+100) \in \Theta(\lg n)$, $2^{2n} = (2^2)^n \in \Theta(4^n)$, $0.001n^4 + 3n^3 + 1 \in \Theta(n^4)$, $\ln^2 n \in \Theta(\log^2 n)$, $\sqrt[3]{n} \in \Theta(n^{\frac{1}{3}})$, $3^n \in \Theta(3^n)$. The list of these functions ordered in increasing order of growth looks as follows:

$$5 \lg(n+100)^{10}, \quad \ln^2 n, \quad \sqrt[3]{n}, \quad 0.001n^4 + 3n^3 + 1, \quad 3^n, \quad 2^{2n}, \quad (n-2)!$$

6. a. $\lim_{n \rightarrow \infty} \frac{p(n)}{n^k} = \lim_{n \rightarrow \infty} \frac{a_k n^k + a_{k-1} n^{k-1} + \dots + a_0}{n^k} = \lim_{n \rightarrow \infty} \left(a_k + \frac{a_{k-1}}{n} + \dots + \frac{a_0}{n^k} \right) = a_k > 0.$

Hence $p(n) \in \Theta(n^k)$.

b.

$$\lim_{n \rightarrow \infty} \frac{a_1^n}{a_2^n} = \lim_{n \rightarrow \infty} \left(\frac{a_1}{a_2} \right)^n = \begin{cases} 0 & \text{if } a_1 < a_2 \Leftrightarrow a_1^n \in o(a_2^n) \\ 1 & \text{if } a_1 = a_2 \Leftrightarrow a_1^n \in \Theta(a_2^n) \\ \infty & \text{if } a_1 > a_2 \Leftrightarrow a_2^n \in o(a_1^n) \end{cases}$$

7. a. The assertion should be correct because it states that if the order of growth of $t(n)$ is smaller than or equal to the order of growth of $g(n)$, then

the order of growth of $g(n)$ is larger than or equal to the order of growth of $t(n)$. The formal proof is immediate, too:

$$t(n) \leq cg(n) \quad \text{for all } n \geq n_0, \text{ where } c > 0,$$

implies

$$\left(\frac{1}{c}\right)t(n) \leq g(n) \quad \text{for all } n \geq n_0.$$

b. The assertion that $\Theta(\alpha g(n)) = \Theta(g(n))$ should be true because $\alpha g(n)$ and $g(n)$ differ just by a positive constant multiple and, hence, by the definition of Θ , must have the same order of growth. The formal proof has to show that $\Theta(\alpha g(n)) \subseteq \Theta(g(n))$ and $\Theta(g(n)) \subseteq \Theta(\alpha g(n))$. Let $f(n) \in \Theta(\alpha g(n))$; we'll show that $f(n) \in \Theta(g(n))$. Indeed,

$$f(n) \leq c\alpha g(n) \quad \text{for all } n \geq n_0 \text{ (where } c > 0)$$

can be rewritten as

$$f(n) \leq c_1 g(n) \quad \text{for all } n \geq n_0 \text{ (where } c_1 = c\alpha > 0),$$

i.e., $f(n) \in \Theta(g(n))$.

Let now $f(n) \in \Theta(g(n))$; we'll show that $f(n) \in \Theta(\alpha g(n))$ for $\alpha > 0$. Indeed, if $f(n) \in \Theta(g(n))$,

$$f(n) \leq cg(n) \quad \text{for all } n \geq n_0 \text{ (where } c > 0)$$

and therefore

$$f(n) \leq \frac{c}{\alpha} \alpha g(n) = c_1 \alpha g(n) \quad \text{for all } n \geq n_0 \text{ (where } c_1 = \frac{c}{\alpha} > 0),$$

i.e., $f(n) \in \Theta(\alpha g(n))$.

c. The assertion is obviously correct (similar to the assertion that $a = b$ if and only if $a \leq b$ and $a \geq b$). The formal proof should show that $\Theta(g(n)) \subseteq O(g(n)) \cap \Omega(g(n))$ and that $O(g(n)) \cap \Omega(g(n)) \subseteq \Theta(g(n))$, which immediately follow from the definitions of O , Ω , and Θ .

d. The assertion is false. The following pair of functions can serve as a counterexample

$$t(n) = \begin{cases} n & \text{if } n \text{ is even} \\ n^2 & \text{if } n \text{ is odd} \end{cases} \quad \text{and} \quad g(n) = \begin{cases} n^2 & \text{if } n \text{ is even} \\ n & \text{if } n \text{ is odd} \end{cases}$$

8. a. We need to prove that if $t_1(n) \in \Omega(g_1(n))$ and $t_2(n) \in \Omega(g_2(n))$, then $t_1(n) + t_2(n) \in \Omega(\max\{g_1(n), g_2(n)\})$.

Proof Since $t_1(n) \in \Omega(g_1(n))$, there exist some positive constant c_1 and some nonnegative integer n_1 such that

$$t_1(n) \geq c_1 g_1(n) \quad \text{for all } n \geq n_1.$$

Since $t_2(n) \in \Omega(g_2(n))$, there exist some positive constant c_2 and some nonnegative integer n_2 such that

$$t_2(n) \geq c_2 g_2(n) \quad \text{for all } n \geq n_2.$$

Let us denote $c = \min\{c_1, c_2\}$ and consider $n \geq \max\{n_1, n_2\}$ so that we can use both inequalities. Adding the two inequalities above yields the following:

$$\begin{aligned} t_1(n) + t_2(n) &\geq c_1 g_1(n) + c_2 g_2(n) \\ &\geq c g_1(n) + c g_2(n) = c[g_1(n) + g_2(n)] \\ &\geq c \max\{g_1(n), g_2(n)\}. \end{aligned}$$

Hence $t_1(n) + t_2(n) \in \Omega(\max\{g_1(n), g_2(n)\})$, with the constants c and n_0 required by the Ω definition being $\min\{c_1, c_2\}$ and $\max\{n_1, n_2\}$, respectively.

b. The proof follows immediately from the theorem proved in the text (the Ω part), the assertion proved in part (a) of this exercise (the Ω part), and the definition of Θ (see Exercise 7c).

9. a. Since the running time of the sorting part of the algorithm will still dominate the running time of the second, it's the former that will determine the time efficiency of the entire algorithm. Formally, it follows from equality

$$\Theta(n \log n) + O(n) = \Theta(n \log n),$$

whose validity is easy to prove in the same manner as that of the section's theorem.

b. Since the second part of the algorithm will use no extra space, the space efficiency class will be determined by that of the first (sorting) part. Therefore, it will be in $\Theta(n)$.

10. The key idea here is to walk intermittently right and left going each time exponentially farther from the initial position. A simple implementation of this idea is to do the following until the door is reached: For $i = 0, 1, \dots$, make 2^i steps to the right, return to the initial position, make 2^i steps to

the left, and return to the initial position again. Let $2^{k-1} < n \leq 2^k$. The number of steps this algorithm will need to find the door can be estimated above as follows:

$$\sum_{i=0}^{k-1} 4 \cdot 2^i + 3 \cdot 2^k = 4(2^k - 1) + 3 \cdot 2^k < 7 \cdot 2^k = 14 \cdot 2^{k-1} < 14n.$$

Hence the number of steps made by the algorithm is in $O(n)$. (Note: It is not difficult to improve the multiplicative constant with a better algorithm.)

Exercises 2.3

1. Compute the following sums.

a. $1 + 3 + 5 + 7 + \dots + 999$

b. $2 + 4 + 8 + 16 + \dots + 1024$

c. $\sum_{i=3}^{n+1} 1$ d. $\sum_{i=3}^{n+1} i$ e. $\sum_{i=0}^{n-1} i(i+1)$

f. $\sum_{j=1}^n 3^{j+1}$ g. $\sum_{i=1}^n \sum_{j=1}^n ij$ h. $\sum_{i=0}^{n-1} 1/i(i+1)$

2. Find the order of growth of the following sums.

a. $\sum_{i=0}^{n-1} (i^2+1)^2$ b. $\sum_{i=2}^{n-1} \lg i^2$

c. $\sum_{i=1}^n (i+1)2^{i-1}$ d. $\sum_{i=0}^{n-1} \sum_{j=0}^{i-1} (i+j)$

Use the $\Theta(g(n))$ notation with the simplest function $g(n)$ possible.

3. The sample variance of n measurements x_1, x_2, \dots, x_n can be computed as

$$\frac{\sum_{i=1}^n (x_i - \bar{x})^2}{n-1} \text{ where } \bar{x} = \frac{\sum_{i=1}^n x_i}{n}$$

or

$$\frac{\sum_{i=1}^n x_i^2 - (\sum_{i=1}^n x_i)^2/n}{n-1}.$$

Find and compare the number of divisions, multiplications, and additions/subtractions (additions and subtractions are usually bunched together) that are required for computing the variance according to each of these formulas.

4. Consider the following algorithm.

```
Algorithm Mystery( $n$ )  
//Input: A nonnegative integer  $n$   
 $S \leftarrow 0$   
for  $i \leftarrow 1$  to  $n$  do  
     $S \leftarrow S + i * i$   
return  $S$ 
```

- a. What does this algorithm compute?
- b. What is its basic operation?
- c. How many times is the basic operation executed?

d. What is the efficiency class of this algorithm?

e. Suggest an improvement or a better algorithm altogether and indicate its efficiency class. If you cannot do it, try to prove that, in fact, it cannot be done.

5. Consider the following algorithm.

```
Algorithm Secret( $A[0..n-1]$ )  
//Input: An array  $A[0..n-1]$  of  $n$  real numbers  
 $minval \leftarrow A[0]$ ;  $maxval \leftarrow A[0]$   
for  $i \leftarrow 1$  to  $n-1$  do  
    if  $A[i] < minval$   
         $minval \leftarrow A[i]$   
    if  $A[i] > maxval$   
         $maxval \leftarrow A[i]$   
return  $maxval - minval$ 
```

Answer questions a–e of Problem 4 about this algorithm.

6. Consider the following algorithm.

```
Algorithm Enigma( $A[0..n-1, 0..n-1]$ )  
//Input: A matrix  $A[0..n-1, 0..n-1]$  of real numbers  
for  $i \leftarrow 0$  to  $n-2$  do  
    for  $j \leftarrow i+1$  to  $n-1$  do  
        if  $A[i, j] \neq A[j, i]$   
            return false  
return true
```

Answer the questions a–e of Problem 4 about this algorithm.

7. Improve the implementation of the matrix multiplication algorithm (see Example 3) by reducing the number of additions made by the algorithm. What effect will this change have on the algorithm's efficiency?
8. Determine the asymptotic order of growth for the total number of times all the doors are toggled in the Locker Doors puzzle (Problem 11 in Exercises 1.1).
9. Prove the formula

$$\sum_{i=1}^n i = 1 + 2 + \dots + n = \frac{n(n+1)}{2}$$

either by mathematical induction or by following the insight of a 10-year old schoolboy named Karl Friedrich Gauss (1777–1855) who grew up to become one of the greatest mathematicians of all times.

10. Consider the following version of an important algorithm that we will study later in the book.

Algorithm $GE(A[0..n-1, 0..n])$

//Input: An n -by- $n+1$ matrix $A[0..n-1, 0..n]$ of real numbers

for $i \leftarrow 0$ **to** $n-2$ **do**

for $j \leftarrow i+1$ **to** $n-1$ **do**


for $k \leftarrow i$ **to** n **do**

$A[j, k] \leftarrow A[j, k] - A[i, k] * A[j, i] / A[i, i]$


a.▷ Find the time efficiency class of this algorithm.

b.▷ What glaring inefficiency does this pseudocode contain and how can it be eliminated to speed the algorithm up?


11. *von Neumann's neighborhood* How many one-by-one squares are generated by the algorithm that starts with a single square and on each of its n iterations adds new squares all round the outside. How many one-by-one squares are generated on the n th iteration? [Gar99, p.88] (In the parlance of cellular automata theory, the answer is the number of cells in the von Neumann neighborhood of range n .) The results for $n = 0, 1$, and 2 are illustrated below:



$n = 0$



$n = 1$



$n = 2$

Hints to Exercises 2.3

1. Use the common summation formulas and rules listed in Appendix A. You may need to perform some simple algebraic operations before applying them.
2. Find a sum among those in Appendix A that looks similar to the sum in question and try to transform the latter to the former. Note that you do not have to get a closed-end formula for a sum before establishing its order of growth.
3. Just follow the formulas in question.
4. a. Tracing the algorithm to get its output for a few small values of n (e.g., $n = 1, 2$, and 3) should help if you need it.

b. We faced the same question for the examples discussed in the text. One of them is particularly pertinent here.

c. Follow the plan outlined in this section.

d. As a function of n , the answer should follow immediately from your answer to part c. You may also want to give an answer as a function of the number of bits in the n 's representation (why?).

e. Have you not encountered this sum somewhere?
5. a. Tracing the algorithm to get its output for a few small values of n (e.g., $n = 1, 2$, and 3) should help if you need it.

b. We faced the same question for the examples discussed in the text. One of them is particularly pertinent here.

c. You can either follow the section's plan by setting up and computing a sum or answer the question directly. (Try to do both.)

d. Your answer will immediately follow from the answer to part c).

e. Does the algorithm always have to make two comparisons on each iteration? This idea can be developed further to get a more significant improvement than the obvious one—try to do it for a two-element array and then generalize the insight. But can we hope to find an algorithm with a better than linear efficiency?
6. a. Elements $A[i, j]$ and $A[j, i]$ are symmetric with respect to the main diagonal of the matrix.

b. There is just one candidate here.

- c. You may investigate the worst case only.
 - d. Your answer will immediately follow from the answer to part c.
 - e. Compare the problem the algorithm solves with the way it does this.
7. Computing a sum of n numbers can be done with $n - 1$ additions. How many does the algorithm make in computing each element of the product matrix?
 8. Set up a sum for the number of times all the doors are toggled and find its asymptotic order of growth by using some properties from Appendix A.
 9. For the general step of the proof by induction, use the formula

$$\sum_{i=1}^n i = \sum_{i=1}^{n-1} i + n.$$

The young Gauss computed the sum $1 + 2 + \dots + 99 + 100$ by noticing that it can be computed as the sum of 50 pairs, each with the same sum.

10. a. Setting up a sum should pose no difficulties. Using the standard summation formulas and rules will require more effort than in the previous examples, however.
 - b. Optimize the algorithm's innermost loop.
11. Set up a sum for the number of squares after n iterations of the algorithm and then simplify it to get a closed-form answer.

Solutions to Exercises 2.3

$$1. \text{ a. } 1+3+5+7+\dots+999 = \sum_{i=1}^{500} (2i-1) = \sum_{i=1}^{500} 2i - \sum_{i=1}^{500} 1 = 2 \frac{500 \cdot 501}{2} - 500 = 250,000.$$

(Or by using the formula for the sum of odd integers: $\sum_{i=1}^{500} (2i-1) = 500^2 = 250,000$.)

Or by using the formula for the sum of the arithmetic progression with $a_1 = 1$, $a_n = 999$, and $n = 500$: $\frac{(a_1+a_n)n}{2} = \frac{(1+999)500}{2} = 250,000$.)

$$\text{b. } 2+4+8+16+\dots+1,024 = \sum_{i=1}^{10} 2^i = \sum_{i=0}^{10} 2^i - 1 = (2^{11} - 1) - 1 = 2,046.$$

(Or by using the formula for the sum of the geometric series with $a = 2$, $q = 2$, and $n = 9$: $a \frac{q^{n+1}-1}{q-1} = 2 \frac{2^{10}-1}{2-1} = 2,046$.)

$$\text{c. } \sum_{i=3}^{n+1} 1 = (n+1) - 3 + 1 = n - 1.$$

$$\text{d. } \sum_{i=3}^{n+1} i = \sum_{i=0}^{n+1} i - \sum_{i=0}^2 i = \frac{(n+1)(n+2)}{2} - 3 = \frac{n^2+3n-4}{2}.$$

$$\begin{aligned} \text{e. } \sum_{i=0}^{n-1} i(i+1) &= \sum_{i=0}^{n-1} (i^2 + i) = \sum_{i=0}^{n-1} i^2 + \sum_{i=0}^{n-1} i = \frac{(n-1)n(2n-1)}{6} + \frac{(n-1)n}{2} \\ &= \frac{(n^2-1)n}{3}. \end{aligned}$$

$$\text{f. } \sum_{j=1}^n 3^{j+1} = 3 \sum_{j=1}^n 3^j = 3 \left[\sum_{j=0}^n 3^j - 1 \right] = 3 \left[\frac{3^{n+1}-1}{3-1} - 1 \right] = \frac{3^{n+2}-9}{2}.$$

$$\begin{aligned} \text{g. } \sum_{i=1}^n \sum_{j=1}^n ij &= \sum_{i=1}^n i \sum_{j=1}^n j = \sum_{i=1}^n i \frac{n(n+1)}{2} = \frac{n(n+1)}{2} \sum_{i=1}^n i = \frac{n(n+1)}{2} \frac{n(n+1)}{2} \\ &= \frac{n^2(n+1)^2}{4}. \end{aligned}$$

$$\text{h. } \sum_{i=1}^n 1/i(i+1) = \sum_{i=1}^n \left(\frac{1}{i} - \frac{1}{i+1} \right)$$

$$= \left(\frac{1}{1} - \frac{1}{2} \right) + \left(\frac{1}{2} - \frac{1}{3} \right) + \dots + \left(\frac{1}{n-1} - \frac{1}{n} \right) + \left(\frac{1}{n} - \frac{1}{n+1} \right) = 1 - \frac{1}{n+1} = \frac{n}{n+1}.$$

(This is a special case of the so-called *telescoping series*—see Appendix

$$A - \sum_{i=l}^u (a_i - a_{i-1}) = a_u - a_{l-1}.)$$

$$\begin{aligned} 2. \text{ a. } \sum_{i=0}^{n-1} (i^2 + 1)^2 &= \sum_{i=0}^{n-1} (i^4 + 2i^2 + 1) = \sum_{i=0}^{n-1} i^4 + 2 \sum_{i=0}^{n-1} i^2 + \sum_{i=0}^{n-1} 1 \\ &\in \Theta(n^5) + \Theta(n^3) + \Theta(n) = \Theta(n^5) \text{ (or just } \sum_{i=0}^{n-1} (i^2 + 1)^2 \approx \sum_{i=0}^{n-1} i^4 \in \Theta(n^5)). \end{aligned}$$

$$\begin{aligned} \text{b. } \sum_{i=2}^{n-1} \log_2 i^2 &= \sum_{i=2}^{n-1} 2 \log_2 i = 2 \sum_{i=2}^{n-1} \log_2 i = 2 \sum_{i=1}^n \log_2 i - 2 \log_2 n \\ &\in 2\Theta(n \log n) - \Theta(\log n) = \Theta(n \log n). \end{aligned}$$

$$\begin{aligned}
\text{c. } \sum_{i=1}^n (i+1)2^{i-1} &= \sum_{i=1}^n i2^{i-1} + \sum_{i=1}^n 2^{i-1} = \frac{1}{2} \sum_{i=1}^n i2^i + \sum_{j=0}^{n-1} 2^j \\
&\in \Theta(n2^n) + \Theta(2^n) = \Theta(n2^n) \text{ (or } \sum_{i=1}^n (i+1)2^{i-1} \approx \frac{1}{2} \sum_{i=1}^n i2^i \in \Theta(n2^n)). \\
\text{d. } \sum_{i=0}^{n-1} \sum_{j=0}^{i-1} (i+j) &= \sum_{i=0}^{n-1} [\sum_{j=0}^{i-1} i + \sum_{j=0}^{i-1} j] = \sum_{i=0}^{n-1} [i^2 + \frac{(i-1)i}{2}] = \sum_{i=0}^{n-1} [\frac{3}{2}i^2 - \frac{1}{2}i] \\
&= \frac{3}{2} \sum_{i=0}^{n-1} i^2 - \frac{1}{2} \sum_{i=0}^{n-1} i \in \Theta(n^3) - \Theta(n^2) = \Theta(n^3).
\end{aligned}$$

3. For the first formula: $D(n) = 2$, $M(n) = n$, $A(n) + S(n) = [(n-1) + (n-1)] + (n+1) = 3n-1$.

For the second formula: $D(n) = 2$, $M(n) = n+1$, $A(n) + S(n) = [(n-1) + (n-1)] + 2 = 2n$.

4. a. Computes $S(n) = \sum_{i=1}^n i^2$.

b. Multiplication (or, if multiplication and addition are assumed to take the same amount of time, either of the two).

c. $C(n) = \sum_{i=1}^n 1 = n$.

d. $C(n) = n \in \Theta(n)$. Since the number of bits $b = \lfloor \log_2 n \rfloor + 1 \approx \log_2 n$ and hence $n \approx 2^b$, $C(n) \approx 2^b \in \Theta(2^b)$.

e. Use the formula $\sum_{i=1}^n i^2 = \frac{n(n+1)(2n+1)}{6}$ to compute the sum in $\Theta(1)$ time (which assumes that the time of arithmetic operations stay constant irrespective of the size of the operations' operands).

5. a. Computes the range, i.e., the difference between the array's largest and smallest elements.

b. An element comparison.

c. $C(n) = \sum_{i=1}^{n-1} 2 = 2(n-1)$.

d. $\Theta(n)$.

e. An obvious improvement for some inputs (but not for the worst case) is to replace the two if-statements by the following one:

if $A[i] < \text{minval}$ $\text{minval} \leftarrow A[i]$

else if $A[i] > \text{maxval}$ $\text{maxval} \leftarrow A[i]$.

Another improvement, both more subtle and substantial, is based on the observation that it is more efficient to update the minimum and maximum values seen so far not for each element but for a pair of two consecutive elements. If two such elements are compared with each other first, the updates will require only two more comparisons for the total of three comparisons per pair. Note that the same improvement can be obtained by a divide-and-conquer algorithm (see Chapter 4).

6. a. The algorithm returns “true” if its input matrix is symmetric and “false” if it is not.

- b. Comparison of two matrix elements.

$$\begin{aligned} \text{c. } C_{\text{worst}}(n) &= \sum_{i=0}^{n-2} \sum_{j=i+1}^{n-1} 1 = \sum_{i=0}^{n-2} [(n-1) - (i+1) + 1] \\ &= \sum_{i=0}^{n-2} (n-1-i) = (n-1) + (n-2) + \dots + 1 = \frac{(n-1)n}{2}. \end{aligned}$$

- d. Quadratic: $C_{\text{worst}}(n) \in \Theta(n^2)$ (or $C(n) \in O(n^2)$).

- e. The algorithm is optimal because any algorithm that solves this problem must, in the worst case, compare $(n-1)n/2$ elements in the upper-triangular part of the matrix with their symmetric counterparts in the lower-triangular part, which is all this algorithm does.

7. Replace the body of the j loop by the following fragment:

```

C[i, j] ← A[i, 0] * B[0, j]
for k ← 1 to n - 1 do
    C[i, j] ← C[i, j] + A[i, k] * B[k, j]

```

This will decrease the number of additions from n^3 to $n^3 - n^2$, but the number of multiplications will still be n^3 . The algorithm’s efficiency class will remain cubic.

8. Let $T(n)$ be the total number of times all the doors are toggled. The problem statement implies that

$$T(n) = \sum_{i=1}^n \lfloor n/i \rfloor.$$

Since $x - 1 < \lfloor x \rfloor \leq x$ and $\sum_{i=1}^n 1/i \approx \ln n + \gamma$, where $\gamma = 0.5772\dots$ (see

Appendix A),

$$T(n) \leq \sum_{i=1}^n n/i = n \sum_{i=1}^n 1/i \approx n(\ln n + \gamma) \in \Theta(n \log n).$$

Similarly,

$$T(n) > \sum_{i=1}^n (n/i - 1) = n \sum_{i=1}^n 1/i - \sum_{i=1}^n 1 \approx n(\ln n + \gamma) - n \in \Theta(n \log n).$$

This implies that $T(n) \in \Theta(n \log n)$.

Note: Alternatively, we could use the formula for approximating sums by definite integrals (see Appendix A):

$$T(n) \leq \sum_{i=1}^n n/i = n(1 + \sum_{i=2}^n 1/i) \leq n(1 + \int_1^n \frac{1}{x} dx) = n(1 + \ln n) \in \Theta(n \log n)$$

and

$$T(n) > \sum_{i=1}^n (n/i - 1) = n \sum_{i=1}^n 1/i - \sum_{i=1}^n 1 \geq n \int_1^{n+1} \frac{1}{x} dx - n = n \ln(n+1) - n \in \Theta(n \log n)$$

9. Here is a proof by mathematical induction that $\sum_{i=1}^n i = \frac{n(n+1)}{2}$ for every positive integer n .

(i) Basis step: For $n = 1$, $\sum_{i=1}^n i = \sum_{i=1}^1 i = 1$ and $\frac{n(n+1)}{2} \Big|_{n=1} = \frac{1(1+1)}{2} = 1$.

(ii) Inductive step: Assume that $\sum_{i=1}^n i = \frac{n(n+1)}{2}$ for a positive integer n .

We need to show that then $\sum_{i=1}^{n+1} i = \frac{(n+1)(n+2)}{2}$. This is obtained as follows:

$$\sum_{i=1}^{n+1} i = \sum_{i=1}^n i + (n+1) = \frac{n(n+1)}{2} + (n+1) = \frac{n(n+1) + 2(n+1)}{2} = \frac{(n+1)(n+2)}{2}.$$

The young Gauss computed the sum

$$1 + 2 + \dots + 99 + 100$$

by noticing that it can be computed as the sum of 50 pairs, each with the sum 101:

$$1 + 100 = 2 + 99 = \dots = 50 + 51 = 101.$$

Hence the entire sum is equal to $50 \cdot 101 = 5,050$. (The well-known historic anecdote claims that his teacher gave this assignment to a class to keep

the class busy.) The Gauss idea can be easily generalized to an arbitrary n by adding

$$S(n) = 1 + 2 + \dots + (n - 1) + n$$

and

$$S(n) = n + (n - 1) + \dots + 2 + 1$$

to obtain

$$2S(n) = (n + 1)n \text{ and hence } S(n) = \frac{n(n + 1)}{2}.$$

10. a. The number of multiplications $M(n)$ and the number of divisions $D(n)$ made by the algorithm are given by the same sum:

$$\begin{aligned} M(n) &= D(n) = \sum_{i=0}^{n-2} \sum_{j=i+1}^{n-1} \sum_{k=i}^n 1 = \sum_{i=0}^{n-2} \sum_{j=i+1}^{n-1} (n - i + 1) = \\ &= \sum_{i=0}^{n-2} (n - i + 1)(n - 1 - (i + 1) + 1) = \sum_{i=0}^{n-2} (n - i + 1)(n - i - 1) \\ &= (n + 1)(n - 1) + n(n - 2) + \dots + 3 * 1 \\ &= \sum_{j=1}^{n-1} (j + 2)j = \sum_{j=1}^{n-1} j^2 + \sum_{j=1}^{n-1} 2j = \frac{(n - 1)n(2n - 1)}{6} + 2\frac{(n - 1)n}{2} \\ &= \frac{n(n - 1)(2n + 5)}{6} \approx \frac{1}{3}n^3 \in \Theta(n^3). \end{aligned}$$

b. The inefficiency is the repeated evaluation of the ratio $A[j, i] / A[i, i]$ in the algorithm's innermost loop, which, in fact, does not change with the loop variable k . Hence, this *loop invariant* can be computed just once before entering this loop: $temp \leftarrow A[j, i] / A[i, i]$; the innermost loop is then changed to

$$A[j, k] \leftarrow A[j, k] - A[i, k] * temp.$$

This change eliminates the most expensive operation of the algorithm, the division, from its innermost loop. The running time gain obtained by this change can be estimated as follows:

$$\frac{T_{old}(n)}{T_{new}(n)} \approx \frac{c_M \frac{1}{3}n^3 + c_D \frac{1}{3}n^3}{c_M \frac{1}{3}n^3} = \frac{c_M + c_D}{c_M} = \frac{c_D}{c_M} + 1,$$

where c_D and c_M are the time for one division and one multiplication, respectively.

11. The answer can be obtained by a straightforward evaluation of the sum

$$2 \sum_{i=1}^n (2i - 1) + (2n + 1) = 2n^2 + 2n + 1.$$

(One can also get the closed-form answer by noting that the cells on the alternating diagonals of the von Neumann neighborhood of range n compose two squares of sizes $n + 1$ and n , respectively.)

Exercises 2.4

1. Solve the following recurrence relations.

a. $x(n) = x(n-1) + 5$ for $n > 1$, $x(1) = 0$

b. $x(n) = 3x(n-1)$ for $n > 1$, $x(1) = 4$

c. $x(n) = x(n-1) + n$ for $n > 0$, $x(0) = 0$

d. $x(n) = x(n/2) + n$ for $n > 1$, $x(1) = 1$ (solve for $n = 2^k$)

e. $x(n) = x(n/3) + 1$ for $n > 1$, $x(1) = 1$ (solve for $n = 3^k$)

2. Set up and solve a recurrence relation for the number of calls made by $F(n)$, the recursive algorithm for computing $n!$.
3. Consider the following recursive algorithm for computing the sum of the first n cubes: $S(n) = 1^3 + 2^3 + \dots + n^3$.

Algorithm $S(n)$

//Input: A positive integer n

//Output: The sum of the first n cubes

if $n = 1$ **return** 1

else return $S(n-1) + n * n * n$

- a. Set up and solve a recurrence relation for the number of times the algorithm's basic operation is executed.
 - b. How does this algorithm compare with the straightforward nonrecursive algorithm for computing this function?
4. Consider the following recursive algorithm.

Algorithm $Q(n)$

//Input: A positive integer n

if $n = 1$ **return** 1

else return $Q(n-1) + 2 * n - 1$

- a. Set up a recurrence relation for this function's values and solve it to determine what this algorithm computes.
- b. Set up a recurrence relation for the number of multiplications made by this algorithm and solve it.
- c. Set up a recurrence relation for the number of additions/subtractions made by this algorithm and solve it.

5. a. In the original version of the Tower of Hanoi puzzle, as it was published by Edouard Lucas, a French mathematician, in the 1890s, the world will end after 64 disks have been moved from a mystical Tower of Brahma. Estimate the number of years it will take if monks could move one disk per minute. (Assume that monks do not eat, sleep, or die.)
 - b. How many moves are made by the i th largest disk ($1 \leq i \leq n$) in this algorithm?
 - c.► Design a nonrecursive algorithm for the Tower of Hanoi puzzle.
6. a. Prove that the exact number of additions made by the recursive algorithm *BinRec*(n) for an arbitrary positive integer n is $\lfloor \log_2 n \rfloor$.
 - b. Set up a recurrence relation for the number of additions made by the nonrecursive version of this algorithm (see Section 2.3, Example 4) and solve it.
7. a. Design a recursive algorithm for computing 2^n for any nonnegative integer n that is based on the formula: $2^n = 2^{n-1} + 2^{n-1}$.
 - b. Set up a recurrence relation for the number of additions made by the algorithm and solve it.
 - c. Draw a tree of recursive calls for this algorithm and count the number of calls made by the algorithm.
 - d. Is it a good algorithm for solving this problem?
8. Consider the following recursive algorithm.

Algorithm *Min1*($A[0..n-1]$)
 //Input: An array $A[0..n-1]$ of real numbers
if $n = 1$ **return** $A[0]$
else $temp \leftarrow \text{Min1}(A[0..n-2])$
 if $temp \leq A[n-1]$ **return** $temp$
 else return $A[n-1]$

- a. What does this algorithm compute?
 - b. Set up a recurrence relation for the algorithm's basic operation count and solve it.
9. Consider another algorithm for solving the same problem as the one in Problem 8 which recursively divides an array into two halves:

call $\text{Min2}(A[0..n-1])$ where

Algorithm $\text{Min2}(A[l..r])$
if $l = r$ **return** $A[l]$
else $\text{temp1} \leftarrow \text{Min2}(A[l.. \lfloor (l+r)/2 \rfloor])$
 $\text{temp2} \leftarrow \text{Min2}(A[\lfloor (l+r)/2 \rfloor + 1..r])$
if $\text{temp1} \leq \text{temp2}$ **return** temp1
else return temp2

a. Set up a recurrence relation for the algorithm's basic operation and solve it.

b. Which of the algorithms Min1 or Min2 is faster? Can you suggest an algorithm for the problem they solve that would be more efficient than either of them?

10. The determinant of an n -by- n matrix

$$A = \begin{bmatrix} a_{11} & a_{1n} \\ a_{21} & a_{2n} \\ \vdots & \vdots \\ a_{n1} & a_{nn} \end{bmatrix},$$

denoted $\det A$, can be defined as a_{11} for $n = 1$ and, for $n > 1$, by the recursive formula

$$\det A = \sum_{j=1}^n s_j a_{1j} \det A_j,$$

where s_j is $+1$ if j is odd and -1 if j is even, a_{1j} is the element in row 1 and column j , and A_j is the $(n-1)$ -by- $(n-1)$ matrix obtained from matrix A by deleting its row 1 and column j .

a.▷ Set up a recurrence relation for the number of multiplications made by the algorithm implementing this recursive definition.

b.▷ Without solving the recurrence, what can you say about the solution's order of growth as compared to $n!$?

9. *von Neumann neighborhood revisited* Find the number of cells in the von Neumann neighborhood of range n (see Problem 11 in Exercises 2.3) by setting up and solving a recurrence relation.

Hints to Exercises 2.4

1. Each of these recurrences can be solved by the method of backward substitutions.
2. The recurrence relation in question is almost identical to the recurrence relation for the number of multiplications, which was set up and solved in the section.
3. a. The question is similar to that about the efficiency of the recursive algorithm for computing $n!$.

b. Write a pseudocode for the nonrecursive algorithm and determine its efficiency.
4. a. Note that you are asked here about a recurrence for the function's values, not about a recurrence for the number of times its operation is executed. Just follow the pseudocode to set it up. It is easier to solve this recurrence by *forward* substitutions (see Appendix B).

b. This question is very similar to one we have already discussed.

c. You may want to include the subtraction needed to decrease n .
5. a. Use the formula for the number of disk moves derived in the section.

b. Solve the problem for 3 disks to investigate the number of moves made by each of the disks. Then generalize the observations and prove their validity for the general case of n disks.

c. If you fail, do not feel discouraged: though a nonrecursive algorithm for this problem is not complicated, it is not easy to discover. As a consolation, find a solution on the Web.
6. a. Consider separately the cases of even and odd values of n and show that for both of them $\lfloor \log_2 n \rfloor$ satisfies the recurrence relation and its initial condition.

b. Just follow the algorithm's pseudocode.
7. a. Use the formula $2^n = 2^{n-1} + 2^{n-1}$ without simplifying it; do not forget to provide a condition for stopping your recursive calls.

b. A similar algorithm was investigated in Section 2.4.

c. A similar question was investigated in Section 2.4.

d. A bad efficiency class of an algorithm by itself does not mean that

the algorithm is bad. For example, the classic algorithm for the Tower of Hanoi puzzle is optimal despite its exponential-time efficiency. Therefore a claim that a particular algorithm is not good requires a reference to a better one.

8. a. Tracing the algorithm for $n = 1$ and $n = 2$ should help.
b. It is very similar to one of the examples discussed in this section.
9. Tracing the algorithm for $n = 1$ and $n = 4$ should help.
10. a. Use the definition's formula to get the recurrence relation for the number of multiplications made by the algorithm.
b. Investigate the right-hand side of the recurrence relation. Computing the first few values of $M(n)$ may be helpful, too.
11. You might want to use the neighborhood's symmetry to obtain a simple formula for the number of squares added to the neighborhood on the n th iteration of the algorithm.

Solutions to Exercises 2.4

1. a. $x(n) = x(n-1) + 5$ for $n > 1$, $x(1) = 0$

$$\begin{aligned}
 x(n) &= x(n-1) + 5 \\
 &= [x(n-2) + 5] + 5 = x(n-2) + 5 \cdot 2 \\
 &= [x(n-3) + 5] + 5 \cdot 2 = x(n-3) + 5 \cdot 3 \\
 &= \dots \\
 &= x(n-i) + 5 \cdot i \\
 &= \dots \\
 &= x(1) + 5 \cdot (n-1) = 5(n-1).
 \end{aligned}$$

Note: The solution can also be obtained by using the formula for the n term of the arithmetical progression:

$$x(n) = x(1) + d(n-1) = 0 + 5(n-1) = 5(n-1).$$

b. $x(n) = 3x(n-1)$ for $n > 1$, $x(1) = 4$

$$\begin{aligned}
 x(n) &= 3x(n-1) \\
 &= 3[3x(n-2)] = 3^2x(n-2) \\
 &= 3^2[3x(n-3)] = 3^3x(n-3) \\
 &= \dots \\
 &= 3^i x(n-i) \\
 &= \dots \\
 &= 3^{n-1}x(1) = 4 \cdot 3^{n-1}.
 \end{aligned}$$

Note: The solution can also be obtained by using the formula for the n term of the geometric progression:

$$x(n) = x(1)q^{n-1} = 4 \cdot 3^{n-1}.$$

c. $x(n) = x(n-1) + n$ for $n > 0$, $x(0) = 0$

$$\begin{aligned}
 x(n) &= x(n-1) + n \\
 &= [x(n-2) + (n-1)] + n = x(n-2) + (n-1) + n \\
 &= [x(n-3) + (n-2)] + (n-1) + n = x(n-3) + (n-2) + (n-1) + n \\
 &= \dots \\
 &= x(n-i) + (n-i+1) + (n-i+2) + \dots + n \\
 &= \dots \\
 &= x(0) + 1 + 2 + \dots + n = \frac{n(n+1)}{2}.
 \end{aligned}$$

d. $x(n) = x(n/2) + n$ for $n > 1$, $x(1) = 1$ (solve for $n = 2^k$)

$$\begin{aligned}
x(2^k) &= x(2^{k-1}) + 2^k \\
&= [x(2^{k-2}) + 2^{k-1}] + 2^k = x(2^{k-2}) + 2^{k-1} + 2^k \\
&= [x(2^{k-3}) + 2^{k-2}] + 2^{k-1} + 2^k = x(2^{k-3}) + 2^{k-2} + 2^{k-1} + 2^k \\
&= \dots \\
&= x(2^{k-i}) + 2^{k-i+1} + 2^{k-i+2} + \dots + 2^k \\
&= \dots \\
&= x(2^{k-k}) + 2^1 + 2^2 + \dots + 2^k = 1 + 2^1 + 2^2 + \dots + 2^k \\
&= 2^{k+1} - 1 = 2 \cdot 2^k - 1 = 2n - 1.
\end{aligned}$$

e. $x(n) = x(n/3) + 1$ for $n > 1$, $x(1) = 1$ (solve for $n = 3^k$)

$$\begin{aligned}
x(3^k) &= x(3^{k-1}) + 1 \\
&= [x(3^{k-2}) + 1] + 1 = x(3^{k-2}) + 2 \\
&= [x(3^{k-3}) + 1] + 2 = x(3^{k-3}) + 3 \\
&= \dots \\
&= x(3^{k-i}) + i \\
&= \dots \\
&= x(3^{k-k}) + k = x(1) + k = 1 + \log_3 n.
\end{aligned}$$

2. $C(n) = C(n-1) + 1$, $C(0) = 1$ (there is a call but no multiplications when $n = 0$).

$$\begin{aligned}
C(n) &= C(n-1) + 1 = [C(n-2) + 1] + 1 = C(n-2) + 2 = \dots \\
&= C(n-i) + i = \dots = C(0) + n = 1 + n.
\end{aligned}$$

3. a. Let $M(n)$ be the number of multiplications made by the algorithm. We have the following recurrence relation for it:

$$M(n) = M(n-1) + 2, \quad M(1) = 0.$$

We can solve it by backward substitutions:

$$\begin{aligned}
 M(n) &= M(n-1) + 2 \\
 &= [M(n-2) + 2] + 2 = M(n-2) + 2 + 2 \\
 &= [M(n-3) + 2] + 2 + 2 = M(n-3) + 2 + 2 + 2 \\
 &= \dots \\
 &= M(n-i) + 2i \\
 &= \dots \\
 &= M(1) + 2(n-1) = 2(n-1).
 \end{aligned}$$

b. Here is a pseudocode for the nonrecursive option:

Algorithm *NonrecS*(n)

//Computes the sum of the first n cubes nonrecursively

//Input: A positive integer n

//Output: The sum of the first n cubes.

$S \leftarrow 1$

for $i \leftarrow 2$ **to** n **do**

$S \leftarrow S + i * i * i$

return S

The number of multiplications made by this algorithm will be

$$\sum_{i=2}^n 2 = 2 \sum_{i=2}^n 1 = 2(n-1).$$

This is exactly the same number as in the recursive version, but the nonrecursive version doesn't carry the time and space overhead associated with the recursion's stack.

4. a. $Q(n) = Q(n-1) + 2n - 1$ for $n > 1$, $Q(1) = 1$.

Computing the first few terms of the sequence yields the following:

$$Q(2) = Q(1) + 2 \cdot 2 - 1 = 1 + 2 \cdot 2 - 1 = 4;$$

$$Q(3) = Q(2) + 2 \cdot 3 - 1 = 4 + 2 \cdot 3 - 1 = 9;$$

$$Q(4) = Q(3) + 2 \cdot 4 - 1 = 9 + 2 \cdot 4 - 1 = 16.$$

Thus, it appears that $Q(n) = n^2$. We'll check this hypothesis by substituting this formula into the recurrence equation and the initial condition. The left hand side yields $Q(n) = n^2$. The right hand side yields

$$Q(n-1) + 2n - 1 = (n-1)^2 + 2n - 1 = n^2.$$

The initial condition is verified immediately: $Q(1) = 1^2 = 1$.

b. $M(n) = M(n-1) + 1$ for $n > 1$, $M(1) = 0$. Solving it by backward substitutions (it's almost identical to the factorial example—see Example 1 in the section) or by applying the formula for the n th term of an arithmetical progression yields $M(n) = n - 1$.

c. Let $C(n)$ be the number of additions and subtractions made by the algorithm. The recurrence for $C(n)$ is $C(n) = C(n-1) + 3$ for $n > 1$, $C(1) = 0$. Solving it by backward substitutions or by applying the formula for the n th term of an arithmetical progression yields $C(n) = 3(n-1)$.

Note: If we don't include in the count the subtractions needed to decrease n , the recurrence will be $C(n) = C(n-1) + 2$ for $n > 1$, $C(1) = 0$. Its solution is $C(n) = 2(n-1)$.

5. a. The number of moves is given by the formula: $M(n) = 2^n - 1$. Hence

$$\frac{2^{64} - 1}{60 \cdot 24 \cdot 365} \approx 3.5 \cdot 10^{13} \text{ years}$$

vs. the age of the Universe estimated to be about $13 \cdot 10^9$ years.

b. Observe that for every move of the i th disk, the algorithm first moves the tower of all the disks smaller than it to another peg (this requires one move of the $(i+1)$ st disk) and then, after the move of the i th disk, this smaller tower is moved on the top of it (this again requires one move of the $(i+1)$ st disk). Thus, for each move of the i th disk, the algorithm moves the $(i+1)$ st disk exactly twice. Since for $i = 1$, the number of moves is equal to 1, we have the following recurrence for the number of moves made by the i th disk:

$$m(i+1) = 2m(i) \quad \text{for } 1 \leq i < n, \quad m(1) = 1.$$

Its solution is $m(i) = 2^{i-1}$ for $i = 1, 2, \dots, n$. (The easiest way to obtain this formula is to use the formula for the generic term of a geometric progression.) Note that the answer agrees nicely with the formula for the total number of moves:

$$M(n) = \sum_{i=1}^n m(i) = \sum_{i=1}^n 2^{i-1} = 1 + 2 + \dots + 2^{n-1} = 2^n - 1.$$

c. Here is a description of the well-known nonrecursive algorithm by P. Buneman and L. Levy as presented in [Har92], p.115. Arrange the three pegs in a circle and do the following repeatedly until, prior to Step 2, all disks are correctly piled on some other peg:

Step 1: Move the smallest disk from its current peg to the next peg in clockwise direction.

Step 2: Make the only move possible that does not involve the smallest disk.

Note: There are other iterative algorithms for the Tower of Hanoi puzzle (see, in particular, <http://hanoitower.mkolar.org/algo.html> for examples of such algorithms and http://www.cs.wm.edu/~pkstoc/h_papers.html for an extensive bibliography on this problem).

6. a. We'll verify by substitution that $A(n) = \lfloor \log_2 n \rfloor$ satisfies the recurrence for the number of additions

$$A(n) = A(\lfloor n/2 \rfloor) + 1 \quad \text{for every } n > 1.$$

Let n be even, i.e., $n = 2k$.

The left-hand side is:

$$A(n) = \lfloor \log_2 n \rfloor = \lfloor \log_2 2k \rfloor = \lfloor \log_2 2 + \log_2 k \rfloor = (1 + \lfloor \log_2 k \rfloor) = \lfloor \log_2 k \rfloor + 1.$$

The right-hand side is:

$$A(\lfloor n/2 \rfloor) + 1 = A(\lfloor 2k/2 \rfloor) + 1 = A(k) + 1 = \lfloor \log_2 k \rfloor + 1.$$

Let n be odd, i.e., $n = 2k + 1$.

The left-hand side is:

$$\begin{aligned} A(n) &= \lfloor \log_2 n \rfloor = \lfloor \log_2(2k + 1) \rfloor = \text{using } \lfloor \log_2 x \rfloor = \lceil \log_2(x + 1) \rceil - 1 \\ &= \lceil \log_2(2k + 2) \rceil - 1 = \lceil \log_2 2(k + 1) \rceil - 1 \\ &= \lceil \log_2 2 + \log_2(k + 1) \rceil - 1 = 1 + \lceil \log_2(k + 1) \rceil - 1 = \lfloor \log_2 k \rfloor + 1. \end{aligned}$$

The right-hand side is:

$$A(\lfloor n/2 \rfloor) + 1 = A(\lfloor (2k + 1)/2 \rfloor) + 1 = A(\lfloor k + 1/2 \rfloor) + 1 = A(k) + 1 = \lfloor \log_2 k \rfloor + 1.$$

The initial condition is verified immediately: $A(1) = \lfloor \log_2 1 \rfloor = 0$.

- b. The recurrence relation for the number of additions is identical to the one for the recursive version:

$$A(n) = A(\lfloor n/2 \rfloor) + 1 \quad \text{for } n > 1, \quad A(1) = 0,$$

with the solution $A(n) = \lfloor \log_2 n \rfloor + 1$.

7. a. **Algorithm** $Power(n)$

//Computes 2^n recursively by the formula $2^n = 2^{n-1} + 2^{n-1}$

//Input: A nonnegative integer n

//Output: Returns 2^n


if $n = 0$ **return** 1

else return $Power(n - 1) + Power(n - 1)$

b. $A(n) = 2A(n-1) + 1, \quad A(0) = 0.$

$$\begin{aligned}
 A(n) &= 2A(n-1) + 1 \\
 &= 2[2A(n-2) + 1] + 1 = 2^2 A(n-2) + 2 + 1 \\
 &= 2^2 [2A(n-3) + 1] + 2 + 1 = 2^3 A(n-3) + 2^2 + 2 + 1 \\
 &= \dots \\
 &= 2^i A(n-i) + 2^{i-1} + 2^{i-2} + \dots + 1 \\
 &= \dots \\
 &= 2^n A(0) + 2^{n-1} + 2^{n-2} + \dots + 1 = 2^{n-1} + 2^{n-2} + \dots + 1 = 2^n - 1.
 \end{aligned}$$

c. The tree of recursive calls for this algorithm looks as follows:



Note that it has one extra level compared to the similar tree for the Tower of Hanoi puzzle.

d. It's a very bad algorithm because it is vastly inferior to the algorithm that simply multiplies an accumulator by 2 n times, not to mention much more efficient algorithms discussed later in the book. Even if only additions are allowed, adding two 2^{n-1} times is better than this algorithm.

8. a. The algorithm computes the value of the smallest element in a given array.

b. The recurrence for the number of key comparisons is

$$C(n) = C(n-1) + 1 \quad \text{for } n > 1, \quad C(1) = 0.$$

Solving it by backward substitutions yields $C(n) = n - 1$.

9. a. The recurrence for the number of key comparisons is

$$C(n) = C(\lceil n/2 \rceil) + C(\lfloor n/2 \rfloor) + 1 \quad \text{for } n > 1, \quad C(1) = 0.$$

Solving it for $n = 2^k$ by backward substitutions yields the following:

$$\begin{aligned}
C(2^k) &= 2C(2^{k-1}) + 1 \\
&= 2[2C(2^{k-2}) + 1] + 1 = 2^2C(2^{k-2}) + 2 + 1 \\
&= 2^2[2C(2^{k-3}) + 1] + 2 + 1 = 2^3C(2^{k-3}) + 2^2 + 2 + 1 \\
&= \dots \\
&= 2^iC(2^{k-i}) + 2^{i-1} + 2^{i-2} + \dots + 1 \\
&= \dots \\
&= 2^kC(2^{k-k}) + 2^{k-1} + 2^{k-2} + \dots + 1 = 2^k - 1 = n - 1.
\end{aligned}$$

One can verify that this solution is, in fact, valid for every (positive integer) value of n by checking it separately for the even ($n = 2k$) and odd ($n = 2k + 1$) cases.

b. A simple standard scan through the array in question requires the same number of key comparisons while avoiding the overhead associated with recursive calls. (It is clear, however, that any algorithm for this problem must be in $\Omega(n)$.)

10. a. Let $M(n)$ be the number of multiplications made by the algorithm based on the formula $\det A = \sum_{j=1}^n s_j a_{1j} \det A_j$. If we don't include multiplications by s_j , which are just ± 1 , then

$$M(n) = \sum_{j=1}^n (M(n-1) + 1),$$

i.e.,

$$M(n) = n(M(n-1) + 1) \quad \text{for } n > 1 \quad \text{and} \quad M(1) = 0.$$

b. Since $M(n) = nM(n-1) + n$, the sequence $M(n)$ grows to infinity at least as fast as the factorial function defined by $F(n) = nF(n-1)$.

8. The number of squares added on the n th iteration to each of the four symmetric sides of the von Neumann neighborhood is equal to n . Hence we obtain the following recurrence for $S(n)$, the total number of squares in the neighborhood after the n th iteration:

$$S(n) = S(n-1) + 4n \quad \text{for } n > 0 \quad \text{and} \quad S(0) = 1.$$

Solving the recurrence by backward substitutions yields the following:

$$\begin{aligned}
 S(n) &= S(n-1) + 4n \\
 &= [S(n-2) + 4(n-1)] + 4n = S(n-2) + 4(n-1) + 4n \\
 &= [S(n-3) + 4(n-2)] + 4(n-1) + 4n = S(n-3) + 4(n-2) + 4(n-1) + 4n \\
 &= \dots \\
 &= S(n-i) + 4(n-i+1) + 4(n-i+2) + \dots + 4n \\
 &= \dots \\
 &= S(0) + 4 \cdot 1 + 4 \cdot 2 + \dots + 4n = 1 + 4(1 + 2 + \dots + n) = 1 + 4n(n+1)/2 = 2n^2 -
 \end{aligned}$$

Exercises 2.5

1. Find a Web site dedicated to applications of the Fibonacci numbers and study it.
2. Check by direct substitutions that the function $\frac{1}{\sqrt{5}}(\phi^n - \hat{\phi}^n)$ indeed satisfies recurrence (2.6) and initial conditions (2.7).
3. The maximum values of the Java primitive types `int` and `long` are $2^{31} - 1$ and $2^{63} - 1$, respectively. Find the smallest n for which the n th Fibonacci number is not going to fit in a memory allocated for
 - a. the type `int`.
 - b. the type `long`.
4. *Climbing stairs* Find the number of different ways to climb an n -stair staircase if each step is either one or two stairs. (For example, a 3-stair staircase can be climbed three ways: 1-1-1, 1-2, and 2-1.) [TUC80]
5. Consider the recursive definition-based algorithm for computing the n th Fibonacci number $F(n)$. Let $C(n)$ and $Z(n)$ be the number of times $F(1)$ and $F(0)$, respectively, are computed. Prove that
 - a. $C(n) = F(n)$.
 - b. $Z(n) = F(n - 1)$.
6. Improve algorithm *Fib* of the text so that it requires only $\Theta(1)$ space.
7. Prove the equality


$$\begin{bmatrix} F(n-1) & F(n) \\ F(n) & F(n+1) \end{bmatrix} = \begin{bmatrix} 0 & 1 \\ 1 & 1 \end{bmatrix}^n \quad \text{for } n \geq 1.$$

8. \triangleright How many modulo divisions are made by Euclid's algorithm on two consecutive Fibonacci numbers $F(n)$ and $F(n-1)$ as the algorithm's input?
9. \triangleright a. Prove **Cassini's identity**:

$$F(n+1)F(n-1) - [F(n)]^2 = (-1)^n \quad \text{for } n \geq 1.$$

b. *Disappearing square* Consider the following paradox, which is based on Cassini's identity. Take an 8-by-8 chessboard (more generally, any $F(n)$ -by- $F(n)$ board divided into $[F(n)]^2$ squares). Cut it into two trapezoids and two triangles as shown in the left portion of the figure. Then reassemble it as shown in the right portion of the figure. The area of the left rectangle is $8 \times 8 = 64$ squares, while the area of the right triangle is

$13 \times 5 = 65$ squares. Explain the paradox.



10. In the language of your choice, implement two algorithms for computing the last five digits of the n th Fibonacci number that are based on (a) the recursive definition-based algorithm $F(n)$; (b) the iterative definition-based algorithm $Fib(n)$. Perform an experiment to find the largest value of n for which your programs run under 1 minute on your computer.

Hints to Exercises 2.5

1. Use a search engine.
2. It is easier to substitute ϕ^n and $\hat{\phi}^n$ into the recurrence equation separately. Why will this suffice?
3. Use an approximate formula for $F(n)$ to find the smallest values of n to exceed the numbers given.
4. There are several ways to solve this problem. The most elegant of them makes it possible to put the problem in this section.
5. Set up the recurrence relations for $C(n)$ and $Z(n)$, with appropriate initial conditions, of course.
6. All the information needed on each iteration of the algorithm is the values of the last two consecutive Fibonacci numbers. Modify the algorithm to take advantage of this fact.
7. Prove it by mathematical induction.
8. Consider first a small example such as computing $\gcd(13, 8)$.
9. a. Prove it by mathematical induction.

b. The renowned authors of “Concrete Mathematics”—R. Graham, D. Knuth, and O. Patashnik—note in their well-known book on discrete mathematics [Gra94] that this was one of Lewis Carroll’s favorite puzzles. Unfortunately, they also note that “The paradox is explained because...well, magic tricks aren’t supposed to be explained.” Given the prominence of these scientists, I have no choice but to obey and refrain from giving you a hint.
10. The last k digits of an integer N can be obtained by computing $N \bmod 10^k$. Performing all operations of your algorithms modulo 10^k (see Appendix A) will enable you to circumvent the exponential growth of the Fibonacci numbers. Also note that Section 2.6 is devoted to a general discussion of the empirical analysis of algorithms.

Solutions to Exercises 2.5

1. n/a

2. On substituting ϕ^n into the left-hand side of the equation, we obtain $F(n) - F(n-1) - F(n-2) = \phi^n - \phi^{n-1} - \phi^{n-2} = \phi^{n-2}(\phi^2 - \phi - 1) = 0$ because ϕ is one of the roots of the characteristic equation $r^2 - r - 1 = 0$. The verification of $\hat{\phi}^n$ works out for the same reason. Since the equation $F(n) - F(n-1) - F(n-2) = 0$ is homogeneous and linear, any linear combination of its solutions ϕ^n and $\hat{\phi}^n$, i.e., any sequence of the form $\alpha\phi^n + \beta\hat{\phi}^n$ will also be a solution to $F(n) - F(n-1) - F(n-2) = 0$. In particular, it will be the case for the Fibonacci sequence $\frac{1}{\sqrt{5}}\phi^n - \frac{1}{\sqrt{5}}\hat{\phi}^n$. Both initial conditions are checked out in a quite straightforward manner (but, of course, not individually for ϕ^n and $\hat{\phi}^n$).

3. a. The question is to find the smallest value of n such that $F(n) > 2^{31} - 1$. Using the formula $F(n) = \frac{1}{\sqrt{5}}\phi^n$ rounded to the nearest integer, we get (approximately) the following inequality:

$$\frac{1}{\sqrt{5}}\phi^n > 2^{31} - 1 \quad \text{or} \quad \phi^n > \sqrt{5}(2^{31} - 1).$$

After taking natural logarithms of both hand sides, we obtain

$$n > \frac{\ln(\sqrt{5}(2^{31} - 1))}{\ln \phi} \approx 46.3.$$

Thus, the answer is $n = 47$.

- b. Similarly, we have to find the smallest value of n such that $F(n) > 2^{63} - 1$. Thus,

$$\frac{1}{\sqrt{5}}\phi^n > 2^{63} - 1, \quad \text{or} \quad \phi^n > \sqrt{5}(2^{63} - 1)$$

or, after taking natural logarithms of both hand sides,

$$n > \frac{\ln(\sqrt{5}(2^{63} - 1))}{\ln \phi} \approx 92.4.$$

Thus, the answer is $n = 93$.

4. Let $W(n)$ be the number of different ways to climb an n -stair staircase. $W(n-1)$ of them start with a one-stair climb and $W(n-2)$ of them start with a two-stair climb. Thus,

$$W(n) = W(n-1) + W(n-2) \quad \text{for } n \geq 3, \quad W(1) = 1, \quad W(2) = 2.$$

Solving this recurrence either “from scratch” or better yet noticing that the solution runs one step ahead of the canonical Fibonacci sequence $F(n)$, we obtain $W(n) = F(n+1)$ for $n \geq 1$.

5. Since $F(n)$ is computed recursively by the formula $F(n) = F(n-1) + F(n-2)$, the recurrence equations for $C(n)$ and $Z(n)$ will be the same as the recurrence for $F(n)$. The initial conditions will be:

$$C(0) = 0, \quad C(1) = 1 \quad \text{and} \quad Z(0) = 1, \quad Z(1) = 0$$

for $C(n)$ and $Z(n)$, respectively. Therefore, since both the recurrence equation and the initial conditions for $C(n)$ and $F(n)$ are the same, $C(n) = F(n)$. As to the assertion that $Z(n) = F(n-1)$, it is easy to see that it should be the case since the sequence $Z(n)$ looks as follows:

$$1, 0, 1, 1, 2, 3, 5, 8, 13, 21, \dots,$$

i.e., it is the same as the Fibonacci numbers shifted one position to the right. This can be formally proved by checking that the sequence $F(n-1)$ (in which $F(-1)$ is defined as 1) satisfies the recurrence relation

$$Z(n) = Z(n-1) + Z(n-2) \quad \text{for } n > 1 \quad \text{and} \quad Z(0) = 1, \quad Z(1) = 0.$$

It can also be proved either by mathematical induction or by deriving an explicit formula for $Z(n)$ and showing that this formula is the same as the value of the explicit formula for $F(n)$ with n replaced by $n-1$.

6. **Algorithm** *Fib2*(n)

//Computes the n -th Fibonacci number using just two variables

//Input: A nonnegative integer n

//Output: The n -th Fibonacci number

$u \leftarrow 0; \quad v \leftarrow 1$

for $i \leftarrow 2$ **to** n **do**

$v \leftarrow v + u$

$u \leftarrow v - u$

if $n = 0$ **return** 0

else return v

7. (i) The validity of the equality for $n = 1$ follows immediately from the definition of the Fibonacci sequence.
(ii) Assume that

$$\begin{bmatrix} F(n-1) & F(n) \\ F(n) & F(n+1) \end{bmatrix} = \begin{bmatrix} 0 & 1 \\ 1 & 1 \end{bmatrix}^n \quad \text{for a positive integer } n.$$

We need to show that then

$$\begin{bmatrix} F(n) & F(n+1) \\ F(n+1) & F(n+2) \end{bmatrix} = \begin{bmatrix} 0 & 1 \\ 1 & 1 \end{bmatrix}^{n+1}.$$

Indeed,

$$\begin{aligned} \begin{bmatrix} 0 & 1 \\ 1 & 1 \end{bmatrix}^{n+1} &= \begin{bmatrix} 0 & 1 \\ 1 & 1 \end{bmatrix} \begin{bmatrix} 0 & 1 \\ 1 & 1 \end{bmatrix}^n \\ &= \begin{bmatrix} 0 & 1 \\ 1 & 1 \end{bmatrix} \begin{bmatrix} F(n-1) & F(n) \\ F(n) & F(n+1) \end{bmatrix} = \begin{bmatrix} F(n) & F(n+1) \\ F(n+1) & F(n+2) \end{bmatrix}. \end{aligned}$$

8. The principal observation here is the fact that Euclid's algorithm replaces two consecutive Fibonacci numbers as its input by another pair of consecutive Fibonacci numbers, namely:

$$\gcd(F(n), F(n-1)) = \gcd(F(n-1), F(n-2)) \quad \text{for every } n \geq 4.$$

Indeed, since $F(n-2) < F(n-1)$ for every $n \geq 4$,

$$F(n) = F(n-1) + F(n-2) < 2F(n-1).$$

Therefore for every $n \geq 4$, the quotient and remainder of division of $F(n)$ by $F(n-1)$ are 1 and $F(n) - F(n-1) = F(n-2)$, respectively. This is exactly what we asserted at the beginning of the solution. In turn, this leads to the following recurrence for the number of divisions $D(n)$:

$$D(n) = D(n-1) + 1 \quad \text{for } n \geq 4, \quad D(3) = 1,$$

whose initial condition $D(3) = 1$ is obtained by tracing the algorithm on the input pair $F(3), F(2)$, i.e., 2,1. The solution to this recurrence is:

$$D(n) = n - 2 \quad \text{for every } n \geq 3.$$

(One can also easily find directly that $D(2) = 1$ and $D(1) = 0$.)

9. a. (i) For $n = 1$: $F(2)F(0) - [F(1)]^2 = 1 \cdot 0 - 1^2 = (-1)^1$.
(ii) Let $F(n+1)F(n-1) - [F(n)]^2 = (-1)^n$ for some positive integer n . We'll show that

$$F(n+2)F(n) - [F(n+1)]^2 = (-1)^{n+1}.$$

Indeed,

$$\begin{aligned} &F(n+2)F(n) - [F(n+1)]^2 \\ &= [F(n+1) + F(n)]F(n) - [F(n+1)]^2 \\ &= F(n+1)F(n) + [F(n)]^2 - [F(n+1)]^2 \\ &= F(n+1)F(n) + [F(n+1)F(n-1) - (-1)^n] - [F(n+1)]^2 \\ &= F(n+1)[F(n) + F(n-1) - F(n+1)] + (-1)^{n+1} = (-1)^{n+1}. \end{aligned}$$

b. The paradox is based on a bit of visual mischief: the four pieces don't actually form the rectangle because the sides of the triangle and trapezoid don't form a straight line. If they did, the smaller and larger right triangles would be similar, and hence the ratios of their sides, which are 3:8 and 5:13, respectively, would have to be the same.

10. n/a

Exercises 2.6

1. Consider the following well-known sorting algorithm (we shall study it more closely later in the book) with a counter inserted to count the number of key comparisons.

Algorithm *SortAnalysis*($A[0..n-1]$)
 //Input: An array $A[0..n-1]$ of n orderable elements
 //Output: The total number of key comparisons made
 $count \leftarrow 0$
for $i \leftarrow 1$ **to** $n-1$ **do**
 $v \leftarrow A[i]$
 $j \leftarrow i-1$
 while $j \geq 0$ **and** $A[j] > v$ **do**
 $count \leftarrow count + 1$
 $A[j+1] \leftarrow A[j]$
 $j \leftarrow j-1$
 $A[j+1] \leftarrow v$

Is the comparison counter inserted in the right place? If you believe it is, prove it; if you believe it is not, make an appropriate correction.

2. a. Run the program of Problem 1, with a properly inserted counter (or counters) for the number of key comparisons, on 20 random arrays of sizes 1000, 1500, 2000, 2500, ..., 9000, 9500.
 b. Analyze the data obtained to form a hypothesis about the algorithm's average-case efficiency.
 c. Estimate the number of key comparisons one should expect for a randomly generated array of size 10,000 sorted by the same algorithm.
3. Repeat Problem 2 by measuring the program's running time in milliseconds.
4. Hypothesize a likely efficiency class of an algorithm based on the following empirical observations of its basic operation's count:

size	1000	2000	3000	4000	5000	6000	7000	8000	9000
count	11,966	24,303	39,992	53,010	67,272	78,692	91,274	113,063	129,799

5. What scale transformation will make a logarithmic scatterplot look like a linear one?
6. How can we distinguish a scatterplot for an algorithm in $\Theta(\lg \lg n)$ from a scatterplot for an algorithm in $\Theta(\lg n)$?

7. a. Find empirically the largest number of divisions made by Euclid's algorithm for computing $\gcd(m, n)$ for $1 \leq n \leq m \leq 100$.
- b. For each positive integer k , find empirically the smallest pair of integers $1 \leq n \leq m \leq 100$ for which Euclid's algorithm needs to make k divisions in order to find $\gcd(m, n)$.

8. The average-case efficiency of Euclid's algorithm on inputs of size n can be measured by the average number of divisions $D_{avg}(n)$ made by the algorithm in computing $\gcd(n, 1)$, $\gcd(n, 2)$, ..., $\gcd(n, n)$. For example,

$$D_{avg}(5) = \frac{1}{5}(1 + 2 + 3 + 2 + 1) = 1.8.$$

Produce a scatterplot of $D_{avg}(n)$ and indicate a likely average-case efficiency class of the algorithm.

9. Run an experiment to ascertain the efficiency class of the sieve of Eratosthenes (see Section 1.1).
10. Run a timing experiment for the three algorithms for computing $\gcd(m, n)$ presented in Section 1.1.

Hints to Exercises 2.6

1. Does it return a correct comparison count for every array of size 2?
2. Debug your comparison counting and random input generating for small array sizes first.
3. On a reasonably fast desktop, you may well get zero time at least for smaller sizes in your sample. Section 2.6 mentions a trick for overcoming this difficulty.
4. Check how fast the count values grow with doubling the size.
5. A similar question was discussed in the section.
6. Compare the values of the functions $\lg \lg n$ and $\lg n$ for $n = 2^k$.
7. Insert the division counter into a program implementing the algorithm and run it for the input pairs in the range indicated.
8. Get the empirical data for random values of n in a range of between, say, 10^2 and 10^4 or 10^5 and plot the data obtained. (You may want to use different scales for the axes of your coordinate system.)
9. n/a
10. n/a

Solutions to Exercises 2.6

1. It doesn't count the comparison $A[j] > v$ when the comparison fails (and, hence, the body of the while loop is not executed). If the language implies that the second comparison will always be executed even if the first clause of the conjunction fails, the count should be simply incremented by one either right before the **while** statement or right after the **while** statement's end. If the second clause of the conjunction is not executed after the first clause fails, we should add the line

if $j \geq 0$ $count \leftarrow count + 1$

right after the **while** statement's end.

2. a. One should expect numbers very close to $n^2/4$ (the approximate theoretical number of key comparisons made by insertion sort on random arrays).

b. The closeness of the ratios $C(n)/n^2$ to a constant suggests the $\Theta(n^2)$ average-case efficiency. The same conclusion can also be drawn by observing the four-fold increase in the number of key comparisons in response to doubling the array's size.

c. $C(10,000)$ can be estimated either as $10,000^2/4$ or as $4C(5,000)$.
3. See answers to Exercise 2. Note, however, that the timing data is inherently much less accurate and volatile than the counting data.
4. The data exhibits a behavior indicative of an $n \lg n$ algorithm.
5. If $M(n) \approx c \log n$, then the transformation $n = a^k$ ($a > 1$) will yield $M(a^k) \approx (c \log a)k$.
6. The function $\lg \lg n$ grows much more slowly than the slow-growing function $\lg n$. Also, if we transform the plots by substitution $n = 2^k$, the plot of the former would look logarithmic while the plot of the latter would appear linear.
7. a. 9 (for $m = 89$ and $n = 55$)

b. Two consecutive Fibonacci numbers— $m = F_{k+2}$, $n = F_{k+1}$ —are the smallest pair of integers $m \geq n > 0$ that requires k comparisons for every $k \geq 2$. (This is a well-known theoretical fact established by G. Lamé (e.g., [KnuII].) For $k = 1$, the answer is F_{k+1} and F_k , which are both equal to 1.

8. The experiment should confirm the known theoretical result: the average-case efficiency of Euclid's algorithm is in $\Theta(\lg n)$. For a slightly different metric $T(n)$ investigated by D. Knuth, $T(n) \approx \frac{12 \ln 2}{\pi^2} \ln n \approx 0.843 \ln n$ (see [KnuII], Section 4.5.3).
9. n/a
10. n/a

This file contains the exercises, hints, and solutions for Chapter 3 of the book "Introduction to the Design and Analysis of Algorithms," 2nd edition, by A. Levitin. The problems that might be challenging for at least some students are marked by \triangleright ; those that might be difficult for a majority of students are marked by \blacktriangleright .

Exercises 3.1


1. a. Give an example of an algorithm that should not be considered an application of the brute-force approach.
- b. Give an example of a problem that cannot be solved by a brute-force algorithm.
2. a. What is the efficiency of the brute-force algorithm for computing a^n as a function of n ? As a function of the number of bits in the binary representation of n ?
- b. If you are to compute $a^n \bmod m$ where $a > 1$ and n is a large positive integer, how would you circumvent the problem of a very large magnitude of a^n ?
3. For each of the algorithms in Problems 4, 5, and 6 of Exercises 2.3, tell whether or not the algorithm is based on the brute-force approach.
4. a. Design a brute-force algorithm for computing the value of a polynomial

$$p(x) = a_n x^n + a_{n-1} x^{n-1} + \dots + a_1 x + a_0$$

at a given point x_0 and determine its worst-case efficiency class.

- b. If the algorithm you designed is in $\Theta(n^2)$, design a linear algorithm for this problem.
- c. Is it possible to design an algorithm with a better than linear efficiency for this problem?
5. Sort the list E, X, A, M, P, L, E in alphabetical order by selection sort.
6. Is selection sort stable? (The definition of a stable sorting algorithm was given in Section 1.3.)
7. Is it possible to implement selection sort for linked lists with the same $\Theta(n^2)$ efficiency as the array version?
8. Sort the list E, X, A, M, P, L, E in alphabetical order by bubble sort.
9. a. Prove that if bubble sort makes no exchanges on its pass through a list, the list is sorted and the algorithm can be stopped.

- b. Write a pseudocode of the method that incorporates this improvement.
 - c. Prove that the worst-case efficiency of the improved version is quadratic.
10. Is bubble sort stable?
11. *Alternating disks* You have a row of $2n$ disks of two colors, n dark and n light. They alternate: dark, light, dark, light, and so on. You want to get all the dark disks to the right-hand end, and all the light disks to the left-hand end. The only moves you are allowed to make are those which interchange the positions of two neighboring disks.



Design an algorithm for solving this puzzle and determine the number of moves it makes. [Gar99], p.75

Hints to Exercises 3.1

1. a. Think of algorithms that have impressed you with their efficiency and/or sophistication. Neither characteristic is indicative of a brute-force algorithm.

b. Surprisingly, it is not a very easy question to answer. Mathematical problems (including those you have studied in your secondary school and college courses) are a good source of such examples.
2. a. The first question was all but answered in the section. Expressing the answer as a function of the number of bits can be done by using the formula relating the two metrics.

b. How can we compute $(ab) \bmod m$?
3. It helps to have done the exercises in question.
4. a. The most straightforward algorithm, which is based on substituting x_0 into the formula, is quadratic.

b. Analyzing what unnecessary computations the quadratic algorithm does should lead you to a better (linear) algorithm.

c. How many coefficients does a polynomial of degree n have? Can one compute its value at an arbitrary point without processing all of them?
5. Just trace the algorithm on the input given. (It was done for another input in the section.)
6. Although the majority of elementary sorting algorithms are stable, do not rush with your answer. A general remark about stability made in Section 1.3, where the notion of stability is introduced, could be helpful, too.
7. Generally speaking, implementing an algorithm for a linked list poses problems if the algorithm requires accessing the list's elements not in a sequential order.
8. Just trace the algorithm on the input given. (See an example in the section.)
9. a. A list is sorted if and only if all its adjacent elements are in a correct order. Why?

b. Add a boolean flag to register the presence or absence of switches.

c. Identify worst-case inputs first.
10. Can bubble sort change the order of two equal elements in its input?

11. Thinking about the puzzle as a sorting-like problem may and may not lead you to the most simple and efficient solution.

Solutions to Exercises 3.1

1. a. Euclid's algorithm and the standard algorithm for finding the binary representation of an integer are examples from the algorithms previously mentioned in this book. There are, of course, many more examples in its other chapters.

b. Solving nonlinear equations or computing definite integrals are examples of problems that cannot be solved exactly (except for special instances) by any algorithm.

2. a. $M(n) = n \approx 2^b$ where $M(n)$ is the number of multiplications made by the brute-force algorithm in computing a^n and b is the number of bits in the n 's binary representation. Hence, the efficiency is linear as a function of n and exponential as a function of b .

- b. Perform all the multiplications modulo m , i.e.,

$$a^i \bmod m = (a^{i-1} \bmod m \cdot a \bmod m) \bmod m \quad \text{for } i = 1, \dots, n.$$

3. Problem 4 (computes $\sum_1^n i^2$): yes

Problem 5 (computes the range of an array's values): yes

Problem 6 (checks whether a matrix is symmetric): yes

4. a. Here is a pseudocode of the most straightforward version:

Algorithm *BruteForcePolynomialEvaluation*($P[0..n], x$)
//The algorithm computes the value of polynomial P at a given point x
//by the “highest-to-lowest term” brute-force algorithm
//Input: Array $P[0..n]$ of the coefficients of a polynomial of degree n ,
// stored from the lowest to the highest and a number x
//Output: The value of the polynomial at the point x
 $p \leftarrow 0.0$
for $i \leftarrow n$ **downto** 0 **do**
 $power \leftarrow 1$
 for $j \leftarrow 1$ **to** i **do**
 $power \leftarrow power * x$
 $p \leftarrow p + P[i] * power$
return p

We will measure the input's size by the polynomial's degree n . The basic operation of this algorithm is a multiplication of two numbers; the number of multiplications $M(n)$ depends on the polynomial's degree only.

Although it is not difficult to find the total number of multiplications in this algorithm, we can count just the number of multiplications in the algorithm's inner-most loop to find the algorithm's efficiency class:

$$M(n) = \sum_{i=0}^n \sum_{j=1}^i 1 = \sum_{i=0}^n i = \frac{n(n+1)}{2} \in \Theta(n^2).$$

b. The above algorithm is very inefficient: we recompute powers of x again and again as if there were no relationship among them. Thus, the obvious improvement is based on computing consecutive powers more efficiently. If we proceed from the highest term to the lowest, we could compute x^{i-1} by using x^i but this would require a division and hence a special treatment for $x = 0$. Alternatively, we can move from the lowest term to the highest and compute x^i by using x^{i-1} . Since the second alternative uses multiplications instead of divisions and does not require any special treatment for $x = 0$, it is both more efficient and cleaner. It leads to the following algorithm:

Algorithm *BetterBruteForcePolynomialEvaluation*($P[0..n], x$)
 //The algorithm computes the value of polynomial P at a given point x
 //by the “lowest-to-highest term” algorithm
 //Input: Array $P[0..n]$ of the coefficients of a polynomial of degree n ,
 // from the lowest to the highest, and a number x
 //Output: The value of the polynomial at the point x
 $p \leftarrow P[0]; \text{ power} \leftarrow 1$
for $i \leftarrow 1$ **to** n **do**
 $\text{power} \leftarrow \text{power} * x$
 $p \leftarrow p + P[i] * \text{power}$
return p

The number of multiplications here is

$$M(n) = \sum_{i=1}^n 2 = 2n$$

(while the number of additions is n), i.e., we have a linear algorithm.

Note: Horner's Rule discussed in Section 6.5 needs only n multiplications (and n additions) to solve this problem.

c. No, because any algorithm for evaluating an arbitrary polynomial of degree n at an arbitrary point x must process all its $n + 1$ coefficients. (Note that even when $x = 1$, $p(x) = a_n + a_{n-1} + \dots + a_1 + a_0$, which needs at least n additions to be computed correctly for arbitrary a_n, a_{n-1}, \dots, a_0 .)

5.

	E	X	A	M	P	L	E
A		X	E	M	P	L	E
A	E		X	M	P	L	E
A	E	E		M	P	L	X
A	E	E	L		P	M	X
A	E	E	L	M		P	X
A	E	E	L	M	P		X

6. Selection sort is not stable: In the process of exchanging elements that are not adjacent to each other, the algorithm can reverse an ordering of equal elements. The list 2', 2'', 1 is such an example.

7. Yes. Both operations—finding the smallest element and swapping it—can be done as efficiently with a linked list as with an array.

8. E, X, A, M, P, L, E

E	\leftrightarrow	X	\leftrightarrow	A		M		P		L		E
E		A		X	\leftrightarrow	M		P		L		E
E		A		M		X	\leftrightarrow	P		L		E
E		A		M		P		X	\leftrightarrow	L		E
E		A		M		P		L		X	\leftrightarrow	E
E		A		M		P		L		E		X
E	\leftrightarrow	A		M		P		L		E		
A		E	\leftrightarrow	M	\leftrightarrow	P	\leftrightarrow	L		E		
A		E		M		L		P	\leftrightarrow	E		
A		E		M		L		E		P		
A	\leftrightarrow	E	\leftrightarrow	M	\leftrightarrow	L		E				
A		E		L		M	\leftrightarrow	E				
A		E		L		E		M				
A	\leftrightarrow	E	\leftrightarrow	L	\leftrightarrow	E						
A		E		E		L						
A	\leftrightarrow	E	\leftrightarrow	E	\leftrightarrow	L						

The algorithm can be stopped here (see the next question).

9. a. Pass i ($0 \leq i \leq n-2$) of bubble sort can be represented by the following diagram:

$$A_0, \dots, A_j \xleftrightarrow{?} A_{j+1}, \dots, A_{n-i-1} \leq \underset{\text{in their final positions}}{A_{n-i} \leq \dots \leq A_{n-1}}$$

If there are no swaps during this pass, then

$$A_0 \leq A_1 \leq \dots \leq A_j \leq A_{j+1} \leq \dots \leq A_{n-i-1},$$

with the larger (more accurately, not smaller) elements in positions $n - i$ through $n - 1$ being sorted during the previous iterations.

b. Here is a pseudocode for the improved version of bubble sort:

Algorithm *BetterBubbleSort*($A[0..n - 1]$)
 //The algorithm sorts array $A[0..n - 1]$ by improved bubble sort
 //Input: An array $A[0..n - 1]$ of orderable elements
 //Output: Array $A[0..n - 1]$ sorted in ascending order
 $count \leftarrow n - 1$ //number of adjacent pairs to be compared
 $sflag \leftarrow \mathbf{true}$ //swap flag
while $sflag$ **do**
 $sflag \leftarrow \mathbf{false}$
 for $j \leftarrow 0$ **to** $count - 1$ **do**
if $A[j + 1] < A[j]$
 swap $A[j]$ and $A[j + 1]$
 $sflag \leftarrow \mathbf{true}$
 $count \leftarrow count - 1$

c. The worst-case inputs will be strictly decreasing arrays. For them, the improved version will make the same comparisons as the original version, which was shown in the text to be quadratic.

10. Bubble sort is stable. It follows from the fact that it swaps adjacent elements only, provided $A[j + 1] < A[j]$.
5. Here is a simple and efficient (in fact, optimal) algorithm for this problem: Starting with the first and ending with the last light disk, swap it with each of the i ($1 \leq i \leq n$) dark disks to the left of it. The i th iteration of the algorithm can be illustrated by the following diagram, in which 1s and 0s correspond to the dark and light disks, respectively.

$$\underbrace{00..011..11}_{i-1} \underbrace{10}_{i-1} 10..10 \quad \Rightarrow \quad 00..\underbrace{0011}_{i-1}..\underbrace{11}_{i-1} 10..10$$

The total number of swaps made is equal to $\sum_{i=1}^n i = n(n + 1)/2$.

Exercises 3.2

1. Find the number of comparisons made by the sentinel version of sequential search
 - a. in the worst case.
 - b. in the average case if the probability of a successful search is p ($0 \leq p \leq 1$).
2. As shown in Section 2.1, the average number of key comparisons made by sequential search (without a sentinel, under standard assumptions about its inputs) is given by the formula

$$C_{avg}(n) = \frac{p(n+1)}{2} + n(1-p),$$

where p is the probability of a successful search. Determine, for a fixed n , the values of p ($0 \leq p \leq 1$) for which this formula yields the largest value of $C_{avg}(n)$ and the smallest value of $C_{avg}(n)$.

3. *Gadgets testing* A firm wants to determine the highest floor of its n -story headquarters from which a gadget can fall with no impact on the gadget's functionality. The firm has two identical gadgets to experiment with. Design an algorithm in the best efficiency class you can to solve this problem.
4. Determine the number of character comparisons made by the brute-force algorithm in searching for the pattern **GANDHI** in the text

`THERE_IS_MORE_TO_LIFE_THAN_INCREASING_ITS_SPEED`

(Assume that the length of the text—it is 47 characters long—is known before the search starts.)
5. How many comparisons (both successful and unsuccessful) are made by the brute-force string-matching algorithm in searching for each of the following patterns in the binary text of 1000 zeros?
 - a. 00001 b. 10000 c. 01010
6. Give an example of a text of length n and a pattern of length m that constitutes the worst-case input for the brute-force string-matching algorithm. Exactly how many character comparisons are made for such input?
7. Write a visualization program for the brute-force string-matching algorithm.

8. In solving the string-matching problem, would there be any advantage in comparing pattern and text characters right-to-left instead of left-to-right?
9. Consider the problem of counting, in a given text, the number of substrings that start with an A and end with a B. (For example, there are four such substrings in CABAAXBYA.)
 - (a) Design a brute-force algorithm for this problem and determine its efficiency class.
 - (b) Design a more efficient algorithm for this problem [Gin04].
10. *Word Find* A popular diversion in the United States, Word Find, asks the player to find each of a given set of words in a square table filled with single letters. A word can read horizontally (left or right), vertically (up or down), or along a 45 degree diagonal (in any of the four directions), formed by consecutively adjacent cells of the table; it may wrap around the table's boundaries but it must read in the same direction with no zigzagging. The same cell of the table may be used in different words, but, in a given word, the same cell may be used no more than once. Write a computer program for solving this puzzle.
11. *Battleship game* Write a program for playing Battleship (a classic strategy game) on the computer which is based on a version of brute-force pattern matching. The rules of the game are as follows. There are two opponents in the game (in this case, a human player and the computer). The game is played on two identical boards (10-by-10 tables of squares) on which each opponent places his or her ships, not seen by the opponent. Each player has five ships, each of which occupies a certain number of squares on the board: a destroyer (2 squares), a submarine (3 squares), a cruiser (3 squares), a battleship (4 squares), and an aircraft carrier (5 squares). Each ship is placed either horizontally or vertically, with no two ships touching each other. The game is played by the opponents taking turns "shooting" at each other's ships. A result of every shot is displayed as either a hit or a miss. In case of a hit, the player gets to go again and keeps playing until this player misses. The goal is to sink all the opponent's ships before the opponent succeeds in doing it first. (To sink a ship, all squares occupied by the ship must be hit.)

Hints to Exercises 3.2

1. Modify the analysis of the algorithm's version in Section 2.1.
2. As a function of p , what kind of function is C_{avg} ?
3. Solve a simpler problem with a single gadget first. Then design a better than linear algorithm for the problem with two gadgets.
4. The content of this quote from Mahatma Gandhi is more thought provoking than this drill.
5. For each input, one iteration of the algorithm yields all the information you need to answer the question.
6. It will suffice to limit your search for an example to binary texts and patterns.
7. You may use either bit strings or a natural language text for the visualization program. It would be a good idea to implement, as an option, a search for all occurrences of a given pattern in a given text.
8. The answer, surprisingly, is yes.
9.
 - a. For a given occurrence of A in the text, what are the substrings you need to count?
 - b. For a given occurrence of B in the text, what are the substrings you need to count?
10. Test your program thoroughly. Be especially careful about a possibility of words read diagonally with wrapping around the table's border.
11. A (very) brute-force algorithm can simply shoot at adjacent feasible cells starting at, say, one of the corners of the board. Can you suggest a better strategy? (You can investigate relative efficiencies of different strategies by making two programs implementing them play each other.) Is your strategy better than the one that shoots at randomly generated cells of the opponent's board?

Solutions to Exercises 3.2

1. a. $C_{\text{worst}}(n) = n + 1$.

b. $C_{\text{avg}}(n) = \frac{(2-p)(n+1)}{2}$. In the manner almost identical to the analysis in Section 2.1, we obtain

$$\begin{aligned} C_{\text{avg}}(n) &= [1 \cdot \frac{p}{n} + 2 \cdot \frac{p}{n} + \dots + i \cdot \frac{p}{n} + \dots + n \cdot \frac{p}{n}] + (n+1) \cdot (1-p) \\ &= \frac{p}{n} [1 + 2 + \dots + i + \dots + n] + (n+1)(1-p) \\ &= \frac{p}{n} \frac{n(n+1)}{2} + (n+1)(1-p) = \frac{(2-p)(n+1)}{2}. \end{aligned}$$

2. The expression

$$\frac{p(n+1)}{2} + n(1-p) = p \frac{n+1}{2} + n - np = n - p(n - \frac{n+1}{2}) = n - \frac{n-1}{2}p$$

is a linear function of p . Since the p 's coefficient is negative for $n > 1$, the function is strictly decreasing on the interval $0 \leq p \leq 1$ from n to $(n+1)/2$. Hence $p = 0$ and $p = 1$ are its maximum and minimum points, respectively, on this interval. (Of course, this is the answer we should expect: The average number of comparisons should be the largest when the probability of a successful search is 0, and it should be the smallest when the probability of a successful search is 1.)

3. Drop the first gadget from floors $\lceil \sqrt{n} \rceil$, $2\lceil \sqrt{n} \rceil$, and so on until either the floor $i\lceil \sqrt{n} \rceil$ a drop from which makes the gadget malfunction is reached or no such floor in this sequence is encountered before the top of the building is reached. In the former case, the floor to be found is higher than $(i-1)\lceil \sqrt{n} \rceil$ and lower than $i\lceil \sqrt{n} \rceil$. So, drop the second gadget from floors $(i-1)\lceil \sqrt{n} \rceil + 1$, $(i-1)\lceil \sqrt{n} \rceil + 2$, and so on until the first floor a drop from which makes the gadget malfunction is reached. The floor immediately preceding that floor is the floor in question. If no drop in the first-pass sequence resulted in the gadget's failure, the floor in question is higher than $i\lceil \sqrt{n} \rceil$, the last tried floor of that sequence. Hence, continue the successive examination of floors $i\lceil \sqrt{n} \rceil + 1$, $i\lceil \sqrt{n} \rceil + 2$, and so on until either a failure is registered or the last floor is reached. The number of times the two gadgets are dropped doesn't exceed $\lceil \sqrt{n} \rceil + \lceil \sqrt{n} \rceil$, which puts it in $O(\sqrt{n})$.

4. 43 comparisons.

The algorithm will make $47 - 6 + 1 = 42$ trials: In the first one, the G of the pattern will be aligned against the first T of the text; in the last one, it will be aligned against the last space. On each but one trial, the algorithm will make one unsuccessful comparison; on one trial—when the G of the pattern is aligned against the G of the text—it will make two

comparisons. Thus, the total number of character comparisons will be $41 \cdot 1 + 1 \cdot 2 = 43$.

5. a. For the pattern 00001, the algorithm will make four successful and one unsuccessful comparison on each of its trials and then shift the pattern one position to the right:

0 0 0 0 0 0	0 0 0 0 0
0 0 0 0 1	
0 0 0 0 1	
etc.	

0 0 0 0 1

The total number of character comparisons will be $C = 5 \cdot 996 = 4980$.

- b. For the pattern 10000, the algorithm will make one unsuccessful comparison on each of its trials and then shift the pattern one position to the right:

0 0 0 0 0 0	0 0 0 0 0
1 0 0 0 0	
1 0 0 0 0	
etc.	

1 0 0 0 0

The total number of character comparisons will be $C = 1 \cdot 996 = 996$.

- c. For the pattern 01010, the algorithm will make one successful and one unsuccessful comparison on each of its trials and then shift the pattern one position to the right:

0 0 0 0 0 0	0 0 0 0 0
0 1 0 1 0	
0 1 0 1 0	
etc.	

0 1 0 1 0

The total number of character comparisons will be $C = 2 \cdot 996 = 1,992$.

6. The text composed of n zeros and the pattern $\underbrace{0 \dots 0}_{m-1}1$ is an example of the worst-case input. The algorithm will make $m(n - m + 1)$ character comparisons on such input.

7. n/a

8. Comparing pairs of the pattern and text characters right-to-left can allow farther pattern shifts after a mismatch. This is the main insight the two string matching algorithms discussed in Section 7.2 are based on. (As a specific example, consider searching for the pattern 11111 in the text of one thousand zeros.)

9. a. Note that the number of desired substrings that starts with an A at a given position i ($0 \leq i < n - 1$) in the text is equal to the number of B's to the right of that position. This leads to the following simple algorithm:

Initialize the count of the desired substrings to 0. Scan the text left to right doing the following for every character except the last one: If an A is encountered, count the number of all the B's following it and add this number to the count of desired substrings. After the scan ends, return the last value of the count.

For the worst case of the text composed of n A's, the total number of character comparisons is

$$n + (n - 1) + \dots + 2 = n(n + 1)/2 - 1 \in \Theta(n^2).$$

- b. Note that the number of desired substrings that ends with a B at a given position i ($0 < i \leq n - 1$) in the text is equal to the number of A's to the left of that position. This leads to the following algorithm:

Initialize the count of the desired substrings and the count of A's encountered to 0. Scan the text left to right until the text is exhausted and do the following. If an A is encountered, increment the A's count; if a B is encountered, add the current value of the A's count to the desired substring count. After the text is exhausted, return the last value of the desired substring count.

Since the algorithm makes a single pass through a given text spending constant time on each of its characters, the algorithm is linear.

10. n/a

11. n/a

Exercises 3.3

1. Can you design a more efficient algorithm than the one based on the brute-force strategy to solve the closest-pair problem for n points x_1, \dots, x_n on the real line?
2. Let $x_1 < x_2 < \dots < x_n$ be real numbers representing coordinates of n villages located along a straight road. A post office needs to be built in one of these villages.
 - a. Design an efficient algorithm to find the post-office location minimizing the average distance between the villages and the post office.
 - b. Design an efficient algorithm to find the post-office location minimizing the maximum distance from a village to the post office.

3. a.▷ There are several alternative ways to define a distance between two points $P_1 = (x_1, y_1)$ and $P_2 = (x_2, y_2)$. In particular, the **Manhattan distance** is defined as

$$d_M(P_1, P_2) = |x_1 - x_2| + |y_1 - y_2|.$$

Prove that d_M satisfies the following axioms that every distance function must satisfy:

(i) $d_M(P_1, P_2) \geq 0$ for any two points P_1 and P_2 and $d_M(P_1, P_2) = 0$ if and only if $P_1 = P_2$;

(ii) $d_M(P_1, P_2) = d_M(P_2, P_1)$;

(iii) $d_M(P_1, P_2) \leq d_M(P_1, P_3) + d_M(P_3, P_2)$ for any P_1, P_2 , and P_3 .

b. Sketch all the points in the x, y coordinate plane whose Manhattan distance to the origin $(0,0)$ is equal to 1. Do the same for the Euclidean distance.

c.▷ True or false: A solution to the closest-pair problem does not depend on which of the two metrics— d_E (Euclidean) or d_M (Manhattan)—is used.

4. *Odd pie fight* There are $n \geq 3$ people positioned in a field (Euclidean plane) so that each has a unique nearest neighbor. Each person has a cream pie. At a signal, everybody hurls his or her pie at the nearest neighbor. Assuming that n is odd and that nobody can miss his or her target, true or false: There always remains at least one person not hit by a pie? [Car79].
5. The closest-pair problem can be posed in k -dimensional space in which the Euclidean distance between two points $P' = (x'_1, \dots, x'_k)$ and $P'' =$

(x''_1, \dots, x''_k) is defined as

$$d(P', P'') = \sqrt{\sum_{s=1}^k (x'_s - x''_s)^2}.$$

What is the time-efficiency class of the brute-force algorithm for the k -dimensional closest-pair problem?

6. Find the convex hulls of the following sets and identify their extreme points (if they have any).
 - a. a line segment
 - b. a square
 - c. the boundary of a square
 - d. a straight line
7. Design a linear-time algorithm to determine two extreme points of the convex hull of a given set of $n > 1$ points in the plane.
8. What modification needs to be made in the brute-force algorithm for the convex-hull problem to handle more than two points on the same straight line?
9. Write a program implementing the brute-force algorithm for the convex-hull problem.
10. Consider the following small instance of the linear programming problem:

$$\begin{array}{ll}\text{maximize} & 3x + 5y \\ \text{subject to} & x + y \leq 4 \\ & x + 3y \leq 6 \\ & x \geq 0, \ y \geq 0\end{array}$$

- a. Sketch, in the Cartesian plane, the problem's **feasible region** defined as the set of points satisfying all the problem's constraints.
- b. Identify the region's extreme points.
- c. Solve the optimization problem given by using the following theorem:
A linear programming problem with a nonempty bounded feasible region always has a solution, which can be found at one of the extreme points of its feasible region.

Hints to Exercises 3.3

1. Sorting n real numbers can be done in $O(n \log n)$ time.
2. a. Solving the problem for $n = 2$ and $n = 3$ should lead you to the critical insight.

b. Where would you put the post office if it would not have to be at one of the village locations?
3. a. Check requirements (i)–(iii) by using basic properties of absolute values.

b. For the Manhattan distance, the points in question are defined by equation $|x - 0| + |y - 0| = 1$. You can start by sketching the points in the positive quadrant of the coordinate system (i.e., the points for which $x, y \geq 0$) and then sketch the rest by using the symmetries.

c. The assertion is false. You can choose, say, $P_1(0, 0)$, $P_2(1, 0)$ and find P_3 to complete a counterexample.
4. True; prove it by mathematical induction.
5. Your answer should be a function of two parameters: n and k . A special case of this problem (for $k = 2$) was solved in the text.
6. Review the examples given in the section.
7. Some of the extreme points of a convex hull are easier to find than others.
8. If there are other points of a given set on the straight line through P_i and P_j , which of all these points need to be preserved for further processing?
9. Your program should work for any set of n distinct points, including sets with many colinear points.
10. a. The set of points satisfying inequality $ax + by \leq c$ is the half-plane of the points on one side of the straight line $ax + by = c$, including all the points on the line itself. Sketch such a half-plane for each of the inequalities and find their intersection.

b. The extreme points are the vertices of the polygon obtained in part a.

c. Compute and compare the values of the objective function at the extreme points.

Solutions to Exercises 3.3

- Sort the numbers in ascending order, compute the differences between adjacent numbers in the sorted list, and find the smallest such difference. If sorting is done in $O(n \log n)$ time, the running time of the entire algorithm will be in

$$O(n \log n) + \Theta(n) + \Theta(n) = O(n \log n).$$

- If we put the post office at location x_i , the average distance between it and all the points $x_1 < x_2 < \dots < x_n$ is given by the formula $\frac{1}{n} \sum_{j=1}^n |x_j - x_i|$. Since the number of points n stays the same, we can ignore the multiple $\frac{1}{n}$ and minimize $\sum_{j=1}^n |x_j - x_i|$. We'll have to consider the cases of even and odd n separately.

Let n be even. Consider first the case of $n = 2$. The sum $|x_1 - x| + |x_2 - x|$ is equal to $x_2 - x_1$, the length of the interval with the endpoints at x_1 and x_2 , for any point x of this interval (including the endpoints), and it is larger than $x_2 - x_1$ for any point x outside of this interval. This implies that for any even n , the sum

$$\sum_{j=1}^n |x_j - x| = [|x_1 - x| + |x_n - x|] + [|x_2 - x| + |x_{n-1} - x|] + \dots + [|x_{n/2} - x| + |x_{n/2+1} - x|]$$

is minimized when x belongs to each of the intervals $[x_1, x_n] \supset [x_2, x_{n-1}] \supset \dots \supset [x_{n/2}, x_{n/2+1}]$. If x must be one of the points given, either $x_{n/2}$ or $x_{n/2+1}$ solves the problem.

Let $n > 1$ be odd. Then, the sum $\sum_{j=1}^n |x_j - x|$ is minimized when $x = x_{\lceil n/2 \rceil}$, the point for which the number of the given points to the left of it is equal to the number of the given points to the right of it.

Note that the point $x_{\lceil n/2 \rceil}$ —the $\lceil n/2 \rceil$ th smallest called the *median*—solves the problem for even n 's as well. For a sorted list implemented as an array, the median can be found in $\Theta(1)$ time by simply returning the $\lceil n/2 \rceil$ th element of the array. (Section 5.6 provides a more general discussion of algorithms for computing the median.)

- Assuming that the points x_1, x_2, \dots, x_n are given in increasing order, the answer is the point x_i that is the closest to $m = (x_1 + x_n)/2$, the middle point between x_1 and x_n . (The middle point would be the obvious solution if the post-office didn't have to be at one of the given locations.) Indeed, if we put the post office at any location x_i to the left of m , the longest distance from a village to the post office would be $x_n - x_i$; this distance is minimal for the rightmost among such points. If we put the post office at any location x_i to the right of m , the longest distance from a village to the post office would be $x_i - x_1$; this distance is minimal for the leftmost among such points.

Algorithm *PostOffice1*(P)

//Input: List P of n ($n \geq 2$) points x_1, x_2, \dots, x_n in increasing order


```

//Output: Point  $x_i$  that minimizes  $\max_{1 \leq j \leq n} |x_j - x_i|$  among all  $x_1, x_2, \dots, x_n$ 
 $m \leftarrow (x_1 + x_n)/2$ 
 $i \leftarrow 1$ 
while  $x_i < m$  do
     $i \leftarrow i + 1$ 
if  $x_i - x_1 < x_n - x_{i-1}$ 
    return  $x_i$ 
else return  $x_{i-1}$ 

```

The time efficiency of this algorithm is $O(n)$.

3. a. For $d_M(P_1, P_2) = |x_1 - x_2| + |y_1 - y_2|$, we have the following:

(i) $d_M(P_1, P_2) = |x_1 - x_2| + |y_1 - y_2| \geq 0$ and $d_M(P_1, P_2) = 0$ if and only if both $x_1 = x_2$ and $y_1 = y_2$, i.e., P_1 and P_2 coincide.

(ii) $d_M(P_1, P_2) = |x_1 - x_2| + |y_1 - y_2| = |x_2 - x_1| + |y_2 - y_1| = d_M(P_2, P_1)$.

(iii) $d_M(P_1, P_2) = |x_1 - x_2| + |y_1 - y_2|$
 $= |(x_1 - x_3) + (x_3 - x_2)| + |(y_1 - y_3) + (y_3 - y_2)|$
 $\leq |x_1 - x_3| + |x_3 - x_2| + |y_1 - y_3| + |y_3 - y_2| = d(P_1, P_3) + d(P_3, P_2)$.

b. For the Manhattan distance, the points in question are defined by the equation

$$|x - 0| + |y - 0| = 1, \text{ i.e., } |x| + |y| = 1.$$

The graph of this equation is the boundary of the square with its vertices at $(1, 0)$, $(0, 1)$, $(-1, 0)$, and $(0, -1)$.

For the Euclidean distance, the points in question are defined by the equation

$$\sqrt{(x - 0)^2 + (y - 0)^2} = 1, \text{ i.e., } x^2 + y^2 = 1.$$

The graph of this equation is the circumference of radius 1 and the center at $(0, 0)$.

c. False. Consider points $P_1(0, 0)$, $P_2(1, 0)$, and, say, $P_3(\frac{1}{2}, \frac{3}{4})$. Then

$$d_E(P_1, P_2) = 1 \text{ and } d_E(P_3, P_1) = d_E(P_3, P_2) = \sqrt{\left(\frac{1}{2}\right)^2 + \left(\frac{3}{4}\right)^2} < 1.$$

Therefore, for the Euclidean distance, the two closest points are either P_1 and P_3 or P_2 and P_3 . For the Manhattan distance, we have

$$d_M(P_1, P_2) = 1 \text{ and } d_M(P_3, P_1) = d_M(P_3, P_2) = \frac{1}{2} + \frac{3}{4} = \frac{5}{4} > 1.$$

Therefore, for the Manhattan distance, the two closest points are P_1 and P_2 .

4. We'll prove by induction that there will always remain at least one person not hit by a pie. The basis step is easy: If $n = 3$, the two persons with the smallest pairwise distance between them throw at each other, while the third person throws at one of them (whoever is closer). Therefore, this third person remains "unharmd".

For the inductive step, assume that the assertion is true for odd $n \geq 3$, and consider $n + 2$ persons. Again, the two persons with the smallest pairwise distance between them (the closest pair) throw at each other. Consider two possible cases as follows. If the remaining n persons all throw at one another, at least one of them remains "unharmd" by the inductive assumption. If at least one of the remaining n persons throws at one of the closest pair, among the remaining $n - 1$ persons, at most $n - 1$ pies are thrown at one another, and hence at least one person must remain "unharmd" because there is not enough pies to hit everybody in that group. This completes the proof.

5. The number of squarings will be


$$\begin{aligned} C(n, k) &= \sum_{i=1}^{n-1} \sum_{j=i+1}^n \sum_{s=1}^k 1 = \sum_{i=1}^{n-1} \sum_{j=i+1}^n k = k \sum_{i=1}^{n-1} (n - i) \\ &= k[(n - 1) + (n - 2) + \dots + 1] = \frac{k(n - 1)n}{2} \in \Theta(kn^2). \end{aligned}$$

6. a. The convex hull of a line segment is the line segment itself; its extreme points are the endpoints of the segment.
- b. The convex hull of a square is the square itself; its extreme points are the four vertices of the square.
- c. The convex hull of the boundary of a square is the region comprised of the points within that boundary and on the boundary itself; its extreme points are the four vertices of the square.
- d. The convex hull of a straight line is the straight line itself. It doesn't have any extreme points.
7. Find the point with the smallest x coordinate; if there are several such points, find the one with the smallest y coordinate among them. Similarly, find the point with the largest x coordinate; if there are several such points, find the one with the largest y coordinate among them.

8. If there are other points of a given set on the straight line through P_i and P_j (while all the other points of the set lie on the same side of the line), a line segment of the convex hull's boundary will have its end points at the two farthest set points on the line. All the other points on the line can be eliminated from further processing.

9. n/a

10. a. Here is a sketch of the feasible region in question:



- b. The extreme points are: $(0, 0)$, $(4, 0)$, $(3, 1)$, and $(0, 2)$.

c.

Extreme point	Value of $3x + 5y$
$(0, 0)$	0
$(4, 0)$	12
$(3, 1)$	14
$(0, 2)$	10

So, the optimal solution is $(3, 1)$, with the maximal value of $3x + 5y$ equal to 14. (Note: This instance of the linear programming problem is discussed further in Section 10.1.)

Exercises 3.4

1. a. Assuming that each tour can be generated in constant time, what will be the efficiency class of the exhaustive-search algorithm outlined in the text for the traveling salesman problem?

b. If this algorithm is programmed on a computer that makes one billion additions per second, estimate the maximum number of cities for which the problem can be solved in (i) one hour; (ii) 24-hours; (iii) one year; (iv) one century.
2. Outline an exhaustive-search algorithm for the Hamiltonian circuit problem.
3. Outline an algorithm to determine whether a connected graph represented by its adjacency matrix has an Eulerian circuit. What is the efficiency class of your algorithm?
4. Complete the application of exhaustive search to the instance of the assignment problem started in the text.
5. Give an example of the assignment problem whose optimal solution does not include the smallest element of its cost matrix.
6. Consider the ***partition problem***: given n positive integers, partition them into two disjoint subsets with the same sum of their elements. (Of course, the problem does not always have a solution.) Design an exhaustive search algorithm for this problem. Try to minimize the number of subsets the algorithm needs to generate.
7. Consider the ***clique problem***: given a graph G and a positive integer k , determine whether the graph contains a ***clique*** of size k , i.e., a complete subgraph of k vertices. Design an exhaustive-search algorithm for this problem.
8. Explain how exhaustive search can be applied to the sorting problem and determine the efficiency class of such an algorithm.
9. A magic square of order n is an arrangement of the numbers from 1 to n^2 in an n -by- n matrix, with each number occurring exactly once, so that each row, each column, and each main diagonal has the same sum.
 - a. Prove that if a magic square of order n exists, the sum in question must be equal to $n(n^2 + 1)/2$.
 - b. Design an exhaustive search algorithm for generating all magic squares of order n .
 - c. Go to the Internet or your library and find a better algorithm for

generating magic squares.

d. Implement the two algorithms—the exhaustive search and the one you have found—and run an experiment to determine the largest value of n for which each of the algorithms is able to find a magic square of order n in less than one minute of your computer’s time.

10. *Famous alphametic* A puzzle in which the digits in a correct mathematical expression, such as a sum, are replaced by letters is called **cryptarithm**; if, in addition, the puzzle’s words make sense, it is said to be an **alphametic**. The most well-known alphametic was published by the renowned British puzzlist H. E. Dudeney (1857-1930):

$$\begin{array}{r} \text{S E N D} \\ + \text{M O R E} \\ \hline \text{M O N E Y} \end{array}$$

Two conditions are assumed: First, the correspondence between letters and digits is one-to-one, that is each letter represents one digit only and different letters represent different digits. Second, the digit zero does not appear as the left-most digit in any of the numbers. To solve an alphametic means to find which digit each letter represents. Note that a solution’s uniqueness cannot be assumed and has to be verified by the solver.

- a. Write a program for solving cryptarithms by exhaustive search. Assume that a given cryptarithm is a sum of two words.
- b. Solve Dudeney’s puzzle the way it was expected to be solved when it was first published in 1924.

Hints to Exercises 3.4

1. a. Identify the algorithm's basic operation and count the number of times it will be executed.

b. For each of the time amounts given, find the largest value of n for which this limit will not be exceeded.
2. How different is the traveling salesman problem from the problem of finding a Hamiltonian circuit?
3. Your algorithm should check the well-known conditions that are both necessary and sufficient for the existence of a Eulerian circuit in a connected graph.
4. Generate the remaining $4! - 6 = 18$ possible assignments, compute their costs, and find the one with the smallest cost.
5. Make the size of your counterexample as small as possible.
6. Rephrase the problem so that the sum of elements in one subset, rather than two, needs to be checked on each try of a possible partition.
7. Follow the definitions of a clique and of an exhaustive search algorithm.
8. Try all possible orderings of the elements given.
9. a. Add all the elements in the magic square in two different ways.

b. What combinatorial objects do you have to generate here?
10. a. For testing, you may use alphametic collections available on the Internet.

b. Given the absence of electronic computers in 1924, you must refrain here from using the Internet.

Solutions to Exercises 3.4

1. a. $\Theta(n!)$

For each tour (a sequence of $n+1$ cities), one needs n additions to compute the tour's length. Hence, the total number of additions $A(n)$ will be n times the total number of tours considered, i.e., $n \cdot \frac{1}{2}(n-1)! = \frac{1}{2}n! \in \Theta(n!)$.

b. (i) $n_{\max} = 15$; (ii) $n_{\max} = 16$; (iii) $n_{\max} = 18$; (iv) $n_{\max} = 20$.

Given the answer to part a, we have to find the largest value of n such that

$$\frac{1}{2}n!10^{-9} \leq t$$

where t is the time available (in seconds). Thus, for $t = 1\text{hr} = 3.6 \cdot 10^3 \text{sec}$, we get the inequality

$$n! \leq 2 \cdot 10^9 t = 7.2 \cdot 10^{12}.$$

The largest value of n for which this inequality holds is 15 (since $15! \approx 1.3 \cdot 10^{12}$ and $16! \approx 2.1 \cdot 10^{13}$).

For the other given values of t , the answers can be obtained in the same manner.

2. The problem of finding a Hamiltonian circuit is very similar to the traveling salesman problem. Generate permutations of n vertices that start and end with, say, the first vertex, and check whether every pair of successive vertices in a current permutation are connected by an edge. If it's the case, the current permutation represents a Hamiltonian circuit, otherwise, a next permutation needs to be generated.
3. A connected graph has a Eulerian circuit if and only if all its vertices have even degrees. An algorithm should check this condition until either an odd vertex is encountered (then a Eulerian circuit doesn't exist) or all the vertices turn out to be even (then a Eulerian circuit must exist). For a graph (with no loops) represented by its n -by- n adjacency matrix, the degree of a vertex is the number of ones in the vertex's row. Thus, computing its degree will take the $\Theta(n)$ time, checking whether it's even will take $\Theta(1)$ time, and it will be done between 1 and n times. Hence, the algorithm's efficiency will be in $O(n^2)$.

4. The following assignments were generated in the chapter's text:

$C = \begin{bmatrix} 9 & 2 & 7 & 8 \\ 6 & 4 & 3 & 7 \\ 5 & 8 & 1 & 8 \\ 7 & 6 & 9 & 4 \end{bmatrix}$	1, 2, 3, 4	cost = $9+4+1+4 = 18$	etc.
	1, 2, 4, 3	cost = $9+4+8+9 = 30$	
	1, 3, 2, 4	cost = $9+3+8+4 = 24$	
	1, 3, 4, 2	cost = $9+3+8+6 = 26$	
	1, 4, 2, 3	cost = $9+7+8+9 = 33$	
	1, 4, 3, 2	cost = $9+7+1+6 = 23$	

The remaining ones are

2, 1, 3, 4	cost = 2+6+1+4 = 13
2, 1, 4, 3	cost = 2+6+8+9 = 25
2, 3, 1, 4	cost = 2+3+5+4 = 14
2, 3, 4, 1	cost = 2+3+8+7 = 20
2, 4, 1, 3	cost = 2+7+5+9 = 23
2, 4, 3, 1	cost = 2+7+1+7 = 17
3, 1, 2, 4	cost = 7+6+8+4 = 25
3, 1, 4, 2	cost = 7+6+8+6 = 27
3, 2, 1, 4	cost = 7+4+5+4 = 20
3, 2, 4, 1	cost = 7+4+8+7 = 26
3, 4, 1, 2	cost = 7+7+5+6 = 25
3, 4, 2, 1	cost = 7+7+8+7 = 29
4, 1, 2, 3	cost = 8+6+8+9 = 31
4, 1, 3, 2	cost = 8+6+1+6 = 21
4, 2, 1, 3	cost = 8+4+5+9 = 26
4, 2, 3, 1	cost = 8+4+1+7 = 20
4, 3, 1, 2	cost = 8+3+5+6 = 22
4, 3, 2, 1	cost = 8+3+8+7 = 26

The optimal solution is: Person 1 to Job 2, Person 2 to Job 1, Person 3 to Job 3, and Person 4 to Job 4, with the total (minimal) cost of the assignment being 13.

5. Here is a very simple example:

$$\begin{bmatrix} 1 & 2 \\ 2 & 9 \end{bmatrix}$$

6. Start by computing the sum S of the numbers given. If S is odd, stop because the problem doesn't have a solution. If S is even, generate the subsets until either a subset whose elements' sum is $S/2$ is encountered or no more subsets are left. Note that it will suffice to generate only subsets with no more than $n/2$ elements.
7. Generate a subset of k vertices and check whether every pair of vertices in the subset is connected by an edge. If it's true, stop (the subset is a clique); otherwise, generate the next subset.
8. Generate a permutation of the elements given and check whether they are ordered as required by comparing values of its consecutive elements. If they are, stop; otherwise, generate the next permutation. Since the

number of permutations of n items is equal to $n!$ and checking a permutation requires up to $n - 1$ comparisons, the algorithm's efficiency class is in $O(n!(n - 1)) = O((n + 1)!)$.

9. a. Let s be the sum of the numbers in each row of an n -by- n magic square. Let us add all the numbers in rows 1 through n . We will get the following equality:

$$sn = 1 + 2 + \dots + n^2, \text{ i.e., } sn = \frac{n^2(n^2 + 1)}{2}, \text{ which implies that } s = \frac{n(n^2 + 1)}{2}.$$

b. Number positions in an n -by- n matrix from 1 through n^2 . Generate a permutation of the numbers 1 through n^2 , put them in the corresponding positions of the matrix, and check the magic-square equality (proved in part (a)) for every row, every column, and each of the two main diagonals of the matrix.

c. n/a

d. n/a

10. a. Since the letter-digit correspondence must be one-to-one and there are only ten distinct decimal digits, the exhaustive search needs to check $P(10, k) = 10!/(10 - k)!$ possible substitutions, where k is the number of distinct letters in the input. (The requirement that the first letter of a word cannot represent 0 can be used to reduce this number further.) Thus a program should run in a quite reasonable amount of time on today's computers. Note that rather than checking two cases—with and without a “1-carry”—for each of the decimal positions, the program can check just one equality, which stems from the definition of the decimal number system. For Dudeney's alphametic, for example, this equality is $1000(S+M) + 100(E+O) + 10(N+R) + (D+E) = 10000M + 1000O + 100N + 10E + Y$

b. Here is a “computerless” solution to this classic problem. First, notice that M must be 1. (Since both S and M are not larger than 9, their sum, even if increased by 1 because of the carry from the hundred column, must be less than 20.) We will have to rely on some further insights into specifics of the problem. The leftmost digits of the addends imply one of the two possibilities: either $S + M = 10 + O$ (if there was no carry from the hundred column) or $1 + S + M = 10 + O$ (if there was such a carry). First, let us pursue the former of the two possibilities. Since $M = 1$, $S \leq 9$ and $O \geq 0$, the equation $S + 1 = 10 + O$ has only one solution: $S = 9$ and $O = 0$. This leaves us with

$$\begin{array}{r}
 \text{E N D} \\
 + \text{O R E} \\
 \hline
 \text{N E Y}
 \end{array}$$

Since we deal here with the case of no carry from the hundreds and E and N must be distinct, the only possibility is a carry from the tens: $1 + E = N$ and either $N + R = 10 + E$ (if there was no carry from the rightmost column) or $1 + N + R = 10 + E$ (if there was such a carry). The first combination leads to a contradiction: Substituting $1 + E$ for N into $N + R = 10 + E$, we obtain $R = 9$, which is incompatible with the same digit already represented by S. The second combination of $1 + E = N$ and $1 + N + R = 10 + E$ implies, after substituting the first of these equations into the second one, $R = 8$. Note that the only remaining digit values still unassigned are 2, 3, 4, 5, 6, and 7. Finally, for the rightmost column, we have the equation $D + E = 10 + Y$. But $10 + Y \geq 12$, because the smallest unassigned digit value is 2 while $D + E \leq 12$ because the two largest unassigned digit values are 6 and 7 and $E < N$. Thus, $D + E = 10 + Y = 12$. Hence $Y = 2$ and $D + E = 12$. The only pair of still unassigned digit values that add up to 12, 5 and 7, must be assigned to E and D, respectively, since doing this the other way ($E = 7$, $D = 5$) would imply $N = E + 1 = 8$, which is already represented by R. Thus, we found the following solution to the puzzle:

$$\begin{array}{r}
 9\ 5\ 6\ 7 \\
 + 1\ 0\ 8\ 5 \\
 \hline
 1\ 0\ 6\ 5\ 2
 \end{array}$$


Is this the only solution? To answer this question, we should pursue the carry possibility from the hundred column to the thousand column (see above). Then $1 + S + M = 10 + O$ or, since $M = 1$, $S = 8 + O$. But $S \leq 9$, while $8 + O \geq 10$ since $O \geq 2$. Hence the last equation has no solutions in our domain. This proves that the puzzle has no other solutions.

This file contains the exercises, hints, and solutions for Chapter 4 of the book "Introduction to the Design and Analysis of Algorithms," 2nd edition, by A. Levitin. The problems that might be challenging for at least some students are marked by \triangleright ; those that might be difficult for a majority of students are marked by \blacktriangleright .

Exercises 4.1

1.
 - a. Write a pseudocode for a divide-and-conquer algorithm for finding the position of the largest element in an array of n numbers.
 - b. What will be your algorithm's output for arrays with several elements of the largest value?
 - c. Set up and solve a recurrence relation for the number of key comparisons made by your algorithm.
 - d. How does this algorithm compare with the brute-force algorithm for this problem?
2.
 - a. Write a pseudocode for a divide-and-conquer algorithm for finding values of both the largest and smallest elements in an array of n numbers.
 - b. Set up and solve (for $n = 2^k$) a recurrence relation for the number of key comparisons made by your algorithm.
 - c. How does this algorithm compare with the brute-force algorithm for this problem?
3.
 - a. Write a pseudocode for a divide-and-conquer algorithm for the exponentiation problem of computing a^n where $a > 0$ and n is a positive integer.
 - b. Set up and solve a recurrence relation for the number of multiplications made by this algorithm.
 - c. How does this algorithm compare with the brute-force algorithm for this problem?
4. As mentioned in Chapter 2, logarithm bases are irrelevant in most contexts arising in analyzing an algorithm's efficiency class. Is it true for both assertions of the Master Theorem that include logarithms?
5. Find the order of growth for solutions of the following recurrences.
 - a. $T(n) = 4T(n/2) + n$, $T(1) = 1$

- b. $T(n) = 4T(n/2) + n^2$, $T(1) = 1$
- c. $T(n) = 4T(n/2) + n^3$, $T(1) = 1$
6. Apply mergesort to sort the list E, X, A, M, P, L, E in alphabetical order.
7. Is mergesort a stable sorting algorithm?
8. a. Solve the recurrence relation for the number of key comparisons made by mergesort in the worst case. (You may assume that $n = 2^k$.)
- b. Set up a recurrence relation for the number of key comparisons made by mergesort on best-case inputs and solve it for $n = 2^k$.
- c. Set up a recurrence relation for the number of key moves made by the version of mergesort given in Section 4.1. Does taking the number of key moves into account change the algorithm's efficiency class?
9. Let $A[0..n-1]$ be an array of n distinct real numbers. A pair $(A[i], A[j])$ is said to be an ***inversion*** if these numbers are out of order, i.e., $i < j$ but $A[i] > A[j]$. Design an $O(n \log n)$ algorithm for counting the number of inversions.
10. One can implement mergesort without a recursion by starting with merging adjacent elements of a given array, then merging sorted pairs, and so on. Implement this bottom-up version of mergesort in the language of your choice.
11. *Tromino puzzle* A tromino is an L-shaped tile formed by adjacent 1-by-1 squares. The problem is to cover any 2^n -by- 2^n chessboard with one missing square (anywhere on the board) with trominoes. Trominoes should cover all the squares of the board except the missing one with no overlaps.



Design a divide-and-conquer algorithm for this problem.

Hints to Exercises 4.1

1. In more than one respect, this question is similar to the divide-and-conquer computation of the sum of n numbers. Also, you were asked to analyze an almost identical algorithm in Exercises 2.4.
2. Unlike Problem 1, a divide-and-conquer algorithm for this problem can be more efficient by a constant factor than the brute-force algorithm.
3. How would you compute a^8 by solving two exponentiation problems of size 4? How about a^9 ?
4. Look at the notations used in the theorem's statement.
5. Apply the Master Theorem.
6. Trace the algorithm as it was done for another input in the section.
7. How can mergesort reverse a relative ordering of two elements?
8.
 - a. Use backward substitutions, as usual.
 - b. What inputs minimize the number of key comparisons made by mergesort? How many comparisons are made by mergesort on such inputs during the merging stage?
 - c. Do not forget to include key moves made both before the split and during the merging.
9. Modify mergesort to solve the problem.
10. n/a
11. A divide-and-conquer algorithm works by reducing a problem's instance to several smaller instances of the *same* problem.

Solutions to Exercises 4.1

1. a. Call **Algorithm** *MaxIndex*($A[0..n-1]$) where

```

Algorithm MaxIndex( $A[l..r]$ )
//Input: A portion of array  $A[0..n-1]$  between indices  $l$  and  $r$  ( $l \leq r$ )
//Output: The index of the largest element in  $A[l..r]$ 
if  $l = r$  return  $l$ 
else  $temp1 \leftarrow \text{MaxIndex}(A[l..\lfloor(l+r)/2\rfloor])$ 
       $temp2 \leftarrow \text{MaxIndex}(A[\lfloor(l+r)/2\rfloor + 1..r])$ 
      if  $A[temp1] \geq A[temp2]$ 
        return  $temp1$ 
      else return  $temp2$ 

```

- b. This algorithm returns the index of the leftmost largest element.
c. The recurrence for the number of element comparisons is

$$C(n) = C(\lceil n/2 \rceil) + C(\lfloor n/2 \rfloor) + 1 \quad \text{for } n > 1, \quad C(1) = 0.$$

Solving it by backward substitutions for $n = 2^k$ yields the following:

$$\begin{aligned}
 C(2^k) &= 2C(2^{k-1}) + 1 \\
 &= 2[2C(2^{k-2}) + 1] + 1 = 2^2C(2^{k-2}) + 2 + 1 \\
 &= 2^2[2C(2^{k-3}) + 1] + 2 + 1 = 2^3C(2^{k-3}) + 2^2 + 2 + 1 \\
 &= \dots \\
 &= 2^iC(2^{k-i}) + 2^{i-1} + 2^{i-2} + \dots + 1 \\
 &= \dots \\
 &= 2^kC(2^{k-k}) + 2^{k-1} + 2^{k-2} + \dots + 1 = 2^k - 1 = n - 1.
 \end{aligned}$$

We can verify that $C(n) = n - 1$ satisfies, in fact, the recurrence for every value of $n > 1$ by substituting it into the recurrence equation and considering separately the even ($n = 2i$) and odd ($n = 2i + 1$) cases. Let $n = 2i$, where $i > 0$. Then the left-hand side of the recurrence equation is $n - 1 = 2i - 1$. The right-hand side is

$$\begin{aligned}
 C(\lceil n/2 \rceil) + C(\lfloor n/2 \rfloor) + 1 &= C(\lceil 2i/2 \rceil) + C(\lfloor 2i/2 \rfloor) + 1 \\
 &= 2C(i) + 1 = 2(i - 1) + 1 = 2i - 1,
 \end{aligned}$$

which is the same as the left-hand side.

Let $n = 2i + 1$, where $i > 0$. Then the left-hand side of the recurrence equation is $n - 1 = 2i$. The right-hand side is

$$\begin{aligned}
 C(\lceil n/2 \rceil) + C(\lfloor n/2 \rfloor) + 1 &= C(\lceil (2i + 1)/2 \rceil) + C(\lfloor (2i + 1)/2 \rfloor) + 1 \\
 &= C(i + 1) + C(i) + 1 = (i + 1 - 1) + (i - 1) + 1 = 2i,
 \end{aligned}$$

which is the same as the left-hand side in this case, too.

d. A simple standard scan through the array in question requires the same number of key comparisons but avoids the overhead associated with recursive calls.

2. a. Call **Algorithm** *MinMax*($A[0..n-1]$, *minval*, *maxval*) where

Algorithm *MinMax*($A[l..r]$, *minval*, *maxval*)
 //Finds the values of the smallest and largest elements in a given subarray
 //Input: A portion of array $A[0..n-1]$ between indices l and r ($l \leq r$)
 //Output: The values of the smallest and largest elements in $A[l..r]$
 //assigned to *minval* and *maxval*, respectively
if $r = l$
 $\text{minval} \leftarrow A[l]; \quad \text{maxval} \leftarrow A[l]$
else if $r - l = 1$
 if $A[l] \leq A[r]$
 $\text{minval} \leftarrow A[l]; \quad \text{maxval} \leftarrow A[r]$
 else $\text{minval} \leftarrow A[r]; \quad \text{maxval} \leftarrow A[l]$
else $//r - l > 1$
 $\text{MinMax}(A[l..\lfloor (l+r)/2 \rfloor], \text{minval}, \text{maxval})$
 $\text{MinMax}(A[\lfloor (l+r)/2 \rfloor + 1..r], \text{minval2}, \text{maxval2})$
 if $\text{minval2} < \text{minval}$
 $\text{minval} \leftarrow \text{minval2}$
 if $\text{maxval2} > \text{maxval}$
 $\text{maxval} \leftarrow \text{maxval2}$

b. Assuming for simplicity that $n = 2^k$, we obtain the following recurrence for the number of element comparisons $C(n)$:

$$C(n) = 2C(n/2) + 2 \quad \text{for } n > 2, \quad C(2) = 1, \quad C(1) = 0.$$

Solving it by backward substitutions for $n = 2^k$, $k \geq 1$, yields the following:

$$\begin{aligned} C(2^k) &= 2C(2^{k-1}) + 2 \\ &= 2[2C(2^{k-2}) + 2] + 2 = 2^2C(2^{k-2}) + 2^2 + 2 \\ &= 2^2[2C(2^{k-3}) + 2] + 2^2 + 2 = 2^3C(2^{k-3}) + 2^3 + 2^2 + 2 \\ &= \dots \\ &= 2^iC(2^{k-i}) + 2^i + 2^{i-1} + \dots + 2 \\ &= \dots \\ &= 2^{k-1}C(2) + 2^{k-1} + \dots + 2 = 2^{k-1} + 2^k - 2 = \frac{3}{2}n - 2. \end{aligned}$$

c. This algorithm makes about 25% fewer comparisons— $1.5n$ compared to $2n$ —than the brute-force algorithm. (Note that if we didn't stop recursive calls when $n = 2$, we would've lost this gain.) In fact, the algorithm is

optimal in terms of the number of comparisons made. As a practical matter, however, it might not be faster than the brute-force algorithm because of the recursion-related overhead. (As noted in the solution to Problem 5 of Exercises 2.3, a nonrecursive scan of a given array that maintains the minimum and maximum values seen so far and updates them not for each element but for a pair of two consecutive elements makes the same number of comparisons as the divide-and-conquer algorithm but doesn't have the recursion's overhead.)

3. a. The following divide-and-conquer algorithm for computing a^n is based on the formula $a^n = a^{\lfloor n/2 \rfloor} a^{\lceil n/2 \rceil}$:

Algorithm *DivConqPower*(a, n)
 //Computes a^n by a divide-and-conquer algorithm
 //Input: A positive number a and a positive integer n
 //Output: The value of a^n
if $n = 1$ **return** a
else return *DivConqPower*($a, \lfloor n/2 \rfloor$) * *DivConqPower*($a, \lceil n/2 \rceil$)

- b. The recurrence for the number of multiplications is

$$M(n) = M(\lfloor n/2 \rfloor) + M(\lceil n/2 \rceil) + 1 \quad \text{for } n > 1, \quad M(1) = 0.$$

The solution to this recurrence (solved above for Problem 1) is $n - 1$.

- c. Though the algorithm makes the same number of multiplications as the brute-force method, it has to be considered inferior to the latter because of the recursion overhead.

4. For the second case, where the solution's class is indicated as $\Theta(n^d \log n)$, the logarithm's base could change the function by a constant multiple only and, hence, is irrelevant. For the third case, where the solution's class is $\Theta(n^{\log_b a})$, the logarithm is in the function's exponent and, hence, must be indicated since functions n^α have different orders of growth for different values of α .


5. The applications of the Master Theorem yield the following.

a. $T(n) = 4T(n/2) + n$. Here, $a = 4$, $b = 2$, and $d = 1$. Since $a > b^d$, $T(n) \in \Theta(n^{\log_2 4}) = \Theta(n^2)$.

b. $T(n) = 4T(n/2) + n^2$. Here, $a = 4$, $b = 2$, and $d = 2$. Since $a = b^d$, $T(n) \in \Theta(n^2 \log n)$.

c. $T(n) = 4T(n/2) + n^3$. Here, $a = 4$, $b = 2$, and $d = 3$. Since $a < b^d$, $T(n) \in \Theta(n^3)$.

6. Here is a trace of mergesort applied to the input given:



7. Mergesort is stable, provided its implementation employs the comparison \leq in merging. Indeed, assume that we have two elements of the same value in positions i and j , $i < j$, in a subarray before its two (sorted) halves are merged. If these two elements are in the same half of the subarray, their relative ordering will stay the same after the merging because the elements of the same half are processed by the merging operation in the FIFO fashion. Consider now the case when $A[i]$ is in the first half while $A[j]$ is in the second half. $A[j]$ is placed into the new array either after the first half becomes empty (and, hence, $A[i]$ has been already copied into the new array) or after being compared with some key $k > A[j]$ of the first half. In the latter case, since the first half is sorted before the merging begins, $A[i] = A[j] < k$ cannot be among the unprocessed elements of the first half. Hence, by the time of this comparison, $A[i]$ has been already copied into the new array and therefore will precede $A[j]$ after the merging operation is completed.

8. a. The recurrence for the number of comparisons in the worst case, which was given in Section 4.1, is

$$C_w(n) = 2C_w(n/2) + n - 1 \text{ for } n > 1 \text{ (and } n = 2^k), \quad C_w(1) = 0.$$

Solving it by backward substitutions yields the following:

$$\begin{aligned} C_w(2^k) &= 2C_w(2^{k-1}) + 2^k - 1 \\ &= 2[2C_w(2^{k-2}) + 2^{k-1} - 1] + 2^k - 1 = 2^2C_w(2^{k-2}) + 2 \cdot 2^k - 2 - 1 \\ &= 2^2[2C_w(2^{k-3}) + 2^{k-2} - 1] + 2 \cdot 2^k - 2 - 1 = 2^3C_w(2^{k-3}) + 3 \cdot 2^k - 2^2 - 2 - 1 \\ &= \dots \\ &= 2^iC_w(2^{k-i}) + i2^k - 2^{i-1} - 2^{i-2} - \dots - 1 \\ &= \dots \\ &= 2^kC_w(2^{k-k}) + k2^k - 2^{k-1} - 2^{k-2} - \dots - 1 = k2^k - (2^k - 1) = n \log n - n + 1 \end{aligned}$$

- b. The recurrence for the number of comparisons on best-case inputs (lists sorted in ascending or descending order) is

$$C_b(n) = 2C_b(n/2) + n/2 \text{ for } n > 1 \text{ (and } n = 2^k), \quad C_b(1) = 0.$$

Thus,

$$\begin{aligned} C_b(2^k) &= 2C_b(2^{k-1}) + 2^{k-1} \\ &= 2[2C_b(2^{k-2}) + 2^{k-2}] + 2^{k-1} = 2^2C_b(2^{k-2}) + 2^{k-1} + 2^{k-1} \\ &= 2^2[2C_b(2^{k-3}) + 2^{k-3}] + 2^{k-1} + 2^{k-1} = 2^3C_b(2^{k-3}) + 2^{k-1} + 2^{k-1} + 2^{k-1} \\ &= \dots \\ &= 2^iC_b(2^{k-i}) + i2^{k-1} \\ &= \dots \\ &= 2^kC_b(2^{k-k}) + k2^{k-1} = k2^{k-1} = \frac{1}{2}n \log n. \end{aligned}$$

- c. If $n > 1$, the algorithm copies $\lfloor n/2 \rfloor + \lceil n/2 \rceil = n$ elements first and then makes n more moves during the merging stage. This leads to the following recurrence for the number of moves $M(n)$:

$$M(n) = 2M(n/2) + 2n \text{ for } n > 1, \quad M(1) = 0.$$


According to the Master Theorem, its solution is in $\Theta(n \log n)$ —the same class established by the analysis of the number of key comparisons only.

5. Let *ModifiedMergesort* be a mergesort modified to return the number of inversions in its input array $A[0..n-1]$ in addition to sorting it. Obviously, for an array of size 1, *ModifiedMergesort*($A[0]$) should return 0. Let

i_{left} and i_{right} be the number of inversions returned by *ModifiedMerge-sort*($A[0..mid - 1]$) and *ModifiedMerge-sort*($A[mid..n - 1]$), respectively, where mid is the index of the middle element in the input array $A[0..n - 1]$. The total number of inversions in $A[0..n - 1]$ can then be computed as $i_{left} + i_{right} + i_{merge}$, where i_{merge} , the number of inversions involving elements from both halves of $A[0..n - 1]$, is computed during the merging as follows. Let $A[i]$ and $A[j]$ be two elements from the left and right half of $A[0..n - 1]$, respectively, that are compared during the merging. If $A[i] < A[j]$, we output $A[i]$ to the sorted list without incrementing i_{merge} because $A[i]$ cannot be a part of an inversion with any of the remaining elements in the second half, which are greater than $A[j]$. If, on the other hand, $A[i] > A[j]$, we output $A[j]$ and increment i_{merge} by $mid - i$, the number of remaining elements in the first half, because all those elements (and only they) form an inversion with $A[j]$.

10. n/a

11. For $n > 1$, we can always place one L-tromino at the center of the $2^n \times 2^n$ chessboard with one missing square to reduce the problem to four subproblems of tiling $2^{n-1} \times 2^{n-1}$ boards, each with one missing square too. The orientation of this centrally placed piece is determined by the board's quarter with the missing square as shown by the example below.



Then each of the four smaller problems can be solved recursively until a trivial case of a 2×2 board with a missing square is reached.

Exercises 4.2

1. Apply quicksort to sort the list

E, X, A, M, P, L, E

in alphabetical order. Draw the tree of the recursive calls made.

2. For the partitioning procedure outlined in Section 4.2:

- a. Prove that if the scanning indices stop while pointing to the same element, i.e., $i = j$, the value they are pointing to must be equal to p .

- b. Prove that when the scanning indices stop, j cannot point to an element more than one position to the left of the one pointed to by i .

- c. Why is it worth stopping the scans after encountering an element equal to the pivot?

3. Is quicksort a stable sorting algorithm?

4. Give an example of an array of n elements for which the sentinel mentioned in the text is actually needed. What should be its value? Also explain why a single sentinel suffices for any input.

5. For the version of quicksort given in the text:

- a. Are arrays made up of all equal elements the worst-case input, the best-case input, or neither?

- b. Are strictly decreasing arrays the worst-case input, the best-case input, or neither?

6. a. For quicksort with the median-of-three pivot selection, are increasing arrays the worst-case input, the best-case input, or neither?

- b. Answer the same question for decreasing arrays.

7. ► Solve the average-case recurrence for quicksort.

8. Design an algorithm to rearrange elements of a given array of n real numbers so that all its negative elements precede all its positive elements. Your algorithm should be both time- and space-efficient.

9. ▷ The **Dutch flag problem** is to rearrange an array of characters R , W , and B (red, white, and blue are the colors of the Dutch national flag) so that all the R 's come first, the W 's come next, and the B 's come last. Design a linear in-place algorithm for this problem.

10. Implement quicksort in the language of your choice. Run your program on a sample of inputs to verify the theoretical assertions about the algorithm's efficiency.
11. \triangleright *Nuts and bolts* You are given a collection of n bolts of different widths and n corresponding nuts. You are allowed to try a nut and bolt together, from which you can determine whether the nut is larger than the bolt, smaller than the bolt, or matches the bolt exactly. However, there is no way to compare two nuts together or two bolts together. The problem is to match each bolt to its nut. Design an algorithm for this problem with average-case efficiency in $\Theta(n \log n)$. [Raw91]

Hints to Exercises 4.2

1. We traced the algorithm on an another instance in the section.
2. a. Use the rules for stopping the scans.

b. Use the rules for stopping the scans.

c. Consider an array whose all elements are the same.
3. The definition of *stability* of a sorting algorithm was given in Section 1.3. Generally speaking, algorithms that can exchange elements far apart are not stable.
4. Trace the algorithm to see on which inputs index i gets out of bounds.
5. Study what the text's version of quicksort does on such arrays. You should base your answers on the number of key comparisons, of course.
6. Where will splits occur on the inputs in question?
7. This requires several standard tricks for solving more sophisticated recurrence relations. A solution can be found in most books on the design and analysis of algorithms.
8. Use the partition idea.
9. You may want to solve first the two-color flag problem, i.e., rearrange efficiently an array of R 's and B 's. (A similar problem is Problem 8 in these exercises.)
10. n/a
11. Use the partition idea.

2. a. Let $i = j$ be the coinciding values of the scanning indices. According to the rules for stopping the i (left-to-right) and j (right-to-left) scans, $A[i] \geq p$ and $A[j] \leq p$ where p is the pivot's value. Hence, $A[i] = A[j] = p$.
 b. Let i be the value of the left-to-right scanning index after it stopped. Since $A[i - 1] \leq p$, the right-to-left scanning index will have to stop no later than reaching $i - 1$.
 c. Stopping the scans after encountering an element equal to the pivot tends to yield better (i.e., more equal) splits. For example, if we did otherwise for an array of n equal elements, we would have gotten a split into subarrays of sizes $n - 1$ and 0 .
 3. Quicksort is not stable. As a counterexample, consider its performance on a two-element array of equal values.
 4. With the pivot being the leftmost element, the left-to-right scan will get out of bounds if and only if the pivot is larger than all the other elements. Appending a sentinel of value equal $A[0]$ (or larger than $A[0]$) after the array's last element will stop the index of the left-to-right scan of $A[0..n-1]$ from going beyond position n . A single sentinel will suffice by the following reason. In quicksort, when $Partition(A[l..r])$ is called for $r < n - 1$, all the elements to the right of position r are greater than or equal to all the elements in $A[l..r]$. Hence, $A[r + 1]$ will automatically play the role of a sentinel to stop index i going beyond position $r + 1$.
 5. a. Arrays composed of all equal elements constitute the best case because all the splits will happen in the middle of corresponding subarrays.
 b. Strictly decreasing arrays constitute the worst case because all the splits will yield one empty subarray. (Note that we need to show this to be the case on two consecutive iterations of the algorithm because the first iteration does not yield a decreasing array of size $n - 1$.)
 6. The best case for both questions. For either an increasing or decreasing subarray, the median of the first, last, and middle values will be the median of the entire subarray. Using it as a pivot will split the subarray in the middle. This will cause the total number of key comparisons be the smallest.
 7. Here is a solution that follows [Sed88], p. 121:

$$C(n) = \frac{1}{n} \sum_{s=0}^{n-1} [(n+1) + C(s) + C(n-1-s)]$$

can be rewritten as

$$C(n) = (n+1) + \frac{1}{n} \sum_{s=0}^{n-1} [C(s) + C(n-1-s)].$$

Since $\sum_{s=0}^{n-1} C(s) = \sum_{s=0}^{n-1} C(n-1-s)$, the equation can be reduced to

$$C(n) = (n+1) + \frac{2}{n} \sum_{s=0}^{n-1} C(s)$$

or

$$nC(n) = n(n+1) + 2 \sum_{s=0}^{n-1} C(s).$$

Substituting $n-1$ for n in the last equation yields

$$(n-1)C(n-1) = (n-1)n + 2 \sum_{s=0}^{n-2} C(s).$$

Subtracting the last equation from the one before yields, after obvious simplifications, the recurrence

$$nC(n) = (n+1)C(n-1) + 2n,$$

which, after dividing both hand sides by $n(n+1)$, becomes

$$\frac{C(n)}{n+1} = \frac{C(n-1)}{n} + \frac{2}{n+1}.$$

Substituting $B(n) = \frac{C(n)}{n+1}$, we obtain the following recurrence relation:

$$B(n) = B(n-1) + \frac{2}{n+1} \quad \text{for } n \geq 2, \quad B(1) = B(0) = 0.$$

The latter can be solved either by backward substitutions (or by “telescoping”) to obtain

$$B(n) = 2 \sum_{k=3}^{n+1} \frac{1}{k}.$$

Hence

$$B(n) = 2H_{n+1} - 3, \quad \text{where } H_{n+1} = \sum_{k=1}^{n+1} \frac{1}{k} \approx \ln(n+1) \quad (\text{see Appendix A}).$$

Thus,

$$C(n) = (n+1)B(n) \approx 2(n+1) \ln(n+1) \approx 2n \ln n.$$

8. The following algorithm uses the partition idea similar to that of quicksort, although it's implemented somewhat differently. Namely, on each iteration the algorithm maintains three sections (possibly empty) in a given array: all the elements in $A[0..i-1]$ are negative, all the elements in $A[i..j]$ are unknown, and all the elements in $A[j+1..n]$ are nonnegative:

$A[0]$...	$A[i-1]$	$A[i]$...	$A[j]$	$A[j+1]$...	$A[n-1]$
all are < 0			unknown			all are ≥ 0		

On each iteration, the algorithm shrinks the size of the unknown section by one element either from the left or from the right.

Algorithm *NegBeforePos*($A[0..n-1]$)

//Puts negative elements before positive (and zeros, if any) in an array

//Input: Array $A[0..n-1]$ of real numbers

//Output: Array $A[0..n-1]$ in which all its negative elements precede nonnegative

$i \leftarrow 0; \quad j \leftarrow n-1$

while $i \leq j$ **do** // $i < j$ would suffice

if $A[i] < 0$ //shrink the unknown section from the left

$i \leftarrow i+1$

else //shrink the unknown section from the right

 swap($A[i], A[j]$)

$j \leftarrow j-1$

Note: If we want all the zero elements placed after all the negative elements but before all the positive ones, the problem becomes the Dutch flag problem (see Problem 9 in these exercises).

9. The following algorithm uses the partition idea similar to that of quicksort. (See also a simpler 2-color version of this problem in Problem 8 in these exercises.) On each iteration, the algorithm maintains four sections (possibly empty) in a given array: all the elements in $A[0..r-1]$ are filled with R's, all the elements in $A[r..w-1]$ are filled with W's, all the elements in $A[w..b]$ are unknown, and all the elements in $A[b+1..n-1]$ are filled with B's.

$A[0]$...	$A[r-1]$	$A[r]$...	$A[w-1]$	$A[w]$...	$A[b]$	$A[b+1]$...	$A[n-1]$
all are filled with R's			all are filled with W's			unknown			all are filled with B's		

On each iteration, the algorithm shrinks the size of the unknown section by one element either from the left or from the right.

Algorithm *DutchFlag*($A[0..n-1]$)

//Sorts an array with values in a three-element set

//Input: An array $A[0..n-1]$ of characters from {'R', 'W', 'B'}

```

//Output: Array  $A[0..n-1]$  in which all its  $R$  elements precede
//          all its  $W$  elements that precede all its  $B$  elements
 $r \leftarrow 0$ ;  $w \leftarrow 0$ ;  $b \leftarrow n-1$ 
while  $w \leq b$  do
    if  $A[w] = 'R'$ 
         $\text{swap}(A[r], A[w]);$   $r \leftarrow r+1$ ;  $w \leftarrow w+1$ 
    else if  $A[w] = 'W'$ 
         $w \leftarrow w+1$ 
    else  $//A[w] = 'B'$ 
         $\text{swap}(A[w], A[b]);$   $b \leftarrow b-1$ 

```

10. n/a

11. Randomly select a nut and try each of the bolts for it to find the matching bolt and separate the bolts that are smaller and larger than the selected nut into two disjoint sets. Then try each of the unmatched nuts against the matched bolt to separate those that are larger from those that are smaller than the bolt. As a result, we've identified a matching pair and partitioned the remaining nuts and bolts into two smaller independent instances of the same problem. The average number of nut-bolt comparisons $C(n)$ is defined by the recurrence very similar to the one for quicksort in Section 4.2:

$$C(n) = \frac{1}{n} \sum_{s=0}^{n-1} [(2n-1) + C(s) + C(n-1-s)], \quad C(1) = 0, \quad C(0) = 0.$$

The solution to this recurrence can be shown to be in $\Theta(n \log n)$ by repeating the steps outlined in the solution to Problem 7.

Note: See a $O(n \log n)$ deterministic algorithm for this problem in the paper by Janos Komlos, Yuan Ma and Endre Szemerédi "Matching Nuts and Bolts in $O(n \log n)$ Time," SIAM J. Discrete Math. 11, No.3, 347-372 (1998).

Exercises 4.3

1. a. What is the largest number of key comparisons made by binary search in searching for a key in the following array?

3	14	27	31	39	42	55	70	74	81	85	93	98
---	----	----	----	----	----	----	----	----	----	----	----	----

- b. List all the keys of this array that will require the largest number of key comparisons when searched for by binary search.
 - c. Find the average number of key comparisons made by binary search in a successful search in this array. (Assume that each key is searched for with the same probability.)
 - d. Find the average number of key comparisons made by binary search in an unsuccessful search in this array. (Assume that searches for keys in each of the 14 intervals formed by the array's elements are equally likely.)
2. Solve the recurrence $C_{worst}(n) = C_{worst}(\lfloor n/2 \rfloor) + 1$ for $n > 1$, $C_{worst}(1) = 1$, for $n = 2^k$ by backward substitutions.
3. a.▷ Prove the equality

$$\lfloor \log_2 n \rfloor + 1 = \lceil \log_2(n+1) \rceil \quad \text{for } n \geq 1.$$


- b.▷ Prove that $C_{worst}(n) = \lfloor \log_2 n \rfloor + 1$ satisfies equation (4.2) for every positive integer n .
4. Estimate how many times faster an average successful search will be in a sorted array of 100,000 elements if it is done by binary search versus sequential search.
5. Sequential search can be used with about the same efficiency whether a list is implemented as an array or as a linked list. Is it also true for binary search? (Of course, we assume that a list is sorted for binary search.)
6. How can one use binary search for range searching, i.e., for finding all the elements in a sorted array whose values fall between two given values L and U (inclusively), $L \leq U$? What is the worst-case efficiency of this algorithm?
7. Write a pseudocode for a recursive version of binary search.
8. Design a version of binary search that uses only two-way comparisons such as \leq and $=$. Implement your algorithm in the language of your choice and carefully debug it (such programs are notorious for being prone to bugs).
9. Analyze the time efficiency of the two-way comparison version designed in Problem 8.

10. A version of the popular problem-solving task involves presenting people with an array of 42 pictures—seven rows of six pictures each—and asking them to identify the target picture by asking questions that can be answered yes or no. Further, people are then required to identify the picture with as few questions as possible. Suggest the most efficient algorithm for this problem and indicate the largest number of questions that may be necessary.

Hints to Exercises 4.3

1. a. Take advantage of the formula that gives the immediate answer.

(b)–(d) The most efficient prop for answering such questions is a binary tree that mirrors the algorithm's operations in searching for an arbitrary search key. The first three nodes of such a tree for the instance in question will look as follows:



(The first number inside a node is the index m of the array's element being compared with a search key; the number in the parentheses is the value of the element itself, i.e., $A[m]$.)


2. If you need to refresh your memory, look up Section 2.4, where we solved an almost identical recurrence relation, and Appendix B.
3. a. Use the fact that n is bounded below and above by some consecutive powers of 2, i.e., $2^k \leq n < 2^{k+1}$.

b. The case of an even n ($n = 2i$) was considered in the section. For an odd n ($n = 2i + 1$), substitute the function into both hand sides of the equation and show their equality. The formula of part (a) might be useful. Do not forget to verify the initial condition, too.
4. Estimate the ratio of the average number of key comparisons made in a successful search by sequential search to that for binary search.
5. How would you reach the middle element in a linked list?
6. Find separately the elements that are greater than or equal to L and those that are smaller than or equal to U . Do not forget that neither L nor U have to be among the array values.
7. You may find the diagram of binary search in the text helpful.
8. Use the comparison $K \leq A[m]$ where $m \leftarrow \lfloor (l + r)/2 \rfloor$ until $l = r$. Then check whether the search is successful or not.
9. The analysis is almost identical to that of the text's version of binary search.
10. Number the pictures and use this numbering in your questions.

Solutions to Exercises 4.3

1. a. According to formula (4.4), $C_{worst}(13) = \lceil \log_2(13 + 1) \rceil = 4$.

b. In the comparison tree below, the first number indicates the element's index, the second one is its value:



The searches for each of the elements on the last level of the tree, i.e., the elements in positions 1(14), 3(31), 5(42), 8(74), 10(85), and 12(98) will require the largest number of key comparisons.

$$c. C_{avg}^{yes} = \frac{1}{13} \cdot 1 \cdot 1 + \frac{1}{13} \cdot 2 \cdot 2 + \frac{1}{13} \cdot 3 \cdot 4 + \frac{1}{13} \cdot 4 \cdot 6 = \frac{41}{13} \approx 3.2.$$

$$d. C_{avg}^{no} = \frac{1}{14} \cdot 3 \cdot 2 + \frac{1}{14} \cdot 4 \cdot 12 = \frac{54}{14} \approx 3.9.$$

2. For $n = 2^k$ (and omitting the subscript to simplify the notation), we get the recurrence relation

$$C(2^k) = C(2^{k-1}) + 1 \quad \text{for } k > 0, \quad C(1) = 1.$$

By making backward substitutions, we obtain the following:

$$\begin{aligned}
 C(2^k) &= C(2^{k-1}) + 1 && \text{substitute } C(2^{k-1}) = C(2^{k-2}) + 1 \\
 &= [C(2^{k-2}) + 1] + 1 = C(2^{k-2}) + 2 && \text{substitute } C(2^{k-2}) = C(2^{k-3}) + 1 \\
 &= [C(2^{k-3}) + 1] + 2 = C(2^{k-3}) + 3 && \dots \\
 &\dots && \dots \\
 &= C(2^{k-i}) + i \\
 &\dots \\
 &= C(2^{k-k}) + k.
 \end{aligned}$$

Thus, we end up with

$$C(2^k) = C(1) + k = 1 + k$$

or, after returning to the original variable $n = 2^k$ and hence $k = \log_2 n$,

$$C(n) = \log_2 n + 1.$$

3. a. Every positive integer n is bounded below and above by some consecutive powers of 2, i.e.,

$$2^k \leq n < 2^{k+1},$$

where k is a nonnegative integer (uniquely defined by the value of n). Taking the base-2 logarithms, we obtain

$$k \leq \log_2 n < k + 1.$$

Hence $k \leq \lfloor \log_2 n \rfloor$ and, since $\lfloor \log_2 n \rfloor \leq \log_2 n < k + 1$, we have the inequality

$$k \leq \lfloor \log_2 n \rfloor < k + 1,$$

which implies that $\lfloor \log_2 n \rfloor = k$. Similarly, $2^k < n + 1 \leq 2^{k+1}$, and an argument analogous to the one just used to show that $\lfloor \log_2 n \rfloor = k$, we obtain that $\lceil \log_2(n + 1) \rceil = k + 1$. Hence, $\lceil \log_2(n + 1) \rceil = \lfloor \log_2 n \rfloor + 1$.

- b. The case of even n was considered in the text of Chapter 4.3. Let $n > 1$ be odd, i.e., $n = 2i + 1$, where $i > 0$. The left-hand side is:

$$\begin{aligned} C_w(n) &= \lfloor \log_2 n \rfloor + 1 = \lfloor \log_2(2i + 1) \rfloor + 1 = \lceil \log_2(2i + 2) \rceil = \lceil \log_2 2(i + 1) \rceil \\ &= \lceil \log_2 2 + \log_2(i + 1) \rceil = 1 + \lceil \log_2(i + 1) \rceil = 1 + (\lfloor \log_2 i \rfloor + 1) = \lfloor \log_2 i \rfloor + 2. \end{aligned}$$

The right-hand side is:

$$\begin{aligned} C_w(\lfloor n/2 \rfloor) + 1 &= C_w(\lfloor (2i + 1)/2 \rfloor) + 1 = C_w(\lfloor i + 1/2 \rfloor) + 1 = C_w(i) + 1 = \\ &= (\lfloor \log_2 i \rfloor + 1) + 1 = \lfloor \log_2 i \rfloor + 2, \end{aligned}$$

which is the same as the left-hand side.

The initial condition is verified immediately: $C_w(1) = \lfloor \log_2 1 \rfloor + 1 = 1$.

4. The ratio in question can be estimated as follows:

$$\frac{C_{avg}^{seq.}(n)}{C_{avg}^{bin.}(n)} \approx \frac{n/2}{\log_2 n} = (\text{for } n = 10^5) \frac{10^5/2}{\log_2 10^5} = \frac{1}{2} \cdot \frac{10^5}{5 \log_2 10} = \frac{10^4}{\log_2 10} \approx 3000.$$

5. Unlike an array, where any element can be accessed in constant time, reaching the middle element in a linked list is a $\Theta(n)$ operation. Hence, though implementable in principle, binary search would be a horribly inefficient algorithm for searching in a (sorted) linked list.

6. Step 1: Check whether $A[0] \leq U$ and $A[n-1] \leq L$. If this is not true, stop: there are no such elements.

Step 2: Search for L using the text's version of binary search. If the search was successful, record the value m of the index returned by the algorithm; if the search was unsuccessful, record the value of l on the exit from the algorithm.

Step 3: Search for U using the text's version of binary search. If the search was successful, record the value m of the index returned by the algorithm; if the search was unsuccessful, record the value of r on the exit from the algorithm.

The final answer (if the problem has a solution) is the range of the array indices between l and r (inclusively), where l and r are the values recorded in Steps 2 and 3, respectively.

7. Call $BSR(A[0..n-1], K)$ where

Algorithm $BSR(A[l..r], K)$
 //Implements binary search recursively.
 //Input: A sorted (sub)array $A[l..r]$ and a search key K
 //Output: An index of the array's element equal to K
 // or -1 if there is no such element.
if $l > r$ **return** -1
else $m \leftarrow \lfloor (l+r)/2 \rfloor$
 if $K = A[m]$ **return** m
 else if $K < A[m]$ **return** $BSR(A[l..m-1], K)$
 else return $BSR(A[m+1..r], K)$

8. **Algorithm** $TwoWayBinarySearch(A[0..n-1], K)$

//Implements binary search with two-way comparisons
 //Input: A sorted array $A[0..n-1]$ and a search key K
 //Output: An index of the array's element equal to K
 // or -1 if there is no such element.
 $l \leftarrow 0; \quad r \leftarrow n-1$
while $l < r$ **do**
 $m \leftarrow \lfloor (l+r)/2 \rfloor$
 if $K \leq A[m]$
 $r \leftarrow m$
 else $l \leftarrow m+1$
if $K = A[l]$ **return** l
else return -1

9. Algorithm $TwoWayBinarySearch$ makes $\lceil \log_2 n \rceil + 1$ two-way comparisons in the worst case, which is obtained by solving the recurrence $C_w(n) = C_w(\lceil n/2 \rceil) + 1$ for $n > 1$, $C_w(1) = 1$. Also note that the best-case efficiency of this algorithm is not in $\Theta(1)$ but in $\Theta(\log n)$.

10. Apply a two-way comparison version of binary search using the picture numbering. That is, assuming that pictures are numbered from 1 to 42, start with a question such as “Is the picture’s number > 21 ?”. The largest number of questions that may be required is 6. (Because the search can be assumed successful, one less comparison needs to be made than in *TwoWayBinarySearch*, yielding here $\lceil \log_2 42 \rceil = 6$.)

Exercises 4.4

1. Design a divide-and-conquer algorithm for computing the number of levels in a binary tree. What is the efficiency class of your algorithm?
2. The following algorithm seeks to compute the number of leaves in a binary tree.

Algorithm *LeafCounter*(T)

//Computes recursively the number of leaves in a binary tree

//Input: A binary tree T


//Output: The number of leaves in T

if $T = \emptyset$ **return** 0

else return *LeafCounter*(T_L) + *LeafCounter*(T_R)

Is this algorithm correct? If it is, prove it; if it is not, make an appropriate correction.

3. Prove equality (4.5) by mathematical induction.
4. Traverse the following binary tree
 - a. in preorder.
 - b. in inorder.
 - c. in postorder.



5. Write a pseudocode for one of the classic traversal algorithms (preorder, inorder, and postorder) for binary trees. Assuming that your algorithm is recursive, find the number of recursive calls made.
6. Which of the three classic traversal algorithms yields a sorted list if applied to a binary search tree? Prove this property.
7. a. Draw a binary tree with ten nodes labeled 0, 1, 2, ..., 9 in such a way that the inorder and postorder traversals of the tree yield the following lists: 9, 3, 1, 0, 4, 2, 7, 6, 8, 5 (inorder) and 9, 1, 4, 0, 3, 6, 7, 5, 8, 2 (postorder).
 - b. Give an example of two permutations of the same n labels 0, 1, 2, ..., $n-1$ that cannot be inorder and postorder traversal lists of the same binary tree.
 - c. Design an algorithm that constructs a binary tree for which two given lists of n labels 0, 1, 2, ..., $n-1$ are generated by the inorder and postorder traversals of the tree. Your algorithm should also identify inputs for which the problem has no solution.

8. \triangleright The *internal path length* I of an extended binary tree is defined as the sum of the lengths of the paths—taken over all internal nodes—from the root to each internal node. Similarly, the *external path length* E of an extended binary tree is defined as the sum of the lengths of the paths—taken over all external nodes—from the root to each external node. Prove that $E = I + 2n$ where n is the number of internal nodes in the tree.
9. Write a program for computing the internal path length of a binary search tree. Use it to investigate empirically the average number of key comparisons for searching in a randomly generated binary search tree.
10. *Chocolate bar puzzle* Given an n -by- m chocolate bar, you need to break it into nm 1-by-1 pieces. You can break a bar only in a straight line, and only one bar can be broken at a time. Design an algorithm that solves the problem with the minimum number of bar breaks. What is this minimum number? Justify your answer by using properties of a binary tree.

Hints to Exercises 4.4

1. The problem is almost identical to the one discussed in this section.
2. Trace the algorithm on a small input.
3. Use strong induction on the number of internal nodes.
4. This is a standard exercise that you have probably done in your data structures course. With the traversal definitions given at the end of this section, you should be able to trace them even if you have never seen them before.
5. The pseudocodes can simply mirror the traversals' descriptions.
6. If you do not know the answer to this important question, you may want to check the results of the traversals on a small binary search tree. For a proof, answer the question: what can be said about two nodes with keys k_1 and k_2 if $k_1 < k_2$?
7. Find the root's label of the binary tree first and then identify the labels of the nodes in its left and right subtrees.
8. Use strong induction on the number of internal nodes.
9. n/a
10. Breaking the chocolate bar can be represented by a binary tree.

Solutions to Exercises 4.4

1. **Algorithm** *Levels*(T)

//Computes recursively the number of levels in a binary tree

//Input: Binary tree T

//Output: Number of levels in T

if $T = \emptyset$ **return** 0

else return $\max\{\text{Levels}(T_L), \text{Levels}(T_R)\} + 1$

This is a $\Theta(n)$ algorithm, by the same reason *Height*(T) discussed in the section is.

2. The algorithm is incorrect because it returns 0 instead of 1 for the one-node binary tree. Here is a corrected version:

Algorithm *LeafCounter*(T)

//Computes recursively the number of leaves in a binary tree

//Input: A binary tree T

//Output: The number of leaves in T

if $T = \emptyset$ **return** 0

//empty tree

else if $T_L = \emptyset$ **and** $T_R = \emptyset$ **return** 1

//one-node tree

else return $\text{LeafCounter}(T_L) + \text{LeafCounter}(T_R)$ //general case

3. Here is a proof of equality (4.5) by strong induction on the number of internal nodes $n \geq 0$. The basis step is true because for $n = 0$ we have the empty tree whose extended tree has 1 external node by definition. For the inductive step, let us assume that

$$x = k + 1$$

for any extended binary tree with $0 \leq k < n$ internal nodes. Let T be a binary tree with n internal nodes and let n_L and x_L be the numbers of internal and external nodes in the left subtree of T , respectively, and let n_R and x_R be the numbers of internal and external nodes in the right subtree of T , respectively. Since $n > 0$, T has a root, which is its internal node, and hence

$$n = n_L + n_R + 1.$$

Since both $n_L < n$ and $n_R < n$, we can use equality (4.5), assumed to be correct for the left and right subtree of T , to obtain the following:

$$x = x_L + x_R = (n_L + 1) + (n_R + 1) = (n_L + n_R + 1) + 1 = n + 1,$$

which completes the proof.

4. a. Preorder: $a\ b\ d\ e\ c\ f$

b. Inorder: $d\ b\ e\ a\ c\ f$

c. Postorder: $d\ e\ b\ f\ c\ a$

5. Here is a pseudocode of the preorder traversal:

Algorithm *Preorder*(T)

//Implements the preorder traversal of a binary tree

//Input: Binary tree T (with labeled vertices)

//Output: Node labels listed in preorder

if $T \neq \emptyset$

 print label of T 's root

Preorder(T_L) // T_L is the root's left subtree

Preorder(T_R) // T_R is the root's right subtree


The number of calls, $C(n)$, made by the algorithm is equal to the number of nodes, both internal and external, in the extended tree. Hence, according to the formula in the section,

$$C(n) = 2n + 1.$$

6. The inorder traversal yields a sorted list of keys of a binary search tree. In order to prove it, we need to show that if $k_1 < k_2$ are two keys in a binary search tree then the inorder traversal visits the node containing k_1 before the node containing k_2 . Let k_3 be the key at their nearest common ancestor. (Such a node is uniquely defined for any pair of nodes in a binary tree. If one of the two nodes at hand is an ancestor of the other, their nearest common ancestor coincides with the ancestor.) If the k_3 's node differ from both k_1 's node and k_2 's node, the definition of a binary search tree implies that k_1 and k_2 are in the left and right subtrees of k_3 , respectively. If k_3 's node coincides with k_2 's node (k_1 's node), k_1 's node (k_2 's node) is in the left (right) subtree rooted at k_2 's node (k_1 's node). In each of these cases, the inorder traversal visits k_1 's node before k_2 's node.

7. a. The root's label is listed last in the postorder tree: hence, it is 2. The labels preceding 2 in the order list—9,3,1,0,4—form the inorder traversal list of the left subtree; the corresponding postorder list for the left subtree traversal is given by the first four labels of the postorder list: 9,1,4,0,3. Similarly, for the right subtree, the inorder and postorder lists are, respectively, 7,6,8,5 and 6,7,5,8. Applying the same logic recursively to each of

the subtrees yields the following binary tree:



b. There is no such example for $n = 2$. For $n = 3$, lists 0,1,2 (inorder) and 2,0,1 (postorder) provide one.

c. The problem can be solved by a recursive algorithm based on the following observation: There exists a binary tree with inorder traversal list i_0, i_1, \dots, i_{n-1} and postorder traversal list p_0, p_1, \dots, p_{n-1} if and only if $p_{n-1} = i_k$ (the root's label), the sets formed by the first k labels in both lists are the same: $\{i_0, i_1, \dots, i_{k-1}\} = \{p_0, p_1, \dots, p_{k-1}\}$ (the labels of the nodes in the left subtree) and the sets formed by the other $n - k - 1$ labels excluding the root are the same: $\{i_{k+1}, i_{k+2}, \dots, i_{n-1}\} = \{p_k, p_{k+1}, \dots, p_{n-2}\}$ (the labels of the nodes in the right subtree).

Algorithm $Tree(i_0, i_1, \dots, i_{n-1}, p_0, p_1, \dots, p_{n-1})$

//Construct recursively the binary tree based on the inorder and postorder traversal lists

//Input: Lists i_0, i_1, \dots, i_{n-1} and p_0, p_1, \dots, p_{n-1} of inorder and postorder traversals, respectively

//Output: Binary tree T , specified in preorder, whose inorder and postorder traversals yield the lists given or

// -1 if such a tree doesn't exist

Find element i_k in the inorder list that is equal to the last element p_{n-1} of the postorder list.

if the previous search was unsuccessful **return** -1

else $print(i_k)$

$Tree(i_0, i_1, \dots, i_{k-1}, p_0, p_1, \dots, p_{k-1})$

$Tree(i_{k+1}, i_{k+2}, \dots, i_{n-1}, p_k, p_{k+1}, \dots, p_{n-2})$

8. We can prove equality $E = I + 2n$, where E and I are, respectively, the external and internal path lengths in an extended binary tree with n internal nodes by induction on n . The basis case, for $n = 0$, holds because both E and I are equal to 0 in the extended tree of the empty binary tree.

For the general case of induction, we assume that

$$E = I + 2k$$

for any extended binary tree with $0 \leq k < n$ internal nodes. To prove the equality for an extended binary tree T with n internal nodes, we are going to use this equality for T_L and T_R , the left and right subtrees of T . (Since $n > 0$, the root of the tree is an internal node, and hence the number of internal nodes in both the left and right subtree is less than n .) Thus,

$$E_L = I_L + 2n_L,$$

where E_L and I_L are external and internal paths, respectively, in the left subtree T_L , which has n_L internal and x_L external nodes, respectively. Similarly,

$$E_R = I_R + 2n_R,$$

where E_R and I_R are external and internal paths, respectively, in the right subtree T_R , which has n_R internal and x_R external nodes, respectively. Since the length of the simple path from the root of $T_L(T_R)$ to a node in $T_L(T_R)$ is one less than the length of the simple path from the root of T to that node, we have

$$\begin{aligned} E &= (E_L + x_L) + (E_R + x_R) \\ &= (I_L + 2n_L + x_L) + (I_R + 2n_R + x_R) \\ &= [(I_L + n_L) + (I_R + n_R)] + (n_L + n_R) + (x_L + x_R) \\ &= I + (n - 1) + x, \end{aligned}$$

where x is the number of external nodes in T . Since $x = n + 1$ (see Section 4.4), we finally obtain the desired equality:

$$E = I + (n - 1) + x = I + 2n.$$

9. n/a
10. We can represent operations of any algorithm solving the problem by a full binary tree in which parental nodes represent breakable pieces and leaves represent 1-by-1 pieces of the original bar. The number of the latter is nm ; and the number of the former, which is equal to the number of the bar breaks, is one less, i.e., $nm - 1$, according to equation (4.5) in Section 4.4. (Note: This elegant solution was suggested to the author by Simon Berkovich, one of the book's reviewers.)

Alternatively, we can reason as follows: Since only one bar can be broken at a time, any break increases the number of pieces by 1. Hence, $nm - 1$

breaks are needed to get from a single n -by- m piece to nm one-by-one pieces, which is obtained by *any* sequence of $nm - 1$ allowed breaks. (The same argument can be made more formally by mathematical induction.)

Exercises 4.5

1. What are the smallest and largest numbers of digits the product of two decimal n -digit integers can have?
2. Compute $2101 * 1130$ by applying the divide-and-conquer algorithm outlined in the text.
3. a. Prove the equality $a^{\log_b c} = c^{\log_b a}$, which was used twice in Section 4.5.
 b. Why is $n^{\log_2 3}$ better than $3^{\log_2 n}$ as a closed-form formula for $M(n)$?
4. a. Why did we not include multiplications by 10^n in the multiplication count $M(n)$ of the large-integer multiplication algorithm?
 b. In addition to assuming that n is a power of 2, we made, for the sake of simplicity, another, more subtle, assumption in setting up a recurrence relation for $M(n)$ which is not always true (it does not change the final answer, however.) What is this assumption?
5. How many one-digit additions are made by the pen-and-pencil algorithm in multiplying two n -digit integers? (You may disregard potential carries.)
6. Verify the formulas underlying Strassen's algorithm for multiplying 2-by-2 matrices.
7. Apply Strassen's algorithm to compute

$$\begin{bmatrix} 1 & 0 & 2 & 1 \\ 4 & 1 & 1 & 0 \\ 0 & 1 & 3 & 0 \\ 5 & 0 & 2 & 1 \end{bmatrix} * \begin{bmatrix} 0 & 1 & 0 & 1 \\ 2 & 1 & 0 & 4 \\ 2 & 0 & 1 & 1 \\ 1 & 3 & 5 & 0 \end{bmatrix}$$

exiting the recursion when $n = 2$, i.e., computing the products of 2-by-2 matrices by the brute-force algorithm.

8. Solve the recurrence for the number of additions required by Strassen's algorithm. (Assume that n is a power of 2.)
9. V. Pan [Pan78] has discovered a divide-and-conquer matrix multiplication algorithm that is based on multiplying two 70-by-70 matrices using 143,640 multiplications. Find the asymptotic efficiency of Pan's algorithm (you can ignore additions) and compare it with that of Strassen's algorithm.
10. Practical implementations of Strassen's algorithm usually switch to the brute-force method after matrix sizes become smaller than some "crossover point". Run an experiment to determine such crossover point on your computer system.

Hints to Exercises 4.5

1. You might want to answer the question for $n = 2$ first and then generalize it.
2. Trace the algorithm on the input given. You will have to use it again in order to compute products of two-digit numbers as well.
3. a. Take logarithms of both sides of the equality.

b. What did we use the closed-form formula for?
4. a. How do we multiply by powers of 10?

b. Try to repeat the argument for, say, $98 * 76$.
5. Counting the number of one-digit additions made by the pen-and-pencil algorithm in multiplying, say, two four-digit numbers, should help answering the general question.
6. Check the formulas by simple algebraic manipulations.
7. Trace Strassen's algorithm on the input given. (It takes some work, but it would have been much more of it if you were asked to stop the recursion when $n = 1$.) It is a good idea to check your answer by multiplying the matrices by the brute-force (i.e., definition-based) algorithm, too.
8. Use the method of backward substitutions to solve the recurrence given in the text.
9. The recurrence for the number of multiplications in Pan's algorithm is similar to that for Strassen's algorithm. Use the Master Theorem to find the order of growth of its solution.
10. n/a

Solutions to Exercises 4.5

1. The smallest decimal n -digit positive integer is $\underbrace{10\dots 0}_{n-1}$, i. e., 10^{n-1} . The product of two such numbers is $10^{n-1} \cdot 10^{n-1} = 10^{2n-2}$, which has $2n-1$ digits (1 followed by $2n-2$ zeros).

The largest decimal n -digit integer is $\underbrace{9\dots 9}_n$, i.e., $10^n - 1$. The product of two such numbers is $(10^n - 1)(10^n - 1) = 10^{2n} - 2 \cdot 10^n + 1$, which has $2n$ digits (because 10^{2n-1} and $10^{2n} - 1$ are the smallest and largest numbers with $2n$ digits, respectively, and $10^{2n} - 1 < 10^{2n} - 2 \cdot 10^n + 1 < 10^{2n} - 1$).

2. For $2101 * 1130$:

$$\begin{aligned} c_2 &= 21 * 11 \\ c_0 &= 01 * 30 \\ c_1 &= (21 + 01) * (11 + 30) - (c_2 + c_0) = 22 * 41 - 21 * 11 - 01 * 30. \end{aligned}$$

For $21 * 11$:

$$\begin{aligned} c_2 &= 2 * 1 = 2 \\ c_0 &= 1 * 1 = 1 \\ c_1 &= (2 + 1) * (1 + 1) - (2 + 1) = 3 * 2 - 3 = 3. \\ \text{So, } 21 * 11 &= 2 \cdot 10^2 + 3 \cdot 10^1 + 1 = 231. \end{aligned}$$

For $01 * 30$:

$$\begin{aligned} c_2 &= 0 * 3 = 0 \\ c_0 &= 1 * 0 = 0 \\ c_1 &= (0 + 1) * (3 + 0) - (0 + 0) = 1 * 3 - 0 = 3. \\ \text{So, } 01 * 30 &= 0 \cdot 10^2 + 3 \cdot 10^1 + 0 = 30. \end{aligned}$$

For $22 * 41$:

$$\begin{aligned} c_2 &= 2 * 4 = 8 \\ c_0 &= 2 * 1 = 2 \\ c_1 &= (2 + 2) * (4 + 1) - (8 + 2) = 4 * 5 - 10 = 10. \\ \text{So, } 22 * 41 &= 8 \cdot 10^2 + 10 \cdot 10^1 + 2 = 902. \end{aligned}$$

Hence

$$2101 * 1130 = 231 \cdot 10^4 + (902 - 231 - 30) \cdot 10^2 + 30 = 2,374,130.$$

3. a. Taking the base- b logarithms of both hand sides of the equality $a^{\log_b c} = c^{\log_b a}$ yields $\log_b c \log_b a = \log_b a \log_b c$. Since two numbers are equal if and only if their logarithms to the same base are equal, the equality in question is proved.
- b. It is easier to compare $n^{\log_2 3}$ with n^2 (the number of digit multiplications made by the classic algorithm) than $3^{\log_2 n}$ with n^2 .
4. a. When working with decimal integers, multiplication by a power of 10 can be done by a shift.
- b. In the formula for c_1 , the sum of two $n/2$ -digit integers can have not $n/2$ digits, as it was assumed, but $n/2 + 1$.
5. Let a and b be two n -digit integers such that the product of each pair of their digits is a one-digit number. Then the result of the pen-and-pencil algorithm will look as follows:

$a :$	a_{n-1}	\dots	a_1	a_0
$b :$	b_{n-1}	\dots	b_1	b_0
	$d_{n-1,0}$	\dots	$d_{1,0}$	$d_{0,0}$
$d_{n-1,1}$	$d_{n-2,1}$		$d_{0,1}$	
	$d_{n-2,n-1}$	\dots		
$d_{n-1,n-1}$	$d_{n-2,n-1}$	\dots	$d_{1,n-1}$	$d_{0,n-1}$
# additions	0	1	\dots	$n-2$
			$n-1$	\dots
				1
				0

Hence, the total number of additions without carries will be

$$\begin{aligned}
& 1 + 2 + \dots + (n-1) + (n-2) + \dots + 1 \\
&= [1 + 2 + \dots + (n-1)] + [(n-2) + \dots + 1] \\
&= \frac{(n-1)n}{2} + \frac{(n-2)(n-1)}{2} = (n-1)^2.
\end{aligned}$$

6. $m_1 + m_4 - m_5 + m_7 =$
 $(a_{00} + a_{11})(b_{00} + b_{11}) + a_{11}(b_{10} - b_{00}) - (a_{00} + a_{01})b_{11} + (a_{01} - a_{11})(b_{10} + b_{11}) =$
 $a_{00}b_{00} + a_{11}b_{00} + a_{00}b_{11} + a_{11}b_{11} + a_{11}b_{10} - a_{11}b_{00} - a_{00}b_{11} - a_{01}b_{11} + a_{01}b_{10} -$
 $a_{11}b_{10} + a_{01}b_{11} - a_{11}b_{11}$
 $= a_{00}b_{00} + a_{01}b_{10}$
- $m_3 + m_5 = a_{00}(b_{01} - b_{11}) + (a_{00} + a_{01})b_{11} = a_{00}b_{01} - a_{00}b_{11} + a_{00}b_{11} + a_{01}b_{11} =$
 $a_{00}b_{01} + a_{01}b_{11}$
- $m_2 + m_4 = (a_{10} + a_{11})b_{00} + a_{11}(b_{10} - b_{00}) = a_{10}b_{00} + a_{11}b_{00} + a_{11}b_{10} - a_{11}b_{00} =$

$$a_{10}b_{00} + a_{11}b_{10}$$

$$\begin{aligned} m_1 + m_3 - m_2 + m_6 &= (a_{00} + a_{11})(b_{00} + b_{11}) + a_{00}(b_{01} - b_{11}) - (a_{10} + a_{11})b_{00} + (a_{10} - a_{00})(b_{00} + b_{01}) = \\ &= a_{00}b_{00} + a_{11}b_{00} + a_{00}b_{11} + a_{11}b_{11} + a_{00}b_{01} - a_{00}b_{11} - a_{10}b_{00} - a_{11}b_{00} + a_{10}b_{00} - \\ &= a_{10}b_{01} + a_{11}b_{11}. \end{aligned}$$

7. For the matrices given, Strassen's algorithm yields the following:

$$C = \left[\begin{array}{c|c} C_{00} & C_{01} \\ \hline C_{10} & C_{11} \end{array} \right] = \left[\begin{array}{c|c} A_{00} & A_{01} \\ \hline A_{10} & A_{11} \end{array} \right] \left[\begin{array}{c|c} B_{00} & B_{01} \\ \hline B_{10} & B_{11} \end{array} \right]$$

where

$$A_{00} = \begin{bmatrix} 1 & 0 \\ 4 & 1 \end{bmatrix}, \quad A_{01} = \begin{bmatrix} 2 & 1 \\ 1 & 0 \end{bmatrix}, \quad A_{10} = \begin{bmatrix} 0 & 1 \\ 5 & 0 \end{bmatrix}, \quad A_{11} = \begin{bmatrix} 3 & 0 \\ 2 & 1 \end{bmatrix},$$

$$B_{00} = \begin{bmatrix} 0 & 1 \\ 2 & 1 \end{bmatrix}, \quad B_{01} = \begin{bmatrix} 0 & 1 \\ 0 & 4 \end{bmatrix}, \quad B_{10} = \begin{bmatrix} 2 & 0 \\ 1 & 3 \end{bmatrix}, \quad B_{11} = \begin{bmatrix} 1 & 1 \\ 5 & 0 \end{bmatrix}.$$

Therefore,

$$M_1 = (A_{00} + A_{11})(B_{00} + B_{11}) = \begin{bmatrix} 4 & 0 \\ 6 & 2 \end{bmatrix} \begin{bmatrix} 1 & 2 \\ 7 & 1 \end{bmatrix} = \begin{bmatrix} 4 & 8 \\ 20 & 14 \end{bmatrix},$$

$$M_2 = (A_{10} + A_{11})B_{00} = \begin{bmatrix} 3 & 1 \\ 7 & 1 \end{bmatrix} \begin{bmatrix} 0 & 1 \\ 2 & 1 \end{bmatrix} = \begin{bmatrix} 2 & 4 \\ 2 & 8 \end{bmatrix},$$

$$M_3 = A_{00}(B_{01} - B_{11}) = \begin{bmatrix} 1 & 0 \\ 4 & 1 \end{bmatrix} \begin{bmatrix} -1 & 0 \\ -5 & 4 \end{bmatrix} = \begin{bmatrix} -1 & 0 \\ -9 & 4 \end{bmatrix},$$

$$M_4 = A_{11}(B_{10} - B_{00}) = \begin{bmatrix} 3 & 0 \\ 2 & 1 \end{bmatrix} \begin{bmatrix} 2 & -1 \\ -1 & 2 \end{bmatrix} = \begin{bmatrix} 6 & -3 \\ 3 & 0 \end{bmatrix},$$

$$M_5 = (A_{00} + A_{01})B_{11} = \begin{bmatrix} 3 & 1 \\ 5 & 1 \end{bmatrix} \begin{bmatrix} 1 & 1 \\ 5 & 0 \end{bmatrix} = \begin{bmatrix} 8 & 3 \\ 10 & 5 \end{bmatrix},$$

$$M_6 = (A_{10} - A_{00})(B_{00} + B_{01}) = \begin{bmatrix} -1 & 1 \\ 1 & -1 \end{bmatrix} \begin{bmatrix} 0 & 2 \\ 2 & 5 \end{bmatrix} = \begin{bmatrix} 2 & 3 \\ -2 & -3 \end{bmatrix},$$

$$M_7 = (A_{01} - A_{11})(B_{10} + B_{11}) = \begin{bmatrix} -1 & 1 \\ -1 & -1 \end{bmatrix} \begin{bmatrix} 3 & 1 \\ 6 & 3 \end{bmatrix} = \begin{bmatrix} 3 & 2 \\ -9 & -4 \end{bmatrix}.$$

Accordingly,

$$\begin{aligned}
C_{00} &= M_1 + M_4 - M_5 + M_7 \\
&= \begin{bmatrix} 4 & 8 \\ 20 & 14 \end{bmatrix} + \begin{bmatrix} 6 & -3 \\ 3 & 0 \end{bmatrix} - \begin{bmatrix} 8 & 3 \\ 10 & 5 \end{bmatrix} + \begin{bmatrix} 3 & 2 \\ -9 & -4 \end{bmatrix} = \begin{bmatrix} 5 & 4 \\ 4 & 5 \end{bmatrix}, \\
C_{01} &= M_3 + M_5 \\
&= \begin{bmatrix} -1 & 0 \\ -9 & 4 \end{bmatrix} + \begin{bmatrix} 8 & 3 \\ 10 & 5 \end{bmatrix} = \begin{bmatrix} 7 & 3 \\ 1 & 9 \end{bmatrix}, \\
C_{10} &= M_2 + M_4 \\
&= \begin{bmatrix} 2 & 4 \\ 2 & 8 \end{bmatrix} + \begin{bmatrix} 6 & -3 \\ 3 & 0 \end{bmatrix} = \begin{bmatrix} 8 & 1 \\ 5 & 8 \end{bmatrix}, \\
C_{11} &= M_1 + M_3 - M_2 + M_6 \\
&= \begin{bmatrix} 4 & 8 \\ 20 & 14 \end{bmatrix} + \begin{bmatrix} -1 & 0 \\ -9 & 4 \end{bmatrix} - \begin{bmatrix} 2 & 4 \\ 2 & 8 \end{bmatrix} + \begin{bmatrix} 2 & 3 \\ -2 & -3 \end{bmatrix} = \begin{bmatrix} 3 & 7 \\ 7 & 7 \end{bmatrix}.
\end{aligned}$$

That is,

$$C = \begin{bmatrix} 5 & 4 & 7 & 3 \\ 4 & 5 & 1 & 9 \\ 8 & 1 & 3 & 7 \\ 5 & 8 & 7 & 7 \end{bmatrix}.$$

8. For $n = 2^k$, the recurrence $A(n) = 7A(n/2) + 18(n/2)^2$ for $n > 1$, $A(1) = 0$, becomes

$$A(2^k) = 7A(2^{k-1}) + \frac{9}{2}4^k \quad \text{for } k > 1, \quad A(1) = 0.$$

Solving it by backward substitutions yields the following:

$$\begin{aligned}
A(2^k) &= 7A(2^{k-1}) + \frac{9}{2}4^k \\
&= 7[7A(2^{k-2}) + \frac{9}{2}4^{k-1}] + \frac{9}{2}4^k = 7^2A(2^{k-2}) + 7 \cdot \frac{9}{2}4^{k-1} + \frac{9}{2}4^k \\
&= 7^2[7A(2^{k-3}) + \frac{9}{2}4^{k-2}] + 7 \cdot \frac{9}{2}4^{k-1} + \frac{9}{2}4^k \\
&= 7^3A(2^{k-3}) + 7^2 \cdot \frac{9}{2}4^{k-2} + 7 \cdot \frac{9}{2}4^{k-1} + \frac{9}{2}4^k \\
&= \dots \\
&= 7^k A(2^{k-k}) + \frac{9}{2} \sum_{i=0}^{k-1} 7^i 4^{k-i} = 7^k \cdot 0 + \frac{9}{2}4^k \sum_{i=0}^{k-1} \left(\frac{7}{4}\right)^i \\
&= \frac{9}{2}4^k \frac{(7/4)^k - 1}{(7/4) - 1} = 6(7^k - 4^k).
\end{aligned}$$

Returning back to the variable $n = 2^k$, we obtain

$$A(n) = 6(7^{\log_2 n} - 4^{\log_2 n}) = 6(n^{\log_2 7} - n^2).$$

(Note that the number of additions in Strassen's algorithm has the same order of growth as the number of multiplications: $\Theta(n^s)$ where $s = n^{\log_2 7} \approx n^{2.807}$.)

9. The recurrence for the number of multiplications in Pan's algorithm is

$$M(n) = 143640M(n/70) \quad \text{for } n > 1, \quad M(1) = 1.$$

Solving it for $n = 70^k$ or applying the Master Theorem yields $M(n) \in \Theta(n^p)$ where

$$p = \log_{70} 143640 = \frac{\ln 143640}{\ln 70} \approx 2.795.$$

This number is slightly smaller than the exponent of Strassen's algorithm

$$s = \log_2 7 = \frac{\ln 7}{\ln 2} \approx 2.807.$$

10. n/a

Exercises 4.6

1. a. For the one-dimensional version of the closest-pair problem, i.e., for the problem of finding two closest numbers among a given set of n real numbers, design an algorithm that is directly based on the divide-and-conquer technique and determine its efficiency class.

b. Is it a good algorithm for this problem?
2. Consider the version of the divide-and-conquer two-dimensional closest-pair algorithm in which we simply sort each of the two sets C_1 and C_2 in ascending order of their y coordinates on each recursive call. Assuming that sorting is done by mergesort, set up a recurrence relation for the running time in the worst case and solve it for $n = 2^k$.
3. Implement the divide-and-conquer closest-pair algorithm, outlined in this section, in the language of your choice.
4. Find a visualization of an algorithm for the closest-pair problem on the Web. What algorithm does this visualization represent?
5. The **Voronoi polygon** for a point P of a set S of points in the plane is defined to be the perimeter of the set of all points in the plane closer to P than to any other point in S . The union of all the Voronoi polygons of the points in S is called the **Voronoi diagram** of S .

a. What is the Voronoi diagram for a set of three points?

b. Find a visualization of an algorithm for generating the Voronoi diagram on the Web and study a few examples of such diagrams. Based on your observations, can you tell how the solution to the previous question is generalized to the general case?
6. Explain how one can find point P_{\max} in the quickhull algorithm analytically.
7. What is the best-case efficiency of quickhull?
8. Give a specific example of inputs that make the quickhull algorithm run in quadratic time.
9. Implement the quickhull algorithm in the language of your choice.
10. *Shortest path around* There is a fenced area in the two-dimensional Euclidean plane in the shape of a convex polygon with vertices at points $P_1(x_1, y_1)$, $P_2(x_2, y_2)$, ..., $P_n(x_n, y_n)$ (not necessarily in this order). There are two more points, $A(x_A, y_A)$ and $B(x_B, y_B)$, such that $x_A < \min\{x_1, x_2, \dots, x_n\}$ and $x_B > \max\{x_1, x_2, \dots, x_n\}$. Design a reasonably efficient algorithm for computing the length of the shortest path between A and B . [ORo98], p.68

Hints to Exercises 4.6

1. a. How many points need to be considered in the combining-solutions stage of the algorithm?

b. Design a simpler algorithm in the same efficiency class.
2. Recall (see Section 4.1) that the number of comparisons made by mergesort in the worst case is $C_{worst}(n) = n \log_2 n - n + 1$ (for $n = 2^k$). You may use just the highest-order term of this formula in the recurrence you need to set up.
3. n/a
4. n/a
5. The answer to part (a) comes directly from a textbook on plane geometry.
6. Use the formula relating the value of a determinant with the area of a triangle.
7. It must be in $\Omega(n)$, of course. (Why?)
8. Design a sequence of n points for which the algorithm decreases the problem's size just by one on each of its recursive calls.
9. n/a
10. The path cannot cross inside the fenced area but it can go along the fence.

Solutions to Exercises 4.6

1. a. Assuming that the points are sorted in increasing order, we can find the closest pair (or, for simplicity, just the distance between two closest points) by comparing three distances: the distance between the two closest points in the first half of the sorted list, the distance between the two closest points in its second half, and the distance between the rightmost point in the first half and the leftmost point in the second half. Therefore, after sorting the numbers of a given array $P[0..n-1]$ in increasing order, we can call $ClosestNumbers(P[0..n-1])$, where

Algorithm $ClosestNumbers(P[l..r])$

// A divide-and-conquer alg. for the one-dimensional closest-pair problem

// Input: A subarray $P[l..r]$ ($l \leq r$) of a given array $P[0..n-1]$

// of real numbers sorted in nondecreasing order

// Output: The distance between the closest pair of numbers

if $r = l$ **return** ∞

else if $r - l = 1$ **return** $P[r] - P[l]$

else return $\min\{ClosestNumbers(P[l..\lfloor(l+r)/2\rfloor]),$
 $ClosestNumbers(P[\lceil(l+r)/2\rceil+1..r]),$
 $P[\lfloor(l+r)/2\rfloor+1] - P[\lceil(l+r)/2\rceil]\}$

For $n = 2^k$, the recurrence for the running time $T(n)$ of this algorithm is

$$T(n) = 2T(n/2) + c.$$

Its solution, according to the Master Theorem, is in $\Theta(n^{\log_2 2}) = \Theta(n)$. If sorting of the input's numbers is done with a $\Theta(n \log n)$ algorithm such as mergesort, the overall running time will be in $\Theta(n \log n) + \Theta(n) = \Theta(n \log n)$.

b. A simpler algorithm can sort the numbers given (e.g., by mergesort) and then compare the distances between the adjacent elements in the sorted list. The resulting algorithm has the same $\Theta(n \log n)$ efficiency but it is arguably simpler than the divide-and-conquer algorithms above.

Note: In fact, any algorithm that solves this problem must be in $\Omega(n \log n)$ (see Problem 11 in Exercises 11.1).

2. $T(n) = 2T(n/2) + 2\frac{n}{2} \log_2 \frac{n}{2}$ for $n > 2$ (and $n = 2^k$), $T(2) = 1$.


Thus, $T(2^k) = 2T(2^{k-1}) + 2^k(k-1)$. Solving it by backward substitutions yields the following:

$$\begin{aligned}
T(2^k) &= 2T(2^{k-1}) + 2^k(k-1) \\
&= 2[2T(2^{k-2}) + 2^{k-1}(k-2)] + 2^k(k-1) = 2^2T(2^{k-2}) + 2^k(k-2) + 2^k(k-1) \\
&= 2^2[2T(2^{k-3}) + 2^{k-2}(k-3)] + 2^k(k-2) + 2^k(k-1) = 2^3T(2^{k-3}) + 2^k(k-3) + 2^k(k-2) + 2^k(k-1) \\
&\dots \\
&= 2^iT(2^{k-i}) + 2^k(k-i) + 2^k(k-i+1) + \dots + 2^k(k-1) \\
&\dots \\
&= 2^{k-1}T(2^1) + 2^k + 2^k2 + \dots + 2^k(k-1) \\
&= 2^{k-1} + 2^k(1 + 2 + \dots + (k-1)) = 2^{k-1} + 2^k \frac{(k-1)k}{2} \\
&= 2^{k-1}(1 + (k-1)k) = \frac{n}{2}(1 + (\log_2 n - 1) \log_2 n) \in \Theta(n \log^2 n).
\end{aligned}$$

3. n/a

4. n/a

5. a. The Voronoi diagram of three points not on the same line is formed by the perpendicular bisectors of the sides of the triangle with vertices at P_1 , P_2 , and P_3 :



(If P_1 , P_2 , and P_3 lie on the same line, with P_2 between P_1 and P_3 , the Voronoi diagram is formed by the perpendicular bisectors of the segments with the endpoints at P_1 and P_2 and at P_2 and P_3 .)

- b. The Voronoi polygon of a set of points is made up of perpendicular bisectors; a point of their intersection has at least three of the set's points nearest to it.

6. Since all the points in question serve as the third vertex for triangles with the same base P_1P_n , the farthest point is the one that maximizes the area of such a triangle. The area of a triangle, in turn, can be computed as one half of the magnitude of the determinant

$$\begin{vmatrix} x_1 & y_1 & 1 \\ x_2 & y_2 & 1 \\ x_3 & y_3 & 1 \end{vmatrix} = x_1y_2 + x_3y_1 + x_2y_3 - x_3y_2 - x_2y_1 - x_1y_3.$$

In other words, P_{\max} is a point whose coordinates (x_3, y_3) maximize the absolute value of the above expression in which (x_1, y_1) and (x_2, y_2) are the coordinates of P_1 and P_n , respectively.

7. If all n points lie on the same line, both S_1 and S_2 will be empty and the convex hull (a line segment) will be found in linear time, assuming that the input points have been already sorted before the algorithm begins. Note: Any algorithm that finds the convex hull for a set of n points must be in $\Omega(n)$ because all n points must be processed before the convex hull is found.
8. Among many possible answers, one can take two endpoints of the horizontal diameter of some circumference as points P_1 and P_n and obtain the other points P_i , $i = 2, \dots, n-1$, of the set in question by placing them successively in the middle of the circumference's upper arc between P_{i-1} and P_n .
9. n/a
10. Find the upper and lower hulls of the set $\{A, B, P_1, \dots, P_n\}$ (e.g., by quickhull), compute their lengths (by summing up the lengths of the line segments making up the polygonal chains) and return the smaller of the two.

This file contains the exercises, hints, and solutions for Chapter 5 of the book "Introduction to the Design and Analysis of Algorithms," 2nd edition, by A. Levitin. The problems that might be challenging for at least some students are marked by \triangleright ; those that might be difficult for a majority of students are marked by \blacktriangleright .

Exercises 5.1

1. *Ferrying soldiers* A detachment of n soldiers must cross a wide and deep river with no bridge in sight. They notice two 12-year-old boys playing in a rowboat by the shore. The boat is so tiny, however, that it can only hold two boys or one soldier. How can the soldiers get across the river and leave the boys in joint possession of the boat? How many times need the boat pass from shore to shore?
2. \triangleright *Alternating glasses* There are $2n$ glasses standing next to each other in a row, the first n of them filled with a soda drink while the remaining n glasses are empty. Make the glasses alternate in a filled-empty-filled-empty pattern in the minimum number of glass moves. [Gar78], p.7.
3. Design a decrease-by-one algorithm for generating the power set of a set of n elements. (The power set of a set S is the set of all the subsets of S , including the empty set and S itself.)
4. Apply insertion sort to sort the list E, X, A, M, P, L, E in alphabetical order.
5. a. What sentinel should be put before the first element of an array being sorted in order to avoid checking the in-bound condition $j \geq 0$ on each iteration of the inner loop of insertion sort?

b. Will the version with the sentinel be in the same efficiency class as the original version?
6. Is it possible to implement insertion sort for sorting linked lists? Will it have the same $O(n^2)$ efficiency as the array version?
7. Consider the following version of insertion sort.

```

Algorithm InsertSort2( $A[0..n-1]$ )
for  $i \leftarrow 1$  to  $n-1$  do
     $j \leftarrow i-1$ 
    while  $j \geq 0$  and  $A[j] > A[j+1]$  do
        swap( $A[j], A[j+1]$ )
         $j \leftarrow j-1$ 

```

What is its time efficiency? How is it compared to that of the version given in the text?

8. Let $A[0..n-1]$ be an array of n sortable elements. (For simplicity, you can assume that all the elements are distinct.) Recall that a pair of its elements $(A[i], A[j])$ is called an ***inversion*** if $i < j$ and $A[i] > A[j]$.

a. What arrays of size n have the largest number of inversions and what is this number? Answer the same questions for the smallest number of inversions.

b.► Show that the average-case number of key comparisons in insertion sort is given by the formula

$$C_{avg}(n) \approx \frac{n^2}{4}.$$

9. ▷ Binary insertion sort uses binary search to find an appropriate position to insert $A[i]$ among the previously sorted $A[0] \leq \dots \leq A[i-1]$. Determine the worst-case efficiency class of this algorithm.

10. Shellsort (more accurately Shell's sort) is an important sorting algorithm which works by applying insertion sort to each of several interleaving sublists of a given list. On each pass through the list, the sublists in question are formed by stepping through the list with an increment h_i taken from some predefined decreasing sequence of step sizes, $h_1 > \dots > h_i > \dots > 1$, which must end with 1. (The algorithm works for any such sequence, though some sequences are known to yield a better efficiency than others. For example, the sequence 1, 4, 13, 40, 121, ... , used, of course, in reverse, is known to be among the best for this purpose.)

a. Apply shellsort to the list

$S, H, E, L, L, S, O, R, T, I, S, U, S, E, F, U, L$

b. Is shellsort a stable sorting algorithm?

c. Implement shellsort, straight insertion sort, binary insertion sort, mergesort, and quicksort in the language of your choice and compare their performance on random arrays of sizes 10^2 , 10^3 , 10^4 , and 10^5 as well as on increasing and decreasing arrays of these sizes.

Hints to Exercises 5.1

1. Solve the problem for $n = 1$.
2. You may consider pouring soda from a filled glass into an empty glass as one move.
3. Use the fact that all the subsets of an n -element set $S = \{a_1, \dots, a_n\}$ can be divided into two groups: those that contain a_n and those that do not.
4. Trace the algorithm as we did in the text for another input (see Fig. 5.4).
5. a. The sentinel should stop the smallest element from moving beyond the first position in the array.

b. Repeat the analysis performed in the text for the sentinel version.
6. Recall that we can access elements of a singly linked list only sequentially.
7. Since the only difference between the two versions of the algorithm is in the inner loop's operations, you should estimate the difference in the running times of one repetition of this loop.
8. a. Answering the questions for an array of three elements should lead to the general answers.

b. Assume for simplicity that all elements are distinct and that inserting $A[i]$ in each of the $i + 1$ possible positions among its predecessors is equally likely. Analyze the sentinel version of the algorithm first.
9. The order of growth of the worst-case number of key comparisons made by binary insertion sort can be obtained from formulas in Section 4.3 and Appendix A. For this algorithm, however, a key comparison is not the operation that determines the algorithm's efficiency class. Which operation does?
10. a. Note that it is more convenient to sort sublists in parallel, i.e., compare $A[0]$ with $A[h_i]$, then $A[1]$ with $A[1 + h_i]$, and so on.

b. Recall that, generally speaking, sorting algorithms that can exchange elements far apart are not stable.

Solutions to Exercises 5.1

1. First, the two boys take the boat to the other side, after which one of them returns with the boat. Then a soldier takes the boat to the other side and stays there while the other boy returns the boat. These four trips reduce the problem's instance of size n (measured by the number of soldiers to be ferried) to the instance of size $n - 1$. Thus, if this four-trip procedure repeated n times, the problem will be solved after the total of $4n$ trips.
2. Assuming that the glasses are numbered left to right from 1 to $2n$, pour soda from glass 2 into glass $2n - 1$. This makes the first and last pair of glasses alternate in the required pattern and hence reduces the problem to the same problem with $2(n - 2)$ middle glasses. If n is even, the number of times this operation needs to be repeated is equal to $n/2$; if n is odd, it is equal to $(n - 1)/2$. The formula $\lfloor n/2 \rfloor$ provides a closed-form answer for both cases. Note that this can also be obtained by solving the recurrence $M(n) = M(n - 2) + 1$ for $n > 2$, $M(2) = 1$, $M(1) = 0$, where $M(n)$ is the number of moves made by the decrease-by-two algorithm described above. Since any algorithm for this problem must move at least one filled glass for each of the $\lfloor n/2 \rfloor$ nonoverlapping pairs of the filled glasses, $\lfloor n/2 \rfloor$ is the least number of moves needed to solve the problem.
3. Here is a general outline of a recursive algorithm that create list $L(n)$ of all the subsets of $\{a_1, \dots, a_n\}$ (see a more detailed discussion in Section 5.4):

if $n = 0$ **return** list $L(0)$ containing the empty set as its only element
else create recursively list $L(n - 1)$ of all the subsets of $\{a_1, \dots, a_{n-1}\}$
 append a_n to each element of $L(n - 1)$ to get list T
 return $L(n)$ obtained by concatenation of $L(n - 1)$ and T
4. Sorting the list E, X, A, M, P, L, E in alphabetical order with insertion sort:

E	X	A	M	P	L	E
E		X				
E	X		A			
A	E	X		M		
A	E	M	X		P	
A	E	M	P	X		L
A	E	L	M	P	X	E
A	E	E	L	M	P	X

5. a. $-\infty$ or, more generally, any value less than or equal to every element in the array.

b. Yes, the efficiency class will stay the same. The number of key comparisons for strictly decreasing arrays (the worst-case input) will be

$$C_{worst}(n) = \sum_{i=1}^{n-1} \sum_{j=-1}^{i-1} 1 = \sum_{i=1}^{n-1} (i+1) = \sum_{i=1}^{n-1} i + \sum_{i=1}^{n-1} 1 = \frac{(n-1)n}{2} + (n-1) \in \Theta(n^2).$$

6. Yes, but we will have to scan the sorted part left to right while inserting $A[i]$ to get the same $O(n^2)$ efficiency as the array version.

7. The efficiency classes of both versions will be the same. The inner loop of *InsertionSort* consists of one key assignment and one index decrement; the inner loop of *InsertionSort2* consists of one key swap (i.e., three key assignments) and one index decrement. If we disregard the time spent on the index decrements, the ratio of the running times should be estimated as $3c_a/c_a = 3$; if we take into account the time spent on the index decrements, the ratio's estimate becomes $(3c_a + c_d)/(c_a + c_d)$, where c_a and c_d are the times of one key assignment and one index decrement, respectively.

8. a. The largest number of inversions for $A[i]$ ($0 \leq i \leq n-1$) is $n-1-i$; this happens if $A[i]$ is greater than all the elements to the right of it. Therefore, the largest number of inversions for an entire array happens for a strictly decreasing array. This largest number is given by the sum:

$$\sum_{i=0}^{n-1} (n-1-i) = (n-1) + (n-2) + \dots + 1 + 0 = \frac{(n-1)n}{2}.$$

The smallest number of inversions for $A[i]$ ($0 \leq i \leq n-1$) is 0; this happens if $A[i]$ is smaller than or equal to all the elements to the right of it. Therefore, the smallest number of inversions for an entire array will be 0 for nondecreasing arrays.

b. Assuming that all elements are distinct and that inserting $A[i]$ in each of the $i+1$ possible positions among its predecessors is equally likely, we obtain the following for the expected number of key comparisons on the i th iteration of the algorithm's sentinel version:

$$\frac{1}{i+1} \sum_{j=1}^{i+1} j = \frac{1}{i+1} \frac{(i+1)(i+2)}{2} = \frac{i+2}{2}.$$

Hence for the average number of key comparisons, $C_{avg}(n)$, we have

$$C_{avg}(n) = \sum_{i=1}^{n-1} \frac{i+2}{2} = \frac{1}{2} \sum_{i=1}^{n-1} i + \sum_{i=1}^{n-1} 1 = \frac{1}{2} \frac{(n-1)n}{2} + n-1 \approx \frac{n^2}{4}.$$

For the no-sentinel version, the number of key comparisons to insert $A[i]$ before and after $A[0]$ will be the same. Therefore the expected number of key comparisons on the i th iteration of the no-sentinel version is:

$$\frac{1}{i+1} \sum_{j=1}^i j + \frac{i}{i+1} = \frac{1}{i+1} \frac{i(i+1)}{2} + \frac{i}{i+1} = \frac{i}{2} + \frac{i}{i+1}.$$

Hence, for the average number of key comparisons, $C_{avg}(n)$, we have

$$C_{avg}(n) = \sum_{i=1}^{n-1} \left(\frac{i}{2} + \frac{i}{i+1} \right) = \frac{1}{2} \sum_{i=1}^{n-1} i + \sum_{i=1}^{n-1} \frac{i}{i+1}.$$

We have a closed-form formula for the first sum:

$$\frac{1}{2} \sum_{i=1}^{n-1} i = \frac{1}{2} \frac{(n-1)n}{2} = \frac{n^2 - n}{4}.$$

The second sum can be estimated as follows:

$$\sum_{i=1}^{n-1} \frac{i}{i+1} = \sum_{i=1}^{n-1} \left(1 - \frac{1}{i+1} \right) = \sum_{i=1}^{n-1} 1 - \sum_{i=1}^{n-1} \frac{1}{i+1} = n-1 - \sum_{j=2}^n \frac{1}{j} = n - H_n,$$

where $H_n = \sum_{j=1}^n 1/j \approx \ln n$ according to a well-known formula quoted in Appendix A. Hence, for the no-sentinel version of insertion sort too, we have

$$C_{avg}(n) \approx \frac{n^2 - n}{4} + n - H_n \approx \frac{n^2}{4}.$$

9. The largest number of key comparisons will be, in particular, for strictly increasing or decreasing arrays:

$$C_{\max}(n) = \sum_{i=1}^{n-1} ([\log_2 i] + 1) = \sum_{i=1}^{n-1} [\log_2 i] + \sum_{i=1}^{n-1} 1 \in \Theta(n \log n) + \Theta(n) = \Theta(n \log n).$$

It is the number of key *moves*, however, that will dominate the number of key comparisons in the worst case of strictly decreasing arrays. The number of key moves will be exactly the same as for the classic insertion sort, putting the algorithm's worst-case efficiency in $\Theta(n^2)$.

10. a. Applying shellsort to the list $S_1, H, E_1, L_1, L_2, S_2, O, R, T, I, S_3, U_1, S_4, E_2, F, U_2, L_3$ with the step-sizes 13, 4, and 1 yields the following:


0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
S_1	H	E_1	L_1	L_2	S_2	O	R	T	I	S_3	U_1	S_4	E_2	F	U_2
S_1													E_2		
E_2													S_1		
	H													F	
	F													H	
		E_1													U_2
			L_1												
E_2				L_2											
	F				S_2										
		E_1				O									
			L_1				R								
				L_2				T							
					S_2				I						
					I				S_2						
						O				S_3					
							R				U_1				
								T				S_4			
								S_4					T		
									S_2					S_1	
										S_3					H
										H				S_3	
											U_1				U_2
												T			
												L_3			
E_2	F	E_1	L_1	L_2	I	O	R	S_4	S_2	H	U_1	L_3	S_1	S_3	U_2

The final pass—sorting the last array by insertion sort—is omitted from the solution because of its simplicity. Note that since relatively few elements in the last array are out of order as a result of the work done on the preceding passes of shellsort, insertion sort will need significantly fewer comparisons to finish the job than it would have needed if it were applied to the initial array.

b. Shellsort is not stable. As a counterexample for shellsort with the sequence of step-sizes 4 and 1, consider, say, the array 5, 1, 2, 3, 1. The first pass with the step-size of 4 will exchange 5 with the last 1, changing the relative ordering of the two 1's in the array. The second pass with the step-size of 1, which is insertion sort, will not make any exchanges because the array is already sorted.


Exercises 5.2

1. Consider the graph




- a. Write down the adjacency matrix and adjacency lists specifying this graph. (Assume that the matrix rows and columns and vertices in the adjacency lists follow in the alphabetical order of the vertex labels.)
- b. Starting at vertex a and resolving ties by the vertex alphabetical order, traverse the graph by depth-first search and construct the corresponding depth-first search tree. Give the order in which the vertices were reached for the first time (pushed onto the traversal stack) and the order in which the vertices became dead ends (popped off the stack).
2. If we define sparse graphs as graphs for which $|E| \in O(|V|)$, which implementation of DFS will have a better time efficiency for such graphs, the one that uses the adjacency matrix or the one that uses the adjacency lists?
3. Let G be a graph with n vertices and m edges.
- a. True or false: All its DFS forests (for traversals starting at different vertices) will have the same number of trees?
- b. True or false: All its DFS forests will have the same number of tree edges and the same number of back edges?
4. Traverse the graph of Problem 1 by breadth-first search and construct the corresponding breadth-first search tree. Start the traversal at vertex a and resolve ties by the vertex alphabetical order.
5. Prove that a cross edge in a BFS tree of an undirected graph can connect vertices only on either the same level or on two adjacent levels of a BFS tree.
6. a. Explain how one can check a graph's acyclicity by using breadth-first search.
- b. Does either of the two traversals—DFS or BFS—always find a cycle faster than the other? If you answer yes, indicate which of them is better and explain why it is the case; if you answer no, give two examples supporting your answer.

7. Explain how one can identify connected components of a graph by using
- a depth-first search.
 - a breadth-first search.
8. A graph is said to be **bipartite** if all its vertices can be partitioned into two disjoint subsets X and Y so that every edge connects a vertex in X with a vertex in Y . (We can also say that a graph is bipartite if its vertices can be colored in two colors so that every edge has its vertices colored in different colors; such graphs are also called **2-colorable**). For example, graph (i) is bipartite while graph (ii) is not.




(i)



(ii)

- Design a DFS-based algorithm for checking whether a graph is bipartite.
 - Design a BFS-based algorithm for checking whether a graph is bipartite.
9. Write a program that, for a given graph, outputs
- vertices of each connected component;
 - its cycle or a message that the graph is acyclic.
10. One can model a maze by having a vertex for a starting point, a finishing point, dead ends, and all the points in the maze where more than one path can be taken, and then connecting the vertices according to the paths in the maze.
- Construct such a graph for the following maze.



- Which traversal—DFS or BFS—would you use if you found yourself in a maze and why?

Hints to Exercises 5.2

1. a. Use the definitions of the adjacency matrix and adjacency lists given in Section 1.4.

b. Perform the DFS traversal the same way it is done for another graph in the text (see Fig. 5.5).
2. Compare the efficiency classes of the two versions of DFS for sparse graphs.
3. a. What is the number of such trees equal to?

b. Answer this question for connected graphs first.
4. Perform the BFS traversal the same way it is done in the text (see Fig. 5.6).
5. You may use the fact that the level of a vertex in a BFS tree indicates the number of edges in the shortest (minimum-edge) path from the root to that vertex.
6. a. What property of a BFS forest indicates a cycle's presence? (The answer is similar to the one for a DFS forest.)

b. The answer is no. Find two examples supporting this answer.
7. Given the fact that both traversals can reach a new vertex if and only if it is adjacent to one of the previously visited vertices, which vertices will be visited by the time either traversal halts (i.e., its stack or queue becomes empty)?
8. Use a DFS forest and a BFS forest for parts (a) and (b), respectively.
9. Use either DFS or BFS.
10. a. Follow the instructions of the problem's statement.

b. Trying both traversals should lead you to a correct answer very fast.


Solutions to Exercises 5.2

1. a. Here are the adjacency matrix and adjacency lists for the graph in question:

	<i>a</i>	<i>b</i>	<i>c</i>	<i>d</i>	<i>e</i>	<i>f</i>	<i>g</i>
<i>a</i>	0	1	1	1	1	0	0
<i>b</i>	1	0	0	1	0	1	0
<i>c</i>	1	0	0	0	0	0	1
<i>d</i>	1	1	0	0	0	1	0
<i>e</i>	1	0	0	0	0	0	1
<i>f</i>	0	1	0	1	0	0	0
<i>g</i>	0	0	1	0	1	0	0

<i>a</i>	→ <i>b</i> → <i>c</i> → <i>d</i> → <i>e</i>
<i>b</i>	→ <i>a</i> → <i>d</i> → <i>f</i>
<i>c</i>	→ <i>a</i> → <i>g</i>
<i>d</i>	→ <i>a</i> → <i>b</i> → <i>f</i>
<i>e</i>	→ <i>a</i> → <i>g</i>
<i>f</i>	→ <i>b</i> → <i>d</i>
<i>g</i>	→ <i>c</i> → <i>e</i>

- b. See below: (i) the graph; (ii) the traversal's stack (the first subscript number indicates the order in which the vertex was visited, i.e., pushed onto the stack, the second one indicates the order in which it became a dead-end, i.e., popped off the stack); (iii) the DFS tree (with the tree edges shown with solid lines and the back edges shown with dashed lines).



2. The time efficiency of DFS is $\Theta(|V|^2)$ for the adjacency matrix representation and $\Theta(|V| + |E|)$ for the adjacency lists representation, respectively. If $|E| \in O(|V|)$, the former remains $\Theta(|V|^2)$ while the latter becomes $\Theta(|V|)$. Hence, for sparse graphs, the adjacency lists version of DFS is more efficient than the adjacency matrix version.

3. a. The number of DFS trees is equal to the number of connected components of the graph. Hence, it will be the same for all DFS traversals of the graph.

- b. For a connected (undirected) graph with $|V|$ vertices, the number of tree edges $|E^{(tree)}|$ in a DFS tree will be $|V| - 1$ and, hence, the number of

back edges $|E^{(back)}|$ will be the total number of edges minus the number of tree edges: $|E| - (|V| - 1) = |E| - |V| + 1$. Therefore, it will be independent from a particular DFS traversal of the same graph. This observation can be extended to an arbitrary graph with $|C|$ connected components by applying this reasoning to each of its connected components:


$$|E^{(tree)}| = \sum_{c=1}^{|C|} |E_c^{(tree)}| = \sum_{c=1}^{|C|} (|V_c| - 1) = \sum_{c=1}^{|C|} |V_c| - \sum_{c=1}^{|C|} 1 = |V| - |C|$$

and

$$|E^{(back)}| = |E| - |E^{(tree)}| = |E| - (|V| - |C|) = |E| - |V| + |C|,$$

where $|E_c^{(tree)}|$ and $|V_c|$ are the numbers of tree edges and vertices in the c th connected component, respectively.

4. Here is the result of the BFS traversal of the graph of Problem 1:




(i) the graph; (ii) the traversal's queue; (iii) the tree (the tree and cross edges are shown with solid and dotted lines, respectively).

5. We'll prove the assertion in question by contradiction. Assume that a BFS tree of some undirected graph has a cross edge connecting two vertices u and v such that $level[u] \geq level[v] + 2$. But $level[u] = d[u]$ and $level[v] = d[v]$, where $d[u]$ and $d[v]$ are the lengths of the minimum-edge paths from the root to vertices u and v , respectively. Hence, we have $d[u] \geq d[v] + 2$. The last inequality contradicts the fact that $d[u]$ is the length of the minimum-edge path from the root to vertex u because the minimum-edge path of length $d[v]$ from the root to vertex v followed by edge (v, u) has fewer edges than $d[u]$.

6. a. A graph has a cycle if and only if its BFS forest has a cross edge.

b. Both traversals, DFS and BFS, can be used for checking a graph's

acyclicity. For some graphs, a DFS traversal discovers a back edge in its DFS forest sooner than a BFS traversal discovers a cross edge (see example (i) below); for others the exactly opposite is the case (see example (ii) below).




7. Start a DFS (or BFS) traversal at an arbitrary vertex and mark the visited vertices with 1. By the time the traversal's stack (queue) becomes empty, all the vertices in the same connected component as the starting vertex, and only they, will have been marked with 1. If there are unvisited vertices left, restart the traversal at one of them and mark all the vertices being visited with 2, and so on until no unvisited vertices are left.

8. a. Let F be a DFS forest of a graph. It is not difficult to see that F is 2-colorable if and only if there is no back edge connecting two vertices both on odd levels or both on even levels. It is this property that a DFS traversal needs to verify. Note that a DFS traversal can mark vertices as even or odd when it reaches them for the first time.

- b. Similarly to part (a), a graph is 2-colorable if and only if its BFS forest has no cross edge connecting vertices on the same level. Use a BFS traversal to check whether or not such a cross edge exists.

9. n/a


10. a. Here is the maze and a graph representing it:




b. DFS is much more convenient for going through a maze than BFS. When DFS moves to a next vertex, it is connected to a current vertex by an edge (i.e., “close nearby” in the physical maze), which is not generally the case for BFS. In fact, DFS can be considered a generalization of an ancient right-hand rule for maze traversal: go through the maze in such a way so that your right hand is always touching a wall.

Exercises 5.3

1. Apply the DFS-based algorithm to solve the topological sorting problem for the following digraphs:



(a)



(b)

2. a. Prove that the topological sorting problem has a solution for a digraph if and only if it is a dag.
b. For a digraph with n vertices, what is the largest number of distinct solutions the topological sorting problem can have?
3. a. What is the time efficiency of the DFS-based algorithm for topological sorting?
b. How can one modify the DFS-based algorithm to avoid reversing the vertex ordering generated by DFS?
4. Can one use the order in which vertices are pushed onto the DFS stack (instead of the order they are popped off it) to solve the topological sorting problem?
5. Apply the source-removal algorithm to the digraphs of Problem 1.
6. a. Prove that a dag must have at least one source.
b. How would you find a source (or determine that such a vertex does not exist) in a digraph represented by its adjacency matrix? What is the time efficiency of this operation?
c. How would you find a source (or determine that such a vertex does not exist) in a digraph represented by its adjacency lists? What is the time efficiency of this operation?
7. \triangleright Can you implement the source-removal algorithm for a digraph represented by its adjacency lists so that its running time is in $O(|V| + |E|)$?
8. Implement the two topological sorting algorithms in the language of your choice. Run an experiment to compare their running times.
9. A digraph is called **strongly connected** if for any pair of two distinct vertices u and v , there exists a directed path from u to v and a directed path

from v to u . In general, a digraph's vertices can be partitioned into disjoint maximal subsets of vertices that are mutually accessible via directed paths of the digraph; these subsets are called ***strongly connected components***. There are two DFS-based algorithms for identifying strongly connected components. Here is the simpler (but somewhat less efficient) one of the two:


Step 1. Do a DFS traversal of the digraph given and number its vertices in the order that they become dead ends.

Step 2. Reverse the directions of all the edges of the digraph.

Step 3. Do a DFS traversal of the new digraph by starting (and, if necessary, restarting) the traversal at the highest numbered vertex among still unvisited vertices.

The strongly connected components are exactly the subsets of vertices in each DFS tree obtained during the last traversal.

a. Apply this algorithm to the following digraph to determine its strongly connected components.



b. What is the time efficiency class of this algorithm? Give separate answers for the adjacency matrix representation and adjacency list representation of an input graph.

c. How many strongly connected components does a dag have?

10. *Celebrity problem* A celebrity among a group of n people is a person who knows nobody but is known by everybody else. The task is to identify a celebrity by only asking questions to people of the form: "Do you know him/her?" Design an efficient algorithm to identify a celebrity or determine that the group has no such person. How many questions does your algorithm need in the worst case?

Hints to Exercises 5.3

1. Trace the algorithm as it is done in the text for another digraph (see Fig. 5.10).
2. a. You need to prove two assertions: (i) if a digraph has a directed cycle, then the topological sorting problem does not have a solution; (ii) if a digraph has no directed cycles, the problem has a solution.

b. Consider an extreme type of a digraph.
3. a. How does it relate to the time efficiency of DFS?

b. Do you know the length of the list to be generated by the algorithm? Where should you put, say, the first vertex being popped off a DFS traversal stack for the vertex to be in its final position?
4. Try to do this for a small example or two.
5. Trace the algorithm on the instances given as it is done in the section (see Fig. 5.11).
6. a. Use a proof by contradiction.


b. If you have difficulty answering the question, consider an example of a digraph with a vertex with no incoming edges and write down its adjacency matrix.

c. The answer follows from the definitions of the source and adjacency lists.
7. For each vertex, store the number of edges entering the vertex in the remaining subgraph. Maintain a queue of the source vertices.
8. n/a
9. a. Trace the algorithm on the input given by following the steps of the algorithm as indicated.

b. Determine the efficiency for each of the three principal steps of the algorithm and then determine the overall efficiency. Of course, the answers will depend on whether a graph is represented by its adjacency matrix or by its adjacency lists.
10. Solve first a simpler version in which a celebrity must be present.

Solutions to Exercises 5.3

1. a. The digraph and the stack of its DFS traversal that starts at vertex a are given below:



f
 $e \quad g$
 $b \quad c$
 $a \quad d$


The vertices are popped off the stack in the following order:

$e \ f \ g \ b \ c \ a \ d.$

The topological sorting order obtained by reversing the list above is

$d \ a \ c \ b \ g \ f \ e.$

- b. The digraph below is not a dag. Its DFS traversal that starts at a encounters a back edge from e to a :



e
 g
 d
 c
 b
 a

2. a. Let us prove by contradiction that if a digraph has a directed cycle, then the topological sorting problem does not have a solution. Assume that v_{i_1}, \dots, v_{i_n} is a solution to the topological sorting problem for a digraph with a directed cycle. Let v_{i_k} be the leftmost vertex of this cycle on the list v_{i_1}, \dots, v_{i_n} . Since the cycle's edge entering v_{i_k} goes right to left, we have a contradiction that proves the assertion.

If a digraph has no directed cycles, a solution to the topological sorting problem is fetched by either of the two algorithms discussed in the section. (The correctness of the DFS-based algorithm was explained there; the correctness of the source removal algorithm stems from the assertion of Problem 6a.)

- b. For a digraph with n vertices and no edges, any permutation of its vertices solves the topological sorting problem. Hence, the answer to the question is $n!$.

3. a. Since reversing the order in which vertices have been popped off the DFS traversal stack is in $\Theta(|V|)$, the running time of the algorithm will be the same as that of DFS (except for the fact that it can stop before processing the entire digraph if a back edge is encountered). Hence, the running time of the DFS-based algorithm is in $O(|V|^2)$ for the adjacency matrix representation and in $O(|V| + |E|)$ for the adjacency lists representation.

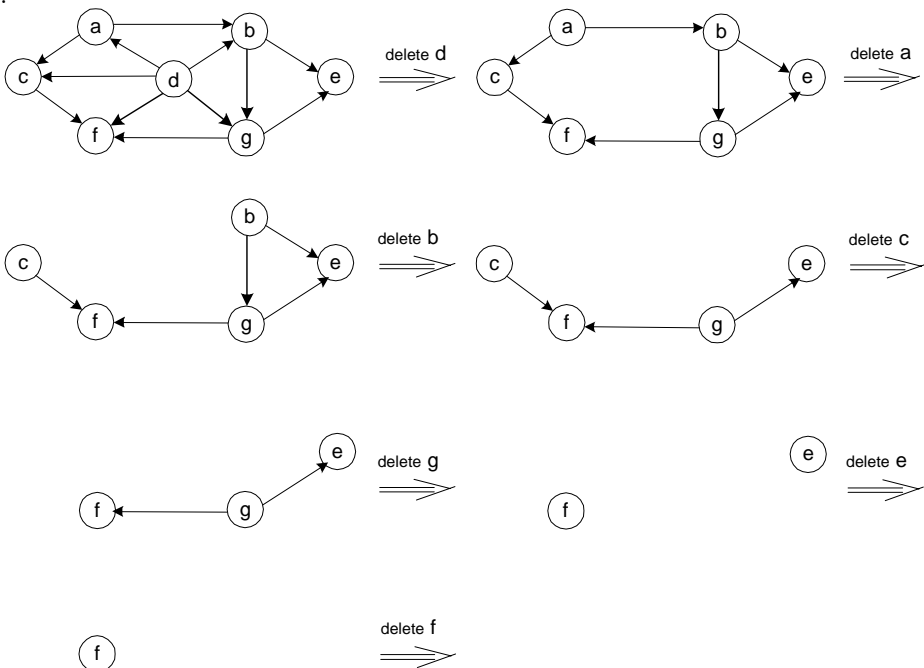
- b. Fill the array of length $|V|$ with vertices being popped off the DFS traversal stack right to left.

4. The answer is no. Here is a simple counterexample:




The DFS traversal that starts at a pushes the vertices on the stack in the order a, b, c , and neither this ordering nor its reversal solves the topological sorting problem correctly.

5. a.



The topological ordering obtained is $d \ a \ b \ c \ g \ e \ f$.



The topological sorting is impossible.

6. a. Assume that, on the contrary, there exists a dag with every vertex having an incoming edge. Reversing all its edges would yield a dag with every vertex having an outgoing edge. Then, starting at an arbitrary vertex and following a chain of such outgoing edges, we would get a directed cycle no later than after $|V|$ steps. This contradiction proves the assertion.


b. A vertex of a dag is a source if and only if its column in the adjacency matrix contains only 0's. Looking for such a column is a $O(|V|^2)$ operation.

c. A vertex of a dag is a source if and only if this vertex appears in none of the dag's adjacency lists. Looking for such a vertex is a $O(|V| + |E|)$ operation.

7. The answer to this well-known problem is yes (see, e.g., [Knu1], pp. 264-265).

8. n/a

9. a. The digraph given is




The stack of the first DFS traversal, with a as its starting vertex, will look as follows:


f_1
 g_2 h_6
 b_3 d_5 e_7
 a_4 c_8

(The numbers indicate the order in which the vertices are popped off the stack.)

The digraph with the reversed edges is



The stack and the DFS trees (with only tree edges shown) of the DFS traversal of the second digraph will be as follows:



The strongly connected components of the given digraph are:

$$\{c, h, e\}, \quad \{d\}, \quad \{a, f, g, b\}.$$

b. If a graph is represented by its adjacency matrix, then the efficiency of the first DFS traversal will be in $\Theta(|V|^2)$. The efficiency of the edge-reversal step (set $B[j, i]$ to 1 in the adjacency matrix of the new digraph if $A[i, j] = 1$ in the adjacency matrix of the given digraph and to 0 otherwise) will also be in $\Theta(|V|^2)$. The time efficiency of the last DFS traversal of the new graph will be in $\Theta(|V|^2)$, too. Hence, the efficiency of the entire algorithm will be in $\Theta(|V|^2) + \Theta(|V|^2) + \Theta(|V|^2) = \Theta(|V|^2)$.

The answer for a graph represented by its adjacency lists will be, by similar reasoning (with a necessary adjustment for the middle step), in $\Theta(|V| + |E|)$.

10. The problem can be solved by a recursive algorithm based on the decrease-by-one strategy. Indeed, by asking just one question, we can eliminate the number of people who can be a celebrity by 1, solve the problem for the remaining group of $n - 1$ people recursively, and then verify the returned solution by asking no more than two questions. Here is a more detailed description of this algorithm:

If $n = 1$, return that one person as a celebrity. If $n > 1$, proceed as follows:

Step 1 Select two people from the group given, say, A and B, and ask A whether A knows B. If A knows B, remove A from the remaining people who can be a celebrity; if A doesn't know B, remove B from this group.

Step 2 Solve the problem recursively for the remaining group of $n - 1$ people who can be a celebrity.

Step 3 If the solution returned in Step 2 indicates that there is no celebrity among the group of $n - 1$ people, the larger group of n people cannot contain a celebrity either. If Step 2 identified as a celebrity a person other than either A or B, say, C, ask whether C knows the person removed in Step 1 and, if the answer is no, whether the person removed in Step 1 knows C. If the answer to the second question is yes, return C as a celebrity and "no celebrity" otherwise. If Step 2 identified B as a celebrity, just ask whether B knows A: return B as a celebrity if the answer is no and "no celebrity" otherwise. If Step 2 identified A as a celebrity, ask whether B knows A: return A as a celebrity if the answer is yes and "no celebrity" otherwise.

The recurrence for $Q(n)$, the number of questions needed in the worst case, is as follows:

$$Q(n) = Q(n - 1) + 3 \quad \text{for } n > 2, \quad Q(2) = 2, \quad Q(1) = 0.$$

Its solution is $Q(n) = 2 + 3(n - 2)$ for $n > 1$ and $Q(1) = 0$.

Note: A computer implementation of this algorithm can be found, e.g., in Manber's *Introduction to Algorithms: A Creative Approach*. Addison-Wesley, 1989.

Exercises 5.4

1. Is it realistic to implement an algorithm that requires generating all permutations of a 25-element set on your computer? What about all the subsets of such a set?
2. Generate all permutations of $\{1, 2, 3, 4\}$ by
 - a. the bottom-up minimal-change algorithm.
 - b. the Johnson-Trotter algorithm.
 - c. the lexicographic-order algorithm.
3. Write a program for generating permutations in lexicographic order.
4. ► Consider a simple implementation of the following algorithm for generating permutations discovered by B. Heap [Hea63].

Algorithm *HeapPermute*(n)

//Implements Heap's algorithm for generating permutations

//Input: A positive integer n and a global array $A[1..n]$

//Output: All permutations of elements of A

if $n = 1$

write A

else

for $i \leftarrow 1$ **to** n **do**

HeapPermute($n - 1$)

if n is odd

 swap $A[1]$ and $A[n]$

else swap $A[i]$ and $A[n]$

- a. Trace the algorithm by hand for $n = 2, 3$, and 4.
 - b. Prove correctness of Heap's algorithm.
 - c. What is the time efficiency of this algorithm?
5. Generate all the subsets of a four-element set $A = \{a_1, a_2, a_3, a_4\}$ by each of the two algorithms outlined in this section.
 6. What simple trick would make the bit string-based algorithm generate subsets in squashed order?
 7. Write a pseudocode for a recursive algorithm for generating all 2^n bit strings of length n .
 8. Write a nonrecursive algorithm for generating 2^n bit strings of length n that implements bit strings as arrays and does not use binary additions.

9. a. Use the decrease-by-one technique to generate the binary reflected Gray code for $n = 4$.
- b.▷ Design a general decrease-by-one algorithm for generating the binary reflected Gray code of order n .
10. ► Design a decrease-and-conquer algorithm for generating all combinations of k items chosen from n , i.e., all k -element subsets of a given n -element set. Is your algorithm a minimal-change algorithm?
11. *Gray code and the Tower of Hanoi*
- (a) ▷ Show that the disk moves made in the classic recursive algorithm for the Tower-of-Hanoi puzzle can be used for generating the binary reflected Gray code.
- (b) ► Show how the binary reflected Gray code can be used for solving the Tower-of-Hanoi puzzle.

Hints to Exercises 5.4

1. Use standard formulas for the numbers of these combinatorial objects. For the sake of simplicity, you may assume that generating one combinatorial object takes the same time as, say, one assignment.
2. We traced the algorithms on smaller instances in the section.
3. See an outline of this algorithm in the section.
4. a. Trace the algorithm for $n = 2$; take advantage of this trace in tracing the algorithm for $n = 3$ and then use the latter for $n = 4$.

b. Show that the algorithm generates $n!$ permutations and that all of them are distinct. Use mathematical induction.

c. Set up a recurrence relation for the number of swaps made by the algorithm. Find its solution and the solution's order of growth. You may need the formula: $e \approx \sum_{i=0}^n \frac{1}{i!}$.
5. We traced both algorithms on smaller instances in the section.
6. Tricks become boring after they have been given away.
7. This is not a difficult exercise because of the obvious way of getting bit strings of length n from bit strings of length $n - 1$.
8. You may still mimic the binary addition without using it explicitly.
9. A Gray code for $n = 3$ is given at the end of the section. It is not difficult to see how to use it to generate a Gray code for $n = 4$. Gray codes have a useful geometric interpretation based on mapping its bit strings to vertices of the n -dimensional cube. Find such a mapping for $n = 1, 2$, and 3 . This geometric interpretation might help you with designing a general algorithm for generating a Gray code of order n .
10. There are several decrease-and-conquer algorithms for this problem. They are more subtle than one might expect. Generating combinations in a pre-defined order (increasing, decreasing, lexicographic) helps with both a design and a correctness proof. The following simple property is very helpful. Assuming with no loss of generality that the underlying set is $\{1, 2, \dots, n\}$, there are $\binom{n-i}{k-1}$ k -subsets whose smallest element is i , $i = 1, 2, \dots, n - k + 1$.
11. Represent the disk movements by flipping bits in a binary n -tuple.

Solutions to Exercises 5.4

1. Since $25! \approx 1.5 \cdot 10^{25}$, it would take an unrealistically long time to generate this number of permutations even on a supercomputer. On the other hand, $2^{25} \approx 3.3 \cdot 10^7$, which would take about 0.3 seconds to generate on a computer making one hundred million operations per second.
2. a. The permutations of $\{1, 2, 3, 4\}$ generated by the bottom-up minimal-change algorithm:

start	1
insert 2 into 1 right to left	12 21
insert 3 into 12 right to left	123 132 312
insert 3 into 21 left to right	321 231 213
insert 4 into 123 right to left	1234 1243 1423 4123
insert 4 into 132 left to right	4132 1432 1342 1324
insert 4 into 312 right to left	3124 3142 3412 4312
insert 4 into 321 left to right	4321 3421 3241 3214
insert 4 into 231 right to left	2314 2341 2431 4231
insert 4 into 213 left to right	4213 2413 2143 2134

- b. The permutations of $\{1, 2, 3, 4\}$ generated by the Johnson-Trotter algorithm. (Read horizontally; the largest mobile element is shown in bold.)

$\overleftarrow{\overleftarrow{\overleftarrow{\overleftarrow{1}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{2}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{3}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{4}}}}$	$\overleftarrow{\overleftarrow{\overleftarrow{\overleftarrow{1}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{2}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{4}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{3}}}}$	$\overleftarrow{\overleftarrow{\overleftarrow{\overleftarrow{1}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{4}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{2}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{3}}}}$	$\overleftarrow{\overleftarrow{\overleftarrow{\overleftarrow{4}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{1}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{2}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{3}}}}$
$\overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{4}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{1}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{3}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{2}}}}$	$\overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{1}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{4}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{3}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{2}}}}$	$\overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{1}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{3}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{4}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{2}}}}$	$\overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{1}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{3}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{2}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{4}}}}$
$\overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{3}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{1}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{2}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{4}}}}$	$\overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{3}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{1}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{4}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{2}}}}$	$\overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{3}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{4}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{1}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{2}}}}$	$\overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{4}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{3}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{1}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{2}}}}$
$\overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{4}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{3}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{2}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{1}}}}$	$\overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{3}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{4}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{2}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{1}}}}$	$\overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{3}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{2}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{4}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{1}}}}$	$\overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{3}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{2}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{1}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{4}}}}$
$\overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{2}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{3}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{1}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{4}}}}$	$\overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{2}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{4}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{3}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{1}}}}$	$\overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{2}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{4}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{1}}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{3}}}}$	$\overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{4}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{2}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{3}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{1}}}}$
$\overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{4}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{2}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{1}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{3}}}}$	$\overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{2}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{4}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{1}}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{3}}}}$	$\overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{2}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{1}}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{4}}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{3}}}}$	$\overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{2}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{1}}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{3}}}}} \overrightarrow{\overrightarrow{\overrightarrow{\overrightarrow{4}}}}$

- c. The permutations of $\{1, 2, 3, 4\}$ generated in lexicographic order. (Read horizontally.)

1234	1243	1324	1342	1423	1432
2134	2143	2314	2341	2413	2431
3124	3142	3214	3241	3412	3421
4123	4132	4213	4231	4312	4321

3. n/a

4. a. For $n = 2$:

12 21

For $n = 3$ (read along the rows):

123 213

312 132

231 321

For $n = 4$ (read along the rows):

1234 2134 3124 1324 2314 3214

4231 2431 3421 4321 2341 3241

4132 1432 3412 4312 1342 3142

4123 1423 2413 4213 1243 2143

b. Let $C(n)$ be the number of times the algorithm writes a new permutation (on completion of the recursive call when $n = 1$). We have the following recurrence for $C(n)$:

$$C(n) = \sum_{i=1}^n C(n-1) \quad \text{or} \quad C(n) = nC(n-1) \quad \text{for } n > 1, \quad C(1) = 1.$$

Its solution (see Section 2.4) is $C(n) = n!$. The fact that all the permutations generated by the algorithm are distinct, can be proved by mathematical induction.

c. We have the following recurrence for the number of swaps $S(n)$:

$$S(n) = \sum_{i=1}^n (S(n-1) + 1) \quad \text{or} \quad S(n) = nS(n-1) + n \quad \text{for } n > 1, \quad S(1) = 0.$$

Although it can be solved by backward substitution, this is easier to do after dividing both hand sides by $n!$

$$\frac{S(n)}{n!} = \frac{S(n-1)}{(n-1)!} + \frac{1}{(n-1)!} \quad \text{for } n > 1, \quad S(1) = 0$$

and substituting $T(n) = \frac{S(n)}{n!}$ to obtain the following recurrence:

$$T(n) = T(n-1) + \frac{1}{(n-1)!} \quad \text{for } n > 1, \quad T(1) = 0.$$

Solving the last recurrence by backward substitutions yields

$$T(n) = T(1) + \sum_{i=1}^{n-1} \frac{1}{i!} = \sum_{i=1}^{n-1} \frac{1}{i!}.$$

On returning to variable $S(n) = n!T(n)$, we obtain

$$S(n) = n! \sum_{i=1}^{n-1} \frac{1}{i!} \approx n!(e - 1 - \frac{1}{n!}) \in \Theta(n!).$$

5. Generate all the subsets of a four-element set $A = \{a_1, a_2, a_3, a_4\}$ bottom up:

n	subsets						
0	\emptyset						
1	\emptyset	$\{a_1\}$					
2	\emptyset	$\{a_1\}$	$\{a_2\}$	$\{a_1, a_2\}$			
3	\emptyset	$\{a_1\}$	$\{a_2\}$	$\{a_1, a_2\}$	$\{a_3\}$	$\{a_1, a_3\}$	$\{a_2, a_3\}$
4	\emptyset	$\{a_1\}$	$\{a_2\}$	$\{a_1, a_2\}$	$\{a_3\}$	$\{a_1, a_3\}$	$\{a_2, a_3\}$
	$\{a_4\}$	$\{a_1, a_4\}$	$\{a_2, a_4\}$	$\{a_1, a_2, a_4\}$	$\{a_3, a_4\}$	$\{a_1, a_3, a_4\}$	$\{a_2, a_3, a_4\}$

Generate all the subsets of a four-element set $A = \{a_1, a_2, a_3, a_4\}$ with bit vectors:

bit strings	0000	0001	0010	0011	0100	0101	0110
subsets	\emptyset	$\{a_4\}$	$\{a_3\}$	$\{a_3, a_4\}$	$\{a_2\}$	$\{a_2, a_4\}$	$\{a_2, a_3\}$
bit strings	1000	1001	1010	1011	1100	1101	1110
subsets	$\{a_1\}$	$\{a_1, a_4\}$	$\{a_1, a_3\}$	$\{a_1, a_3, a_4\}$	$\{a_1, a_2\}$	$\{a_1, a_2, a_4\}$	$\{a_1, a_2, a_3\}$

6. Establish the correspondence between subsets of $A = \{a_1, \dots, a_n\}$ and bit strings $b_1 \dots b_n$ of length n by associating bit i with the presence or absence of element a_{n-i+1} for $i = 1, \dots, n$.

7. **Algorithm** *BitstringsRec*(n)

//Generates recursively all the bit strings of a given length

//Input: A positive integer n

//Output: All bit strings of length n as contents of global array $B[0..n-1]$

if $n = 0$

 print(B)

else

$B[n-1] \leftarrow 0$; *BitstringsRec*($n-1$)

$B[n-1] \leftarrow 1$; *BitstringsRec*($n-1$)

8. **Algorithm** *BitstringsNonrec*(n)

//Generates nonrecursively all the bit strings of a given length


//Input: A positive integer n

```

//Output: All bit strings of length  $n$  as contents of global array  $B[0..n-1]$ 
for  $i \leftarrow 0$  to  $n-1$  do
     $B[i] = 0$ 
repeat
     $\text{print}(B)$ 
     $k \leftarrow n-1$ 
    while  $k \geq 0$  and  $B[k] = 1$ 
         $k \leftarrow k-1$ 
    if  $k \geq 0$ 
         $B[k] \leftarrow 1$ 
        for  $i \leftarrow k+1$  to  $n-1$  do
             $B[i] \leftarrow 0$ 
until  $k = -1$ 

```

9. a. As mentioned in the hint to this problem, binary Gray codes have a useful geometric interpretation based on mapping their bit strings to vertices of the n -dimensional cube. Such a mapping is shown below for $n = 1, 2$, and 3.



The list of bit strings in the binary reflexive Gray code for $n = 3$ given in the section is obtained by traversing the vertices of the three-dimensional cube by starting at 000 and following the arrows shown:

000 001 011 010 110 111 101 100.

We can obtain the binary reflexive Gray code for $n = 4$ as follows. Make two copies of the list of bit strings for $n = 3$; add 0 in front of each bit string in the first copy and 1 in front of each bit string in the second copy and then append the second list to the first in reversed order to obtain:

0000 0001 0011 0010 0110 0111 0101 0100 1100 1101 1111 1110 1010 1011

(Note that the last bit string differs from the first one by a single bit, too.)

- b. The mechanism employed in part (a) can be used for constructing

the binary reflexive Gray code for an arbitrary $n \geq 1$: If $n = 1$, return the list 0, 1. If $n > 1$, generate recursively the list of bit strings of size $n - 1$ and make a copy of this list; add 0 in front of each bit string in the first list and add 1 in front of each bit string in the second list; then append the second list in reversed order to the first list.

Note: The correctness of the algorithm stems from the fact that it generates 2^n bit strings and all of them are distinct. (Both these assertions are very easy to check by mathematical induction.)

10. Here is a recursive algorithm from “Problems on Algorithms” by Ian Parberry [Par95], p.120:

call *Choose*(1, k) where

Algorithm *Choose*(i, k)

//Generates all k -subsets of $\{i, i + 1, \dots, n\}$ stored in global array $A[1..k]$
 //in descending order of their components

if $k = 0$

print(A)

else

for $j \leftarrow i$ **to** $n - k + 1$ **do**

$A[k] \leftarrow j$

Choose($j + 1, k - 1$)

11. a. Number the disks from 1 to n in increasing order of their size. The disk movements will be represented by a tuple of n bits, in which the bits will be counted right to left so that the rightmost bit will represent the movements of the smallest disk and the leftmost bit will represent the movements of the largest disk. Initialize the tuple with all 0's. For each move in the puzzle's solution, flip the i th bit if the move involves the i th disk.

b. Use the correspondence described in part a between bit strings of the binary reflected Gray code and the disk moves in the Tower of Hanoi puzzle with the following additional rule for situations when there is a choice of where to place a disk: When faced with a choice in placing a disk, always place an odd numbered disk on top of an even numbered disk; if an even numbered disk is not available, place the odd numbered disk on an empty peg. Similarly, place an even numbered disk on an odd disk, if available, or else on an empty peg.

Exercises 5.5

1. Design a decrease-by-half algorithm for computing $\lfloor \log_2 n \rfloor$ and determine its time efficiency.
2. Consider **ternary search**—the following algorithm for searching in a sorted array $A[0..n - 1]$: if $n = 1$, simply compare the search key K with the single element of the array; otherwise, search recursively by comparing K with $A[\lfloor n/3 \rfloor]$, and if K is larger, compare it with $A[\lfloor 2n/3 \rfloor]$ to determine in which third of the array to continue the search.
 - a. What design technique is this algorithm based on?
 - b. Set up a recurrence relation for the number of key comparisons in the worst case. (You may assume that $n = 3^k$.)
 - c. Solve the recurrence for $n = 3^k$.
 - d. Compare this algorithm's efficiency with that of binary search.
3.
 - a. Write a pseudocode for the divide-into-three algorithm for the fake-coin problem. (Make sure that your algorithm handles properly all values of n , not only those that are multiples of 3.)
 - b. Set up a recurrence relation for the number of weighings in the divide-into-three algorithm for the fake-coin problem and solve it for $n = 3^k$.
 - c. For large values of n , about how many times faster is this algorithm than the one based on dividing coins into two piles? (Your answer should not depend on n .)
4. Apply multiplication à la russe to compute $26 \cdot 47$.
5.
 - a. From the standpoint of time efficiency, does it matter whether we multiply n by m or m by n by the multiplication à la russe algorithm?
 - b. What is the efficiency class of multiplication à la russe?
6. Write a pseudocode for the multiplication à-la-russe algorithm.
7. Find $J(40)$ —the solution to the Josephus problem for $n = 40$.
8. Prove that the solution to the Josephus problem is 1 for every n that is a power of 2.
9. ► For the Josephus problem,
 - a. compute $J(n)$ for $n = 1, 2, \dots, 15$.

- b. discern a pattern in the solutions for the first fifteen values of n and prove its general validity.
- c. prove the validity of getting $J(n)$ by a one-bit cyclic shift left of the binary representation of n .

Hints to Exercises 5.5

1. If the instance of size n is to compute $\lfloor \log_2 n \rfloor$, what is the instance of size $n/2$? What is the relationship between the two?
2. The algorithm is quite similar to binary search, of course. In the worst case, how many key comparisons does it make on each iteration and what fraction of the array remains to be processed?
3. While it is obvious how one needs to proceed if $n \bmod 3 = 0$ or $n \bmod 3 = 1$, it is somewhat less so if $n \bmod 3 = 2$.
4. Trace the algorithm for the numbers given as it is done in the text for another input (see Figure 5.14b).
5. How many iterations does the algorithm do?
6. You can implement the algorithm either recursively or nonrecursively.
7. The fastest way to the answer is to use the formula that exploits the binary representation of n , which is mentioned at the end of Section 5.5.
8. Use the binary representation of n .
9.
 - a. Use forward substitutions (see Appendix B) into the recurrence equations given in the text.
 - b. On observing the pattern in the first fifteen values of n obtained in part (a), express it analytically. Then prove its validity by mathematical induction.
 - c. Start with the binary representation of n and translate into binary the formula for $J(n)$ obtained in part (b).

Solutions to Exercises 5.5

1. **Algorithm** *LogFloor*(n)
//Input: A positive integer n
//Output: Returns $\lfloor \log_2 n \rfloor$
if $n = 1$ **return** 0
else return *LogFloor*($\lfloor \frac{n}{2} \rfloor$) + 1

The algorithm is almost identical to the algorithm for computing the number of binary digits, which was investigated in Section 2.4. The recurrence relation for the number of additions is

$$A(n) = A(\lfloor n/2 \rfloor) + 1 \quad \text{for } n > 1, \quad A(1) = 0.$$

Its solution is $A(n) = \lfloor \log_2 n \rfloor \in \Theta(\log n)$.

2. a. The algorithm is based on the decrease-by-a constant factor (equal to 3) strategy.

$$\text{b. } C(n) = 2 + C(n/3) \quad \text{for } n = 3^k \ (k > 0), \quad C(1) = 1.$$

$$\begin{aligned} \text{c. } C(3^k) &= 2 + C(3^{k-1}) \quad [\text{sub. } C(3^{k-1}) = 2 + C(3^{k-2})] \\ &= 2 + [2 + C(3^{k-2})] = 2 \cdot 2 + C(3^{k-2}) = [\text{sub. } C(3^{k-2}) = 2 + C(3^{k-3})] \\ &= 2 \cdot 2 + [2 + C(3^{k-3})] = 2 \cdot 3 + C(3^{k-3}) = \dots = 2i + C(3^{k-i}) = \dots = \\ &= 2k + C(3^{k-k}) = 2\log_3 n + 1. \end{aligned}$$

- d. We have to compare this formula with the worst-case number of key comparisons in the binary search, which is about $\log_2 n + 1$. Since

$$2\log_3 n + 1 = 2 \frac{\log_2 n}{\log_2 3} + 1 = \frac{2}{\log_2 3} \log_2 n + 1$$

and $2/\log_2 3 > 1$, binary search has a smaller multiplicative constant and hence is more efficient (by about the factor of $2/\log_2 3$) in the worst case, although both algorithms belong to the same logarithmic class.

3. a. If n is a multiple of 3 (i.e., $n \bmod 3 = 0$), we can divide the coins into three piles of $n/3$ coins each and weigh two of the piles. If $n = 3k + 1$ (i.e., $n \bmod 3 = 1$), we can divide the coins into the piles of sizes k , k , and $k + 1$ or $k + 1$, $k + 1$, and $k - 1$. (We will use the second option.) Finally, if $n = 3k + 2$ (i.e., $n \bmod 3 = 2$), we will divide the coins into the piles of sizes $k + 1$, $k + 1$, and k . The following pseudocode assumes that there is exactly one fake coin among the coins given and that the fake coin is lighter than the other coins.

if $n = 1$ the coin is fake

else divide the coins into three piles of $\lceil n/3 \rceil$, $\lceil n/3 \rceil$, and $n - 2\lceil n/3 \rceil$ coins
 weigh the first two piles
if they weigh the same
 discard all of them and continue with the coins of the third pile
else continue with the lighter of the first two piles

b. The recurrence relation for the number of weighing $W(n)$ needed in the worst case is as follows:

$$W(n) = W(\lceil n/3 \rceil) + 1 \quad \text{for } n > 1, \quad W(1) = 0.$$

For $n = 3^k$, the recurrence becomes $W(3^k) = W(3^{k-1}) + 1$. Solving it by backward substitutions yields $W(3^k) = k = \log_3 n$.

c. The ratio of the numbers of weighings in the worst case can be approximated for large values of n by

$$\frac{\log_2 n}{\log_3 n} = \frac{\log_2 n}{\log_3 2 \log_2 n} = \log_2 3 \approx 1.6.$$

4. Compute $26 \cdot 47$ by the multiplication à la russe algorithm:

n	m	
26	47	
13	94	94
6	188	
3	376	376
1	752	752
		<hr/> 1,222

5. a. The number of divisions multiplication à la russe needs for computing $n \cdot m$ and $m \cdot n$ is $\lfloor \log_2 n \rfloor$ and $\lfloor \log_2 m \rfloor$, respectively.

b. Its time efficiency is in $\Theta(\log n)$ where n is the first factor of the product. As a function of b , the number of binary digits of n , the time efficiency is in $\Theta(b)$.

6. **Algorithm** *Russe*(n, m)

//Implements multiplication à la russe nonrecursively

//Input: Two positive integers n and m

//Output: The product of n and m

$p \leftarrow 0$

```

while  $n \neq 1$  do
    if  $n \bmod 2 = 1$   $p \leftarrow p + m$ 
     $n \leftarrow \lfloor n/2 \rfloor$ 
     $m \leftarrow 2 * m$ 
return  $p + m$ 

```

Algorithm *RusseRec*(n, m)

//Implements multiplication à la russe recursively

//Input: Two positive integers n and m

//Output: The product of n and m

if $n \bmod 2 = 0$ **return** *RusseRec*($n/2, 2m$)

else if $n = 1$ **return** m

else return *RusseRec*(($n - 1$)/2, $2m$) + m

7. Using the fact that $J(n)$ can be obtained by a one-bit left cyclic shift of n , we get the following for $n = 40$:

$$J(40) = J(101000_2) = 10001_2 = 17.$$

8. We can use the fact that $J(n)$ can be obtained by a one-bit left cyclic shift of n . If $n = 2^k$, where k is a nonnegative integer, then $J(2^k) = J(\underbrace{10\dots0}_k) = 1$.

9. a. Using the initial condition $J(1) = 1$ and the recurrences $J(2k) = 2J(k) - 1$ and $J(2k + 1) = 2J(k) + 1$ for even and odd values of n , respectively, we obtain the following values of $J(n)$ for $n = 1, 2, \dots, 15$:

n	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
$J(n)$	1	1	3	1	3	5	7	1	3	5	7	9	11	13	15

- b. On inspecting the values obtained in part (a), it is not difficult to observe that for the n 's values between consecutive powers of 2, i.e., for $2^k \leq n < 2^{k+1}$ ($k = 0, 1, 2, 3$) or $n = 2^k + i$ where $i = 0, 1, \dots, 2^k - 1$, the corresponding values of $J(n)$ run the range of odd numbers from 1 to $2^{k+1} - 1$. This observation can be expressed by the formula

$$J(2^k + i) = 2i + 1 \quad \text{for } i = 0, 1, \dots, 2^k - 1.$$

We'll prove that this formula solves the recurrences of the Josephus problem for any nonnegative integer k by induction on k . For the basis value

$k = 0$, we have $J(2^0 + 0) = 2 \cdot 0 + 1 = 1$ as it should for the initial condition. Assuming that for a given nonnegative integer k and for every $i = 0, 1, \dots, 2^k - 1$, $J(2^k + i) = 2i + 1$, we need to show that

$$J(2^{k+1} + i) = 2i + 1 \quad \text{for } i = 0, 1, \dots, 2^{k+1} - 1.$$

If i is even, it can be represented as $2j$ where $j = 0, 1, \dots, 2^k - 1$. Then we obtain

$$J(2^{k+1} + i) = J(2(2^k + j)) = 2J(2^k + j) - 1$$

and, using the induction's assumption, we can continue as follows

$$2J(2^k + j) - 1 = 2[2j + 1] - 1 = 2i + 1.$$

If i is odd, it can be expressed as $2j + 1$ where $0 \leq j < 2^k$. Then we obtain

$$J(2^{k+1} + i) = J(2^{k+1} + 2j + 1) = J(2(2^k + j) + 1) = 2J(2^k + j) + 1$$

and, using the induction's assumption, we can continue as follows

$$2J(2^k + j) + 1 = 2[2j + 1] + 1 = 2i + 1.$$

c. Let $n = (b_k b_{k-1} \dots b_0)_2$ where the first binary digit b_k is 1. In the n 's representation used in part (b), $n = 2^k + i$, $i = (b_{k-1} \dots b_0)_2$. Further, as proved in part (b),

$$J(n) = 2i + 1 = (b_{k-1} \dots b_0 0)_2 + 1 = (b_{k-1} \dots b_0 1)_2 = (b_{k-1} \dots b_0 b_k)_2,$$

which is a one-bit left cyclic shift of $n = (b_k b_{k-1} \dots b_0)_2$.

Note: The solutions to Problem 9 are from [Gra94].

Exercises 5.6

1. a. If we measure the size of an instance of the problem of computing the greatest common divisor of m and n by the size of the second parameter n , by how much can the size decrease after one iteration of Euclid's algorithm?

b. Prove that an instance size will always decrease at least by a factor of 2 after two successive iterations of Euclid's algorithm.
2. a. Apply the partition-based algorithm to find the median of the list of numbers 9, 12, 5, 17, 20.

b. Show that the worst-case efficiency of the partition-based algorithm for the selection problem is quadratic.
3. a. Write a pseudocode for a nonrecursive implementation of the partition-based algorithm for the selection problem.

b. Write a pseudocode for a recursive implementation of this algorithm.
4. Derive the formula underlying interpolation search.
5. \triangleright Give an example of the worst-case input for interpolation search and show that the algorithm is linear in the worst case.
6. a. Find the smallest value of n for which $\log_2 \log_2 n + 1$ is greater than 6.

b. Determine which, if any, of the following assertions are true:

i. $\log \log n \in o(\log n)$ ii. $\log \log n \in \Theta(\log n)$ iii. $\log \log n \in \Omega(\log n)$.
7. a. Outline an algorithm for finding the largest key in a binary search tree. Would you classify your algorithm as a variable-size-decrease algorithm?

b. What is the time efficiency class of your algorithm in the worst case?
8. a. Outline an algorithm for deleting a key from a binary search tree. Would you classify this algorithm as a variable-size-decrease algorithm?

b. What is the time efficiency class of your algorithm?
9. *Misere one-pile Nim* Consider the so-called ***misere version*** of the one-pile Nim, in which the player taking the last chip loses the game. All the other conditions of the game remain the same, i.e., the pile contains n chips and on each move a player takes at least one but no more than m chips. Identify the winning and losing positions (for the player to move) in this game.

10. ▷a. *Moldy chocolate* Two players take turns by breaking an m -by- n chocolate bar, which has one spoiled 1-by-1 square. Each break must be a single straight line cutting all the way across the bar along the boundaries between the squares. After each break, the player who broke the bar last eats the piece that does not contain the spoiled corner. The player left with the spoiled square loses the game. Is it better to go first or second in this game?
- b. Write an interactive program to play this game with the computer. Your program should make a winning move in a winning position and a random legitimate move in a losing position.
11. ▷*Flipping pancakes* There are n pancakes all of different sizes that are stacked on top of each other. You are allowed to slip a flipper under one of the pancakes and flip over the whole stack above the flipper. The purpose is to arrange pancakes according to their size with the biggest at the bottom. (You can see a visualization of this puzzle on the *Interactive Mathematics Miscellany and Puzzles* site [Bog].) Design an algorithm for solving this puzzle.

Hints to Exercises 5.6

1. a. The answer follows immediately from the formula underlying Euclid's algorithm.

b. Let $r = m \bmod n$. Investigate two cases of r 's value relative to n 's value.
2. a. Trace the algorithm on the input given, as was done in the section for another input.

b. Since the algorithm in question is based on the same partitioning idea as quicksort is, it is natural to expect the worst-case inputs to be similar for these algorithms.
3. You should have difficulties with neither implementation of the algorithm outlined in the text.
4. Write an equation of the straight line through the points $(l, A[l])$ and $(r, A[r])$ and find the x coordinate of the point on this line whose y coordinate is v .
5. Construct an array for which interpolation search decreases the remaining subarray by one element on each iteration.
6. a. Solve the inequality $\log_2 \log_2 n + 1 > 6$.

b. Compute $\lim_{n \rightarrow \infty} \frac{\log \log n}{\log n}$. Note that to within a constant multiple, you can consider the logarithms to be natural, i.e., base e .
7. a. The definition of the binary search tree suggests such an algorithm.

b. What will be the worst-case input for your algorithm? How many key comparisons will it make on such an input?
8. a. Consider separately three cases: (1) the key's node is a leaf; (2) the key's node has one child; (3) the key's node has two children.

b. Assume that you know a location of the key to be deleted.
9. Follow the plan used in Section 5.6 for analyzing the standard version of the game.
10. Play several rounds of the game on the graphed paper to become comfortable with the problem. Considering special cases of the spoiled square's location should help you to solve it.
11. Do yourself a favor: try to design an algorithm on your own. It does not have to be optimal, but it should be reasonably efficient.

Solutions to Exercises 5.6

1. a. Since the algorithm uses the formula $\text{gcd}(m, n) = \text{gcd}(n, m \bmod n)$, the size of the new pair will be $m \bmod n$. Hence it can be any integer between 0 and $n-1$. Thus, the size n can decrease by any number between 1 and n .
- b. Two consecutive iterations of Euclid's algorithm are performed according to the following formulas:

$$\text{gcd}(m, n) = \text{gcd}(n, r) = \text{gcd}(r, n \bmod r) \quad \text{where } r = m \bmod n.$$

We need to show that $n \bmod r \leq n/2$. Consider two cases: $r \leq n/2$ and $n/2 < r < n$. If $r \leq n/2$, then

$$n \bmod r < r \leq n/2.$$

If $n/2 < r < n$, then

$$n \bmod r = n - r < n/2,$$

too.

2. a. Since $n = 5$, $k = \lceil 5/2 \rceil = 3$. For the given list 9, 12, 5, 17, 20, with the first element as the pivot, we obtain the following partition

$$\begin{array}{ccccc} \mathbf{9} & 12 & 5 & 17 & 20 \\ 5 & \mathbf{9} & 12 & 17 & 20 \end{array}$$

Since $s = 2 < k = 3$, we proceed with the right part of the list:

$$\begin{array}{ccc} \mathbf{12} & 17 & 20 \\ \mathbf{12} & 17 & 20 \end{array}$$

Since $s = k = 3$, 12 is the median of the list given.

- b. Consider an instance of the selection problem with $k = n$ and a strictly increasing array. The situation is identical to the worst-case analysis of quicksort (see Section 4.2).
3. a. **Algorithm** *Selection*($A[0..n-1], k$)
//Solves the selection problem by partition-based algorithm
//Input: An array $A[0..n-1]$ of orderable elements and integer k ($1 \leq k \leq n$)
//Output: The value of the k th smallest element in $A[0..n-1]$
 $l \leftarrow 0$; $r \leftarrow n-1$
 $A[n] \leftarrow \infty$ //append sentinel
while $l \leq r$ **do**
 $p \leftarrow A[l]$ //the pivot

```

i ← l;  j ← r + 1
repeat
    repeat i ← i + 1 until A[i] ≥ p
    repeat j ← j − 1 until A[j] ≤ p do
        swap(A[i], A[j])
until i ≥ j
swap(A[i], A[j]) //undo last swap
swap(A[l], A[j]) //partition
if j > k − 1  r ← j − 1
else if j < k − 1  l ← j + 1
else return A[k − 1]

```

b. call *SelectionRec*(*A*[0..*n* − 1], *k*) where

Algorithm *SelectionRec*(*A*[*l*..*r*], *k*)
//Solves the selection problem by recursive partition-based algorithm
//Input: A subarray *A*[*l*..*r*] of orderable elements and
// integer *k* (1 ≤ *k* ≤ *r* − *l* + 1)
//Output: The value of the *k*th smallest element in *A*[*l*..*r*]
s ← *Partition*(*A*[*l*..*r*]) //see Section 4.2; must return *l* if *l* = *r*
if *s* > *l* + *k* − 1 *SelectionRec*(*A*[*l*..*s* − 1], *k*)
else if *s* < *l* + *k* − 1 *SelectionRec*(*A*[*s* + 1..*r*], *k* − 1 − *s*)
else return *A*[*s*]

4. Using the standard form of an equation of the straight line through two given points, we obtain

$$y - A[l] = \frac{A[r] - A[l]}{r - l}(x - l).$$

Substituting a given value *v* for *y* and solving the resulting equation for *x* yields

$$x = l + \lfloor \frac{(v - A[l])(r - l)}{A[r] - A[l]} \rfloor$$

after the necessary round-off of the second term to guarantee index *l* to be an integer.

5. If *v* = *A*[*l*] or *v* = *A*[*r*], formula (5.4) will yield *x* = *l* and *x* = *r*, respectively, and the search for *v* will stop successfully after comparing *v* with *A*[*x*]. If *A*[*l*] < *v* < *A*[*r*],

$$0 < \frac{(v - A[l])(r - l)}{A[r] - A[l]} < r - l;$$

therefore

$$0 \leq \lfloor \frac{(v - A[l])(r - l)}{A[r] - A[l]} \rfloor \leq r - l - 1$$

and

$$l \leq l + \lfloor \frac{(v - A[l])(r - l)}{A[r] - A[l]} \rfloor \leq r - 1.$$

Hence, if interpolation search does not stop on its current iteration, it reduces the size of the array that remains to be investigated at least by one. Therefore, its worst-case efficiency is in $O(n)$. We want to show that it is, in fact, in $\Theta(n)$. Consider, for example, array $A[0..n-1]$ in which $A[0] = 0$ and $A[i] = n-1$ for $i = 1, 2, \dots, n-1$. If we search for $v = n-1.5$ in this array by interpolation search, its k th iteration ($k = 1, 2, \dots, n$) will have $l = 0$ and $r = n - k$. We will prove this assertion by mathematical induction on k . Indeed, for $k = 1$ we have $l = 0$ and $r = n - 1$. For the general case, assume that the assertion is correct for some iteration k ($1 \leq k < n$) so that $l = 0$ and $r = n - k$. On this iteration, we will obtain the following by applying the algorithm's formula

$$x = 0 + \lfloor \frac{((n - 1.5) - 0)(n - k)}{(n - 1) - 0} \rfloor.$$

Since

$$\frac{(n - 1.5)(n - k)}{(n - 1)} = \frac{(n - 1)(n - k) - 0.5(n - k)}{(n - 1)} = (n - k) - 0.5 \frac{(n - k)}{(n - 1)} < (n - k)$$

and

$$\frac{(n - 1.5)(n - k)}{(n - 1)} = (n - k) - 0.5 \frac{(n - k)}{(n - 1)} > (n - k) - \frac{(n - k)}{(n - 1)} \geq (n - k) - 1,$$

$$x = \lfloor \frac{(n - 1.5)(n - k)}{(n - 1) - 0} \rfloor = (n - k) - 1 = n - (k + 1).$$

Therefore $A[x] = A[n - (k + 1)] = n - 1$ (unless $k = n - 1$), implying that $l = 0$ and $r = n - (k + 1)$ on the next $(k + 1)$ iteration. (If $k = n - 1$, the assertion holds true for the next and last iteration, too: $A[x] = A[0] = 0$, implying that $l = 0$ and $r = 0$.)

6. a. We can solve the inequality $\log_2 \log_2 n + 1 > 6$ as follows:

$$\begin{aligned} \log_2 \log_2 n + 1 &> 6 \\ \log_2 \log_2 n &> 5 \\ \log_2 n &> 2^5 \\ n &> 2^{32} (> 4 \cdot 10^9). \end{aligned}$$

b. Using the formula $\log_a n = \log_a e \ln n$, we can compute the limit as follows:

$$\begin{aligned} \lim_{n \rightarrow \infty} \frac{\log_a \log_a n}{\log_a n} &= \lim_{n \rightarrow \infty} \frac{\log_a e \ln(\log_a e \ln n)}{\log_a e \ln n} = \lim_{n \rightarrow \infty} \frac{\ln \log_a e + \ln \ln n}{\ln n} \\ &= \lim_{n \rightarrow \infty} \frac{\ln \log_a e}{\ln n} + \lim_{n \rightarrow \infty} \frac{\ln \ln n}{\ln n} = 0 + \lim_{n \rightarrow \infty} \frac{\ln \ln n}{\ln n}. \end{aligned}$$

The second limit can be computed by using L'Hôpital's rule:

$$\lim_{n \rightarrow \infty} \frac{\ln \ln n}{\ln n} = \lim_{n \rightarrow \infty} \frac{[\ln \ln n]'}{[\ln n]'} = \lim_{n \rightarrow \infty} \frac{(1/\ln n)(1/n)}{1/n} = \lim_{n \rightarrow \infty} (1/\ln n) = 0.$$

Hence, $\log \log n \in o(\log n)$.

7. a. Recursively, go to the right subtree until a node with the empty right subtree is reached; return the key of that node. We can consider this algorithm as a variable-size-decrease algorithm: after each step to the right, we obtain a smaller instance of the same problem (whether we measure a tree's size by its height or by the number of nodes).
- b. The worst-case efficiency of the algorithm is linear; we should expect its average-case efficiency to be logarithmic (see the discussion in Section 5.6).
8. a. This is an important and well-known algorithm. Case 1: If a key to be deleted is in a leaf, make the pointer from its parent to the key's node null. (If it doesn't have a parent, i.e., it is the root of a single-node tree, make the tree empty.) Case 2: If a key to be deleted is in a node with a single child, make the pointer from its parent to the key's node to point to that child. (If the node to be deleted is the root with a single child, make its child the new root.) Case 3: If a key K to be deleted is in a node with two children, its deletion can be done by the following three-stage procedure. First, find the smallest key K' in the right subtree of the K 's node. (K' is the immediate successor of K in the inorder traversal of the given binary tree; it can be also found by making one step to the right from the K 's node and then all the way to the left until a node with no left subtree is reached). Second, exchange K and K' . Third, delete K in its new node by using either Case 1 or Case 2, depending on whether that node is a leaf or has a single child.

This algorithm is not a variable-size-decrease algorithm because it does not work by reducing the problem to that of deleting a key from a smaller binary tree.

- b. Consider, as an example of the worst case input, the task of deleting the root from the binary tree obtained by successive insertions of keys $2, 1, n, n-1, \dots, 3$. Since finding the smallest key in the right subtree requires following a chain of $n-2$ pointers, the worst-case efficiency of the deletion algorithm is in $\Theta(n)$. Since the average height of a binary tree constructed from n random keys is a logarithmic function (see Section 5.6), we should expect the average-case efficiency of the deletion algorithm be logarithmic as well.
9. If $n = 1$, Player 1 (the player to move first) loses by definition of the misere game because s/he has no choice but to take the last chip. If $2 \leq n \leq m+1$, Player 1 wins by taking $n-1$ chips to leave Player 2 with one chip. If $n = m+2 = 1 + (m+1)$, Player 1 loses because any legal move puts Player 2 in a winning position. If $m+3 \leq n \leq 2m+2$ (i.e., $2 + (m+1) \leq n \leq 2(m+1)$), Player 1 can win by taking $(n-1) \bmod (m+1)$ chips to leave Player 2 with $m+2$ chips, which is a losing position for the player to move next. Thus, an instance is a losing position for Player 1 if and only if $n \bmod (m+1) = 1$. Otherwise, Player 1 wins by taking $(n-1) \bmod (m+1)$ chips; any deviation from this winning strategy puts the opponent in a winning position. The formal proof of the solution's correctness is by strong induction.
10. The problem is equivalent to the game of Nim, with the piles represented by the rows and columns of the bar between the spoiled square and the bar's edges. Thus, the Nim's theory outlined in the section identifies both winning positions and winning moves in this game. According to this theory, an instance of Nim is a winning one (for the player to move next) if and only if its binary digital sum contains at least one 1. In such a position, a winning move can be found as follows. Scan left to right the binary digital sum of the bit strings representing the number of chips in the piles until the first 1 is encountered. Let j be the position of this 1. Select a bit string with a 1 in position j —this is the pile from which some chips will be taken in a winning move. To determine the number of chips to be left in that pile, scan its bit string starting at position j and flip its bits to make the new binary digital sum contain only 0's.

Note: Under the name of *Yucky Chocolate*, the special case of this problem—with the spoiled square in the bar's corner—is discussed, for example, by Yan Stuart in "Math Hysteria: Fun and Games with Mathematics," Oxford University Press, 2004. For such instances, the player going first loses if $m = n$, i.e., the bar has the square shape, and wins if $m \neq n$. Here is a proof by strong induction, which doesn't involve binary representations of the pile sizes. If $m = n = 1$, the player moving first loses by the game's definition. Assuming that the assertion is true for every k -by- k square bar for all $k \leq n$, consider the $n+1$ -by- $n+1$ bar. Any move (i.e., a break

of the bar) creates a rectangular bar with one side of size $k \leq n$ and the other side's size remaining $n + 1$. The second player can always follow with a break creating a k -by- k square bar with a spoiled corner, which is a losing instance by the inductive assumption. And if $m \neq n$, the first player can always "even" the bar by creating the square with the side's size $\min\{m, n\}$, putting the second player in a losing position.

11. Here is a decrease-and-conquer algorithm for this problem. Repeat the following until the problem is solved: Find the largest pancake that is out of order. (If there is none, the problem is solved.) If it is not on the top of the stack, slide the flipper under it and flip to put the largest pancake on the top. Slide the flipper under the first-from-the-bottom pancake that is not in its proper place and flip to increase the number of pancakes in their proper place at least by one.

The number of flips needed by this algorithm in the worst case is $W(n) = 2n - 3$, where $n \geq 2$ is the number of pancakes. Here is a proof of this assertion by mathematical induction. For $n = 2$, the assertion is correct: the algorithm makes one flip for a two-pancake stack with a larger pancake on the top, and it makes no flips for a two-pancake stack with a larger pancake at the bottom. Assume now that the worst-case number of flips for some value of $n \geq 2$ is given by the formula $W(n) = 2n - 3$. Consider an arbitrary stack of $n + 1$ pancakes. With two flips or less, the algorithm puts the largest pancake at the bottom of the stack, where it doesn't participate in any further flips. Hence, the total number of flips needed for any stack of $n + 1$ pancakes is bounded above by

$$2 + W(n) = 2 + (2n - 3) = 2(n + 1) - 3.$$

In fact, this upper bound is attained on the stack of $n + 1$ pancakes constructed as follows: flip a worst-case stack of n pancakes upside down and insert a pancake larger than all the others between the top and the next-to-the-top pancakes. (On the new stack, the algorithm will make two flips to reduce the problem to flipping the worst-case stack of n pancakes.) This completes the proof of the fact that

$$W(n + 1) = 2(n + 1) - 3,$$

which, in turn, completes our mathematical induction proof.


Note: The Web site mentioned in the problem's statement contains, in addition to a visualization applet, an interesting discussion of the problem. (Among other facts, it mentions that the only research paper published by Bill Gates was devoted to this problem.)

This file contains the exercises, hints, and solutions for Chapter 6 of the book "Introduction to the Design and Analysis of Algorithms," 2nd edition, by A. Levitin. The problems that might be challenging for at least some students are marked by \triangleright ; those that might be difficult for a majority of students are marked by \blacktriangleright .

Exercises 6.1

1. Recall that the *median* of a list of n numbers is defined as its $\lceil n/2 \rceil$ smallest element. (The median is larger than one half the elements and is smaller than the other half.) Design a presorting-based algorithm for finding the median and determine its efficiency class.
2. Consider the problem of finding the distance between the two closest numbers in an array of n numbers. (The distance between two numbers x and y is computed as $|x - y|$.)
 - a. Design a presorting-based algorithm for solving this problem and determine its efficiency class.
 - b. Compare the efficiency of this algorithm with that of the brute-force algorithm (see Problem 9 in Exercises 1.2).
3. Let $A = \{a_1, \dots, a_n\}$ and $B = \{b_1, \dots, b_m\}$ be two sets of numbers. Consider the problem of finding their intersection, i.e., the set C of all the numbers that are in both A and B .
 - a. Design a brute-force algorithm for solving this problem and determine its efficiency class.
 - b. Design a presorting-based algorithm for solving this problem and determine its efficiency class.
4. Consider the problem of finding the smallest and largest elements in an array of n numbers.
 - a. Design a presorting-based algorithm for solving this problem and determine its efficiency class.
 - b. Compare the efficiency of the three algorithms: (i) the brute-force algorithm, (ii) this presorting-based algorithm, and (iii) the divide-and-conquer algorithm (see Problem 2 in Exercises 4.1).
5. Show that the average-case efficiency of one-time searching by the algorithm that consists of the most efficient comparison-based sorting algorithm followed by binary search is inferior to the average-case efficiency of sequential search.

6. Estimate how many searches will be needed to justify time spend on pre-sorting an array of 10^3 elements if sorting is done by mergesort and searching is done by binary search. (You may assume that all searches are for elements known to be in the array.) What about an array of 10^6 elements?
7. To sort or not to sort? Design a reasonably efficient algorithm for solving each of the following problems and determine its efficiency class.
 - a. You are given n telephone bills and m checks sent to pay the bills ($n \geq m$). Assuming that telephone numbers are written on the checks, find out who failed to pay. (For simplicity, you may also assume that only one check is written for a particular bill and that it covers the bill in full.)
 - b. You have a file of n student records indicating each student's number, name, home address, and date of birth. Find out the number of students from each of the 50 U.S. states.
8. \triangleright Given a set of $n \geq 3$ points in the Cartesian plane, connect them in a simple polygon, i.e., a closed path through all the points so that its line segments (the polygon's edges) do not intersect (except for neighboring edges at their common vertex). For example,



- a. Does the problem always have a solution? Does it always have a unique solution?
 - b. Design a reasonably efficient algorithm for solving this problem and indicate its efficiency class.
9. \triangleright You have an array of n numbers and a number s . Find out whether the array contains two elements whose sum is s . (For example, for the array 5, 9, 1, 3 and $s = 6$, the answer is yes, but for the same array and $s = 7$, the answer is no.) Design an algorithm for this problem with a better than quadratic time efficiency.
10. \triangleright You have a list of n open intervals $(a_1, b_1), (a_2, b_2), \dots, (a_n, b_n)$ on the

real line. (An open interval (a, b) comprises all the points strictly between its endpoints a and b , i.e., $(a, b) = \{x \mid a < x < b\}$.) Find the maximal number of these intervals that have a common point. For example, for the intervals $(1, 4)$, $(0, 3)$, $(-1.5, 2)$, $(3.6, 5)$, this maximal number is 3. Design an algorithm for this problem with a better than quadratic time efficiency.

11. a. Design an efficient algorithm for finding all sets of anagrams in a large file such as a dictionary of English words [Ben00]. For example, *eat*, *ate*, and *tea* belong to one such a set.
- b. Write a program implementing the algorithm.

Hints to Exercises 6.1

1. The algorithm is suggested by the problem's statement. Its analysis is similar to the examples discussed in this section.
2. This problem is similar to one of the examples in this section.
3. a. Compare every element in one set with all the elements in the other.

b. In fact, you can use presorting in three different ways: sort elements of just one of the sets, sort elements of each of the sets separately, and sort elements of the two sets together.
4. a. How do we find the smallest and largest elements in a sorted list?

b. The brute-force algorithm and the divide-and-conquer algorithm are both linear.
5. Use the known results about the average-case efficiencies of the algorithms in this question.
6. Assume that sorting requires about $n \log_2 n$ comparisons. Use the known results about the number of comparisons made, on the average, in a successful search by binary search and by sequential search.
7. a. The problem is similar to one of the preceding problems in these exercises.

b. How would you solve this problem if the student information were written on index cards? Better yet, think how somebody else, who has never taken a course on algorithms but possesses a good dose of common sense, would solve this problem.
8. a. Many problems of this kind have exceptions for one particular configuration of points. As to the question about a solution's uniqueness, you can get the answer by considering a few small "random" instances of the problem.

b. Construct a polygon for a few small "random" instances of the problem. Try to construct polygons in some systematic fashion.
9. It helps to think about real numbers as ordered points on the real line. Considering the special case of $s = 0$, with a given array containing both negative and positive numbers, might be helpful, too.
10. After sorting the a_i 's and b_i 's, the problem can be solved in linear time.
11. Use the presorting idea twice.

Solutions to Exercises 6.1

- Sort the list and then simply return the $\lceil n/2 \rceil$ th element of the sorted list. Assuming the efficiency of the sorting algorithm is in $O(n \log n)$, the time efficiency of the entire algorithm will be in

$$O(n \log n) + \Theta(1) = O(n \log n).$$

- Sort the array first and then scan it to find the smallest difference between two successive elements $A[i]$ and $A[i + 1]$ ($0 \leq i \leq n - 2$).

b. The time efficiency of the brute-force algorithm is in $\Theta(n^2)$ because the algorithm considers $n(n - 1)/2$ pairs of the array's elements. (In the crude version given in Problem 9 of Exercises 1.2, the same pair is considered twice but this doesn't change the efficiency's order, of course.) If the presorting is done with a $O(n \log n)$ algorithm, the running time of the entire algorithm will be in

$$O(n \log n) + \Theta(n) = O(n \log n).$$

- Initialize a list to contain elements of $C = A \cap B$ to empty. Compare every element a_i in A with successive elements of B : if $a_i = b_j$, add this value to the list C and proceed to the next element in A . (In fact, if $a_i = b_j$, b_j need not be compared with the remaining elements in A and may be deleted from B .) In the worst case of input sets with no common elements, the total number of element comparisons will be equal to nm , putting the algorithm's efficiency in $O(nm)$.

b. First solution: Sort elements of one of the sets, say, A , stored in an array. Then use binary search to search for each element of B in the sorted array A : if a match is found, add this value to the list C . If sorting is done with a $O(n \log n)$ algorithm, the total running time will be in

$$O(n \log n) + mO(\log n) = O((m + n) \log n).$$

Note that the efficiency formula implies that it is more efficient to sort the smaller one of the two input sets.

Second solution: Sort the lists representing sets A and B , respectively. Scan the lists in the mergesort-like manner but output only the values common to the two lists. If sorting is done with a $O(n \log n)$ algorithm, the total running time will be in

$$O(n \log n) + O(m \log m) + O(n + m) = O(s \log s) \text{ where } s = \max\{n, m\}.$$

Third solution: Combine the elements of both A and B in a single list and sort it. Then scan this sorted list by comparing pairs of its consecutive elements: if $L_i = L_{i+1}$, add this common value to the list C and increment i by two. If sorting is done with an $n \log n$ algorithm, the total running time will be in

$$O((n+m) \log(n+m)) + \Theta(n+m) = O(s \log s) \text{ where } s = \max\{n, m\}.$$

4. a. Sort the list and return its first and last elements as the values of the smallest and largest elements, respectively. Assuming the efficiency of the sorting algorithm used is in $O(n \log n)$, the time efficiency of the entire algorithm will be in

$$O(n \log n) + \Theta(1) + \Theta(1) = O(n \log n).$$

b. The brute-force algorithm and the divide-and-conquer algorithm are both linear, and, hence, superior to the presorting-based algorithm.

5. Since the average-case efficiency of any comparison-based sorting algorithm is known to be in $\Omega(n \log n)$ while the average-case efficiency of binary search is in $\Theta(\log n)$, the average-case efficiency of the searching algorithm combining the two will be in

$$\Omega(n \log n) + \Theta(\log n) = \Omega(n \log n).$$

This is inferior to the average-case efficiency of sequential search, which is linear (see Section 2.1).

6. Let k be the smallest number of searches needed for the sort–binary search algorithm to make fewer comparisons than k searches by sequential search (for average successful searches). Assuming that a sorting algorithm makes about $n \log n$ comparisons on the average and using the formulas for the average number of key comparisons for binary search (about $\log_2 n$) and sequential search (about $n/2$), we get the following inequality

$$n \log_2 n + k \log_2 n \leq kn/2.$$

Thus, we need to find the smallest value of k so that

$$k \geq \frac{n \log_2 n}{n/2 - \log_2 n}.$$

Substituting $n = 10^3$ into the right-hand side yields $k_{\min} = 21$; substituting $n = 10^6$ yields $k_{\min} = 40$.

Note: For large values of n , we can simplify the last inequality by eliminating the relatively insignificant term $\log_2 n$ from the denominator of the right-hand side to obtain

$$k \geq \frac{n \log_2 n}{n/2} \quad \text{or} \quad k \geq 2 \log_2 n.$$

This inequality would yield the answers of 20 and 40 for $n = 10^3$ and $n = 10^6$, respectively.

7. a. The following algorithm will beat the brute-force comparisons of the telephone numbers on the bills and the checks: Using an efficient sorting algorithm, sort the bills and sort the checks. (In both cases, sorting has to be done with respect to their telephone numbers, say, in increasing order.) Then do a merging-like scan of the two sorted lists by comparing the telephone numbers b_i and c_j on the current bill and check, respectively: if $b_i < c_j$, add b_i to the list of unpaid telephone numbers and increment i ; if $b_i > c_j$, increment j ; if $b_i = c_j$, increment both i and j . Stop as soon as one of the two lists becomes empty and append all the remaining telephone numbers on the bill list, if any, to the list of the unpaid ones.


The time efficiency of this algorithm will be in

$$O(n \log n) + O(m \log m) + O(n + m) \underset{n \geq m}{=} O(n \log n).$$

This is superior to the $O(nm)$ efficiency of the brute-force algorithm (but inferior, for the average case, to solving this problem with hashing discussed in Section 7.3).

- b. Initialize 50 state counters to zero. Scan the list of student records and, for a current student record, increment the corresponding state counter. The algorithm's time efficiency will be in $\Theta(n)$, which is superior to any algorithm that uses presorting of student records by a comparison-based algorithm.
8. a. The problem has a solution if and only if all the points don't lie on the same line. And, if a solution exists, it may not be unique.
 - b. Find the lowest point P^* , i.e., the one with the smallest y coordinate. in the set. (If there is a tie, take, say, the leftmost among them, i.e., the one with the smallest x coordinate.) For each of the other $n - 1$ points, compute its angle in the polar coordinate system with the origin at P^* and sort the points in increasing order of these angles, breaking ties in favor of a point closer to P^* . (Instead of the angles, you can use the

line slopes with respect to the horizontal line through P^* .) Connect the points in the order generated, adding the last segment to return to P^* .



Finding the lowest point P^* is in $\Theta(n)$, computing the angles (slopes) is in $\Theta(n)$, sorting the points according to their angles can be done with a $O(n \log n)$ algorithm. The efficiency of the entire algorithm is in $O(n \log n)$.

9. Assume first that $s = 0$. Then $A[i] + A[j] = 0$ if and only if $A[i] = -A[j]$, i.e., these two elements have the same absolute value but opposite signs. We can check for presence of such elements in a given array in several different ways. If all the elements are known to be distinct, we can simply replace each element $A[i]$ by its absolute value $|A[i]|$ and solve the element uniqueness problem for the array of the absolute values in $O(n \log n)$ time with the presorting based algorithm. If a given array can have equal elements, we can modify our approach as follows: We can sort the array in nondecreasing order of their absolute values (e.g., -6, 3, -3, 1, 3 becomes 1, 3, -3, 3, -6), and then scan the array sorted in this fashion to check whether it contains a consecutive pair of elements with the same absolute value and opposite signs (e.g., 1, 3, -3, 3, -6 does). If such a pair of elements exists, the algorithm returns yes, otherwise, it returns no.

The case of an arbitrary value of s is reduced to the case of $s = 0$ by the following substitution: $A[i] + A[j] = s$ if and only if $(A[i] - s/2) + (A[j] - s/2) = 0$. In other words, we can start the algorithm by subtracting $s/2$ from each element and then proceed as described above.

(Note that we took advantage of the instance simplification idea twice: by reducing the problem's instance to one with $s = 0$ and by presorting the array.)

10. Sort all the a_i 's and b_j 's by a $O(n \log n)$ algorithm in a single nondecreasing list, treating b_j as if it were smaller than a_i in case of the tie $a_i = b_j$. Scan the list left to right computing the running difference D between the number of a_i 's and b_j 's seen so far. In other words, initialize D to 0

and then increment or decrement it by 1 depending on whether the next element on the list is a left endpoint a_i or a right endpoint b_j , respectively. The maximum value of D is the number in question.

Note 1: One can also implement this algorithm by sorting a_i 's and b_j 's separately and then computing the running difference D by merging-like processing of the two sorted lists.

Note 2: This solution is suggested by D. Ginat in his paper "Algorithmic Pattern and the Case of the Sliding Delta," *SIGCSE Bulletin*, vol 36, no. 2, June 2004, pp. 29-33.

11. First, attach to every word in the file—as another field of the word's record, for example—its signature defined as the string of the word's letters in alphabetical order. (Obviously, words belong to the same anagram set if and only if they have the same signature.) Sort the records in alphabetical order of their signatures. Scan the list to identify contiguous subsequences, of length greater than one, of records with the same signature.

Note: Jon Bentley describes a real system where a similar problem occurred in [Ben00], p.17, Problem 6.

Exercises 6.2

1. Solve the following system by Gaussian elimination.

$$\begin{aligned}x_1 + x_2 + x_3 &= 2 \\2x_1 + x_2 + x_3 &= 3 \\x_1 - x_2 + 3x_3 &= 8.\end{aligned}$$

2. a. Solve the system of the previous question by the *LU* decomposition method.

b. From the standpoint of general algorithm design techniques, how would you classify the *LU*-decomposition method?
3. Solve the system of Problem 1 by computing the inverse of its coefficient matrix and then multiplying it by the right-hand side vector.
4. Would it be correct to get the efficiency class of the elimination stage of Gaussian elimination as follows?

$$\begin{aligned}C(n) &= \sum_{i=1}^{n-1} \sum_{j=i+1}^n \sum_{k=i}^{n+1} 1 = \sum_{i=1}^{n-1} (n+2-i)(n-i) \\&= \sum_{i=1}^{n-1} [(n+2)n - i(2n+2) + i^2] \\&= \sum_{i=1}^{n-1} (n+2)n - \sum_{i=1}^{n-1} (2n+2)i + \sum_{i=1}^{n-1} i^2.\end{aligned}$$

Since $s_1(n) = \sum_{i=1}^{n-1} (n+2)n \in \Theta(n^3)$, $s_2(n) = \sum_{i=1}^{n-1} (2n+2)i \in \Theta(n^3)$, and $s_3(n) = \sum_{i=1}^{n-1} i^2 \in \Theta(n^3)$, $s_1(n) - s_2(n) + s_3(n) \in \Theta(n^3)$.

5. Write a pseudocode for the back-substitution stage of Gaussian elimination and show that its running time is in $\Theta(n^2)$.
6. Assuming that division of two real numbers takes three times longer than their multiplication, estimate how much faster *BetterGaussElimination* is than *GaussElimination*. (Of course, you should also assume that a compiler is not going to eliminate the inefficiency in *GaussElimination*.)
7. a. Give an example of a system of two linear equations in two unknowns that has a unique solution and solve it by Gaussian elimination.

b. Give an example of a system of two linear equations in two unknowns that has no solution and apply Gaussian elimination to it.

c. Give an example of a system of two linear equations in two unknowns that has infinitely many solutions and apply Gaussian elimination to it.

8. The ***Gauss-Jordan elimination*** method differs from Gaussian elimination in that the elements above the main diagonal of the coefficient matrix are made zero at the same time and by the same use of a pivot row as the elements below the main diagonal.
 - a. Apply the Gauss-Jordan method to the system of Problem 1 of these exercises.
 - b. What general design technique is this algorithm based on?
 - c.▷ In general, how many multiplications are made by this method while solving a system of n equations in n unknowns? How does this compare with the number of multiplications made by the Gaussian elimination method in both its elimination and its back-substitution stages?
9. A system $Ax = b$ of n linear equations in n unknowns has a unique solution if and only if $\det A \neq 0$. Is it a good idea to check this condition before applying Gaussian elimination to a system?
10. a. Apply Cramer's rule to solve the system of Problem 1 of these exercises.
 - b. Estimate how many times longer it will take to solve a system of n linear equations in n unknowns by Cramer's rule than by Gaussian elimination. (Assume that all the determinants in Cramer's rule formulas are computed independently by Gaussian elimination.)
11. ► *Lights out* This one-person game is played on an n -by- n board composed of 1-by-1 light panels. Each panel has a switch that can be turned on and off, thereby toggling the on/off state of this and four vertically and horizontally adjacent panels. (Of course, toggling a corner square affects the total of three panels, and toggling a noncorner panel on the board's border affects the total of four squares.) Given an initial subset of lighted squares, the goal is to turn all the lights off.
 - (a) Show that an answer can be found by solving a system of linear equations with 0/1 coefficients and right-hand sides using modulo 2 arithmetic.
 - (b) Use Gaussian elimination to solve the 2-by-2 "all-ones" instance of this problem, where all the panels of the 2-by-2 board are initially lit.
 - (c) Use Gaussian elimination to solve the 3-by-3 "all-ones" instance of this problem, where all the panels of the 3-by-3 board are initially lit.

Hints to Exercises 6.2

1. Trace the algorithm as we did in solving another system in the section.
2. a. Use the Gaussian elimination results as explained in the text.

b. It is one of the varieties of the transform-and-conquer technique. Which one?
3. To find the inverse, you can either solve the system with three simultaneous right-hand side vectors representing the columns of the 3-by-3 identity matrix or use the LU decomposition of the system's coefficient matrix found in Problem 2.
4. Though the final answer is correct, its derivation contains an error you have to find.
5. The pseudocode of this algorithm is quite straightforward. If you are in doubt, see the section's example tracing the algorithm. The order of growth of the algorithm's running time can be estimated by following the standard plan for the analysis of nonrecursive algorithms.
6. Estimate the ratio of the algorithm running times, by using the approximate formulas for the number of divisions and the number of multiplications in both algorithms.
7. a. This is a "normal" case: one of the two equations should not be proportional to the other.

b. The coefficients of one equation should be the same or proportional to the corresponding coefficients of the other equation while the right-hand sides should not.

c. The two equations should be either the same or proportional to each other (including the right-hand sides).
8. a. Manipulate the matrix rows above a pivot row the same way the rows below the pivot row are changed.

b. Are the Gauss-Jordan method and Gaussian elimination based on the same algorithm design technique or on different ones?

c. Derive the formula for the number of multiplications in the Gauss-Jordan method the same way it was done for Gaussian elimination in Section 6.2.
9. How long will it take to compute the determinant compared to the time needed to apply Gaussian elimination to the system?

10. a. Apply Cramer's rule to the system given.
- b. How many distinct determinants are there in the Cramer's rule formulas?
11. a. If x_{ij} is the number of times the panel in the i th row and j th column needs to be toggled in a solution, what can be said about x_{ij} ? After you answer this question, show that the binary matrix representing an initial state of the board can be represented as a linear combination (in modulo 2 arithmetic) of n^2 binary matrices each representing the affect of toggling an individual panel.
- b. Set up a system of four equations in four unknowns (see part a) and solve it by Gaussian elimination performing all operations in modulo-2 arithmetic.
- c. If you believe that a system of nine equations in nine unknowns is too large to solve by hand, write a program to solve the problem.

Solutions to Exercises 6.2

1. a. Solve the following system by Gaussian elimination

$$\begin{aligned}x_1 + x_2 + x_3 &= 2 \\2x_1 + x_2 + x_3 &= 3 \\x_1 - x_2 + 3x_3 &= 8\end{aligned}$$

$$\left[\begin{array}{cccc} 1 & 1 & 1 & 2 \\ 2 & 1 & 1 & 3 \\ 1 & -1 & 3 & 8 \end{array} \right] \begin{array}{l} \\ \text{row 2} - \frac{2}{1}\text{row 1} \\ \text{row 3} - \frac{1}{1}\text{row 1} \end{array}$$

$$\left[\begin{array}{cccc} 1 & 1 & 1 & 2 \\ 0 & -1 & -1 & -1 \\ 0 & -2 & 2 & 6 \end{array} \right] \begin{array}{l} \\ \\ \text{row 3} - \frac{-2}{-1}\text{row 2} \end{array}$$

$$\left[\begin{array}{cccc} 1 & 1 & 1 & 2 \\ 0 & -1 & -1 & -1 \\ 0 & 0 & 4 & 8 \end{array} \right]$$

Then, by backward substitutions, we obtain the solution as follows:

$$x_3 = 8/4 = 2, \quad x_2 = (-1 + x_3)/(-1) = -1, \quad \text{and} \quad x_1 = (2 - x_3 - x_2)/1 = 1.$$

2. a. Repeating the elimination stage (or using its results obtained in Problem 1), we get the following matrices L and U :

$$L = \begin{bmatrix} 1 & 0 & 0 \\ 2 & 1 & 0 \\ 1 & 2 & 1 \end{bmatrix}, \quad U = \begin{bmatrix} 1 & 1 & 1 \\ 0 & -1 & -1 \\ 0 & 0 & 4 \end{bmatrix}.$$

On substituting $y = Ux$ into $LUx = b$, the system $Ly = b$ needs to be solved first. Here, the augmented coefficient matrix is:

$$\left[\begin{array}{cccc} 1 & 0 & 0 & 2 \\ 2 & 1 & 0 & 3 \\ 1 & 2 & 1 & 8 \end{array} \right]$$

Its solution is

$$y_1 = 2, \quad y_2 = 3 - 2y_1 = -1, \quad y_3 = 8 - y_1 - 2y_2 = 8.$$

Solving now the system $Ux = y$, whose augmented coefficient matrix is

$$\left[\begin{array}{cccc} 1 & 1 & 1 & 2 \\ 0 & -1 & -1 & -1 \\ 0 & 0 & 4 & 8 \end{array} \right],$$

yields the following solution to the system given:

$$x_3 = 2, \quad x_2 = (-1 + x_3)/(-1) = -1, \quad x_1 = 2 - x_3 - x_2 = 1.$$

b. The most fitting answer is the representation change technique.

3. Solving simultaneously the system with the three right-hand side vectors:

$$\begin{bmatrix} 1 & 1 & 1 & 1 & 0 & 0 \\ 2 & 1 & 1 & 0 & 1 & 0 \\ 1 & -1 & 3 & 0 & 0 & 1 \end{bmatrix} \quad \begin{array}{l} \text{row 2} - \frac{2}{1}\text{row 1} \\ \text{row 3} - \frac{1}{1}\text{row 1} \end{array}$$

$$\begin{bmatrix} 1 & 1 & 1 & 1 & 0 & 0 \\ 0 & -1 & -1 & -2 & 1 & 0 \\ 0 & -2 & 2 & -1 & 0 & 1 \end{bmatrix} \quad \text{row 3} - \frac{-2}{-1}\text{row 1}$$

$$\begin{bmatrix} 1 & 1 & 1 & 1 & 0 & 0 \\ 0 & -1 & -1 & -2 & 1 & 0 \\ 0 & 0 & 4 & 3 & -2 & 1 \end{bmatrix}$$

Solving the system with the first right-hand side column

$$\begin{bmatrix} 1 \\ -2 \\ 3 \end{bmatrix}$$

yields the following values of the first column of the inverse matrix:

$$\begin{bmatrix} -1 \\ \frac{5}{4} \\ \frac{3}{4} \end{bmatrix}.$$

Solving the system with the second right-hand side column

$$\begin{bmatrix} 0 \\ 1 \\ -2 \end{bmatrix}$$

yields the following values of the second column of the inverse matrix:

$$\begin{bmatrix} 1 \\ -\frac{1}{2} \\ -\frac{1}{2} \end{bmatrix}.$$

Solving the system with the third right-hand side column

$$\begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}$$

yields the following values of the third column of the inverse matrix:

$$\begin{bmatrix} 0 \\ -\frac{1}{4} \\ \frac{1}{4} \end{bmatrix}.$$

Thus, the inverse of the coefficient matrix is

$$\begin{bmatrix} -1 & 1 & 0 \\ \frac{5}{4} & -\frac{1}{2} & -\frac{1}{4} \\ \frac{3}{4} & -\frac{1}{2} & \frac{1}{4} \end{bmatrix},$$

which leads to the following solution to the original system

$$x = A^{-1}b = \begin{bmatrix} -1 & 1 & 0 \\ \frac{5}{4} & -\frac{1}{2} & -\frac{1}{4} \\ \frac{3}{4} & -\frac{1}{2} & \frac{1}{4} \end{bmatrix} \begin{bmatrix} 2 \\ 3 \\ 8 \end{bmatrix} = \begin{bmatrix} 1 \\ -1 \\ 2 \end{bmatrix}.$$

4. In general, the fact that $f_1(n) \in \Theta(n^3)$, $f_2(n) \in \Theta(n^3)$, and $f_3(n) \in \Theta(n^3)$ does not necessarily imply that $f_1(n) - f_2(n) + f_3(n) \in \Theta(n^3)$, because the coefficients of the highest third-degree terms can cancel each other. As a specific example, consider $f_1(n) = n^3 + n$, $f_2(n) = 2n^3$, and $f_3(n) = n^3$. Each of this functions is in $\Theta(n^3)$, but $f_1(n) - f_2(n) + f_3(n) = n \in \Theta(n)$.

5. **Algorithm** *GaussBackSub*($A[1..n, 1..n+1]$)

//Implements the backward substitution stage of Gaussian elimination
 //by solving a given system with an upper-triangular coefficient matrix
 //Input: Matrix $A[1..n, 1..n+1]$, with the first n columns in the upper-
 //triangular form
 //Output: A solution of the system of n linear equations in n unknowns
 //whose coefficient matrix and right-hand side are the first n columns
 //of A and its $(n+1)$ st column, respectively

for $i \leftarrow n$ **downto** 1 **do**

$temp \leftarrow 0.0$

for $j \leftarrow n$ **downto** $i+1$

$temp \leftarrow temp + A[i, j] * x[j]$

$x[i] \leftarrow (A[i, n+1] - temp) / A[i, i]$

return x

The basic operation is multiplication of two numbers. The number of times it will be executed is given by the sum

$$\begin{aligned} M(n) &= \sum_{i=1}^n \sum_{j=i+1}^n 1 = \sum_{i=1}^n (n - (i+1) + 1) = \sum_{i=1}^n (n - i) \\ &= (n-1) + (n-2) + \dots + 1 = \frac{(n-1)n}{2} \in \Theta(n^2). \end{aligned}$$

6. Let $D^{(G)}(n)$ and $M^{(G)}(n)$ be the numbers of divisions and multiplications made by *GaussElimination*, respectively. Using the count formula derived in Section 6.2, we obtain the following approximate counts:

$$D^{(G)}(n) = M^{(G)}(n) \approx \sum_{i=1}^{n-1} \sum_{j=i+1}^n \sum_{k=i}^{n+1} 1 \approx \frac{1}{3}n^3.$$

Let $D^{(BG)}(n)$ and $M^{(BG)}(n)$ be the numbers of divisions and multiplications made by *BetterGaussElimination*, respectively. We have the following approximations:

$$D^{(BG)}(n) \approx \sum_{i=1}^{n-1} \sum_{j=i+1}^n 1 \approx \frac{1}{2}n^2 \quad \text{and} \quad M^{(BG)}(n) = M^{(G)}(n) \approx \frac{1}{3}n^3.$$

Let c_d and c_m be the times of one division and of one multiplication, respectively. We can estimate the ratio of the running times of the two algorithms as follows:

$$\begin{aligned} \frac{T^{(G)}(n)}{T^{(BG)}(n)} &\approx \frac{c_d D^{(G)}(n) + c_m M^{(G)}(n)}{c_d D^{(BG)}(n) + c_m M^{(BG)}(n)} \approx \frac{c_d \frac{1}{3}n^3 + c_m \frac{1}{3}n^3}{c_d \frac{1}{2}n^2 + c_m \frac{1}{3}n^3} \\ &\approx \frac{c_d \frac{1}{3}n^3 + c_m \frac{1}{3}n^3}{c_m \frac{1}{3}n^3} = \frac{c_d + c_m}{c_m} = \frac{c_d}{c_m} + 1 = 4. \end{aligned}$$

7. a. The elimination stage should yield a 2-by-2 upper-triangular matrix with nonzero coefficients on its main diagonal.
- b. The elimination stage should yield a 2-by-2 matrix whose second row is $0 \ 0 \ \alpha$ where $\alpha \neq 0$.
- c. The elimination stage should yield a 2-by-2 matrix whose second row is $0 \ 0 \ 0$.
8. a. Solve the following system by the Gauss-Jordan method

$$\begin{aligned} x_1 + x_2 + x_3 &= 2 \\ 2x_1 + x_2 + x_3 &= 3 \\ x_1 - x_2 + 3x_3 &= 8 \end{aligned}$$

$$\begin{aligned} \left[\begin{array}{cccc} 1 & 1 & 1 & 2 \\ 2 & 1 & 1 & 3 \\ 1 & -1 & 3 & 8 \end{array} \right] & \begin{array}{l} \text{row 2} - 2\text{row 1} \\ \text{row 3} - \frac{1}{1}\text{row 1} \end{array} \quad \left[\begin{array}{cccc} 1 & 1 & 1 & 2 \\ 0 & -1 & -1 & -1 \\ 0 & -2 & 2 & 6 \end{array} \right] \begin{array}{l} \text{row 1} - \frac{1}{-1}\text{row 2} \\ \text{row 3} - \frac{-2}{-1}\text{row 2} \end{array} \\ \left[\begin{array}{cccc} 1 & 0 & 0 & 1 \\ 0 & -1 & -1 & -1 \\ 0 & 0 & 4 & 8 \end{array} \right] & \begin{array}{l} \text{row 1} - \frac{0}{4}\text{row 3} \\ \text{row 2} - \frac{-1}{4}\text{row 3} \end{array} \quad \left[\begin{array}{cccc} 1 & 0 & 0 & 1 \\ 0 & -1 & 0 & 1 \\ 0 & 0 & 4 & 8 \end{array} \right] \end{aligned}$$

We obtain the solution by dividing the right hand side values by the corresponding elements of the diagonal matrix:

$$x_1 = 1/1 = 1, \quad x_2 = 1/-1 = -1, \quad x_3 = 8/4 = 2.$$

b. The Gauss-Jordan method is also an example of an algorithm based on the instance simplification idea. The two algorithms differ in the kind of a simpler instance to which they transfer a given system: Gaussian elimination transforms a system to an equivalent system with an upper-triangular coefficient matrix whereas the Gauss-Jordan method transforms it to a system with a diagonal matrix.

c. Here is a basic pseudocode for the Gauss-Jordan elimination:

Algorithm *GaussJordan*($A[1..n, 1..n]$, $b[1..n]$)
//Applies Gaussian-Jordan elimination to matrix A of a system's
//coefficients, augmented with vector b of the system's right-hand sides
//Input: Matrix $A[1..n, 1..n]$ and column-vector $b[1..n]$
//Output: An equivalent diagonal matrix in place of A with the
//corresponding right-hand side values in its $(n + 1)$ st column
for $i \leftarrow 1$ **to** n **do** $A[i, n + 1] \leftarrow b[i]$ //augment the matrix
for $i \leftarrow 1$ **to** n **do**
 for $j \leftarrow 1$ **to** n **do**
 if $j \neq i$
 $temp \leftarrow A[j, i] / A[i, i]$ //assumes $A[i, i] \neq 0$
 for $k \leftarrow i$ **to** $n + 1$ **do**
 $A[j, k] \leftarrow A[j, k] - A[i, k] * temp$

The number of multiplications made by the above algorithm can be computed as follows:

$$\begin{aligned} M(n) &= \sum_{i=1}^n \sum_{\substack{j=1 \\ j \neq i}}^n \sum_{k=i}^{n+1} 1 = \sum_{i=1}^n \sum_{\substack{j=1 \\ j \neq i}}^n (n + 1 - i + 1) = \sum_{i=1}^n \sum_{\substack{j=1 \\ j \neq i}}^n (n + 2 - i) \\ &= \sum_{i=1}^n (n + 2 - i)(n - 1) = (n - 1) \sum_{i=1}^n (n + 2 - i) \\ &= (n - 1)[(n + 1) + n + \dots + 2] = (n - 1)[(n + 1) + n + \dots + 1 - 1] \\ &= (n - 1)\left[\frac{(n + 1)(n + 2)}{2} - 1\right] = \frac{(n - 1)n(n + 3)}{2} \approx \frac{1}{2}n^3. \end{aligned}$$

The total number of multiplications made in both elimination and backward substitution stages of the Gaussian elimination method is equal to

$$\frac{n(n - 1)(2n + 5)}{6} + \frac{(n - 1)n}{2} = \frac{(n - 1)n(n + 4)}{3} \approx \frac{1}{3}n^3,$$

which is about 1.5 smaller than in the Gauss-Jordan method.

Note: The Gauss-Jordan method has an important advantage over Gaussian elimination: being more uniform, it is more suitable for efficient implementation on a parallel computer.

9. Since the time needed for computing the determinant of the system's coefficient matrix is about the same as the time needed for solving the system (or detecting that the system does not have a unique solution) by Gaussian elimination, computing the determinant of the coefficient matrix to check whether it is equal to zero is not a good idea from the algorithmic point of view.

10. a. Solve the following system by Cramer's rule:

$$x_1 + x_2 + x_3 = 2$$

$$2x_1 + x_2 + x_3 = 3$$

$$x_1 - x_2 + 3x_3 = 8$$

$$|A| = \begin{vmatrix} 1 & 1 & 1 \\ 2 & 1 & 1 \\ 1 & -1 & 3 \end{vmatrix} = 1 \cdot 1 \cdot 3 + 1 \cdot 1 \cdot 1 + 2 \cdot (-1) \cdot 1 - 1 \cdot 1 \cdot 1 - 2 \cdot 1 \cdot 3 - (-1) \cdot 1 \cdot 1$$

$$|A_1| = \begin{vmatrix} 2 & 1 & 1 \\ 3 & 1 & 1 \\ 8 & -1 & 3 \end{vmatrix} = 2 \cdot 1 \cdot 3 + 1 \cdot 1 \cdot 8 + 3 \cdot (-1) \cdot 1 - 8 \cdot 1 \cdot 1 - 3 \cdot 1 \cdot 3 - (-1) \cdot 1 \cdot 1$$

$$|A_2| = \begin{vmatrix} 1 & 2 & 1 \\ 2 & 3 & 1 \\ 1 & 8 & 3 \end{vmatrix} = 1 \cdot 3 \cdot 3 + 2 \cdot 1 \cdot 1 + 2 \cdot 8 \cdot 1 - 1 \cdot 3 \cdot 1 - 2 \cdot 2 \cdot 3 - 8 \cdot 1 \cdot 1 = 4,$$

$$|A_3| = \begin{vmatrix} 1 & 1 & 2 \\ 2 & 1 & 3 \\ 1 & -1 & 8 \end{vmatrix} = 1 \cdot 1 \cdot 8 + 1 \cdot 3 \cdot 1 + 2 \cdot (-1) \cdot 2 - 1 \cdot 1 \cdot 2 - 2 \cdot 1 \cdot 8 - (-1) \cdot 3 \cdot 1$$

Hence,

$$x_1 = \frac{|A_1|}{|A|} = \frac{-4}{-4} = 1, \quad x_2 = \frac{|A_2|}{|A|} = \frac{4}{-4} = -1, \quad x_3 = \frac{|A_3|}{|A|} = \frac{-8}{-4} = 2.$$

- b. Cramer's rule requires computing $n + 1$ distinct determinants. If each of them is computed by applying Gaussian elimination, it will take about $n + 1$ times longer than solving the system by Gaussian elimination. (The time for the backward substitution stage was not accounted for in the preceding argument because of its quadratic efficiency vs. cubic efficiency of the elimination stage.)

11. a. Any feasible state of the board can be described by an n -by- n binary matrix, in which the element in the i th row and j th column is equal to 1 if and only if the corresponding panel is lit. Let S and F be such matrices representing the initial and final (all-zeros) boards, respectively. The impact of toggling the panel at (i, j) on a board represented by a binary matrix M can be interpreted as the modulo-2 matrix addition $M + A_{ij}$, where A_{ij} is the matrix in which the only entries equal to 1 are those that are in the (i, j) and adjacent to it positions. For example, if M is a 3-by-3 all-ones matrix representing a 3-by-3 board of all-lit panels, then the impact of turning off the $(2, 2)$ panel can be represented as

$$M + A_{22} = \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} + \begin{bmatrix} 0 & 1 & 0 \\ 1 & 1 & 1 \\ 0 & 1 & 0 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 1 \\ 0 & 0 & 0 \\ 1 & 0 & 1 \end{bmatrix}.$$

Let x_{ij} is the number of times the (i, j) panel is toggled in a solution that transforms the board from a starting state S to a final state F . Since the ultimate impact of toggling this panel depends only on whether x_{ij} is even or odd, we can assume with no loss in generality that x_{ij} is either 0 or 1. Then a solution to the puzzle can be expressed by the matrix equation

$$S + \sum_{i,j=1}^n x_{ij} A_{ij} = F,$$

where all the operations are assumed to be performed modulo 2. Taking into account that $F = 0$, the all-zeros n -by- n matrix, the last equation is equivalent to

$$\sum_{i,j=1}^n x_{ij} A_{ij} = S.$$

(The last equation can also be interpreted as transforming the final all-zero board to the initial board S .)

Note: This solution follows Eric W. Weisstein et al. "Lights Out Puzzle" from MathWorld—A Wolfram Web Resource at <http://mathworld.wolfram.com/Light>

- b. The system of linear equations for the instance in question (see part a) is

$$x_{11} \begin{bmatrix} 1 & 1 \\ 1 & 0 \end{bmatrix} + x_{12} \begin{bmatrix} 1 & 1 \\ 0 & 1 \end{bmatrix} + x_{21} \begin{bmatrix} 1 & 0 \\ 1 & 1 \end{bmatrix} + x_{22} \begin{bmatrix} 0 & 1 \\ 1 & 1 \end{bmatrix} = \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix}$$

or

$$\begin{aligned} 1 \cdot x_{11} + 1 \cdot x_{12} + 1 \cdot x_{21} + 0 \cdot x_{22} &= 1 \\ 1 \cdot x_{11} + 1 \cdot x_{12} + 0 \cdot x_{21} + 1 \cdot x_{22} &= 1 \\ 1 \cdot x_{11} + 0 \cdot x_{12} + 1 \cdot x_{21} + 1 \cdot x_{22} &= 1 \\ 0 \cdot x_{11} + 1 \cdot x_{12} + 1 \cdot x_{21} + 1 \cdot x_{22} &= 1. \end{aligned}$$

Solving this system in modulo-2 arithmetic by Gaussian elimination proceeds as follows:


$$\begin{bmatrix} 1 & 1 & 1 & 0 & 1 \\ 1 & 1 & 0 & 1 & 1 \\ 1 & 0 & 1 & 1 & 1 \\ 0 & 1 & 1 & 1 & 1 \end{bmatrix} \quad \begin{bmatrix} 1 & 1 & 1 & 0 & 1 \\ 0 & 0 & 1 & 1 & 0 \\ 0 & 1 & 0 & 1 & 0 \\ 0 & 1 & 1 & 1 & 1 \end{bmatrix} \quad \begin{bmatrix} 1 & 1 & 1 & 0 & 1 \\ 0 & 1 & 0 & 1 & 0 \\ 0 & 0 & 1 & 1 & 0 \\ 0 & 1 & 1 & 1 & 1 \end{bmatrix} \\
 \begin{bmatrix} 1 & 1 & 1 & 0 & 1 \\ 0 & 1 & 0 & 1 & 0 \\ 0 & 0 & 1 & 1 & 0 \\ 0 & 0 & 1 & 0 & 1 \end{bmatrix} \quad \begin{bmatrix} 1 & 1 & 1 & 0 & 1 \\ 0 & 1 & 0 & 1 & 0 \\ 0 & 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 1 & 1 \end{bmatrix}.$$

The backward substitutions yield the solution: $x_{11} = 1$, $x_{12} = 1$, $x_{21} = 1$, $x_{22} = 1$, i.e., each of the four panel switches should be toggled once (in any order).


c. The solution to this instance of the puzzle is $x_{11} = x_{13} = x_{22} = x_{31} = x_{33} = 1$ (with all the other components being 0).

Exercises 6.3


1. Which of the following binary trees are AVL trees?



(a)



(b)



(c)

2. a. For $n = 1, 2, 3, 4$, and 5 , draw all the binary trees with n nodes that satisfy the balance requirement of AVL trees.
- b. Draw a binary tree of height 4 that can be an AVL tree and has the smallest number of nodes among all such trees.
3. Draw diagrams of the single L -rotation and of the double RL -rotation in their general form.
4. For each of the following lists, construct an AVL tree by inserting their elements successively, starting with the empty tree.
 - a. 1, 2, 3, 4, 5, 6
 - b. 6, 5, 4, 3, 2, 1
 - c. 3, 6, 5, 1, 2, 4
5. a. For an AVL tree containing real numbers, design an algorithm for computing the range (i.e., the difference between the largest and smallest numbers in the tree) and determine its worst-case efficiency.
- b.▷ True or false: The smallest and the largest keys in an AVL tree can always be found on either the last level or the next-to-last level?
6. Write a program for constructing an AVL tree for a given list of n distinct integers.
7. a. Construct a 2-3 tree for the list C, O, M, P, U, T, I, N, G. (Use the alphabetical order of the letters and insert them successively starting with the empty tree.)
- b. Assuming that the probabilities of searching for each of the keys (i.e., the letters) are the same, find the largest number and the average number of key comparisons for successful searches in this tree.

8. Let T_B and T_{2-3} be, respectively, a classical binary search tree and a 2-3 tree constructed for the same list of keys inserted in the corresponding trees in the same order. True or false: Searching for the same key in T_{2-3} always takes fewer or the same number of key comparisons as searching in T_B ?
9. For a 2-3 tree containing real numbers, design an algorithm for computing the range (i.e., the difference between the largest and smallest numbers in the tree) and determine its worst-case efficiency.
10. Write a program for constructing a 2-3 tree for a given list of n integers.

Hints to Exercises 6.3

1. Use the definition of AVL trees. Do not forget that an AVL tree is a special case of a binary search tree.
2. For both questions, it is easier to construct the required trees bottom up, i.e., for smaller values of n first.
3. The single L -rotation and the double RL -rotation are the mirror images of the single R -rotation and the double LR -rotation, whose diagrams can be found in this section.
4. Insert the keys one after another doing appropriate rotations the way it was done in the section's example.
5. a. An efficient algorithm immediately follows from the definition of the binary search tree of which the AVL tree is a special case.


b. The correct answer is opposite to the one that immediately comes to mind.
6. n/a
7. a. Trace the algorithm for the input given (see Figure 6.8) for an example.

b. Keep in mind that the number of key comparisons made in searching for a key in a 2-3 tree depends not only on its node's depth but also whether the key is the first or second one in the node.
8. False; find a simple counterexample.
9. Where will the smallest and largest keys be located?
10. n/a


Solutions to Exercises 6.3

- Only (a) is an AVL tree; (b) has a node (in fact, there are two of them: 4 and 6) that violates the balance requirement; (c) is not a binary search tree because 2 is in the right subtree of 3 (and 7 is in the left subtree of 6).


- a. Here are all the binary trees with n nodes (for $n = 1, 2, 3, 4$, and 5) that satisfy the balance requirement of AVL trees.




b. A minimal AVL tree (i.e., a tree with the smallest number of nodes) of height 4 must have its left and right subtrees being minimal AVL trees of heights 3 and 2. Following the same recursive logic further, we will find, as one of the possible examples, the following tree with 12 nodes built bottom up:




3. a. Here is a diagram of the single L -rotation in its general form:




- b. Here is a diagram of the double RL -rotation in its general form:




4. a. Construct an AVL tree for the list 1, 2, 3, 4, 5, 6.



b. Construct an AVL tree for the list 6, 5, 4, 3, 2, 1.




c. Construct an AVL tree for the list 3, 6, 5, 1, 2, 4.



5. a. The simple and efficient algorithm is based on the fact that the smallest and largest keys in a binary search tree are in the leftmost and rightmost nodes of the tree, respectively. Therefore, the smallest key can be found by starting at the root and following the chain of left pointers until a node with the null left pointer is reached: its key is the smallest one in the tree. Similarly, the largest key can be obtained by following the chain of the right pointers. Finally, the range is computed as the difference between the largest and smallest keys found.


In the worst case, the leftmost and rightmost nodes will be on the last level of the tree. Hence, the worst-case efficiency will be in $\Theta(\log n) + \Theta(\log n) + \Theta(1) = \Theta(\log n)$.

b. False. Here is a counterexample in which neither the smallest nor the largest keys are on the last, or the next-to-last, level of an AVL tree:



6. n/a

7. a. Construct a 2-3 tree for the list C, O, M, P, U, T, I, N, G.



b. The largest number of key comparisons in a successful search will be in the searches for O and U; it will be equal to 4. The average number of key comparisons will be given by the following expression:

$$\frac{1}{9}C(C) + \frac{1}{9}C(O) + \frac{1}{9}C(M) + \frac{1}{9}C(P) + \frac{1}{9}C(U) + \frac{1}{9}C(T) + \frac{1}{9}C(I) + \frac{1}{9}C(N) + \frac{1}{9}C(G) \\ = \frac{1}{9} \cdot 3 + \frac{1}{9} \cdot 4 + \frac{1}{9} \cdot 1 + \frac{1}{9} \cdot 2 + \frac{1}{9} \cdot 4 + \frac{1}{9} \cdot 3 + \frac{1}{9} \cdot 3 + \frac{1}{9} \cdot 3 + \frac{1}{9} \cdot 2 = \frac{25}{9} \approx 2.8.$$

8. False. Consider the list B, A. Searching for B in the binary search tree requires 1 comparison while searching for B in the 2-3 tree requires 2 comparisons.

9. The smallest and largest keys in a 2-3 tree are the first key in the leftmost leaf and the second key in the rightmost leaf, respectively. So searching for them requires following the chain of the leftmost pointers from the root to the leaf and of the rightmost pointers from the root to the leaf. Since the height of a 2-3 tree is always in $\Theta(\log n)$, the time efficiency of the algorithm for all cases is in $\Theta(\log n) + \Theta(\log n) + \Theta(1) = \Theta(\log n)$.
10. n/a

Exercises 6.4

1. a. Construct a heap for the list 1, 8, 6, 5, 3, 7, 4 by the bottom-up algorithm.

b. Construct a heap for the list 1, 8, 6, 5, 3, 7, 4 by successive key insertions (top-down algorithm).

c. Is it always true that the bottom-up and top-down algorithms yield the same heap for the same input?
2. Outline an algorithm for checking whether an array $H[1..n]$ is a heap and determine its time efficiency.
3. a. Find the minimum and the maximum number of keys that a heap of height h can contain.

b.▷ Prove that the height of a heap with n nodes is equal to $\lfloor \log_2 n \rfloor$.
4. ▷ Prove the following equation used in Section 6.4

$$\sum_{i=0}^{h-1} 2(h-i)2^i = 2(n - \log_2(n+1)) \quad \text{where } n = 2^{h+1} - 1.$$

5. a. Design an efficient algorithm for finding and deleting an element of the smallest value in a heap and determine its time efficiency.

b. Design an efficient algorithm for finding and deleting an element of a given value v in a given heap H and determine its time efficiency
6. Sort the following lists by heapsort by using the array representation of heaps.
 - a. 1, 2, 3, 4, 5 (in increasing order)
 - b. 5, 4, 3, 2, 1 (in increasing order)
 - c. S, O, R, T, I, N, G (in alphabetical order)
7. Is heapsort a stable sorting algorithm?
8. What variety of the transform-and-conquer technique does heapsort represent?
9. Which sorting algorithm other than heapsort uses a priority queue?

10. Implement three advanced sorting algorithms—mergesort, quicksort, and heapsort—in the language of your choice and investigate their performance on arrays of sizes $n = 10^2, 10^3, 10^4, 10^5$, and 10^6 . For each of these sizes consider:
 - a. randomly generated files of integers in the range $[1..n]$.
 - b. increasing files of integers $1, 2, \dots, n$.
 - c. decreasing files of integers $n, n - 1, \dots, 1$.
11. *Spaghetti sort* Imagine a handful of uncooked spaghetti, individual rods whose lengths represent numbers that need to be sorted.
 - a. Outline a “spaghetti sort”—a sorting algorithm that takes advantage of this unorthodox representation.
 - b. What does this example of computer science folklore (see [Dew93]) have to do with the topic of this chapter in general and heapsort in particular?

Hints to Exercises 6.4

1. a. Trace the two algorithms outlined in the text on the inputs given.

b. Trace the two algorithms outlined in the text on the inputs given.

c. A mathematical fact may not be established by checking its validity for a few examples.
2. For a heap represented by an array, only the parental dominance requirement needs to be checked.
3. a. What structure does a complete tree of height h with the maximum number of nodes have? What about a complete tree with the minimum number of nodes?

b. Use the results established in part (a).
4. First, express the right-hand side as a function of h . Then prove the obtained equality by either using the formula for the sum $\sum i2^i$ given in Appendix A or by mathematical induction on h .
5. a. Where in a heap should we look for its smallest element?

b. Deleting an arbitrary element of a heap can be done by generalizing the algorithm for deleting its root.
6. Trace the algorithm on the inputs given (see Figure 6.14 for an example).
7. As a rule, sorting algorithms that can exchange far apart elements are not stable.
8. One can claim that the answers are different for the two principal representations of heaps.
9. This algorithm is less efficient than heapsort because it uses the array rather than the heap to implement the priority queue.
10. n/a
11. Pick the spaghetti rods up in a bundle and place them end-down (i.e., vertically) onto a tabletop.

Solutions to Exercises 6.4

1. a. Constructing a heap for the list 1, 8, 6, 5, 3, 7, 4 by the bottom-up algorithm (a root of a subtree being heapified is shown in bold):

$$\begin{array}{ccccccc}
 1 & 8 & \mathbf{6} & 5 & 3 & 7 & 4 \\
 1 & \mathbf{8} & 7 & 5 & 3 & 6 & 4 \\
 1 & 8 & 7 & 5 & 3 & 6 & 4
 \end{array}
 \Rightarrow
 \begin{array}{ccccccc}
 1 & 8 & 7 & 5 & 3 & 6 & 4 \\
 8 & 5 & 7 & 1 & 3 & 6 & 4
 \end{array}$$

- b. Constructing a heap for the list 1, 8, 6, 5, 3, 7, 4 by the top-down algorithm (a new element being inserted into a heap is shown in bold):

$$\begin{array}{ccccccc}
 \mathbf{1} & & & & & & \\
 1 & \mathbf{8} & & & & & \\
 8 & 1 & \mathbf{6} & & & & \\
 8 & 1 & 6 & \mathbf{5} & & & \\
 8 & 5 & 6 & 1 & \mathbf{3} & & \\
 8 & 5 & 6 & 1 & 3 & \mathbf{7} & \\
 8 & 5 & 7 & 1 & 3 & 6 & \mathbf{4}
 \end{array}
 \Rightarrow
 \begin{array}{ccccccc}
 8 & 1 & & & & & \\
 8 & 5 & 6 & 1 & & & \\
 8 & 5 & 7 & 1 & 3 & 6 &
 \end{array}$$

- c. False. Although for the input to questions (a) and (b) the constructed heaps are the same, in general, it may not be the case. For example, for the input 1, 2, 3, the bottom-up algorithm yields 3, 2, 1 while the top-down algorithm yields 3, 1, 2.

2. For $i = 1, 2, \dots, \lfloor n/2 \rfloor$, check whether

$$H[i] \geq \max\{H[2i], H[2i + 1]\}.$$

(Of course, if $2i + 1 > n$, just $H[i] \geq H[2i]$ needs to be satisfied.) If the inequality doesn't hold for some i , stop—the array is not a heap; if it holds for every $i = 1, 2, \dots, \lfloor n/2 \rfloor$, the array is a heap.

Since the algorithm makes up to $2\lfloor n/2 \rfloor$ key comparisons, its time efficiency is in $O(n)$.

3. a. A complete binary tree of height h with the minimum number of nodes has the maximum number of nodes on levels 0 through $h - 1$ and one node on the last level. The total number of nodes in such a tree is

$$n_{\min}(h) = \sum_{i=0}^{h-1} 2^i + 1 = (2^h - 1) + 1 = 2^h.$$

A complete binary tree of height h with the maximum number of nodes has the maximum number of nodes on levels 0 through h . The total

number of nodes in such a tree is

$$n_{\max}(h) = \sum_{i=0}^h 2^i = 2^{h+1} - 1.$$

b. The results established in part (a) imply that for any heap with n nodes and height h

$$2^h \leq n < 2^{h+1}.$$

Taking logarithms to base 2 yields

$$h \leq \log_2 n < h + 1.$$

This means that h is the largest integer not exceeding $\log_2 n$, i.e., $h = \lfloor \log_2 n \rfloor$.

4. We are asked to prove that $\sum_{i=0}^{h-1} 2(h-i)2^i = 2(n - \log_2(n+1))$ where $n = 2^{h+1} - 1$.

For $n = 2^{h+1} - 1$, the right-hand side of the equality in question becomes

$$2(2^{h+1} - 1 - \log_2(2^{h+1} - 1 + 1)) = 2(2^{h+1} - 1 - (h + 1)) = 2(2^{h+1} - h - 2).$$

Using the formula $\sum_{i=1}^{h-1} i2^i = (h-2)2^h + 2$ (see Appendix A), the left-hand side can be simplified as follows:

$$\begin{aligned} \sum_{i=0}^{h-1} 2(h-i)2^i &= 2 \sum_{i=0}^{h-1} (h-i)2^i = 2 \left[\sum_{i=0}^{h-1} h2^i - \sum_{i=0}^{h-1} i2^i \right] \\ &= 2[h(2^h - 1) - (h-2)2^h - 2] \\ &= 2(h2^h - h - h2^h + 2^{h+1} - 2) = 2(2^{h+1} - h - 2). \end{aligned}$$

5. a. The parental dominance requirement implies that we can always find the smallest element of a heap $H[1..n]$ among its leaf positions, i.e., among $H[\lfloor n/2 \rfloor + 1], \dots, H[n]$. (One can easily prove this assertion by contradiction.) Therefore, we can find the smallest element by simply scanning sequentially the second half of the array H . Deleting this element can be done by exchanging the found element with the last element $H[n]$, decreasing the heap's size by one, and then, if necessary, sifting up the former $H[n]$ from its new position until it is not larger than its parent.

The time efficiency of searching for the smallest element in the second half of the array is in $\Theta(n)$; the time efficiency of deleting it after it has

been found is in $\Theta(1) + \Theta(1) + O(\log n) = O(\log n)$.

b. Searching for v by sequential search in $H[1..n]$ takes care of the searching part of the question. Assuming that the first matching element is found in position i , the deletion of $H[i]$ can be done with the following three-part procedure (which is similar to the ones used for deleting the root and the smallest element): First, exchange $H[i]$ with $H[n]$; second, decrease n by 1; third, heapify the structure by sifting the former $H[n]$ either up or down depending on whether it is larger than its new parent or smaller than the larger of its new children, respectively.

The time efficiency of searching for an element of a given value is in $O(n)$; the time efficiency of deleting it after it has been found is in $\Theta(1) + \Theta(1) + O(\log n) = O(\log n)$.

6. a. Sort 1, 2, 3, 4, 5 by heapsort

Heap Construction

1 **2** 3 4 5
 1 5 3 4 2
1 5 3 4 2
 5 4 3 1 2

Maximum Deletions

5 4 3 1 2
 2 4 3 1 | **5**
4 2 3 1
 1 2 3 | **4**
3 2 1
 1 2 | **3**
2 1
 1 | **2**
 1

b. Sort 5, 4, 3, 2, 1 (in increasing order) by heapsort

Heap Construction

5 **4** 3 2 1
5 4 3 2 1

Maximum Deletions

5 4 3 2 1
 1 4 3 2 | **5**
4 2 3 1
 1 2 3 | **4**
3 2 1
 1 2 | **3**
2 1
 1 | **2**
 1

c. Sort S, O, R, T, I, N, G (in alphabetic order) by heapsort

Heap Construction							Maximum Deletions						
1	2	3	4	5	6	7	1	2	3	4	5	6	7
S	O	R	T	I	N	G	T	S	R	O	I	N	G
S	O	R	T	I	N	G	G	S	R	O	I	N	T
S	T	R	O	I	N	G	S	O	R	G	I	N	
S	T	R	O	I	N	G	N	O	R	G	I	S	
T	S	R	O	I	N	G	R	O	N	G	I		
							I	O	N	G	R		
							O	I	N	G			
							G	I	N	O			
							N	I	G				
							G	I	N				
							I	G					
							G	I					
							G						

7. Heapsort is not stable. For example, it sorts 1', 1'' into 1'', 1'.
8. If the heap is thought of as a tree, heapsort should be considered a representation-change algorithm; if the heap is thought of as an array with a special property, heapsort should be considered an instance-simplification algorithm.
9. The answer is selection sort. Note that selection sort is less efficient than heapsort because it uses the array, which is an inferior (to the heap) structure for implementing the priority queue.
10. n/a
11. a. After the bunch of spaghetti rods is put in a vertical position on a tabletop, repeatedly take the tallest rod among the remaining ones out until no more rods are left. This will sort the rods in decreasing order of their lengths.

b. The method shares with heapsort its principal idea: represent the items to be sorted in a way that makes finding and deleting the largest item a simple task. From a more general perspective, the spaghetti sort is an example, albeit a rather exotic one, of a representation-change algorithm.

Exercises 6.5

1. Consider the following brute-force algorithm for evaluating a polynomial.

Algorithm *BruteForcePolynomialEvaluation*($P[0..n], x$)

//The algorithm computes the value of polynomial P at a given point x

//by the “highest to lowest term” brute-force algorithm

//Input: An array $P[0..n]$ of the coefficients of a polynomial of degree n ,

// stored from the lowest to the highest and a number x

//Output: The value of the polynomial at the point x

$p \leftarrow 0.0$

for $i \leftarrow n$ **downto** 0 **do**

$power \leftarrow 1$

for $j \leftarrow 1$ **to** i **do**

$power \leftarrow power * x$

$p \leftarrow p + P[i] * power$

return p

Find the total number of multiplications and the total number of additions made by this algorithm.

2. Write a pseudocode for the brute-force polynomial evaluation that stems from substituting a given value of the variable into the polynomial’s formula and evaluating it from the lowest term to the highest one. Determine the number of multiplications and the number of additions made by this algorithm.
3. a. Estimate how much faster Horner’s rule is compared to the “lowest-to-highest term” brute-force algorithm of Problem 2 if (i) the time of one multiplication is significantly larger than the time of one addition; (ii) the time of one multiplication is about the same as the time of one addition.

b. Is Horner’s rule more time efficient at the expense of being less space efficient than the brute-force algorithm?
4. a. Apply Horner’s rule to evaluate the polynomial

$$p(x) = 3x^4 - x^3 + 2x + 5 \text{ at } x = -2.$$

- b. Use the results of the above application of Horner’s rule to find the quotient and remainder of the division of $p(x)$ by $x + 2$.
5. Compare the number of multiplications and additions/subtractions needed by the “long division” of a polynomial $p(x) = a_n x^n + a_{n-1} x^{n-1} + \dots + a_0$ by $x - c$, where c is some constant, with the number of these operations in the “synthetic division.”

6. a. Apply the left-to-right binary exponentiation algorithm to compute a^{17} .
 b. Is it possible to extend the left-to-right binary exponentiation algorithm to work for every nonnegative integer exponent?
7. Apply the right-to-left binary exponentiation algorithm to compute a^{17} .
8. Design a nonrecursive algorithm for computing a^n that mimics the right-to-left binary exponentiation but does not explicitly use the binary representation of n .
9. Is it a good idea to use a general-purpose polynomial evaluation algorithm such as Horner's rule to evaluate the polynomial $p(x) = x^n + x^{n-1} + \dots + x + 1$?
10. According to the corollary of the Fundamental Theorem of Algebra, every polynomial

$$p(x) = a_n x^n + a_{n-1} x^{n-1} + \dots + a_0$$

can be represented in the form

$$p(x) = a_n (x - x_1)(x - x_2) \dots (x - x_n)$$

where x_1, \dots, x_n are the roots of the polynomial (generally, complex and not necessarily distinct). Discuss which of the two representations is more convenient for each of the following operations:

- a. Polynomial evaluation at a given point
 - b. Addition of two polynomials
 - c. Multiplication of two polynomials
11. ► *Polynomial puzzle* Given that a polynomial $p(x)$ of degree eight is such that $p(i) = 1/i$ for $i = 1, 2, \dots, 9$, find $p(10)$.

Hints to Exercises 6.5

1. Set up a sum and simplify it by using the standard formulas and rules for sum manipulation. Do not forget to include the multiplications outside the inner loop.
2. Take advantage of the fact that the value of x^i can be easily computed from the previously computed x^{i-1} .
3. a. Use the formulas for the number of multiplications (and additions) for both algorithms.

b. Does Horner's rule use any extra memory?
4. Apply Horner's rule to the instance given the same way it is applied to another one in the section.
5. If you implement the algorithm for long division by $x - c$ efficiently, the answer might surprise you.
6. a. Trace the left-to-right binary exponentiation algorithm on the instance given the same way it is done for another instance in the section.

b. The answer is yes: the algorithm can be extended to work for the zero exponent as well. How?
7. Trace the right-to-left binary exponentiation algorithm on the instance given the same way it is done for another instance in the section.
8. Compute and use the binary digits of n "on the fly".
9. Use a formula for the sum of the terms of this special kind of a polynomial.
10. Compare the number of operations needed to implement the task in question.
11. Transform $p(x)$ to a polynomial with roots at $x_i = i$ for $i = 1, 2, \dots, 9$ and represent it in the form mentioned in Problem 10.

Solutions to Exercises 6.5

1. The total number of multiplications made by the algorithm can be computed as follows:

$$\begin{aligned} M(n) &= \sum_{i=0}^n \left(\sum_{j=1}^i 1 + 1 \right) = \sum_{i=0}^n (i + 1) = \sum_{i=0}^n i + \sum_{i=0}^n 1 \\ &= \frac{n(n+1)}{2} + (n+1) = \frac{(n+1)(n+2)}{2} \in \Theta(n^2). \end{aligned}$$

The number of additions is obtained as

$$A(n) = \sum_{i=0}^n 1 = n + 1.$$

2. **Algorithm** *BetterBruteForcePolynomialEvaluation*($P[0..n], x$)
 //Computes the value of polynomial P at a given point x
 //by the “lowest-to-highest term” algorithm
 //Input: Array $P[0..n]$ of the coefficients of a polynomial of degree n ,
 // from the lowest to the highest and a number x
 //Output: The value of the polynomial at the point x
 $p \leftarrow P[0]; \text{ power} \leftarrow 1$
for $i \leftarrow 1$ **to** n **do**
 $\text{power} \leftarrow \text{power} * x$
 $p \leftarrow p + P[i] * \text{power}$
return p

The number of multiplications made by this algorithm is

$$M(n) = \sum_{i=1}^n 2 = 2n.$$

The number of additions is

$$A(n) = \sum_{i=1}^n 1 = n.$$

3. a. If only multiplications need to be taken into account, Horner’s rule will be about twice as fast because it makes just n multiplications vs. $2n$ multiplications required by the other algorithm. If one addition takes about the same amount of time as one multiplication, then Horner’s rule will be about $(2n + n)/(n + n) = 1.5$ times faster.
- b. The answer is no, because Horner’s rule doesn’t use any extra memory.

4. a. Evaluate $p(x) = 3x^4 - x^3 + 2x + 5$ at $x = -2$.

coefficients	3	-1	0	2	5
$x = -2$	3	$(-2) \cdot 3 + (-1) = -7$	$(-2) \cdot (-7) + 0 = 14$	$(-2) \cdot 14 + 2 = -26$	$(-2) \cdot (-26) + 5 = 57$

b. The quotient and the remainder of the division of $3x^4 - x^3 + 2x + 5$ by $x + 2$ are $3x^3 - 7x^2 + 14x - 26$ and 57, respectively.

5. The long division by $x - c$ is done as illustrated below

$$\begin{array}{r}
 a_n x^{n-1} + \dots \\
 x - c \overline{) \begin{array}{l} a_n x^n + a_{n-1} x^{n-1} + \dots + a_1 x + a_0 \\ - a_n x^n - c a_n x^{n-1} \\ \hline (c a_n + a_{n-1}) x^{n-1} + \dots + a_1 x + a_0 \end{array}}
 \end{array}$$

This clearly demonstrates that the first iteration—the one needed to get rid of the leading term $a_n x^n$ —requires one multiplication (to get $c a_n$) and one addition (to add a_{n-1}). After this iteration is repeated $n - 1$ more times, the total number of multiplications and the total number of additions will be n each—exactly the same number of operations needed by Horner’s rule. (In fact, it does exactly the same computations as Horner’s algorithm would do in computing the value of the polynomial at $x = c$.) Thus, the long division, though much more cumbersome than the synthetic division for hand-and-pencil computations, is actually not less time efficient from the algorithmic point of view.

6. a. Compute a^{17} by the left-to-right binary exponentiation algorithm. Here, $n = 17 = 10001_2$. So, we have the following table filled left-to-right:

binary digits of n	1	0	0	0	1
product accumulator	a	a^2	$(a^2)^2 = a^4$	$(a^4)^2 = a^8$	$(a^8)^2 \cdot a = a^{17}$

b. Algorithm *LeftRightBinaryExponentiation* will work correctly for $n = 0$ if the variable *product* is initialized to 1 (instead of a) and the loop starts with I (instead of $I - 1$).

7. Compute a^{17} by the right-to-left binary exponentiation algorithm. Here, $n = 17 = 10001_2$. So, we have the following table filled right-to-left:

1	0	0	0	1	binary digits of n
a^{16}	a^8	a^4	a^2	a	terms a^{2^i}
$a \cdot a^{16} = a^{17}$				a	product accumulator

8. **Algorithm** *ImplicitBinaryExponentiation*(a, n)
 //Computes a^n by the implicit right-to-left binary exponentiation
 //Input: A number a and a nonnegative integer n
 //Output: The value of a^n
 $product \leftarrow 1$; $term \leftarrow a$
while $n \neq 0$ **do**
 $b \leftarrow n \bmod 2$; $n \leftarrow \lfloor n/2 \rfloor$
 if $b = 1$
 $product \leftarrow product * term$
 $term \leftarrow term * term$
return $product$

9. Since the polynomial's terms form a geometric series,

$$p(x) = x^n + x^{n-1} + \dots + x + 1 = \begin{cases} \frac{x^{n+1}-1}{x-1} & \text{if } x \neq 1 \\ n+1 & \text{if } x = 1 \end{cases}$$

Its value can be computed faster than with Horner's rule by computing the right-hand side formula with an efficient exponentiation algorithm for evaluating x^{n+1} .

10. a. With Horner's rule, we can evaluate a polynomial in its coefficient form with n multiplications and n additions. The direct substitution of the x value in the factorized form requires the same number of operations, although these may be operations on complex numbers even for a polynomial with real coefficients.

b. Addition of two polynomials is incomparably simpler for polynomials in their coefficient forms, because, in general, knowing the roots of polynomials $p(x)$ and $q(x)$ helps little in deducing the root values of their sum $p(x) + q(x)$.

c. Multiplication of two polynomials is trivial when they are represented in their factorized form. Indeed, if

$$p(x) = a'_n(x - x'_1) \dots (x - x'_n) \quad \text{and} \quad q(x) = a''_m(x - x''_1) \dots (x - x''_m),$$

then

$$p(x)q(x) = a'_n a''_m (x - x'_1) \dots (x - x'_n) (x - x''_1) \dots (x - x''_m).$$

To multiply two polynomials in their coefficient form, we need to multiply out

$$p(x)q(x) = (a'_n x^n + \dots + a'_0)(a''_m x^m + \dots + a''_0)$$

and collect similar terms to get the product represented in the coefficient form as well.

11. Consider $q(x) = xp(x) - 1$. For every $x_i = i$ ($i = 1, 2, \dots, 9$), we have

$$q(x_i) = x_i p(x_i) - 1 = ip(i) - 1 = i(1/i) - 1 = 0.$$

Using the factorized form mentioned in Problem 10,

$$q(x) = c(x-1)(x-2)\dots(x-9).$$

Since, on the one hand, $q(0) = 0p(0) - 1 = -1$, and, on the other, $q(0) = c(0-1)(0-2)\dots(0-9) = -c9!$, we have the following equation to determine the value of constant c :

$$-c9! = -1.$$

Hence $c = 1/9!$ and $q(10) = (1/9!)(10-1)(10-2)\dots(10-9) = 1$. Since

$$q(10) = 10p(10) - 1 = 1,$$

$$p(10) = 2/10.$$

Note: The problem of finding a polynomial of degree not higher than n from its values at $n+1$ points is one of the fundamental problems in numerical analysis, called *interpolation*. Although there exists exactly one such polynomial, there are several different ways to represent it. For example, you may look up *Lagrange's interpolation formula* and apply it to the polynomial in question.

Exercises 6.6

1. a. Prove the equality

$$\text{lcm}(m, n) = \frac{m \cdot n}{\text{gcd}(m, n)}$$

that underlies the algorithm for computing $\text{lcm}(m, n)$.

- b. Euclid's algorithm is known to be in $O(\log n)$. If it is the algorithm that is used for computing $\text{gcd}(m, n)$, what is the efficiency of the algorithm for computing $\text{lcm}(m, n)$?
2. You are given a list of numbers for which you need to construct a min-heap. (A min-heap is a complete binary tree in which every key is less than or equal to the keys in its children.) How would you use an algorithm for constructing a max-heap (a heap as defined in Section 6.4) to construct a min-heap?
3. Prove that the number of different paths of length $k > 0$ from the i th vertex to the j th vertex in a graph (undirected or directed) equals the (i, j) th element of A^k where A is the adjacency matrix of the graph.
4. a. Design an algorithm with a time efficiency better than cubic for checking whether a graph with n vertices contains a cycle of length 3 [Man89].
b. Consider the following algorithm for the same problem. Starting at an arbitrary vertex, traverse the graph by depth-first search and check whether its depth-first search forest has a vertex with a back edge leading to its grandparent. If it does, the graph contains a triangle; if it does not, the graph does not contain a triangle as its subgraph. Is this algorithm correct?
5. Given $n > 3$ points $P_1 = (x_1, y_1), \dots, P_n = (x_n, y_n)$ in the coordinate plane, design an algorithm to check whether all the points lie within a triangle with its vertices at three of the points given. (You can either design an algorithm from scratch or reduce the problem to another one with a known algorithm.)
6. Consider the problem of finding, for a given positive integer n , the pair of integers whose sum is n and whose product is as large as possible. Design an efficient algorithm for this problem and indicate its efficiency class.
7. The assignment problem introduced in Section 3.4 can be stated as follows. There are n people who need to be assigned to execute n jobs, one person per job. (That is, each person is assigned to exactly one job and each job is assigned to exactly one person.) The cost that would accrue if the i th person is assigned to the j th job is a known quantity $C[i, j]$ for each pair $i, j = 1, \dots, n$. The problem is to assign the people to the jobs to minimize

the total cost of the assignment. Express the assignment problem as a 0–1 linear programming problem.

8. Solve the instance of the linear programming problem given in Section 6.6

$$\text{maximize} \quad 0.10x + 0.07y + 0.03z$$

$$\text{subject to} \quad x + y + z = 100$$

$$x \leq \frac{1}{3}y$$

$$z \geq 0.25(x + y)$$

$$x \geq 0, \quad y \geq 0, \quad z \geq 0.$$

9. The graph-coloring problem is usually stated as the vertex-coloring problem: assign the smallest number of colors to vertices of a given graph so that no two adjacent vertices are the same color. Consider the **edge-coloring** problem: assign the smallest number of colors possible to edges of a given graph so that no two edges with the same endpoint are the same color. Explain how the edge-coloring problem can be reduced to a vertex-coloring problem.
10. Consider the two-dimensional **post-office location problem**: Given n points $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$ in the Cartesian plane, find a location (x, y) for a post office that minimizes $\frac{1}{n} \sum_{i=1}^n (|x_i - x| + |y_i - y|)$, the average Manhattan distance from the post office to these points. Explain how the problem can be efficiently solved by the problem reduction technique, provided the post office does not have to be located at one of the input points.
11. *Jealous husbands* There are $n \geq 2$ married couples who need to cross a river. They have a boat that can hold no more than two people at a time. To complicate matters, all the husbands are jealous and will not agree on any crossing procedure that would put a wife on the same bank of the river with another woman's husband without the wife's husband being there too, even if there are other people on the same bank. Can they cross the river under such constraints?
- a. Solve the problem for $n = 2$.
- b.▷ Solve the problem for $n = 3$, which is the classical version of this problem.
- c.▷ Does the problem have a solution for every $n \geq 4$? If it does, explain how and indicate how many river crossings it will take; if it does not, explain why.

Hints to Exercises 6.6

1. a. Use the rules for computing $\text{lcm}(m, n)$ and $\text{gcd}(m, n)$ from the prime factors of m and n .

b. The answer immediately follows from the formula for computing $\text{lcm}(m, n)$.
2. Use a relationship between minimization and maximization problems.
3. Prove the assertion by induction on k .
4. a. Base your algorithm on the following observation: a graph contains a cycle of length 3 if and only if it has two adjacent vertices i and j that are also connected by a path of length 2.

b. Do not jump to a conclusion in answering this question.
5. An easier solution is to reduce the problem to another one with a known algorithm. Since we did not discuss many geometric algorithms in the book, it should not be difficult to figure out to which one this problem needs to be reduced.
6. Express this problem as a maximization problem of a function in one variable.
7. Introduce double-indexed variables x_{ij} to indicate an assignment of the i th person to the j th job.
8. Take advantage of the specific features of this instance to reduce the problem to one with fewer variables.
9. Create a new graph.
10. Solve first the one-dimensional version of this problem (Problem 2(a) in Exercises 3.3).
11. a, b. Create a state-space graph for the problem as it is done for the river-crossing puzzle in the section.

c. Look at the state obtained after the first six river crossings in the solution to part (b).

Solutions to Exercises 6.6

1. a. Since

$$\begin{aligned} \text{lcm}(m, n) = & \text{the product of the common prime factors of } m \text{ and } n \\ & \cdot \text{the product of the prime factors of } m \text{ that are not in } n \\ & \cdot \text{the product of the prime factors of } n \text{ that are not in } m \end{aligned}$$

and

$$\text{gcd}(m, n) = \text{the product of the common prime factors of } m \text{ and } n,$$

the product of $\text{lcm}(m, n)$ and $\text{gcd}(m, n)$ is equal to

$$\begin{aligned} & \text{the product of the common prime factors of } m \text{ and } n \\ & \cdot \text{the product of the prime factors of } m \text{ that are not in } n \\ & \cdot \text{the product of the prime factors of } n \text{ that are not in } m \\ & \cdot \text{the product of the common prime factors of } m \text{ and } n. \end{aligned}$$

Since the product of the first two terms is equal to m and the product of the last two terms is equal to n , we showed that $\text{lcm}(m, n) \cdot \text{gcd}(m, n) = m \cdot n$, and, hence,

$$\text{lcm}(m, n) = \frac{m \cdot n}{\text{gcd}(m, n)}.$$

b. If $\text{gcd}(m, n)$ is computed in $O(\log n)$ time, $\text{lcm}(m, n)$ will also be computed in $O(\log n)$ time, because one extra multiplication and one extra division take only constant time.

2. Replace every key K_i of a given list by $-K_i$ and apply a max-heap construction algorithm to the new list. Then change the signs of all the keys again.
3. The induction basis: For $k = 1$, $A^1[i, j]$ is equal to 1 or 0 depending on whether there is an edge from vertex i to vertex j . In either case, it is also equal to the number of paths of length 1 from i to j . For the general step, assume that for a positive integer k , $A^k[i, j]$ is equal to the number of different paths of length k from vertex i to vertex j . Since $A^{k+1} = A^k A$, we have the following equality for the (i, j) element of A^{k+1} :


$$A^{k+1}[i, j] = A^k[i, 1]A[1, j] + \dots + A^k[i, t]A[t, j] + \dots + A^k[i, n]A[n, j],$$

where $A^k[i, t]$ is equal to the number of different paths of length k from vertex i to vertex t according to the induction hypothesis and $A[t, j]$ is equal to 1 or 0 depending on whether there is an edge from vertex t to

vertex j for $t = 1, \dots, n$. Further, any path of length $k + 1$ from vertex i to vertex j must be made up of a path of length k from vertex i to some intermediate vertex t and an edge from that t to vertex j . Since for different intermediate vertices t we get different paths, the formula above yields the total number of different paths of length $k + 1$ from i to j .

4. a. For the adjacency matrix A of a given graph, compute A^2 with an algorithm whose time efficiency is better than cubic (e.g., Strassen's matrix multiplication discussed in Section 4.5). Check whether there exists a nonzero element $A[i, j]$ in the adjacency matrix such that $A^2[i, j] > 0$: if there is, the graph contains a triangle subgraph, if there is not, the graph does not contain a triangle subgraph.

b. The algorithm is incorrect because the condition is sufficient but not necessary for a graph to contain a cycle of length 3. Consider, as a counterexample, the DFS tree of the traversal that starts at vertex a of the following graph and resolves ties according to alphabetical order of vertices:



It does not contain a back edge to a grandparent of a vertex, but the graph does have a cycle of length 3: $a - b - e - a$.

5. The problem can be reduced to the question about the convex hull of a given set of points: if the convex hull is a triangle, the answer is yes, otherwise, the answer is no. There are several algorithms for finding the convex hull for a set of points; quickhull, which was discussed in section 4.6, is particularly appropriate for this application.
6. Let x be one of the numbers in question; hence, the other number is $n - x$. The problem can be posed as the problem of maximizing $f(x) = x(n - x)$ on the set of all integer values of x . Since the graph of $f(x) = x(n - x)$ is a parabola with the apex at $x = n/2$, the solution is $n/2$ if n is even and $\lfloor n/2 \rfloor$ (or $\lceil n/2 \rceil$) if n is odd. Hence, the numbers in question can be computed as $\lfloor n/2 \rfloor$ and $n - \lfloor n/2 \rfloor$, which works both for even and odd values of n . Assuming that one division by 2 takes a constant time irrespective of n 's size, the algorithm's time efficiency is clearly in $\Theta(1)$.
7. Let x_{ij} be a 0-1 variable indicating an assignment of the i th person to the j th job (or, in terms of the cost matrix C , a selection of the matrix

element from the i th row and the j th column). The assignment problem can then be posed as the following 0–1 linear programming problem:

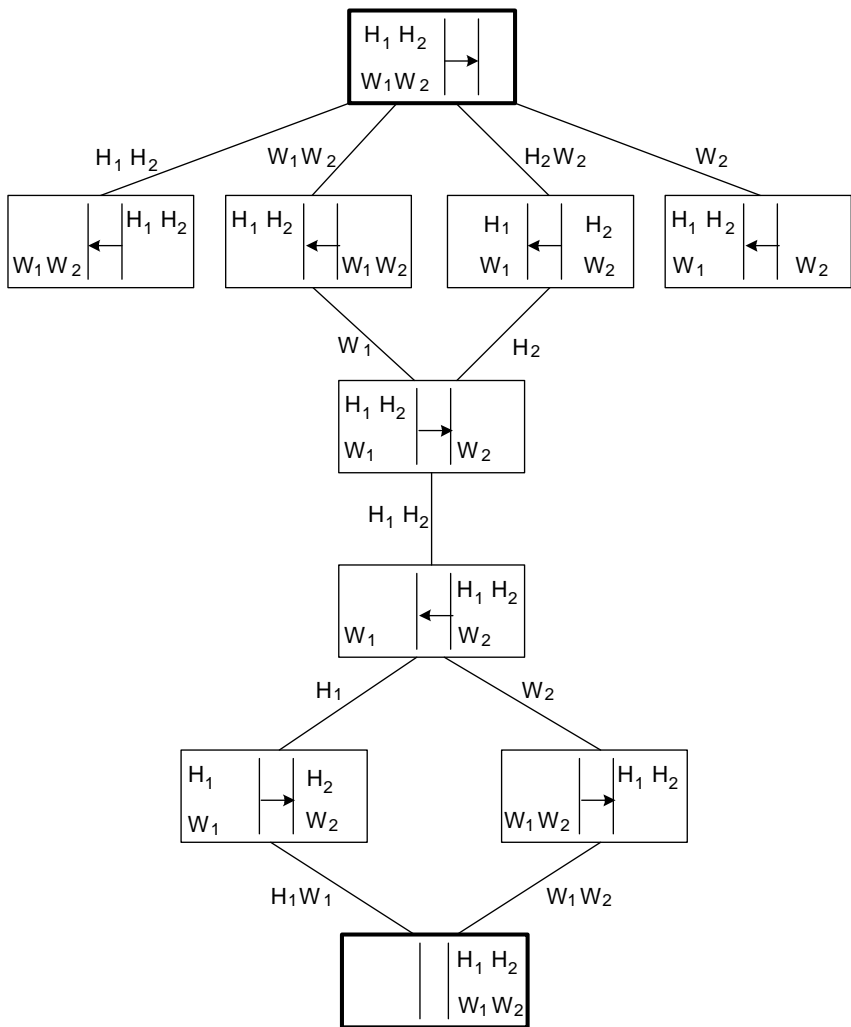
$$\begin{array}{ll} \text{minimize} & \sum_{i=1}^n \sum_{j=1}^n c_{ij} x_{ij} \quad (\text{the total assignment cost}) \\ \text{subject to} & \sum_{j=1}^n x_{ij} = 1 \text{ for } i = 1, \dots, n \text{ (person } i \text{ is assigned to one job)} \\ & \sum_{i=1}^n x_{ij} = 1 \text{ for } j = 1, \dots, n \text{ (job } j \text{ is assigned to one person)} \\ & x_{ij} \in \{0, 1\} \text{ for } i = 1, \dots, n \text{ and } j = 1, \dots, n \end{array}$$

8. We can exploit the specific features of the instance in question to solve it by the following reasoning. Since the expected return from cash is the smallest, the value of cash investment needs to be minimized. Hence, $z = 0.25(x + y)$ in an optimal solution. Substituting $z = 0.25(x + y)$ into $x + y + z = 100$, yields $x + y = 80$ and hence $z = 20$. Similarly, since the expected return on stocks is larger than that of bonds, the amount invested in stocks needs to be maximized. Hence, in an optimal allocation $x = y/3$. Substituting this into $x + y = 80$ yields $y = 60$ and $x = 20$. Thus, the optimal allocation is to put 20 million in stocks, 60 millions in bonds, and 20 million in cash.

Note: This method should not be construed as having a power beyond this particular instance. Generally speaking, we need to use general algorithms such as the simplex method for solving linear programming problems with three or more unknowns. A special technique applicable to instances with only two variables is discussed in Section 10.1 (see also the solution to Problem 10 in Exercises 3.3).

9. Create a new graph whose vertices represent the edges of the given graph and connect two vertices in the new graph by an edge if and only if these vertices represent two edges with a common endpoint in the original graph. A solution of the vertex-coloring problem for the new graph solves the edge-coloring problem for the original graph.
10. The problem is obviously equivalent to minimizing independently $\frac{1}{n} \sum_{i=1}^n |x_i - x|$ and $\frac{1}{n} \sum_{i=1}^n |y_i - y|$. Thus we have two instances of the same problem, whose solution is the median of the numbers defining the instance (see the solution to Problem 2a in Exercises 3.3). Thus, x and y can be found by computing the medians of x_1, x_2, \dots, x_n and y_1, y_2, \dots, y_n , respectively.
11. a. Here is a state-space graph for the two jealous husbands puzzle: H_i, W_i denote the husband and wife of couple i ($i = 1, 2$), respectively; the two bars $||$ denote the river; the arrow indicates the direction of the next trip, which is defined by the boat's location. (For the sake of simplicity, the graph doesn't include crossings that differ by obvious index substitutions such as starting with the first couple $H_1 W_1$ crossing the river instead of

the second one H_2W_2 .) The vertices corresponding to the initial and final states are shown in bold.

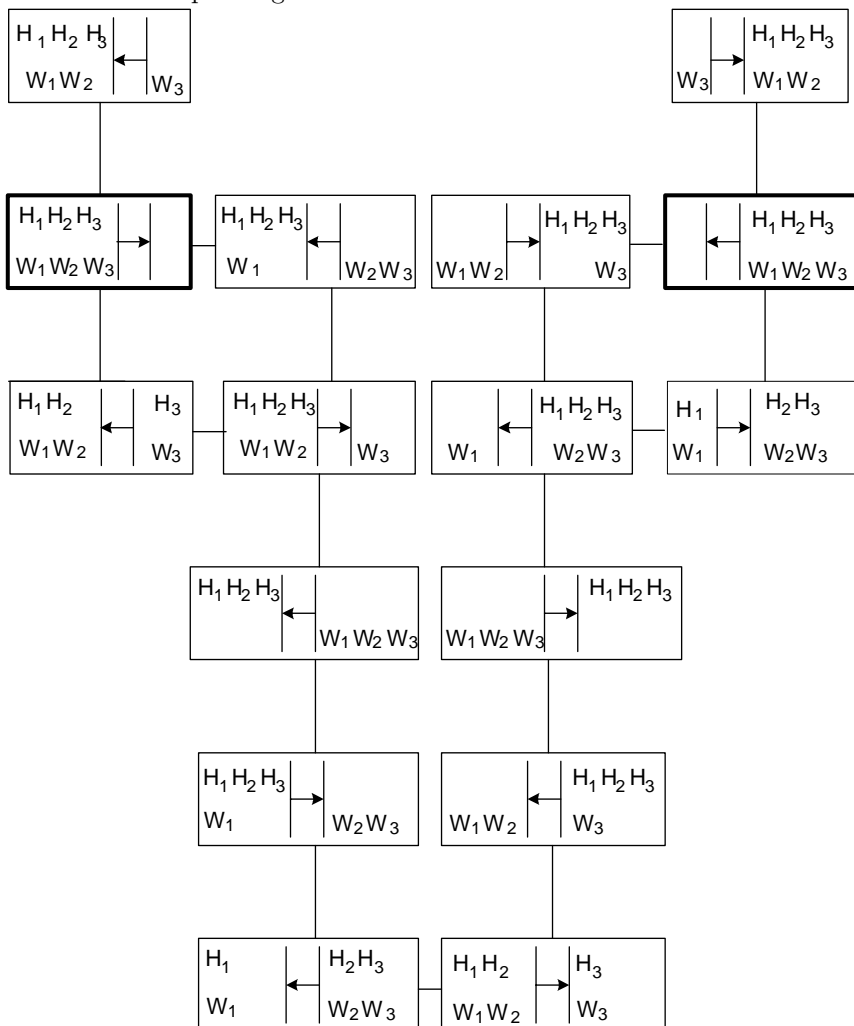


There are four simple paths from the initial-state vertex to the final-state vertex, each five edges long, in this graph. If specified by their edges, they are:

$W_1W_2 \ W_1 \ H_1H_2 \ H_1 \ H_1W_1$
 $W_1W_2 \ W_1 \ H_1H_2 \ W_2 \ W_1W_2$
 $H_2W_2 \ H_2 \ H_1H_2 \ H_1 \ H_1W_1$
 $H_2W_2 \ H_2 \ H_1H_2 \ W_2 \ W_1W_2$

Hence, there are four (to within obvious symmetric substitutions) optimal solutions to this problem, each requiring five river crossings.

b. Here is a state-space graph for the three jealous husbands puzzle: H_i , W_i denote the husband and wife of couple i ($i = 1, 2, 3$), respectively; b stands for the boat; the two bars $|$ denote the river; the arrow indicates the possible direction of the next trip, which is defined by the boat's location. (For the sake of simplicity, the graph doesn't include crossings that differ by obvious index substitutions such as starting with the first or second couple crossing the river instead of the third one H_3W_3 .) The vertices corresponding to the initial and final states are shown in bold.



There are four simple paths from the initial-state vertex to the final-state vertex, each eleven edges long, in this graph. If specified by their edges,

they are:

W_2W_3	W_2	W_1W_2	W_1	H_2H_3	H_2W_2	H_1H_2	W_3	W_2W_3	W_2	W_1W_2
W_2W_3	W_2	W_1W_2	W_1	H_2H_3	H_2W_2	H_1H_2	W_3	W_2W_3	H_1	H_1W_1
H_3W_3	H_3	W_1W_2	W_1	H_2H_3	H_2W_2	H_1H_2	W_3	W_2W_3	W_2	W_1W_2
H_3W_3	H_3	W_1W_2	W_1	H_2H_3	H_2W_2	H_1H_2	W_3	W_2W_3	H_1	H_1W_1

Hence, there are four (to within obvious symmetric substitutions) optimal solutions to this problem, each requiring eleven river crossings.

c. The problem doesn't have a solution for the number of couples $n \geq 4$. If we start with one or more extra (i.e., beyond 3) couples, no new qualitatively different states will result and after the first six river crossings (see the solution to part (b)), we will arrive at the state with $n - 1$ couples and the boat on the original bank and one couple on the other bank. The only allowed transition from that state will be going back to its predecessor by ferrying a married couple to the other side.

This file contains the exercises, hints, and solutions for Chapter 7 of the book "Introduction to the Design and Analysis of Algorithms," 2nd edition, by A. Levitin. The problems that might be challenging for at least some students are marked by \triangleright ; those that might be difficult for a majority of students are marked by \blacktriangleright .

Exercises 7.1

1. Is it possible to exchange numeric values of two variables, say, u and v , without using any extra storage?
2. Will the comparison counting algorithm work correctly for arrays with equal values?
3. Assuming that the set of possible list values is $\{a, b, c, d\}$, sort the following list in alphabetical order by the distribution counting algorithm:

$b, c, d, c, b, a, a, b.$

4. Is the distribution counting algorithm stable?
5. Design a one-line algorithm for sorting any array of size n whose values are n distinct integers from 1 to n .
6. \triangleright The **ancestry problem** asks to determine whether a vertex u is an ancestor of vertex v in a given binary (or, more generally, rooted ordered) tree of n vertices. Design a $O(n)$ input enhancement algorithm that provides sufficient information to solve this problem for any pair of the tree's vertices in constant time.
7. \blacktriangleright The following technique, known as **virtual initialization**, provides a time-efficient way to initialize just some elements of a given array $A[0..n-1]$ so that for each of its elements, we can say in constant time whether it has been initialized and, if it has been, with which value. This is done by utilizing a variable *counter* for the number of initialized elements in A and two auxiliary arrays of the same size, say $B[0..n-1]$ and $C[0..n-1]$, defined as follows. $B[0], \dots, B[\text{counter} - 1]$ contain the indices of the elements of A that were initialized: $B[0]$ contains the index of the element initialized first, $B[1]$ contains the index of the element initialized second, and so on. Furthermore, if $A[i]$ was the k th element ($0 \leq k \leq \text{counter} - 1$) to be initialized, $C[i]$ contains k .

- a. Sketch the state of arrays $A[0..7]$, $B[0..7]$, and $C[0..7]$ after the three assignments

$$A[3] \leftarrow x; \quad A[7] \leftarrow z; \quad A[1] \leftarrow y.$$

- b. In general, how can we check with this scheme whether $A[i]$ has been initialized and, if it has been, with which value?

8. a. Write a program for multiplying two sparse matrices, a p -by- q matrix A and a q -by- r matrix B .

b. Write a program for multiplying two sparse polynomials $p(x)$ and $q(x)$ of degrees m and n , respectively.
9. Write a program that plays the game of tic-tac-toe with the human user by storing all possible positions on the game's 3-by-3 board along with the best move for each of them.

Hints to Exercises 7.1

1. Yes, it is possible. How?
2. Check the algorithm's pseudocode to see what it does upon encountering equal values.
3. Trace the algorithm on the input given (see Figure 7.2 for an example).
4. Check whether the algorithm can reverse a relative ordering of equal elements.
5. Where will $A[i]$ be in the sorted array?
6. Take advantage of the standard traversals of such trees.
7. a. Follow the definitions of the arrays B and C in the description of the method.
b. Find, say, $B[C[3]]$ for the example in part (a).
8. a. Use linked lists to hold nonzero elements of the matrices.
b. Represent each of the given polynomials by a linked list with nodes containing exponent i and coefficient a_i for each nonzero term $a_i x^i$.
9. You may want to take advantage of the board's symmetry to decrease the number of the positions that need to be stored.

Solutions to Exercises 7.1

1. The following operations will exchange values of variables u and v :

$u \leftarrow u + v$ // u holds $u + v$, v holds v

$v \leftarrow u - v$ // u holds $u + v$, v holds u

$u \leftarrow u - v$ // u holds v , v holds u

Note: The same trick is applicable, in fact, to any binary data by employing the “exclusive or” (XOR) operation:

$u \leftarrow u \text{XOR} v$

$v \leftarrow u \text{XOR} v$

$u \leftarrow u \text{XOR} v$

2. Yes, it will work correctly for arrays with equal elements.

3. Input: A: b, c, d, c, b, a, a, b

Frequencies

a	b	c	d
2	3	2	1

 Distribution values

a	b	c	d
2	5	7	8

	$D[a..d]$				$S[0..7]$							
$A[7] = b$	2	5	7	8					b			
$A[6] = a$	2	4	7	8		a						
$A[5] = a$	1	4	7	8	a							
$A[4] = b$	0	4	7	8			b					
$A[3] = c$	0	3	7	8						c		
$A[2] = d$	0	3	6	8							d	
$A[1] = c$	0	3	6	7					c			
$A[0] = b$	0	3	5	7			b					

4. Yes, it is stable because the algorithm scans its input right-to-left and puts equal elements into their section of the sorted array right-to-left as well.

5. **for** $i \leftarrow 0$ **to** $n - 1$ **do** $S[A[i] - 1] \leftarrow A[i]$

6. Vertex u is an ancestor of vertex v in a rooted ordered tree T if and only if the following two inequalities hold

$$preorder(u) \leq preorder(v) \text{ and } postorder(u) \geq postorder(v),$$

where $preorder$ and $postorder$ are the numbers assigned to the vertices by the preorder and postorder traversals of T , respectively. Indeed, preorder traversal visits recursively the root and then the subtrees numbered from left to right. Therefore,

$$preorder(u) \leq preorder(v)$$

if and only if either u is an ancestor of v (i.e., u is on the simple path from the root's tree to v) or u is to the left of v (i.e., u and v are not on the same simple path from the root to a leaf and $T(u)$ is to the left of $T(v)$ where $T(u)$ and $T(v)$ are the subtrees of the nearest common ancestor of u and v , respectively). Similarly, postorder traversal visits recursively the subtrees numbered from left to right and then the root. Therefore,

$$postorder(u) \geq postorder(v)$$


if and only if either u is an ancestor of v or v is to the left of u . Hence,

$$preorder(u) \leq preorder(v) \text{ and } postorder(u) \geq postorder(v)$$

is necessary and sufficient for u to be an ancestor of v .

The time efficiencies of both traversals are in $O(n)$ (Section 4.4); once the preorder and postorder numbers are precomputed, checking the two inequalities takes constant time for any given pair of the vertices.

7. a. The following diagram depicts the results of these assignments (the values of the unspecified elements in the arrays are undefined):



- b. $A[i]$ is initialized if and only if $0 \leq C[i] \leq counter - 1$ and $B[C[i]] = i$. (It is useful to note that the elements of array C define the inverse to the mapping defined by the elements of array B .) Hence, if these two conditions hold, $A[i]$ contains the value it has been initialized with; otherwise, it has not been initialized.

8. n/a

9. n/a

Exercises 7.2

1. Apply Horspool's algorithm to search for the pattern BAOBAB in the text

BESS_KNEW_ABOUT_BAOBABS

2. Consider the problem of searching for genes in DNA sequences using Horspool's algorithm. A DNA sequence consists of a text on the alphabet $\{A, C, G, T\}$ and the gene or gene segment is the pattern.

- a. Construct the shift table for the following gene segment of your chromosome 10:

TCCTATTCTT

- b. Apply Horspool's algorithm to locate the above pattern in the following DNA sequence:

TTATAGATCTCGTATTCTTTTATAGATCTCCTATTCTT

3. How many character comparisons will be made by Horspool's algorithm in searching for each of the following patterns in the binary text of 1000 zeros?
 - a. 00001
 - b. 10000
 - c. 01010
4. For searching in a text of length n for a pattern of length m ($n \geq m$) with Horspool's algorithm, give an example of
 - a. worst-case input.
 - b. best-case input.
5. Is it possible for Horspool's algorithm to make more character comparisons than the brute-force algorithm would make in searching for the same pattern in the same text?
6. If Horspool's algorithm discovers a matching substring, how large a shift should it make to search for a next possible match?
7. How many character comparisons will the Boyer-Moore algorithm make in searching for each of the following patterns in the binary text of 1000 zeros?

- a. 00001
 - b. 10000
 - c. 01010
8. a. Would the Boyer-Moore algorithm work correctly with just the bad-symbol table to guide pattern shifts?
- b. Would the Boyer-Moore algorithm work correctly with just the good-suffix table to guide pattern shifts?
9. a. If the last characters of a pattern and its counterpart in the text do match, does Horspool's algorithm have to check other characters right to left, or can it check them left to right too?
- b. Answer the same question for the Boyer-Moore algorithm.
10. Implement Horspool's algorithm, the Boyer-Moore algorithm, and the brute-force algorithm of Section 3.2 in the language of your choice and run an experiment to compare their efficiencies for matching
- a. random binary patterns in random binary texts.
 - b. random natural language patterns in natural language texts.

Hints to Exercises 7.2

1. Trace the algorithm in the same way it is done in the section for another instance of the string-matching problem.
2. A special alphabet notwithstanding, this application is not different than applications to natural language strings.
3. For each pattern, fill in its shift table and then determine the number of character comparisons (both successful and unsuccessful) on each trial and the total number of trials.
4. Find an example of a binary string of length m and a binary string of length n ($n \geq m$) so that Horspool's algorithm makes
 - a. the largest possible number of character comparisons before making the smallest possible shift.
 - b. the smallest possible number of character comparisons.
5. It is logical to try a worst-case input for Horspool's algorithm.
6. Can the algorithm shift the pattern by more than one position without the possibility of missing another matching substring?
7. For each pattern, fill in the two shift tables and then determine the number of character comparisons (both successful and unsuccessful) on each trial and the total number of trials.
8. Check the description of the Boyer-Moore algorithm.
9. Check the descriptions of the algorithms.
10. n/a

Solutions to Exercises 7.2

1. The shift table for the pattern **BAOBAB** in a text comprised of English letters, the period, and a space will be

<i>c</i>	A	B	C	D	.	.	.	0	.	.	.	Z	.	_
<i>t(c)</i>	1	2	6	6	6			3	6			6	6	6

The actual search will proceed as shown below:

```

    B E S S _ K N E W _ A B O U T _ B A O B A B S
    B A O B A B                               B A O B A B
                B A O B A B                     B A O B A B
                    B A O B A B

```

2. a. For the pattern **TCCTATTCTT** and the alphabet {A, C, G, T}, the shift table looks as follows:

<i>c</i>	A	C	G	T
<i>t(c)</i>	5	2	10	1

- b. Below the text and the pattern, we list the characters of the text that are aligned with the last T of the pattern, along with the corresponding number of character comparisons (both successful and unsuccessful) and the shift size:

the text: **TTATAGATCTCGTATTCTTTTATAGATCTCCTATTCTT**
the pattern: **TCCTATTCTT**

T: 2 comparisons, shift 1
C: 1 comparison, shift 2
T: 2 comparisons, shift 1
A: 1 comparison, shift 5
T: 8 comparisons, shift 1
T: 3 comparisons, shift 1
T: 3 comparisons, shift 1
A: 1 comparison, shift 5
T: 2 comparisons, shift 1
C: 1 comparison, shift 2
C: 1 comparison, shift 2
T: 2 comparisons, shift 1
A: 1 comparison, shift 5
T: 10 comparisons to stop the successful search

3. a. For the pattern **00001**, the shift table is

<i>c</i>	0	1
<i>t(c)</i>	1	5

The algorithm will make one unsuccessful comparison and then shift the pattern one position to the right on each of its trials:

0 0 0 0 0 0	0 0 0 0 0
0 0 0 0 1	
0 0 0 0 1	
etc.	

0 0 0 0 1

The total number of character comparisons will be $C = 1 \cdot 996 = 996$.

b. For the pattern 10000, the shift table is

c	0	1
$t(c)$	1	4

The algorithm will make four successful and one unsuccessful comparison and then shift the pattern one position to the right on each of its trials:

0 0 0 0 0 0	0 0 0 0 0
1 0 0 0 0	
1 0 0 0 0	
etc.	

1 0 0 0 0

The total number of character comparisons will be $C = 5 \cdot 996 = 4980$.

c. For the pattern 01010, the shift table is

c	0	1
$t(c)$	2	1

The algorithm will make one successful and one unsuccessful comparison and then shift the pattern two positions to the right on each of its trials:

0 0 0 0 0 0	0 0 0 0 0 0
0 1 0 1 0	
0 1 0 1 0	
etc.	

0 1 0 1 0

The left end of the pattern in the trials will be aligned against the text's characters in positions 0, 2, 4, ..., 994, which is 498 trials. (We can also get this number by looking at the positions of the right end of the pattern. This leads to finding the largest integer k such that $4 + 2(k - 1) \leq 999$, which is $k = 498$.) Thus, the total number of character comparisons will be $C = 2 \cdot 498 = 996$.

4. a. The worst case: e.g., searching for the pattern $\underbrace{10\dots0}_{m-1}$ in the text of n 0's. $C_w = m(n - m + 1)$.
- b. The best case: e.g., searching for the pattern $\underbrace{0\dots0}_m$ in the text of n 0's. $C_b = m$.
5. Yes: e.g., for the pattern $\underbrace{10\dots0}_{m-1}$ and the text $\underbrace{0\dots0}_n$, $C_{bf} = n - m + 1$ while $C_{Horspool} = m(n - m + 1)$.
6. We can shift the pattern exactly in the same manner as we would in the case of a mismatch, i.e., by the entry $t(c)$ in the shift table for the text's character c aligned against the last character of the pattern.
7. a. For the pattern 00001, the shift tables will be filled as follows:

the bad-symbol table

c	0	1
$t_1(c)$	1	5

the good-suffix table

k	the pattern	d_2
1	0000 1	5
2	0000 1	5
3	0000 1	5
4	0000 1	5

On each of its trials, the algorithm will make one unsuccessful comparison and then shift the pattern by $d_1 = \max\{t_1(0) - 0, 1\} = 1$ position to the right without consulting the good-suffix table:

0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 1 etc.	0 0 0 0 0 0 0 0 0 1
---	------------------------------------

The total number of character comparisons will be $C = 1 \cdot 996 = 996$.

- b. For the pattern 10000, the shift tables will be filled as follows:

the bad-symbol table

c	0	1
$t_1(c)$	1	4

the good-suffix table

k	the pattern	d_2
1	1000 0	3
2	1000 0	2
3	1000 0	1
4	1000 0	5

On each of its trials, the algorithm will make four successful and one

unsuccessful comparison and then shift the pattern by the maximum of $d_1 = \max\{t_1(0) - 4, 1\} = 1$ and $d_2 = t_2(4) = 5$, i.e., by 5 characters to the right:

0 0 0 0 0 0 0 0 0 0	0 0 0 0 0
1 0 0 0 0	
1 0 0 0 0	
etc.	
	1 0 0 0 0

The total number of character comparisons will be $C = 5 \cdot 200 = 1000$.

c. For the pattern 01010, the shift tables will be filled as follows:

the bad-symbol table

c	0	1
$t_1(c)$	2	1

the good-suffix table

k	the pattern	d_2
1	01010	4
2	01010	4
3	01010	2
4	01010	2

On each trial, the algorithm will make one successful and one unsuccessful comparison. The shift's size will be computed as the maximum of $d_1 = \max\{t_1(0) - 1, 1\} = 1$ and $d_2 = t_2(1) = 4$, which is 4. If we count character positions starting with 0, the right end of the pattern in the trials will be aligned against the text's characters in positions 4, 8, 12, ..., with the last term in this arithmetic progression less than or equal to 999. This leads to finding the largest integer k such that $4 + 4(k - 1) \leq 999$, which is $k = 249$.

0 0 0 0 0 0	0 0 0 0 0 0
0 1 0 1 0	
0 1 0 1 0	
etc.	
	0 1 0 1 0

Thus, the total number of character comparisons will be $C = 2 \cdot 249 = 498$.

8. a. Yes, the Boyer-Moore algorithm can get by with just the bad-symbol shift table.
- b. No: The bad-symbol table is necessary because it's the only one used by the algorithm if the first pair of characters does not match.
9. a. Horspool's algorithm can also compare the remaining $m - 1$ characters of the pattern from left to right because it shifts the pattern based only on the text's character aligned with the last character of the pattern.

b. The Boyer-Moore algorithm must compare the remaining $m - 1$ characters of the pattern from right to left because of the good-suffix shift table.

10. n/a

Exercises 7.3

1. For the input 30, 20, 56, 75, 31, 19 and hash function $h(K) = K \bmod 11$
 - a. construct the open hash table.
 - b. find the largest number of key comparisons in a successful search in this table.
 - c. find the average number of key comparisons in a successful search in this table.
2. For the input 30, 20, 56, 75, 31, 19 and hash function $h(K) = K \bmod 11$
 - a. construct the closed hash table.
 - b. find the largest number of key comparisons in a successful search in this table.
 - c. find the average number of key comparisons in a successful search in this table.
3. Why is it not a good idea for a hash function to depend on just one letter (say, the first one) of a natural language word?
4. Find the probability of all n keys being hashed to the same cell of a hash table of size m if the hash function distributes keys evenly among all the cells of the table.
5. ►The **birthday paradox** asks how many people should be in a room so that the chances are better than even that two of them will have the same birthday (month and day). Find the quite unexpected answer to this problem. What implication for hashing does this result have?
6. Answer the following questions for the separate-chaining version of hashing.
 - a. Where would you insert keys if you knew that all the keys in the dictionary are distinct? Which dictionary operations, if any, would benefit from this modification?
 - b. We could keep keys of the same linked list sorted. Which of the dictionary operations would benefit from this modification? How could we take advantage of this if all the keys stored in the entire table need to be sorted?
7. Explain how hashing can be applied to check whether all elements of a list are distinct. What is the time efficiency of this application?

8. Fill in the following table with the average-case efficiency classes for the five implementations of the ADT dictionary:

	unordered array	ordered array	binary search tree	separate chaining	linear probing
search					
insertion					
deletion					

9. We have discussed hashing in the context of techniques based on space-time tradeoffs. But it also takes advantage of another general strategy. Which one?
10. Write a computer program that uses hashing for the following problem. Given a natural language text, generate a list of distinct words with the number of occurrences of each word in the text. Insert appropriate counters in the program to compare the empirical efficiency of hashing with the corresponding theoretical results.

Hints to Exercises 7.3

1. Apply the open hashing (separate chaining) scheme to the input given as it is done in the chapter's text for another input (see Figure 7.5). Then compute the largest number and average number of comparisons for successful searches in the constructed table.
2. Apply the closed hashing (open addressing) scheme to the input given as it is done in the chapter's text for another input (see Figure 7.6). Then compute the largest number and average number of comparisons for successful searches in the constructed table.
3. How many different addresses can such a hash function produce? Would it distribute keys evenly?
4. The question is quite similar to computing the probability of having the same result in n throws of a fair die.
5. Find the probability that n people have different birthdays. As to the hashing connection, what hashing phenomenon deals with coincidences?
6. a. There is no need to insert a new key at the end of the linked list it is hashed to.

b. Which operations are faster in a sorted linked list and why? For sorting, do we have to copy all elements in the nonempty lists in an array and then apply a general purpose sorting algorithm or is there a way to take advantage of the sorted order in each of the nonempty linked lists?
7. After you answer these questions, compare the efficiency of this algorithm with that of the brute-force algorithm (Section 2.3) and of the presorting-based algorithm (Section 6.1).
8. Consider this question as a mini-review: the answers are in Section 7.3 for the last two columns and in the appropriate sections of the book for the others. (Of course, you should use the best algorithms available.)
9. If you need to refresh your memory, check the book's table of contents.
10. n/a

Solutions to Exercises 7.3

1. a.

The list of keys: 30, 20, 56, 75, 31, 19

The hash function: $h(K) = K \bmod 11$

The hash addresses:

K	30	20	56	75	31	19
$h(K)$	8	9	1	9	9	8

The open hash table:

0	1	2	3	4	5	6	7	8	9	10
	↓							↓	↓	
	56							30	20	
								↓	↓	
								19	75	
									↓	
									31	

b. The largest number of key comparisons in a successful search in this table is 3 (in searching for $K = 31$).

c. The average number of key comparisons in a successful search in this table, assuming that a search for each of the six keys is equally likely, is

$$\frac{1}{6} \cdot 1 + \frac{1}{6} \cdot 1 + \frac{1}{6} \cdot 1 + \frac{1}{6} \cdot 2 + \frac{1}{6} \cdot 3 + \frac{1}{6} \cdot 2 = \frac{10}{6} \approx 1.7.$$

2. a.

The list of keys: 30, 20, 56, 75, 31, 19

The hash function: $h(K) = K \bmod 11$

The hash addresses:

K	30	20	56	75	31	19
$h(K)$	8	9	1	9	9	8

0	1	2	3	4	5	6	7	8	9	10
								30		
								30	20	
	56							30	20	
	56							30	20	75
31	56							30	20	75
31	56	19						30	20	75

b. The largest number of key comparisons in a successful search is 6 (when searching for $K = 19$).

c. The average number of key comparisons in a successful search in this table, assuming that a search for each of the six keys is equally likely, is

$$\frac{1}{6} \cdot 1 + \frac{1}{6} \cdot 1 + \frac{1}{6} \cdot 1 + \frac{1}{6} \cdot 2 + \frac{1}{6} \cdot 3 + \frac{1}{6} \cdot 6 = \frac{14}{6} \approx 2.3.$$

3. The number of different values of such a function would be obviously limited by the size of the alphabet. Besides, it is usually not the case that the probability of a word to start with a particular letter is the same for all the letters.
4. The probability of all n keys to be hashed to a particular address is equal to $(\frac{1}{m})^n$. Since there are m different addresses, the answer is $(\frac{1}{m})^n m = \frac{1}{m^{n-1}}$.
5. The probability of n people having different birthdays is $\frac{364}{365} \frac{363}{365} \dots \frac{365-(n-1)}{365}$. The smallest value of n for which this expression becomes less than 0.5 is 23. Sedgewick and Flajolet [SF96] give the following analytical solution to the problem:

$$(1 - \frac{1}{M})(1 - \frac{2}{M}) \dots (1 - \frac{n-1}{M}) \approx \frac{1}{2} \text{ where } M = 365.$$

Taking the natural logarithms of both sides yields

$$\ln(1 - \frac{1}{M})(1 - \frac{2}{M}) \dots (1 - \frac{n-1}{M}) \approx -\ln 2 \text{ or } \sum_{k=1}^{n-1} \ln(1 - \frac{k}{M}) \approx -\ln 2.$$

Using $\ln(1 - x) \approx -x$, we obtain

$$\sum_{k=1}^{n-1} \frac{k}{M} \approx \ln 2 \text{ or } \frac{(n-1)n}{2M} \approx \ln 2. \text{ Hence, } n \approx \sqrt{2M \ln 2} \approx 22.5.$$

The implication for hashing is that we should expect collisions even if the size of a hash table is much larger (by more than a factor of 10) than the number of keys.

6. a. If all the keys are known to be distinct, a new key can always be inserted at the beginning of its linked list; this will make the insertion operation $\Theta(1)$. This will not change the efficiencies of search and deletion, however.
- b. Searching in a sorted list can be stopped as soon as a key larger than the search key is encountered. Both deletion (that must follow a search) and insertion will benefit for the same reason. To sort a dictionary stored in linked lists of a hash table, we can merge the k nonempty lists to get the entire dictionary sorted. (This operation is called the *k-way merge*.) To do this efficiently, it's convenient to arrange the current first elements of the lists in a min-heap.

7. Insert successive elements of the list in a hash table until a matching element is encountered or the list is exhausted. The worst-case efficiency will be in $\Theta(n^2)$: all n distinct keys are hashed to the same address so that the number of key comparisons will be $\sum_{i=1}^n (i-1) \in \Theta(n^2)$. The average-case efficiency, with keys distributed about evenly so that searching for each of them takes $\Theta(1)$ time, will be in $\Theta(n)$.

8.

	unordered array	ordered array	binary search tree	separate chaining	linear probing
search	$\Theta(n)$	$\Theta(\log n)$	$\Theta(\log n)$	$\Theta(1)$	$\Theta(1)$
insertion	$\Theta(1)$	$\Theta(n)$	$\Theta(\log n)$	$\Theta(1)$	$\Theta(1)$
deletion	$\Theta(1)$	$\Theta(n)$	$\Theta(\log n)$	$\Theta(1)$	$\Theta(1)$

9. Representation change—one of the three varieties of transform-and-conquer.

10. n/a

Exercises 7.4

1. Give examples of using an index in real-life applications that do not involve computers.
2. a. Prove the equality

$$1 + \sum_{i=1}^{h-1} 2 \lceil m/2 \rceil^{i-1} (\lceil m/2 \rceil - 1) + 2 \lceil m/2 \rceil^{h-1} = 4 \lceil m/2 \rceil^{h-1} - 1$$

that was used in the derivation of upper bound (7.7) for the height of a B-tree.

- b. Complete the derivation of inequality (7.7).
3. Find the minimum order of the B-tree that guarantees that the number of disk accesses in searching in a file of 100 million records does not exceed 3. Assume that the root's page is stored in main memory.
4. Draw the B-tree obtained after inserting 30 and then 31 in the B-tree in Figure 7.8. Assume that a leaf cannot contain more than three items.
5. Outline an algorithm for finding the largest key in a B-tree.
6. a. A **top-down 2-3-4 tree** is a B-tree of order 4 with the following modification of the *insert* operation. Whenever a search for a leaf for a new key encounters a full node (i.e., a node with three keys), the node is split into two nodes by sending its middle key to the node's parent (or, if the full node happens to be the root, the new root for the middle key is created). Construct a top-down 2-3-4 tree by inserting the following list of keys in the initially empty tree:

10, 6, 15, 31, 20, 27, 50, 44, 18.

- b. What is the principal advantage of this insertion procedure compared with the one described for 2-3 trees in Section 6.3? What is its disadvantage?
7. a. ▷ Write a program implementing a key insertion algorithm in a B-tree.
b. ► Write a program for visualization of a key insertion algorithm in a B-tree.

Hints to Exercises 7.4

1. Thinking about searching for information should lead to a variety of examples.
2. a. Use the standard rules of sum manipulation and, in particular, the geometric series formula.

b. You will need to take the logarithms base $\lceil m/2 \rceil$ in your derivation.
3. Find this value from the inequality in the text that provides the upper-bound of the B-tree's height.
4. Follow the insertion algorithm outlined in this section.
5. The algorithm is suggested by the definition of the B-tree.
6. a. Just follow the description of the algorithm given in the statement of the problem. Note that a new key is always inserted in a leaf and that full nodes are always split on the way down, even though the leaf for the new key may have a room for it.

b. Can a split of a full node cause a cascade of splits through the chain of its ancestors? Can we get a taller search tree than necessary?
7. n/a

Solutions to Exercises 7.4

- Here are a few common examples of using an index: labeling drawers of a file cabinet with, say, a range of letters; an index of a book's terms indicating the page or pages on which the term is defined or mentioned; marking a range of pages in an address book, a dictionary; or an encyclopedia; marking a page of a telephone book or a dictionary with the first and last entry on the page; indexing areas of a geographic map by dividing the map into square regions.

2. a.

$$\begin{aligned}
 & 1 + \sum_{i=1}^{h-1} 2^{\lceil m/2 \rceil^{i-1}} (\lceil m/2 \rceil - 1) + 2^{\lceil m/2 \rceil^{h-1}} \\
 = & 1 + 2(\lceil m/2 \rceil - 1) \sum_{i=1}^{h-1} \lceil m/2 \rceil^{i-1} + 2^{\lceil m/2 \rceil^{h-1}} \\
 = & 1 + 2(\lceil m/2 \rceil - 1) \sum_{j=0}^{h-2} \lceil m/2 \rceil^j + 2^{\lceil m/2 \rceil^{h-1}} \\
 = & 1 + 2(\lceil m/2 \rceil - 1) \frac{\lceil m/2 \rceil^{h-1} - 1}{(\lceil m/2 \rceil - 1)} + 2^{\lceil m/2 \rceil^{h-1}} \\
 = & 1 + 2^{\lceil m/2 \rceil^{h-1}} - 2 + 2^{\lceil m/2 \rceil^{h-1}} \\
 = & 4^{\lceil m/2 \rceil^{h-1}} - 1.
 \end{aligned}$$

b. The inequality

$$n \geq 4^{\lceil m/2 \rceil^{h-1}} - 1$$

is equivalent to

$$\frac{n+1}{4} \geq \lceil m/2 \rceil^{h-1}.$$

Taking the logarithms base $\lceil m/2 \rceil$ of both hand sides yields

$$\log_{\lceil m/2 \rceil} \frac{n+1}{4} \geq \log_{\lceil m/2 \rceil} \lceil m/2 \rceil^{h-1}$$

or

$$\log_{\lceil m/2 \rceil} \frac{n+1}{4} \geq h-1.$$

Hence,

$$h \leq \log_{\lceil m/2 \rceil} \frac{n+1}{4} + 1$$

or, since h is an integer,


$$h \leq \lfloor \log_{\lceil m/2 \rceil} \frac{n+1}{4} \rfloor + 1.$$

3. If the tree's root is stored in main memory, the number of disk accesses will be equal to the number of the levels minus 1, which is exactly the height of the tree. So, we need to find the smallest value of the order m so that the height of the B-tree with $n = 10^8$ keys does not exceed 3. Using the upper bound of the B-tree's height, we obtain the following inequality


$$\lfloor \log_{\lceil m/2 \rceil} \frac{n+1}{4} \rfloor + 1 \leq 3 \quad \text{or} \quad \lfloor \log_{\lceil m/2 \rceil} \frac{n+1}{4} \rfloor \leq 2.$$

By “trial and error,” we can find that the smallest value of m that satisfies this inequality is 585.

4. Since there is enough room for 30 in the leaf for it, the resulting B-tree will look as follows

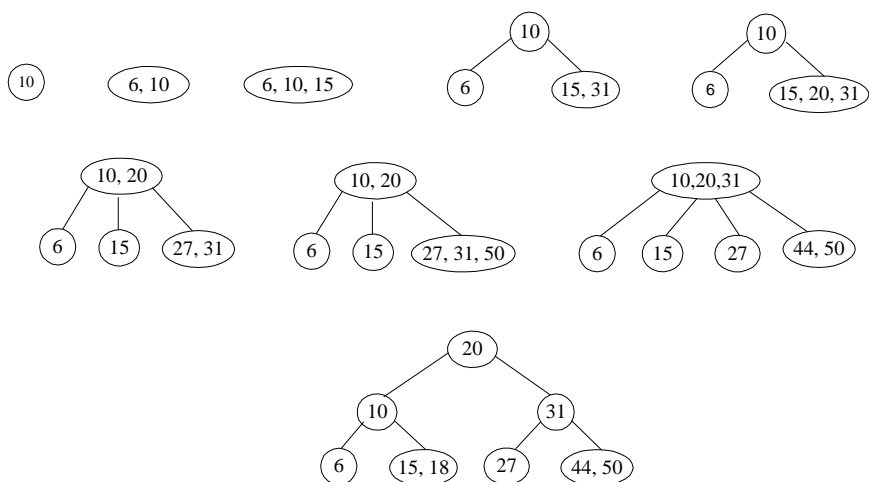


Inserting 31 will require the leaf's split and then its parent's split:



5. Starting at the root, follow the chain of the rightmost pointers to the (rightmost) leaf. The largest key is the last key in that leaf.
6. a. Constructing a top-down 2-3-4 tree by inserting the following list of keys in the initially empty tree:

10, 6, 15, 31, 20, 27, 50, 44, 18.



b. The principal advantage of splitting full nodes (4-nodes with 3 keys) on a way down during insertion of a new key lies in the fact that if the appropriate leaf turns out to be full, its split will never cause a chain reaction of splits because the leaf's parent will always have a room for an extra key. (If the parent is full before the insertion, it is split before the leaf is reached.) This is not the case for the insertion algorithm employed for 2-3 trees (see Section 6.3).

The disadvantage of splitting full nodes on the way down lies in the fact that it can lead to a taller tree than necessary. For the list of part (a), for example, the tree before the last one had a room for key 18 in the leaf containing key 15 and therefore didn't require a split executed by the top-down insertion.

7. n/a

This file contains the exercises, hints, and solutions for Chapter 8 of the book "Introduction to the Design and Analysis of Algorithms," 2nd edition, by A. Levitin. The problems that might be challenging for at least some students are marked by ▷; those that might be difficult for a majority of students are marked by ►.

Exercises 8.1

1. a. What does dynamic programming have in common with divide-and-conquer?
- b. What is a principal difference between the two techniques?
2. a. Compute $C(6, 3)$ by applying the dynamic programming algorithm.
- b. Is it also possible to compute $C(n, k)$ by filling the algorithm's dynamic programming table column by column rather than row by row?
3. Prove the following assertion made in the text while investigating the time efficiency of the dynamic programming algorithm for computing $C(n, k)$:

$$\frac{(k-1)k}{2} + k(n-k) \in \Theta(nk).$$

4. a. What is the space efficiency of *Binomial*, the dynamic programming algorithm for computing $C(n, k)$?
- b. Explain how the space efficiency of this algorithm can be improved. (Try to make as much of an improvement as you can.)
5. a. Find the order of growth of the following functions:
 - i. $C(n, 1)$ ii. $C(n, 2)$ iii. $C(n, n/2)$ for even n 's.

b. What major implication for computing $C(n, k)$ do the answers to the questions in part (a) have?

6. Find the exact number of additions made by the following recursive algorithm based directly on formulas (8.3) and (8.4).

Algorithm *BinomCoeff*(n, k)

if $k = 0$ **or** $k = n$ **return** 1

else return *BinomCoeff*($n - 1, k - 1$) + *BinomCoeff*($n - 1, k$)

7. Which of the following algorithms for computing a binomial coefficient is most efficient?

- a. Use the formula

$$C(n, k) = \frac{n!}{k!(n-k)!}.$$

- b. Use the formula

$$C(n, k) = \frac{n(n-1)\dots(n-k+1)}{k!}.$$

- c. Apply recursively the formula

$$\begin{aligned} C(n, k) &= C(n-1, k-1) + C(n-1, k) \text{ for } n > k > 0, \\ C(n, 0) &= C(n, n) = 1. \end{aligned}$$

- d. Apply the dynamic programming algorithm.

8. Prove that

$$C(n, k) = C(n, n-k) \text{ for } n \geq k \geq 0$$

and explain how this formula can be utilized in computing $C(n, k)$.

9. *Shortest path counting* A chess rook can move horizontally or vertically to any square in the same row or in the same column of a chessboard. Find the number of shortest paths by which a rook can move from one corner of a chessboard to the diagonally opposite corner [Gar78], p.10

(a) by a dynamic programming algorithm.

(b) by using elementary combinatorics.

10. \triangleright *World Series odds* Consider two teams, A and B , playing a series of games until one of the teams wins n games. Assume that the probability of A winning a game is the same for each game and equal to p and the probability of A losing a game is $q = 1 - p$. (Hence, there are no ties.) Let $P(i, j)$ be the probability of A winning the series if A needs i more games to win the series and B needs j more games to win the series.

a. Set up a recurrence relation for $P(i, j)$ that can be used by a dynamic programming algorithm.

b. Find the probability of team A winning a seven-game series if the probability of it winning a game is 0.4.

c. Write a pseudocode of the dynamic programming algorithm for solving this problem and determine its time and space efficiencies.

Hints to Exercises 8.1

1. Compare the definitions of the two techniques.
2. a. Trace the algorithm *Binomial* for $n = 6$ and $k = 3$ by filling the table similar to that in Figure 8.1.

b. Check the algorithm's formula to see which values need to be known to compute $C(n, k)$.

3. Show that there exist positive constants c_1 , c_2 , and a positive integer n_0 such that

$$c_2nk \leq \frac{(k-1)k}{2} + k(n-k) \leq c_1nk$$

for all pairs of integers n, k such that $n \geq n_0$ and $0 \leq k \leq n$.

4. a. The space efficiency can be investigated in almost the same way the time efficiency is investigated in Section 8.1.

b. With some care, one can write a new row of the table over its immediate predecessor.

5. Use an explicit formula for $C(n, k)$ and Stirling's formula if necessary.

6. Set up a recurrence relation for the number of additions $A(n, k)$; solve it by substituting $A(n, k) = B(n, k) - 1$.

7. Find and compare the efficiency classes of these algorithms.

8. The formula can be proved by using either one of the explicit formulas for $C(n, k)$ or by its combinatorial interpretation.

9. a. Relate the number of the rook's shortest paths to the square in the i th row and the j th column of the chessboard to the numbers of the shortest paths to the adjacent squares.

b. Consider one shortest path as 14 consecutive moves to adjacent squares.

10. a. In the situation where teams A and B need i and j games, respectively, to win the series, consider the result of team A winning the game and the result of team A losing the game.

b. Set up a table with five rows ($0 \leq i \leq 4$) and five columns ($0 \leq j \leq 4$) and fill it by using the recurrence derived in part (a).

c. A pseudocode should be guided by the recurrence set up in part (a). The efficiency answers follow immediately from the table's size and the time spent on computing each of its entries.

Solutions to Exercises 8.1

1. a. Both techniques are based on dividing a problem's instance into smaller instances of the same problem.

b. Typically, divide-and-conquer divides an instance into smaller instances with no intersection whereas dynamic programming deals with problems in which smaller instances overlap. Consequently, divide-and-conquer algorithms do not explicitly store solutions to smaller instances and dynamic programming algorithms do.

2. a.

	0	1	2	3
0	1			
1	1	1		
2	1	2	1	
3	1	3	3	1
4	1	4	6	4
5	1	5	10	10
6	1	6	15	20

b. Yes, the table can also be filled column-by-column, with each column filled top-to-bottom starting with 1 on the main diagonal of the table.

3. First, the simple algebra yields

$$\frac{(k-1)k}{2} + k(n-k) = nk - \frac{1}{2}k^2 - \frac{1}{2}k.$$

So, we can obtain an upper bound by eliminating the negative terms:

$$nk - \frac{1}{2}k^2 - \frac{1}{2}k \leq nk \text{ for all } n, k \geq 0.$$

We can get a lower bound by considering $n \geq 2$ and $0 \leq k \leq n$:

$$nk - \frac{1}{2}k^2 - \frac{1}{2}k \geq nk - \frac{1}{2}nk - \frac{1}{2}k \frac{1}{2}n = \frac{1}{4}nk.$$

4. a. The total size of the algorithm's table is $(n+1)(k+1) \in \Theta(nk)$. Considering only the memory cells actually used by the algorithm still yields the same asymptotic class:

$$S(n, k) = \sum_{i=0}^k (i+1) + \sum_{i=k+1}^n (k+1) = \frac{(k+1)(k+2)}{2} + (k+1)(n-k) \in \Theta(nk).$$

b. The following algorithm uses just a one-dimensional array by writing a new row over its immediate predecessor; each row is filled from right to left.

Algorithm *Binomial2*(n, k)

//Computes $C(n, k)$ by the dynamic programming algorithm

//with a one-dimensional table

//Input: A pair of nonnegative integers $n \geq k \geq 0$

//Output: The value of $C(n, k)$

for $i \leftarrow 0$ **to** n **do**

if $i \leq k$ //in the triangular part

$T[i] \leftarrow 1$ //the diagonal element

$u \leftarrow i - 1$ //the rightmost element to be computed

else $u \leftarrow k$ //in the rectangular part

 //overwrite the preceding row moving right to left

for $j \leftarrow u$ **downto** 1 **do**

$T[j] \leftarrow T[j - 1] + T[j]$

return $T[k]$

5. a.

$$\text{i. } C(n, 1) = n \in \Theta(n)$$

$$\text{ii. } C(n, 2) = \frac{n(n-1)}{2} \in \Theta(n^2)$$

$$\begin{aligned} \text{iii. } C(n, \frac{n}{2}) &= \frac{n!}{(n/2)!(n-n/2)!} = \frac{n!}{[(n/2)!]^2} \\ &\approx \frac{\sqrt{2\pi n}(n/e)^n}{[\sqrt{2\pi n/2}(n/2e)^{n/2}]^2} = \frac{2}{\sqrt{2\pi n}} 2^n \in \Theta(2^n n^{-0.5}) \end{aligned}$$

b. A major implication is a very fast growth of $C(n, k)$ if k is close to $n/2$.

6. The recurrence relation for the number of additions made by the algorithm is

$$A(n, k) = A(n-1, k-1) + A(n-1, k) + 1 \quad \text{for } n > k > 0,$$

$$A(n, 0) = A(n, n) = 0 \quad \text{for } n \geq 0.$$

Substituting $A(n, k) = B(n, k) - 1$ yields

$$B(n, k) = B(n-1, k-1) + B(n-1, k) \quad \text{for } n > k > 0,$$

$$B(n, 0) = B(n, n) = 1 \quad \text{for } n \geq 0.$$

Since the last recurrence relation is identical to the recurrence relation for the binomial coefficient, $B(n, k) = C(n, k)$ and hence

$$A(n, k) = B(n, k) - 1 = C(n, k) - 1.$$

7. a. Assuming that $1!$ is computed as $0! \cdot 1$ and, hence, requires one multiplication, the use of the formula

$$C(n, k) = \frac{n!}{k!(n-k)!}$$

requires $n + k + n - k = 2n$ multiplications and one division.

- b. The use of the formula

$$C(n, k) = \frac{n(n-1)\dots(n-k+1)}{k!}$$

requires only $k + k = 2k$ multiplications and one division. (We assumed that the product in the numerator is initialized to 1.) Hence, (b) is more efficient than (a). Also note that the overflow problem is less severe for formula (b) than for formula (a). It can be further alleviated by using the formula

$$\frac{n(n-1)\dots(n-k+1)}{k!} = \frac{n}{k} \frac{n-1}{k-1} \dots \frac{n-(k-1)}{1}.$$

- c. Applying recursively the formula

$$C(n, k) = C(n-1, k-1) + C(n-1, k) \quad \text{for } n > k > 0, \quad C(n, 0) = C(n, n) = 1$$

requires $C(n, k) - 1$ additions (see Problem 6), which is much worse than the dynamic programming algorithm that requires $nk - \frac{1}{2}k^2 - \frac{1}{2}k$ additions (see Section 8.1).

- d. The efficiency of the dynamic programming algorithm, as explained in the section, is in $\Theta(nk)$.

Thus, the choice comes to be between algorithms (b) and (d). Making a choice between two algorithms with different basic operations should be conditioned on their relative speeds. Still, since the time efficiency of algorithm (b) is in $\Theta(k)$ while the time efficiency of algorithm (d) is in $\Theta(nk)$, we have to conclude that the former is more time efficient, at least asymptotically, than the latter. Algorithm (b) is obviously more space efficient than algorithm (d), too.

8. Using the formula $C(n, k) = \frac{n!}{k!(n-k)!}$, we obtain the following equality:

$$C(n, n-k) = \frac{n!}{(n-k)!(n-(n-k))!} = \frac{n!}{(n-k)!k!} = \frac{n!}{k!(n-k)!} = C(n, k).$$

Alternatively, $C(n, k)$ is the number of ways to choose a k -element subset from a set of n elements, which is the same as selecting $n - k$ elements not to be included in a k -element subset.

The formula can be used for speeding up algorithms (b), (c), and (d) of Problem 7: compute $C(n, k)$ if $k \leq n - k$ (i.e., if $k \leq n/2$) and $C(n, n - k)$ otherwise.

9. a. With no loss of generality, we can assume that the rook is initially located in the lower left corner of a chessboard, whose rows and columns are numbered from 1 to 8. Let $P(i, j)$ be the number of the rook's shortest paths from square (1,1) to square (i, j) in the i th row and the j th column, where $1 \leq i, j \leq 8$. Any such path will be composed of vertical and horizontal moves directed toward the goal. Obviously, $P(i, 1) = P(1, j) = 1$ for any $1 \leq i, j \leq 8$. In general, any shortest path to square (i, j) will be reached either from square $(i, j - 1)$ or from square $(i - 1, j)$ square. Hence, we have the following recurrence

$$\begin{aligned} P(i, j) &= P(i, j - 1) + P(i - 1, j) \quad \text{for } 1 < i, j \leq 8, \\ P(i, 1) &= P(1, j) = 1 \quad \text{for } 1 \leq i, j \leq 8. \end{aligned}$$

Using this recurrence, we can compute the values of $P(i, j)$ for each square (i, j) of the board. This can be done either row by row, or column by column, or diagonal by diagonal. (One can also take advantage of the board's symmetry to make the computations only for the squares either on and above or on and below the board's main diagonal.) The results

are given in the diagram below:

1	8	36	120	330	792	1716	3432
1	7	28	84	210	462	924	1716
1	6	21	56	126	252	462	792
1	5	15	35	70	126	210	330
1	4	10	20	35	56	84	120
1	3	6	10	15	21	28	36
1	2	3	4	5	6	7	8
1	1	1	1	1	1	1	1

b. Any shortest path from square (1,1) to square (8,8) can be thought of as 14 consecutive moves to adjacent squares, seven of which being up while the other seven being to the right. For example, the shortest path composed of the vertical move from (1,1) to (8,1) followed by the horizontal move from (8,1) to (8,8) corresponds to the following sequence of 14 one-square moves:

$$(u, u, u, u, u, u, u, r, r, r, r, r, r, r),$$

where u and r stand for a move up and to the right, respectively. Hence, the total number of distinct shortest paths is equal to the number of different ways to choose seven u -positions among the total of 14 possible positions, which is equal to $C(14, 7)$.

10. a. Let $P(i, j)$ be the probability of A winning the series if A needs i more games to win the series and B needs j more games to win the series. If team A wins the game, which happens with probability p , A will need $i - 1$ more wins to win the series while B will still need j wins. If team A loses the game, which happens with probability $q = 1 - p$, A will still need i wins while B will need $j - 1$ wins to win the series. This leads to the recurrence

$$P(i, j) = pP(i - 1, j) + qP(i, j - 1) \text{ for } i, j > 0.$$

The initial conditions follow immediately from the definition of $P(i, j)$:

$$P(0, j) = 1 \text{ for } j > 0, \quad P(i, 0) = 0 \text{ for } i > 0.$$

b. Here is the dynamic programming table in question, with its entries rounded-off to two decimal places. (It can be filled either row-by-row, or column-by-column, or diagonal-by-diagonal.)

$i \setminus j$	0	1	2	3	4
0		1	1	1	1
1	0	0.40	0.64	0.78	0.87
2	0	0.16	0.35	0.52	0.66
3	0	0.06	0.18	0.32	0.46
4	0	0.03	0.09	0.18	0.29

Thus, $P[4, 4] \approx 0.29$.

c. **Algorithm** *WorldSeries*(n, p)

```
//Computes the odds of winning a series of  $n$  games
//Input: A number of wins  $n$  needed to win the series
//       and probability  $p$  of one particular team winning a game
//Output: The probability of this team winning the series
 $q \leftarrow 1 - p$ 
for  $j \leftarrow 1$  to  $n$  do
     $P[0, j] \leftarrow 1.0$ 
for  $i \leftarrow 1$  to  $n$  do
     $P[i, 0] \leftarrow 0.0$ 
    for  $j \leftarrow 1$  to  $n$  do
         $P[i, j] \leftarrow p * P[i - 1, j] + q * P[i, j - 1]$ 
return  $P[n, n]$ 
```

Both the time efficiency and the space efficiency are in $\Theta(n^2)$ because each entry of the $n + 1$ -by- $n + 1$ table (except $P[0, 0]$, which is not computed) is computed in $\Theta(1)$ time.

Exercises 8.2

1. Apply Warshall's algorithm to find the transitive closure of the digraph defined by the following adjacency matrix

$$\begin{bmatrix} 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 \end{bmatrix}$$

2. a. Prove that the time efficiency of Warshall's algorithm is cubic.

b. Explain why the time efficiency of Warshall's algorithm is inferior to that of the traversal-based algorithm for sparse graphs represented by their adjacency lists.
3. Explain how to implement Warshall's algorithm without using extra memory for storing elements of the algorithm's intermediate matrices.
4. Explain how to restructure the innermost loop of the algorithm *Warshall* to make it run faster at least on some inputs.
5. Rewrite the pseudocode of Warshall's algorithm assuming that the matrix rows are represented by bit strings on which the bitwise *or* operation can be performed.
6. a. Explain how Warshall's algorithm can be used to determine whether a given digraph is a dag (directed acyclic graph). Is it a good algorithm for this problem?

b. Is it a good idea to apply Warshall's algorithm to find the transitive closure of an undirected graph?
7. Solve the all-pairs shortest path problem for the digraph with the following weight matrix

$$\begin{bmatrix} 0 & 2 & \infty & 1 & 8 \\ 6 & 0 & 3 & 2 & \infty \\ \infty & \infty & 0 & 4 & \infty \\ \infty & \infty & 2 & 0 & 3 \\ 3 & \infty & \infty & \infty & 0 \end{bmatrix}$$

8. Prove that the next matrix in sequence (8.8) of Floyd's algorithm can be written over its predecessor.
9. Give an example of a graph or a digraph with negative weights for which Floyd's algorithm does not yield the correct result.
10. ► Enhance Floyd's algorithm so that shortest paths themselves, not just their lengths, can be found.

11. *Jack Straws* In the game of Jack Straws, a number of plastic or wooden "straws" are dumped on the table and players try to remove them one-by-one without disturbing the other straws. Here, we are only concerned with if various pairs of straws are connected by a path of touching straws. Given a list of the endpoints for $n > 1$ straws (as if they were dumped on a large piece of graph paper), determine all the pairs of straws that are connected. Note that touching is connecting, but also two straws can be connected indirectly via other connected straws [1994 East-Central Regionals of the ACM International Collegiate Programming Contest].

Hints to Exercises 8.2

1. Apply the algorithm to the adjacency matrix given as it is done in the section for another matrix.
2. a. The answer can be obtained either by considering how many values the algorithm computes or by following the standard plan for analyzing the efficiency of a nonrecursive algorithm (i.e., by setting up a sum to count its basic operation).

b. What is the efficiency class of the traversal-based algorithm for sparse graphs represented by their adjacency lists?
3. Show that we can simply overwrite elements of $R^{(k-1)}$ with elements of $R^{(k)}$ without any other changes in the algorithm.
4. What happens if $R^{(k-1)}[i, k] = 0$?
5. Show first that formula (8.7) (from which the superscripts can be eliminated according to the solution to Problem 3)

$$r_{ij} = r_{ij} \text{ or } r_{ik} \text{ and } r_{kj}.$$

is equivalent to

$$\text{if } r_{ik} \text{ } r_{ij} \leftarrow r_{ij} \text{ or } r_{kj}.$$

6. a. What property of the transitive closure indicates a presence of a directed cycle? Is there a better algorithm for checking this?

b. Which elements of the transitive closure of an undirected graph are equal to 1? Can you find such elements with a faster algorithm?
7. See an example of applying the algorithm to another instance in the section.
8. What elements of matrix $D^{(k-1)}$ does $d_{ij}^{(k)}$, the element in the i th row and the j th column of matrix $D^{(k)}$, depend on? Can these values be changed by the overwriting?
9. Your counterexample must contain a cycle of a negative length.
10. It will suffice to store, in a single matrix P , indices of intermediate vertices k used in updates of the distance matrices. This matrix can be initialized with all zero elements.
11. The problem can be solved by utilizing two well-known algorithms: one from computational geometry, the other dealing with graphs.

Solutions to Exercises 8.2

1. Applying Warshall's algorithm yields the following sequence of matrices (in which newly updated elements are shown in bold):

$$R^{(0)} = \begin{bmatrix} 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 \end{bmatrix} \quad R^{(1)} = \begin{bmatrix} 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 \end{bmatrix}$$

$$R^{(2)} = \begin{bmatrix} 0 & 1 & \mathbf{1} & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 \end{bmatrix} \quad R^{(3)} = \begin{bmatrix} 0 & 1 & 1 & \mathbf{1} \\ 0 & 0 & 1 & \mathbf{1} \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 \end{bmatrix}$$

$$R^{(4)} = \begin{bmatrix} 0 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 \end{bmatrix} = T$$

2. a. For a graph with n vertices, the algorithm computes n matrices $R^{(k)}$ ($k = 1, 2, \dots, n$), each of which has n^2 elements. Hence, the total number of elements to be computed is n^3 . Since computing each element takes constant time, the time efficiency of the algorithm is in $\Theta(n^3)$.
- b. Since one *DFS* or *BFS* traversal of a graph with n vertices and m edges, which is represented by its adjacency lists, takes $\Theta(n + m)$ time, doing this n times takes $n\Theta(n + m) = \Theta(n^2 + nm)$ time. For sparse graphs (i.e., if $m \in O(n)$), $\Theta(n^2 + nm) = \Theta(n^2)$, which is more efficient than the $\Theta(n^3)$ time efficiency of Warshall's algorithm.
3. The algorithm computes the new value to be put in the i th row and the j th column by the formula

$$R^{(k)}[i, j] = R^{(k-1)}[i, j] \text{ or } (R^{(k-1)}[i, k] \text{ and } R^{(k-1)}[k, j]).$$

The formula implies that all the elements in the k th row and all the elements in the k th column never change on the k th iteration, i.e., while computing elements of $R^{(k)}$ from elements of $R^{(k-1)}$. (Substitute k for i and k for j , respectively, into the formula to verify these assertions.) Hence, the new value of the element in the i th row and the j th column, $R^{(k)}[i, j]$, can be written over its old value $R^{(k-1)}[i, j]$ for every i, j .

4. If $R^{(k-1)}[i, k] = 0$,

$$R^{(k)}[i, j] = R^{(k-1)}[i, j] \text{ or } (R^{(k-1)}[i, k] \text{ and } R^{(k-1)}[k, j]) = R^{(k-1)}[i, j],$$

and hence the innermost loop need not be executed. And, since $R^{(k-1)}[i, k]$ doesn't depend on j , its comparison with 0 can be done outside the j loop. This leads to the following implementation of Warshall's algorithm:

Algorithm *Warshall2*($A[1..n, 1..n]$)

//Implements Warshall's algorithm with a more efficient innermost loop

//Input: The adjacency matrix A of a digraph with n vertices

//Output: The transitive closure of the digraph in place of A

for $k \leftarrow 1$ **to** n **do**

for $i \leftarrow 1$ **to** n **do**

if $A[i, k]$

for $j \leftarrow 1$ **to** n **do**

if $A[k, j]$

$A[i, j] \leftarrow 1$

return A

5. First, it is easy to check that

$$r_{ij} \leftarrow r_{ij} \text{ or } (r_{ik} \text{ and } r_{kj}).$$

is equivalent to

$$\text{if } r_{ik} \quad r_{ij} \leftarrow r_{ij} \text{ or } r_{kj}$$

Indeed, if $r_{ik} = 1$ (i.e., **true**),

$$r_{ij} \text{ or } (r_{ik} \text{ and } r_{kj}) = r_{ij} \text{ or } r_{kj},$$

which is exactly the value assigned to r_{ij} by the **if** statement as well. If $r_{ik} = 0$ (i.e., **false**),

$$r_{ij} \text{ or } (r_{ik} \text{ and } r_{kj}) = r_{ij},$$

i.e., the new value of r_{ij} will be the same as its previous value—exactly, the result we obtain by executing the **if** statement.

Here is the algorithm that exploits this observation and the bitwise *or* operation applied to matrix rows:

Algorithm *Warshall3*($A[1..n, 1..n]$)

//Implements Warshall's algorithm for computing the transitive closure

//with the bitwise *or* operation on matrix rows

//Input: The adjacency matrix A of a digraph

//Output: The transitive closure of the digraph in place of A

for $k \leftarrow 1$ **to** n **do**


```

for  $i \leftarrow 1$  to  $n$  do
    if  $A[i, k]$ 
         $row[i] \leftarrow row[i]$  bitwiseor  $row[k]$  //rows of matrix  $A$ 
return  $A$ 

```

6. a. With the book's definition of the transitive closure (which considers only nontrivial paths of a digraph), a digraph has a directed cycle if and only if its transitive closure has a 1 on its main diagonal. The algorithm that finds the transitive closure by applying Warshall's algorithm and then checks the elements of its main diagonal is cubic. This is inferior to the quadratic algorithms for checking whether a digraph represented by its adjacency matrix is a dag, which were discussed in Section 5.3.

b. No. If T is the transitive closure of an undirected graph, $T[i, j] = 1$ if and only if there is a nontrivial path from the i th vertex to the j th vertex. If $i \neq j$, this is the case if and only if the i th vertex and the j th vertex belong to the same connected component of the graph. Thus, one can find the elements outside the main diagonal of the transitive closure that are equal to 1 by a depth-first search or a breadth-first search traversal, which is faster than applying Warshall's algorithm. If $i = j$, $T[i, i] = 1$ if and only if the i th vertex is not isolated, i.e., if it has an edge to at least one other vertex of the graph. Isolated vertices, if any, can be easily identified by the graph's traversal as one-node connected components of the graph.

7. Applying Floyd's algorithm to the given weight matrix generates the following sequence of matrices:

$$\begin{aligned}
 D^{(0)} &= \begin{bmatrix} 0 & 2 & \infty & 1 & 8 \\ 6 & 0 & 3 & 2 & \infty \\ \infty & \infty & 0 & 4 & \infty \\ \infty & \infty & 2 & 0 & 3 \\ 3 & \infty & \infty & \infty & 0 \end{bmatrix} & D^{(1)} &= \begin{bmatrix} 0 & 2 & \infty & 1 & 8 \\ 6 & 0 & 3 & 2 & \mathbf{14} \\ \infty & \infty & 0 & 4 & \infty \\ \infty & \infty & 2 & 0 & 3 \\ 3 & \mathbf{5} & \infty & 4 & 0 \end{bmatrix} \\
 D^{(2)} &= \begin{bmatrix} 0 & 2 & \mathbf{5} & 1 & 8 \\ 6 & 0 & 3 & 2 & 14 \\ \infty & \infty & 0 & 4 & \infty \\ \infty & \infty & 2 & 0 & 3 \\ 3 & 5 & \mathbf{8} & 4 & 0 \end{bmatrix} & D^{(3)} &= \begin{bmatrix} 0 & 2 & 5 & 1 & 8 \\ 6 & 0 & 3 & 2 & 14 \\ \infty & \infty & 0 & 4 & \infty \\ \infty & \infty & 2 & 0 & 3 \\ 3 & 5 & 8 & 4 & 0 \end{bmatrix} \\
 D^{(4)} &= \begin{bmatrix} 0 & 2 & \mathbf{3} & 1 & 4 \\ 6 & 0 & 3 & 2 & \mathbf{5} \\ \infty & \infty & 0 & 4 & \mathbf{7} \\ \infty & \infty & 2 & 0 & 3 \\ 3 & 5 & \mathbf{6} & 4 & 0 \end{bmatrix} & D^{(5)} &= \begin{bmatrix} 0 & 2 & 3 & 1 & 4 \\ 6 & 0 & 3 & 2 & 5 \\ \mathbf{10} & \mathbf{12} & 0 & 4 & 7 \\ \mathbf{6} & \mathbf{8} & 2 & 0 & 3 \\ 3 & 5 & 6 & 4 & 0 \end{bmatrix} = D
 \end{aligned}$$

8. The formula

$$d_{ij}^{(k)} = \min\{d_{ij}^{(k-1)}, d_{ik}^{(k-1)} + d_{kj}^{(k-1)}\}$$

implies that the value of the element in the i th row and the j th column of matrix $D^{(k)}$ is computed from its own value and the values of the elements in the i th row and the k th column and in the k th row and the j th column in the preceding matrix $D^{(k-1)}$. The latter two cannot change their values when the elements of $D^{(k)}$ are computed because of the formulas


$$d_{ik}^{(k)} = \min\{d_{ik}^{(k-1)}, d_{ik}^{(k-1)} + d_{kk}^{(k-1)}\} = \min\{d_{ik}^{(k-1)}, d_{ik}^{(k-1)} + 0\} = d_{ik}^{(k-1)}$$

and

$$d_{kj}^{(k)} = \min\{d_{kj}^{(k-1)}, d_{kk}^{(k-1)} + d_{kj}^{(k-1)}\} = \min\{d_{kj}^{(k-1)}, 0 + d_{kj}^{(k-1)}\} = d_{kj}^{(k-1)}.$$

(Note that we took advantage of the fact that the elements d_{kk} on the main diagonal remain 0's. This can be guaranteed only if the graph doesn't contain a cycle of a negative length.)

9. As a simple counterexample, one can suggest the following digraph:



Floyd's algorithm will yield:

$$D^{(0)} = \begin{bmatrix} 0 & 1 \\ -2 & 0 \end{bmatrix} \quad D^{(1)} = \begin{bmatrix} 0 & 1 \\ -2 & -1 \end{bmatrix} \quad D^{(2)} = \begin{bmatrix} -1 & 0 \\ -3 & -2 \end{bmatrix}$$

None of the four elements of the last matrix gives the correct value of the shortest path, which is, in fact, $-\infty$ because repeating the cycle enough times makes the length of a path arbitrarily small.

Note: Floyd's algorithm can be used for detecting negative-length cycles, but the algorithm should be stopped as soon as it generates a matrix with a negative element on its main diagonal.

10. As pointed out in the hint to this problem, Floyd's algorithm should be enhanced by recording in an n -by- n matrix index k of an intermediate vertex causing an update of the distance matrix. This is implemented in the pseudocode below:

Algorithm *FloydEnhanced*($W[1..n, 1..n]$)

//Input: The weight matrix W of a graph or a digraph

//Output: The distance matrix $D[1..n, 1..n]$ and

```

//          the matrix of intermediate updates  $P[1..n, 1..n]$ 
 $D \leftarrow W$ 
for  $i \leftarrow 1$  to  $n$  do
    for  $j \leftarrow 1$  to  $n$  do
         $P[i, j] \leftarrow 0$  //initial mark
for  $k \leftarrow 1$  to  $n$  do
    for  $i \leftarrow 1$  to  $n$  do
        for  $j \leftarrow 1$  to  $n$  do
            if  $D[i, k] + D[k, j] < D[i, j]$ 
                 $D[i, j] \leftarrow D[i, k] + D[k, j]$ 
                 $P[i, j] \leftarrow k$ 

```

For example, for the digraph in Fig. 8.7 whose vertices are numbered from 1 to 4, matrix P will be as follows:

$$P = \begin{bmatrix} 0 & 3 & 0 & 3 \\ 0 & 0 & 1 & 3 \\ 4 & 0 & 0 & 0 \\ 0 & 3 & 1 & 0 \end{bmatrix}.$$

The list of intermediate vertices on the shortest path from vertex i to vertex j can be then generated by the call to the following recursive algorithm, provided $D[i, j] < \infty$:

Algorithm *ShortestPath*($i, j, P[1..n, 1..n]$)
//The algorithm prints out the list of intermediate vertices of
//a shortest path from the i th vertex to the j th vertex
//Input: Endpoints i and j of the shortest path desired;
// matrix P of updates generated by Floyd's algorithm
//Output: The list of intermediate vertices of the shortest path
// from the i th vertex to the j th vertex
 $k \leftarrow P[i, j]$
if $k \neq 0$
ShortestPath(i, k)
 print(k)
ShortestPath(k, j)

11. First, for each pair of the straws, determine whether the straws intersect. (Although this can be done in $n \log n$ time by a sophisticated algorithm, the quadratic brute-force algorithm would do because of the quadratic efficiency of the subsequent step; both geometric algorithms can be found, e.g., in R. Sedgewick's "Algorithms," Addison-Wesley, 1988.) Record the obtained information in a boolean n -by- n matrix, which must be symmetric. Then find the transitive closure of this matrix in n^2 time by DFS or BFS (see the solution to Problem 6b).

Exercises 8.3

1. Finish the computations started in the section's example of constructing an optimal binary search tree.
2. a. Why is the time efficiency of algorithm *OptimalBST* cubic?
b. Why is the space efficiency of algorithm *OptimalBST* quadratic?
3. Write a pseudocode for a linear-time algorithm that generates the optimal binary search tree from the root table.
4. Devise a way to compute the sums $\sum_{s=i}^j p_s$, which are used in the dynamic programming algorithm for constructing an optimal binary search tree, in constant time (per sum).
5. True or false: The root of an optimal binary search tree always contains the key with the highest search probability?
6. How would you construct an optimal binary search tree for a set of n keys if all the keys are equally likely to be searched for? What will be the average number of comparisons in a successful search in such a tree if $n = 2^k$?
7. a. Show that the number of distinct binary search trees $b(n)$ that can be constructed for a set of n orderable keys satisfies the recurrence relation

$$b(n) = \sum_{k=0}^{n-1} b(k)b(n-1-k) \quad \text{for } n > 0, \quad b(0) = 1.$$

- b. It is known that the solution to this recurrence is given by the Catalan numbers. Verify this assertion for $n = 1, 2, \dots, 5$.
- c. Find the order of growth of $b(n)$. What implication does the answer to this question have for the exhaustive search algorithm for constructing an optimal binary search tree?
8. ► Design a $\Theta(n^2)$ algorithm for finding an optimal binary search tree.
9. ▷ Generalize the optimal binary search algorithm by taking into account unsuccessful searches.
10. **Matrix chain multiplication** Consider the problem of minimizing the total number of multiplications made in computing the product of n matrices

$$A_1 \cdot A_2 \cdot \dots \cdot A_n$$

whose dimensions are d_0 by d_1 , d_1 by d_2 , ..., d_{n-1} by d_n , respectively. (Assume that all intermediate products of two matrices are computed by the

brute-force (definition-based) algorithm.

- a. Give an example of three matrices for which the number of multiplications in $(A_1 \cdot A_2) \cdot A_3$ and $A_1 \cdot (A_2 \cdot A_3)$ differ at least by a factor 1000.
- b.▷ How many different ways are there to compute the chained product of n matrices?
- c.► Design a dynamic programming algorithm for finding an optimal order of multiplying n matrices.

Hints to Exercises 8.3

1. Continue applying formula (8.11) as prescribed by the algorithm.
2. a. The algorithm's time efficiency can be investigated by following the standard plan of analyzing the time efficiency of a nonrecursive algorithm.

b. How much space do the two tables generated by the algorithm use?
3. $k = R[1, n]$ indicates that the root of an optimal tree is the k th key in the list of ordered keys a_1, \dots, a_n . The roots of its left and right subtrees are specified by $R[1, k - 1]$ and $R[k + 1, n]$, respectively.
4. Use a space-for-time tradeoff.
5. If the assertion were true, would we not have a simpler algorithm for constructing an optimal binary search tree?
6. The structure of the tree should simply minimize the average depth of its nodes. Do not forget to indicate a way to distribute the keys among the nodes of the tree.
7. a. Since there is a one-to-one correspondence between binary search trees for a given set of n orderable keys and the total number of binary trees with n nodes (why?), you can count the latter. Consider all the possibilities of partitioning the nodes between the left and right subtrees.

b. Compute the values in question using the two formulas.

c. Use the formula for the n th Catalan number and Stirling's formula for $n!$.
8. Change the bounds of the innermost loop of algorithm *OptimalBST* by exploiting the monotonicity of the root table mentioned at the end of Section 8.3.
9. Assume that a_1, \dots, a_n are distinct keys ordered from the smallest to the largest, p_1, \dots, p_n are the probabilities of searching for them, and q_0, q_1, \dots, q_n are probabilities of unsuccessful searches for keys in intervals $(-\infty, a_1)$, $(a_1, a_2), \dots, (a_n, \infty)$, respectively; $(p_1 + \dots + p_n) + (q_0 + \dots + q_n) = 1$. Set up a recurrence relation similar to recurrence (8.11) for the expected number of key comparisons that takes into account both successful and unsuccessful searches.
10. a. It is easier to find a general formula for the number of multiplications needed for computing $(A_1 \cdot A_2) \cdot A_3$ and $A_1 \cdot (A_2 \cdot A_3)$ for matrices A_1 with dimensions d_0 -by- d_1 , A_2 with dimensions d_1 -by- d_2 , and A_3 with dimensions d_2 -by- d_3 and then choose some specific values for the dimensions to get a required example.

- b. You can get the answer by following the approach used for counting binary trees.
- c. The recurrence relation for the optimal number of multiplications in computing $A_i \cdot \dots \cdot A_j$ is very similar to the recurrence relation for the optimal number of comparisons in searching in a binary tree composed of keys a_i, \dots, a_j .

Solutions to Exercises 8.3

1. The instance of the problem in question is defined by the data

key	A	B	C	D
probability	0.1	0.2	0.4	0.3

The table entries for the dynamic programming algorithm are computed as follows:

$$C[1, 2] = \min \begin{array}{ll} k=1: & C[1, 0] + C[2, 2] + \sum_{s=1}^2 p_s = 0 + 0.2 + 0.3 = 0.5 \\ k=2: & C[1, 1] + C[3, 2] + \sum_{s=1}^2 p_s = 0.1 + 0 + 0.3 = \mathbf{0.4} \end{array} = 0.4$$

$$C[2, 3] = \min \begin{array}{ll} k=2: & C[2, 1] + C[3, 3] + \sum_{s=2}^3 p_s = 0 + 0.4 + 0.6 = 1.0 \\ k=3: & C[2, 2] + C[4, 3] + \sum_{s=2}^3 p_s = 0.2 + 0 + 0.6 = \mathbf{0.8} \end{array} = 0.8$$

$$C[3, 4] = \min \begin{array}{ll} k=3: & C[3, 2] + C[4, 4] + \sum_{s=3}^4 p_s = 0 + 0.3 + 0.7 = \mathbf{1.0} \\ k=4: & C[3, 3] + C[5, 4] + \sum_{s=3}^4 p_s = 0.4 + 0 + 0.7 = 1.1 \end{array} = 1.0$$

$$C[1, 3] = \min \begin{array}{ll} k=1: & C[1, 0] + C[2, 3] + \sum_{s=1}^3 p_s = 0 + 0.8 + 0.7 = 1.5 \\ k=2: & C[1, 1] + C[3, 3] + \sum_{s=1}^3 p_s = 0.1 + 0.4 + 0.7 = 1.2 \\ k=3: & C[1, 2] + C[4, 3] + \sum_{s=1}^3 p_s = 0.4 + 0 + 0.7 = \mathbf{1.1} \end{array} = 1.1$$

$$C[2, 4] = \min \begin{array}{ll} k=2: & C[2, 1] + C[3, 4] + \sum_{s=2}^4 p_s = 0 + 1.0 + 0.9 = 1.9 \\ k=3: & C[2, 2] + C[4, 4] + \sum_{s=2}^4 p_s = 0.2 + 0.3 + 0.9 = \mathbf{1.4} \\ k=4: & C[2, 3] + C[5, 4] + \sum_{s=2}^4 p_s = 0.8 + 0 + 0.9 = 1.7 \end{array} = 1.1$$


$$C[1, 4] = \min \begin{array}{ll} k=1: & C[1, 0] + C[2, 4] + \sum_{s=1}^4 p_s = 0 + 1.4 + 1.0 = 2.4 \\ k=2: & C[1, 1] + C[3, 4] + \sum_{s=1}^4 p_s = 0.1 + 1.0 + 1.0 = 2.1 \\ k=3: & C[1, 2] + C[4, 4] + \sum_{s=1}^4 p_s = 0.4 + 0.3 + 1.0 = \mathbf{1.7} \\ k=4: & C[1, 3] + C[5, 4] + \sum_{s=1}^4 p_s = 1.1 + 0 + 1.0 = 2.1 \end{array} = 1.7$$

the main table

	0	1	2	3	4
1	0	0.1	0.4	1.1	1.7
2		0	0.2	0.8	1.4
3			0	0.4	1.0
4				0	0.3
5					0

the root table

	0	1	2	3	4
1		1	2	3	3
2			2	3	3
3				3	3
4					4
5					



2. a. The number of times the innermost loop is executed is given by the sum

$$\begin{aligned}
\sum_{d=1}^{n-1} \sum_{i=1}^{n-d} \sum_{k=i}^{i+d} 1 &= \sum_{d=1}^{n-1} \sum_{i=1}^{n-d} (i + d - i + 1) = \sum_{d=1}^{n-1} \sum_{i=1}^{n-d} (d + 1) \\
&= \sum_{d=1}^{n-1} (d + 1)(n - d) = \sum_{d=1}^{n-1} (dn + n - d^2 - d) \\
&= \sum_{d=1}^{n-1} nd + \sum_{d=1}^{n-1} n - \sum_{d=1}^{n-1} d^2 - \sum_{d=1}^{n-1} d \\
&= n \frac{(n-1)n}{2} + n(n-1) - \frac{(n-1)n(2n-1)}{6} - \frac{(n-1)n}{2} \\
&= \frac{1}{2}n^3 - \frac{2}{6}n^3 + O(n^2) \in \Theta(n^3).
\end{aligned}$$

- b. The algorithm generates the $(n+1)$ -by- $(n+1)$ table C and the n -by- n table R and fills about one half of each. Hence, the algorithm's space efficiency is in $\Theta(n^2)$.

3. Call *OptimalTree*(1, n) below:

Algorithm *OptimalTree*(i, j)

//Input: Indices i and j of the first and last keys of a sorted list of keys

//composing the tree and table $R[1..n, 1..n]$ obtained by dynamic

//programming

//Output: Indices of nodes of an optimal binary search tree in preorder

if $i \leq j$

$k \leftarrow R[i, j]$

 print(k)

OptimalTree($i, k - 1$)

OptimalTree($k + 1, j$)

4. Precompute $S_k = \sum_{s=1}^k p_s$ for $k = 1, 2, \dots, n$ and set $S_0 = 0$. Then $\sum_{s=i}^j p_s$ can be found as $S_j - S_{i-1}$ for any $1 \leq i \leq j \leq n$.
5. False. Here is a simple counterexample: $A(0.3)$, $B(0.3)$, $C(0.4)$. (The numbers in the parentheses indicate the search probabilities.) The average number of comparisons in a binary search tree with C in its root is $0.3 \cdot 2 + 0.3 \cdot 3 + 0.4 \cdot 1 = 1.9$, while the average number of comparisons in the binary search tree with B in its root is $0.3 \cdot 1 + 0.3 \cdot 2 + 0.4 \cdot 2 = 1.7$.

6. The binary search tree in question should have a maximal number of nodes on each of its levels except the last one. (For simplicity, we can put all the nodes of the last level in their leftmost positions possible to make it complete.) The keys of a given sorted list can be distributed among the nodes of the binary tree by performing its in-order traversal.

Let p/n be the probability of searching for each key, where $0 \leq p \leq 1$. The complete binary tree with 2^k nodes will have 2^i nodes on level i for $i = 0, \dots, k-1$ and one node on level k . Hence, the average number of comparisons in a successful search will be given by the following formula:

$$\begin{aligned}
 C(2^k) &= \sum_{i=0}^{k-1} (p/2^k)(i+1)2^i + (p/2^k)(k+1) \\
 &= (p/2^k) \frac{1}{2} \sum_{i=0}^{k-1} (i+1)2^{i+1} + (p/2^k)(k+1) \\
 &= (p/2^k) \frac{1}{2} \sum_{j=1}^k j2^j + (p/2^k)(k+1) \\
 &= (p/2^k) \frac{1}{2} [(k-1)2^{k+1} + 2] + (p/2^k)(k+1) \\
 &= (p/2^k)[(k-1)2^k + k + 2].
 \end{aligned}$$

7. a. Let $b(n)$ be the number of distinct binary trees with n nodes. If the left subtree of a binary tree with n nodes has k nodes ($0 \leq k \leq n-1$), the right subtree must have $n-1-k$ nodes. The number of such trees is therefore $b(k)b(n-1-k)$. Hence,

$$b(n) = \sum_{k=0}^{n-1} b(k)b(n-1-k) \quad \text{for } n > 0, \quad b(0) = 1.$$

- b. Substituting the first five values of n into the formula above and into the formula for the n th Catalan number yields the following values:

n	0	1	2	3	4	5
$b(n) = c(n)$	1	1	2	5	14	42

c.

$$\begin{aligned}
b(n) &= c(n) = \binom{2n}{n} \frac{1}{n+1} = \frac{(2n)!}{(n!)^2} \frac{1}{n+1} \\
&\approx \frac{\sqrt{2\pi 2n} (2n/e)^{2n}}{[\sqrt{2\pi n} (n/e)^n]^2} \frac{1}{n+1} = \frac{\sqrt{4\pi n} (2n/e)^{2n}}{2\pi n (n/e)^{2n}} \frac{1}{n+1} \\
&\approx \frac{1}{\sqrt{\pi n}} \left(\frac{2n/e}{n/e} \right)^{2n} \frac{1}{n+1} = \frac{1}{\sqrt{\pi n}} 2^{2n} \frac{1}{n+1} \in \Theta(4^n n^{-3/2}).
\end{aligned}$$

This implies that finding an optimal binary search tree by exhaustive search is feasible only for very small values of n and is, in general, vastly inferior to the dynamic programming algorithm.

8. The dynamic programming algorithm finds the root a_{kmin} of an optimal binary search tree for keys a_i, \dots, a_j by minimizing $\{C[i, k-1] + C[k+1, j]\}$ for $i \leq k \leq j$. As pointed out at the end of Section 8.3 (see also [KnuIII], p. 456, Exercise 27), $R[i, j-1] \leq kmin \leq R[i+1, j]$. This observation allows us to change the bounds of the algorithm's innermost loop to the lower bound of $R[i, j-1]$ and the upper bound of $R[i+1, j]$, respectively. The number of times the innermost loop of the modified algorithm will be executed can be estimated as suggested by Knuth (see [KnuIII], p.439):

$$\begin{aligned}
\sum_{d=1}^{n-1} \sum_{i=1}^{n-d} \sum_{k=R[i, j-1]}^{R[i+1, j]} 1 &= \sum_{d=1}^{n-1} \sum_{i=1}^{n-d} \sum_{k=R[i, i+d-1]}^{R[i+1, i+d]} 1 \\
&= \sum_{d=1}^{n-1} \sum_{i=1}^{n-d} (R[i+1, i+d] - R[i, i+d-1] + 1) \\
&= \sum_{d=1}^{n-1} \left(\sum_{i=1}^{n-d} (R[i+1, i+d] - R[i, i+d-1]) + \sum_{i=1}^{n-d} 1 \right).
\end{aligned}$$

By "telescoping" the first sum, we can see that all its terms except the two get cancelled to yield

$$\begin{aligned}
&\sum_{i=1}^{n-d} (R[i+1, i+d] - R[i, i+d-1]) = \\
&= (R[2, 1+d] - R[1, 1+d-1]) \\
&+ (R[3, 2+d] - R[2, 2+d-1]) \\
&+ \dots \\
&+ (R[n-d+1, n-d+d] - R[n-d, n-d+d-1]) \\
&= R[n-d+1, n] - R[1, d].
\end{aligned}$$

Since the second sum $\sum_{i=1}^{n-d} 1$ is equal to $n-d$, we obtain

$$\sum_{i=1}^{n-d} (R[i+1, i+d] - R[i, i+d-1]) + \sum_{i=1}^{n-d} 1 = R[n-d+1, n] - R[1, d] + n-d < 2n.$$

Hence,

$$\sum_{d=1}^{n-1} \sum_{i=1}^{n-d} \sum_{k=R[i,j-1]}^{R[i+1,j]} 1 = \sum_{d=1}^{n-1} (R[n-d+1, n] - R[1, d] + n - d) < \sum_{d=1}^{n-1} 2n < 2n^2.$$

On the other hand, since $R[i+1, i+d] - R[i, i+d-1] \geq 0$,

$$\begin{aligned} \sum_{d=1}^{n-1} \sum_{i=1}^{n-d} \sum_{k=R[i,j-1]}^{R[i+1,j]} 1 &= \sum_{d=1}^{n-1} \sum_{i=1}^{n-d} (R[i+1, i+d] - R[i, i+d-1] + 1) \\ &\geq \sum_{d=1}^{n-1} \sum_{i=1}^{n-d} 1 = \sum_{d=1}^{n-1} (n-d) = \frac{(n-1)n}{2} \geq \frac{1}{4}n^2 \text{ for } n \geq 2. \end{aligned}$$

Therefore, the time efficiency of the modified algorithm is in $\Theta(n^2)$.

9. Let a_1, \dots, a_n be a sorted list of n distinct keys, p_i be a known probability of searching for key a_i for $i = 1, 2, \dots, n$, and q_i be a known probability of searching (unsuccessfully) for a key between a_i and a_{i+1} for $i = 0, 1, \dots, n$ (with q_0 being a probability of searching for a key smaller than a_1 and q_n being a probability of searching for a key greater than a_n). It's convenient to associate unsuccessful searches with external nodes of a binary search tree (see Section 5.4). Repeating the derivation of equation (8.11) for such a tree yields the following recurrence for the expected number of key comparisons

$$C[i, j] = \min_{i \leq k \leq j} \{C[i, k-1] + C[k+1, j]\} + \sum_{s=i}^j p_s + \sum_{s=i-1}^j q_s \quad \text{for } 1 \leq i < j \leq n$$

with the initial condition

$$C[i, i] = p_i + q_{i-1} + q_i \quad \text{for } i = 1, \dots, n.$$

In all other respects, the algorithm remains the same.

10. a. Multiplying two matrices of dimensions α -by- β and β -by- γ by the definition-based algorithm requires $\alpha\beta\gamma$ multiplications. (There are $\alpha\gamma$ elements in the product, each requiring β multiplications to be computed.) If the dimensions of A_1 , A_2 , and A_3 are d_0 -by- d_1 , d_1 -by- d_2 , and d_2 -by- d_3 , respectively, then $(A_1 \cdot A_2) \cdot A_3$ will require

$$d_0 d_1 d_2 + d_0 d_2 d_3 = d_0 d_2 (d_1 + d_3)$$

multiplications, while $A_1 \cdot (A_2 \cdot A_3)$ will need

$$d_1 d_2 d_3 + d_0 d_1 d_3 = d_1 d_3 (d_0 + d_2)$$

multiplications. Here is a simple choice of specific values to make, say, the first of them be 1,000 times larger than the second:

$$d_0 = d_2 = 10^3, \quad d_1 = d_3 = 1.$$

b. Let $m(n)$ be the number of different ways to compute a chain product of n matrices $A_1 \cdot \dots \cdot A_n$. Any parenthesization of the chain will lead to multiplying, as the last operation, some product of the first k matrices ($A_1 \cdot \dots \cdot A_k$) and the last $n - k$ matrices ($A_{k+1} \cdot \dots \cdot A_n$). There are $m(k)$ ways to do the former, and there are $m(n - k)$ ways to do the latter. Hence, we have the following recurrence for the total number of ways to parenthesize the matrix chain of n matrices:

$$m(n) = \sum_{k=1}^{n-1} m(k)m(n-k) \quad \text{for } n > 1, \quad m(1) = 1.$$

Since parenthesizing a chain of n matrices for multiplication is very similar to constructing a binary tree of n nodes, it should come as no surprise that the above recurrence is very similar to the recurrence

$$b(n) = \sum_{k=0}^{n-1} b(k)b(n-1-k) \quad \text{for } n > 1, \quad b(0) = 1,$$

for the number of binary trees mentioned in Section 8.3. Nor is it surprising that their solutions are very similar, too: namely,

$$m(n) = b(n-1) \quad \text{for } n \geq 1,$$

where $b(n)$ is the number of binary trees with n nodes. Let us prove this assertion by mathematical induction. The basis checks immediately: $m(1) = b(0) = 1$. For the general case, let us assume that $m(k) = b(k-1)$ for all positive integers not exceeding some positive integer n (we're using the strong version of mathematical induction); we'll show that the equality holds for $n+1$ as well. Indeed,

$$\begin{aligned} m(n+1) &= \sum_{k=1}^n m(k)m(n+1-k) \\ &= [\text{using the induction's assumption}] \sum_{k=1}^n b(k-1)b(n-k) \\ &= [\text{substituting } l = k-1] \sum_{l=0}^{n-1} b(l)b(n-1-l) \\ &= [\text{see the recurrence for } b(n)] \quad b(n). \end{aligned}$$

c. Let $M[i, j]$ be the optimal (smallest) number of multiplications needed for computing $A_i \cdot \dots \cdot A_j$. If k is an index of the last matrix in the first factor of the last matrix product, then

$$M[i, j] = \max_{1 \leq k \leq j-1} \{M[i, k] + M[k+1, j] + d_{i-1}d_kd_j\} \text{ for } 1 \leq i < j \leq n,$$

$$M[i, i] = 0.$$

This recurrence, which is quite similar to the one for the optimal binary search tree problem, suggests filling the $n+1$ -by- $n+1$ table diagonal by diagonal as in the following algorithm:

Algorithm *MatrixChainMultiplication*($D[0..n]$)
//Solves matrix chain multiplication problem by dynamic programming
//Input: An array $D[0..n]$ of dimensions of n matrices
//Output: The minimum number of multiplications needed to multiply
//a chain of n matrices of the given dimensions and table $T[1..n, 1..n]$
//for obtaining an optimal order of the multiplications
for $i \leftarrow 1$ **to** n **do** $M[i, i] \leftarrow 0$
for $d \leftarrow 1$ **to** $n-1$ **do** //diagonal count
 for $i \leftarrow 1$ **to** $n-d$ **do**
 $j \leftarrow i+d$
 $minval \leftarrow \infty$
 for $k \leftarrow i$ **to** $j-1$ **do**
 $temp \leftarrow M[i, k] + M[k+1, j] + D[i-1] * D[k] * D[j]$
 if $temp < minval$
 $minval \leftarrow temp$
 $kmin \leftarrow k$
 $T[i, j] \leftarrow kmin$
return $M[1, n], T$

To find an optimal order to multiply the matrix chain, call *OptimalMultiplicationOrder*(1, n) below:

Algorithm *OptimalOrder*(i, j)
//Outputs an optimal order to multiply n matrices
//Input: Indices i and j of the first and last matrices in $A_i \dots A_j$ and
// table $T[1..n, 1..n]$ generated by *MatrixChainMultiplication*
//Output: $A_i \dots A_j$ parenthesized for optimal multiplication
if $i = j$
 print(" A_i ")
else
 $k \leftarrow T[i, j]$
 print("(")
 OptimalOrder(i, k)
 OptimalOrder($k+1, j$)
 print(")")

Exercises 8.4

1. a. Apply the bottom-up dynamic programming algorithm to the following instance of the knapsack problem:

item	weight	value
1	3	\$25
2	2	\$20
3	1	\$15
4	4	\$40
5	5	\$50

, capacity $W = 6$.

- b. How many different optimal subsets does the instance of part (a) have?
- c. In general, how can we use the table generated by the dynamic programming algorithm to tell whether there is more than one optimal subset for the knapsack problem's instance?
2. a. Write a pseudocode of the bottom-up dynamic programming algorithm for the knapsack problem.
- b. Write a pseudocode of the algorithm that finds the composition of an optimal subset from the table generated by the bottom-up dynamic programming algorithm for the knapsack problem.
3. For the bottom-up dynamic programming algorithm for the knapsack problem, prove that
 - a. its time efficiency is in $\Theta(nW)$.
 - b. its space efficiency is in $\Theta(nW)$.
 - c. the time needed to find the composition of an optimal subset from a filled dynamic programming table is in $O(n + W)$.
4. a. True or false: A sequence of values in a row of the dynamic programming table for an instance of the knapsack problem is always nondecreasing.
- b. True or false: A sequence of values in a column of the dynamic programming table for an instance of the knapsack problem is always nondecreasing?
5. Apply the memory function method to the instance of the knapsack problem given in Problem 1. Indicate the entries of the dynamic programming table that are: (i) never computed by the memory function method on this instance; (ii) retrieved without a recomputation.

6. Prove that the efficiency class of the memory function algorithm for the knapsack problem is the same as that of the bottom-up algorithm (see Problem 3).
7. \triangleright Write a pseudocode of a memory function for the optimal binary search tree problem. (You may limit your function to finding the smallest number of key comparisons in a successful search.)
8. Give two reasons why the memory function approach is unattractive for the problem of computing a binomial coefficient.
9. \triangleright Design a dynamic programming algorithm for the ***change-making problem***: given an amount n and unlimited quantities of coins of each of the denominations d_1, d_2, \dots, d_m , find the smallest number of coins that add up to n or indicate that the problem does not have a solution.
10. \triangleright Write a research report on one of the following well-known applications of dynamic programming:
 - a. finding the longest common subsequence in two sequences
 - b. optimal string editing
 - c. minimal triangulation of a polygon

Hints to Exercises 8.4

1. a. Use formulas (8.12) and (8.13) to fill in the appropriate table as it is done for another instance of the problem in this section.

b.-c. What would the equality of the two terms in

$$\max\{V[i-1, j], v_i + V[i-1, j-w_i]\}$$

mean?

2. a. Write a pseudocode to fill the table in Fig. 8.11 (say, row by row) by using formulas (8.12) and (8.13).

b. An algorithm for identifying an optimal subset is outlined in the section via an example.
3. How many values does the algorithm compute? How long does it take to compute one value? How many table cells need to be traversed to identify the composition of an optimal subset?
4. Use the definition of $V[i, j]$ to check whether it is always true that
 - a. $V[i, j-1] \leq V[i, j]$ for $1 \leq j \leq W$.
 - b. $V[i-1, j] \leq V[i, j]$ for $1 \leq i \leq n$.
5. Trace the calls of the function *MemoryKnapsack*(i, j) on the instance in question. (An application to another instance can be found in the section.)
6. The algorithm applies formula (8.12) to fill *some* of the table's cells. Why can we still assert that its efficiencies are in $\Theta(nW)$?
7. Modify the bottom-up algorithm of Section 8.3 in the same manner the bottom-up algorithm of Section 8.4 was modified there.
8. One of the reasons deals with the time efficiency, the other deals with the space efficiency.
9. Set up a recurrence for $C[i, j]$, the smallest number of coins in the problem's instance with coins of the first i denominations ($1 \leq i \leq m$) and amount j ($0 \leq j \leq n$). Use $+\infty$ as $C[i, j]$'s value for instances without a solution.
10. n/a

Solutions to Exercises 8.4

1. a.

		capacity j						
	i	0	1	2	3	4	5	6
	0	0	0	0	0	0	0	0
$w_1 = 3, v_1 = 25$	1	0	0	0	25	25	25	25
$w_2 = 2, v_2 = 20$	2	0	0	20	25	25	45	45
$w_3 = 1, v_3 = 15$	3	0	15	20	35	40	45	60
$w_4 = 4, v_4 = 40$	4	0	15	20	35	40	55	60
$w_5 = 5, v_5 = 50$	5	0	15	20	35	40	55	65

The maximal value of a feasible subset is $V[5, 6] = 65$. The optimal subset is {item 3, item 5}.

b.-c. The instance has a unique optimal subset in view of the following general property: An instance of the knapsack problem has a unique optimal solution if and only if the algorithm for obtaining an optimal subset, which retraces backward the computation of $V[n, W]$, encounters no equality between $V[i - 1, j]$ and $v_i + V[i - 1, j - w_i]$ during its operation.

2. a. **Algorithm** *DPKnapsack*($w[1..n], v[1..n], W$)

```
//Solves the knapsack problem by dynamic programming (bottom up)
//Input: Arrays  $w[1..n]$  and  $v[1..n]$  of weights and values of  $n$  items,
//       knapsack capacity  $W$ 
//Output: Table  $V[0..n, 0..W]$  that contains the value of an optimal
//        subset in  $V[n, W]$  and from which the items of an optimal
//        subset can be found
for  $i \leftarrow 0$  to  $n$  do  $V[i, 0] \leftarrow 0$ 
for  $j \leftarrow 1$  to  $W$  do  $V[0, j] \leftarrow 0$ 
for  $i \leftarrow 1$  to  $n$  do
    for  $j \leftarrow 1$  to  $W$  do
        if  $j - w[i] \geq 0$ 
             $V[i, j] \leftarrow \max\{V[i - 1, j], v[i] + V[i - 1, j - w[i]]\}$ 
        else  $V[i, j] \leftarrow V[i - 1, j]$ 
return  $V[n, W], V$ 
```

b. **Algorithm** *OptimalKnapsack*($w[1..n], v[1..n], W$)

```
//Finds the items composing an optimal solution to the knapsack problem
//Input: Arrays  $w[1..n]$  and  $v[1..n]$  of weights and values of  $n$  items,
//       knapsack capacity  $W$ , and table  $V[0..n, 0..W]$  generated by
//       the dynamic programming algorithm
//Output: List  $L[1..k]$  of the items composing an optimal solution
 $k \leftarrow 0$  //size of the list of items in an optimal solution
 $j \leftarrow W$  //unused capacity
```

```

for  $i \leftarrow n$  downto 1 do
  if  $V[i, j] > V[i - 1, j]$ 
     $k \leftarrow k + 1$ ;  $L[k] \leftarrow i$  //include item  $i$ 
     $j \leftarrow j - w[i]$ 
return  $L$ 

```

Note: In fact, we can also stop the algorithm as soon as j , the unused capacity of the knapsack, becomes 0.

3. The algorithm fills a table with $n + 1$ rows and $W + 1$ columns, spending $\Theta(1)$ time to fill one cell (either by applying (8.12) or (8.13)). Hence, its time efficiency and its space efficiency are in $\Theta(nW)$.

In order to identify the composition of an optimal subset, the algorithm repeatedly compares values at no more than two cells in a previous row. Hence, its time efficiency class is in $O(n)$.

4. Both assertions are true:

a. $V[i, j - 1] \leq V[i, j]$ for $1 \leq j \leq W$ is true because it simply means that the maximal value of a subset that fits into a knapsack of capacity $j - 1$ cannot exceed the maximal value of a subset that fits into a knapsack of capacity j .

b. $V[i - 1, j] \leq V[i, j]$ for $1 \leq i \leq n$ is true because it simply means that the maximal value of a subset of the first $i - 1$ items that fits into a knapsack of capacity j cannot exceed the maximal value of a subset of the first i items that fits into a knapsack of the same capacity j .

5. In the table below, the cells marked by a minus indicate the ones for which no entry is computed for the instance in question; the only nontrivial entry that is retrieved without recomputation is $(2, 1)$.

	i	<i>capacity j</i>						
		0	1	2	3	4	5	6
	0	0	0	0	0	0	0	0
$w_1 = 3, v_1 = 25$	1	0	0	0	25	25	25	25
$w_2 = 2, v_2 = 20$	2	0	0	20	-	-	45	45
$w_3 = 1, v_3 = 15$	3	0	15	20	-	-	-	60
$w_4 = 4, v_4 = 40$	4	0	15	-	-	-	-	60
$w_5 = 5, v_5 = 50$	5	0	-	-	-	-	-	65

6. Since some of the cells of a table with $n + 1$ rows and $W + 1$ columns are filled in constant time, both the time and space efficiencies are in

$O(nW)$. But all the entries of the table need to be initialized (unless virtual initialization is used); this puts them in $\Omega(nW)$. Hence, both efficiencies are in $\Theta(nW)$.

7. **Algorithm** *MFOptimalBST*(i, j)

```
//Returns the number of comparisons in a successful search in a BST
//Input: Indices  $i, j$  indicating the range of keys in a sorted list of  $n$  keys
//and an array  $P[1..n]$  of search probabilities used as a global variable
//Uses a global table  $C[1..n + 1, 0..n]$  initialized with a negative number
//Output: The average number of comparisons in the optimal BST
if  $j = i - 1$  return 0
if  $j = i$  return  $P[i]$ 
if  $C[i, j] < 0$ 
     $minval \leftarrow \infty$ 
    for  $k \leftarrow i$  to  $j$  do
         $temp \leftarrow MFOptimalBST(i, k - 1) + MFOptimalBST(k + 1, j)$ 
        if  $temp < minval$ 
             $minval \leftarrow temp$ 
     $sum \leftarrow 0$ ; for  $s \leftarrow i$  to  $j$  do  $sum \leftarrow sum + P[s]$ 
     $C[i, j] \leftarrow minval + sum$ 
return  $C[i, j]$ 
```

Note: The first two lines of this pseudocode can be eliminated by an appropriate initialization of table C .

8. For the problem of computing a binomial coefficient, we know in advance what cells of the table need to be computed. Therefore unnecessary computations can be avoided by the bottom-up dynamic programming algorithm as well. Also, using the memory function method requires $\Theta(nk)$ space, whereas the bottom-up algorithm needs only $\Theta(n)$ because the next row of the table can be written over its immediate predecessor.
9. Let $C[i, j]$ be the smallest number of coins in the problem's instance with coins of the first i denominations ($1 \leq i \leq m$) and amount j ($0 \leq j \leq n$). We will use $+\infty$ as $C[i, j]$'s value for instances without a solution. We have the following recurrence relation for $C[i, j]$:

$$\begin{aligned} C[i, j] &= \min\{[C[i - 1, j], 1 + C[i, j - d_i]]\} \text{ for } 1 \leq i \leq m, 1 \leq j \leq n \\ C[i, 0] &= 0 \text{ for } 1 \leq i \leq m. \end{aligned}$$

(In the recurrence above, we assume that $C[i - 1, j] = +\infty$ if $i = 1$ and $C[i, j - d_i] = +\infty$ if $j - d_i < 0$.) The following pseudocode fills the m -by- $(n + 1)$ dynamic programming table, which will contain the smallest number of coins in $C[m, n]$.

Algorithm *MinimumChange*($D[1..m], n$)
 //Solves the change-making problem by dynamic programming
 //Input: Array $D[1..m]$ of positive numbers (coin denominations) and
 // positive integer n (amount)
 //Output: The smallest number of coins adding up to n
 //Note: Assumes that $\infty + 1 = \infty$
for $i \leftarrow 1$ **to** m **do** $C[i, 0] = 0$
for $i \leftarrow 1$ **to** m **do**
 for $j \leftarrow 1$ **to** n **do**
 if $i = 1$
 if $j < D[1]$ $C[1, j] = \infty$
 else $C[1, j] = 1 + C[1, j - D[1]]$
 else // $i > 1$
 if $j < D[i]$ $C[i, j] = C[i - 1, j]$
 else $C[i, j] = \min\{C[i - 1, j], 1 + C[i, j - D[i]]\}$
return $C[m, n]$ // $C[m, n] = \infty$ signifies solution absence

The composition of an optimal solution can be easily obtained from the dynamic programming table by calling *OptimalCoins*($D[1..m], C[1..m, 0..n], m, n$) below:

Algorithm *OptimalCoins*($D[1..m], C[1..m, 0..n], i, j$)
 //Finds the coins composing an optimal change
 //Input: Arrays $D[1..m]$ of coin denominations, table $C[1..m, 0..n]$
 //generated by the dynamic programming algorithm for amount n ,
 // integers i ($1 \leq i \leq m$) and j ($1 \leq j \leq n$)
 //Output: List $L[1..i]$ of coins composing an optimal solution to
 //the change problem for the first i denominations and amount j
for $k \leftarrow 1$ **to** i **do** $L[k] \leftarrow 0$
if $C[i, j] < \infty$
while $i > 0$ **and** $j > 0$ **do**
 if $C[i, j] = C[i - 1, j]$
 $i \leftarrow i - 1$
 else
 $L[i] \leftarrow L[i] + 1$ //add one coin of denomination $D[i]$
 $j \leftarrow j - D[i]$
return L

10. n/a

This file contains the exercises, hints, and solutions for Chapter 9 of the book "Introduction to the Design and Analysis of Algorithms," 2nd edition, by A. Levitin. The problems that might be challenging for at least some students are marked by \triangleright ; those that might be difficult for a majority of students are marked by \blacktriangleright .

Exercises 9.1

1. Give an instance of the change-making problem for which the greedy algorithm does not yield an optimal solution.
2. Write a pseudocode of the greedy algorithm for the change-making problem, with an amount n and coin denominations $d_1 > d_2 > \dots > d_m$ as its input. What is the time efficiency class of your algorithm?
3. Consider the problem of scheduling n jobs of known durations t_1, \dots, t_n for execution by a single processor. The jobs can be executed in any order, one job at a time. You want to find a schedule that minimizes the total time spent by all the jobs in the system. (The time spent by one job in the system is the sum of the time spent by this job in waiting plus the time spent on its execution.)

Design a greedy algorithm for this problem. \triangleright Does the greedy algorithm always yield an optimal solution?

4. Design a greedy algorithm for the assignment problem (see Section 3.4). Does your greedy algorithm always yield an optimal solution?
5. *Bridge crossing revisited* Consider the generalization of the bridge crossing puzzle (Problem 2 in Exercises 1.2) in which we have $n > 1$ people whose bridge crossing times are t_1, t_2, \dots, t_n . All the other conditions of the problem remain the same: at most two people at the time can cross the bridge (and they move with the speed of the slower of the two) and they must carry with them the only flashlight the group has.


Design a greedy algorithm for this problem and find how long it will take to cross the bridge by using this algorithm. Does your algorithm yields a minimum crossing time for every instance of the problem? If it does—prove it, if it does not—find an instance with the smallest number of people for which this happens.

6. *Bachet-Fibonacci weighing problem* Find an optimal set of n weights $\{w_1, w_2, \dots, w_n\}$ so that it would be possible to weigh on a balance scale any integer load in the largest possible range from 1 to W , provided


a. \triangleright weights can be put only on the free cup of the scale.

b. \blacktriangleright weights can be put on both cups of the scale.

7. a. Apply Prim's algorithm to the following graph. Include in the priority queue all the vertices not already in the tree.



- b. Apply Prim's algorithm to the following graph. Include in the priority queue only the fringe vertices (the vertices not in the current tree which are adjacent to at least one tree vertex).



8. The notion of a minimum spanning tree is applicable to a connected weighted graph. Do we have to check a graph's connectivity before applying Prim's algorithm or can the algorithm do it by itself?
9. a. How can we use Prim's algorithm to find a spanning tree of a connected graph with no weights on its edges?
- b. Is it a good algorithm for this problem?
10. ▷ Prove that any weighted connected graph with distinct weights has exactly one minimum spanning tree.
11. Outline an efficient algorithm for changing an element's value in a min-heap. What is the time efficiency of your algorithm?

Hints to Exercises 9.1

1. As coin denominations for your counterexample, you may use, among a multitude of other possibilities, the ones mentioned in the text: $d_1 = 7$, $d_2 = 5$, $d_3 = 1$.
2. You may use integer divisions in your algorithm.
3. Considering the case of two jobs might help. Of course, after forming a hypothesis, you will have to either prove the algorithm's optimality for an arbitrary input or find a specific counterexample showing that it is not the case.
4. You can apply the greedy approach either to the entire cost matrix or to each of its rows (or columns).
5. Simply apply the greedy approach to the situation at hand. You may assume that $t_1 \leq t_2 \leq \dots \leq t_n$.
6. For both versions of the problem, it is not difficult to get to a hypothesis about the solution's form after considering the cases of $n = 1, 2$, and 3 . It is proving the solutions' optimality that is at the heart of this problem.
7. a. Trace the algorithm for the graph given. An example can be found in the text of the section.

b. After the next fringe vertex is added to the tree, add all the unseen vertices adjacent to it to the priority queue of fringe vertices.
8. Applying Prim's algorithm to a weighted graph that is not connected should help in answering this question.
9. a. Since Prim's algorithm needs weights on a graph's edges, some weights have to be assigned.

b. Do you know other algorithms that can solve this problem?
10. Strictly speaking, the wording of the question asks you to prove two things: the fact that at least one minimum spanning tree exists for any weighted connected graph and the fact that a minimum spanning tree is unique if all the weights are distinct numbers. The proof of the former stems from the obvious observation about finiteness of the number of spanning trees for a weighted connected graph. The proof of the latter can be obtained by repeating the correctness proof of Prim's algorithm with a minor adjustment at the end.
11. Consider two cases: the key's value was decreased (this is the case needed for Prim's algorithm) and the key's value was increased.

Solutions to Exercises 9.1

1. Here is one of many such instances: For the coin denominations $d_1 = 7$, $d_2 = 5$, $d_3 = 1$ and the amount $n = 10$, the greedy algorithm yields one coin of denomination 7 and three coins of denomination 1. The actual optimal solution is two coins of denomination 5.

2. **Algorithm** *Change*($n, D[1..m]$)

```
//Implements the greedy algorithm for the change-making problem
//Input: A nonnegative integer amount  $n$  and
//       a decreasing array of coin denominations  $D$ 
//Output: Array  $C[1..m]$  of the number of coins of each denomination
//        in the change or the "no solution" message
for  $i \leftarrow 1$  to  $m$  do
     $C[i] \leftarrow \lfloor n/D[i] \rfloor$ 
     $n \leftarrow n \bmod D[i]$ 
if  $n = 0$  return  $C$ 
else return "no solution"
```

The algorithm's time efficiency is in $\Theta(m)$. (We assume that integer divisions take a constant time no matter how big dividends are.) Note also that if we stop the algorithm as soon as the remaining amount becomes 0, the time efficiency will be in $O(m)$.

3. a. Sort the jobs in nondecreasing order of their execution times and execute them in that order.
 b. Yes, this greedy algorithm always yields an optimal solution. Indeed, for any ordering (i.e., permutation) of the jobs i_1, i_2, \dots, i_n , the total time in the system is given by the formula

$$t_{i_1} + (t_{i_1} + t_{i_2}) + \dots + (t_{i_1} + t_{i_2} + \dots + t_{i_n}) = nt_{i_1} + (n-1)t_{i_2} + \dots + t_{i_n}.$$

Thus, we have a sum of numbers $n, n-1, \dots, 1$ multiplied by "weights" t_1, t_2, \dots, t_n assigned to the numbers in some order. To minimize such a sum, we have to assign smaller t 's to larger numbers. In other words, the jobs should be executed in nondecreasing order of their execution times.

Here is a more formal proof of this fact. We will show that if jobs are executed in some order i_1, i_2, \dots, i_n , in which $t_{i_k} > t_{i_{k+1}}$ for some k , then the total time in the system for such an ordering can be decreased. (Hence, no such ordering can be an optimal solution.) Let us consider the other job ordering, which is obtained by swapping the jobs k and $k+1$. Obviously, the time in the systems will remain the same for all but these two

jobs. Therefore, the difference between the total time in the system for the new ordering and the one before the swap will be

$$\begin{aligned} & [(\sum_{j=1}^{k-1} t_{i_j} + t_{i_{k+1}}) + (\sum_{j=1}^{k-1} t_{i_j} + t_{i_{k+1}} + t_{i_k})] - [(\sum_{j=1}^{k-1} t_{i_j} + t_{i_k}) + (\sum_{j=1}^{k-1} t_{i_j} + t_{i_k} + t_{i_{k+1}})] \\ &= t_{i_{k+1}} - t_{i_k} < 0. \end{aligned}$$

4. a. The all-matrix version: Repeat the following operation n times. Select the smallest element in the unmarked rows and columns of the cost matrix and then mark its row and column.

The row-by-row version: Starting with the first row and ending with the last row of the cost matrix, select the smallest element in that row which is not in a previously marked column. After such an element is selected, mark its column to prevent selecting another element from the same column.

- b. Neither of the versions always yields an optimal solution. Here is a simple counterexample:

$$C = \begin{bmatrix} 1 & 2 \\ 2 & 100 \end{bmatrix}$$

5. Repeat the following step $n-2$ times: Send to the other side the pair of two fastest remaining persons and then return the flashlight with the fastest person. Finally, send the remaining two people together. Assuming that $t_1 \leq t_2 \leq \dots \leq t_n$, the total crossing time will be equal to

$$(t_2+t_1)+(t_3+t_1)+\dots+(t_{n-1}+t_1)+t_n = \sum_{i=2}^n t_i + (n-2)t_1 = \sum_{i=1}^n t_i + (n-3)t_1.$$

Note: For an algorithm that always yields a minimal crossing time, see Günter Rote, “Crossing the Bridge at Night,” *EATCS Bulletin*, vol. 78 (October 2002), 241–246.

The solution to the instance of Problem 2 in Exercises 1.2 shows that the greedy algorithm doesn’t always yield the minimal crossing time for $n > 3$. No smaller counterexample can be given as a simple exhaustive check for $n = 3$ demonstrates. (The obvious solution for $n = 2$ is the one generated by the greedy algorithm as well.)

6. a. Let's apply the greedy approach to the first few instances of the problem in question. For $n = 1$, we have to use $w_1 = 1$ to balance weight 1. For $n = 2$, we simply add $w_2 = 2$ to balance the first previously unattainable weight of 2. The weights $\{1, 2\}$ can balance every integral weights up to their sum 3. For $n = 3$, in the spirit of greedy thinking, we take the next previously unattainable weight: $w_3 = 4$. The three weights $\{1, 2, 4\}$ allow to weigh any integral load l between 1 and their sum 7, with l 's binary expansion indicating the weights needed for load l :

load l	1	2	3	4	5	6	7
l 's binary expansion	1	10	11	100	101	110	111
weights for load l	1	2	2+1	4	4+1	4+2	4+2+1

Generalizing these observations, we should hypothesize that for any positive integer n the set of consecutive powers of 2 $\{w_i = 2^{i-1}, i = 1, 2, \dots, n\}$ makes it possible to balance every integral load in the largest possible range, which is up to and including $\sum_{i=1}^n 2^{i-1} = 2^n - 1$. The fact that every integral weight l in the range $1 \leq l \leq 2^n - 1$ can be balanced with this set of weights follows immediately from the binary expansion of l , which yields the weights needed for weighing l . (Note that we can obtain the weights needed for a given load l by applying to it the greedy algorithm for the change-making problem with denominations $d_i = 2^{i-1}, i = 1, 2, \dots, n$.)

In order to prove that no set of n weights can cover a larger range of consecutive integral loads, it will suffice to note that there are just $2^n - 1$ nonempty selections of n weights and, hence, no more than $2^n - 1$ sums they yield. Therefore, the largest range of consecutive integral loads they can cover cannot exceed $2^n - 1$.

[Alternatively, to prove that no set of n weights can cover a larger range of consecutive integral loads, we can prove by induction on i that if any multiset of n weights $\{w_i, i = 1, \dots, n\}$ —which we can assume without loss of generality to be sorted in nondecreasing order—can balance every integral load starting with 1, then $w_i \leq 2^{i-1}$ for $i = 1, 2, \dots, n$. The basis checks out immediately: w_1 must be 1, which is equal to 2^{1-1} . For the general case, assume that $w_k \leq 2^{k-1}$ for every $1 \leq k < i$. The largest weight the first $i - 1$ weights can balance is $\sum_{k=1}^{i-1} w_k \leq \sum_{k=1}^{i-1} 2^{k-1} = 2^{i-1} - 1$. If w_i were larger than 2^{i-1} , then this load could have been balanced neither with the first $i - 1$ weights (which are too light even taken together) nor with the weights $w_i \leq \dots \leq w_n$ (which are heavier than 2^{i-1} even individually). Hence, $w_i \leq 2^{i-1}$, which completes the proof by induction. This immediately implies that no n weights can balance every integral load up to the upper limit larger than $\sum_{i=1}^n w_i \leq \sum_{i=1}^n 2^{i-1} = 2^n - 1$, the limit attainable with the consecutive powers of 2 weights.]

- b. If weights can be put on both cups of the scale, then a larger range can

be reached with n weights for $n > 1$. (For $n = 1$, the single weight still needs to be 1, of course.) The weights $\{1, 3\}$ enable weighing of every integral load up to 4; the weights $\{1, 3, 9\}$ enable weighing of every integral load up to 13, and, in general, the weights $\{w_i = 3^{i-1}, i = 1, 2, \dots, n\}$ enable weighing of every integral load up to and including their sum of $\sum_{i=1}^n 3^{i-1} = (3^n - 1)/2$. A load's expansion in the ternary system indicates the weights needed. If the ternary expansion contains only 0's and 1's, the load requires putting the weights corresponding to the 1's on the opposite cup of the balance. If the ternary expansion of load l , $l \leq (3^n - 1)/2$, contains one or more 2's, we can replace each 2 by $(3-1)$ to represent it in the form

$$l = \sum_{i=1}^n \beta_i 3^{i-1}, \text{ where } \beta_i \in \{0, 1, -1\}, \quad n = \lceil \log_3(l+1) \rceil.$$

In fact, every positive integer can be uniquely represented in this form, obtained from its ternary expansion as described above. For example,


$$\begin{aligned} 5 &= 12_3 = 1 \cdot 3^1 + 2 \cdot 3^0 = 1 \cdot 3^1 + (3-1) \cdot 3^0 = 2 \cdot 3^1 - 1 \cdot 3^0 \\ &= (3-1) \cdot 3^1 - 1 \cdot 3^0 = 1 \cdot 3^2 - 1 \cdot 3^1 - 1 \cdot 3^0. \end{aligned}$$

(Note that if we start with the rightmost 2, after a simplification, the new rightmost 2, if any, will be at some position to the left of the starting one. This proves that after a finite number of such replacements, we will be able to eliminate all the 2's.) Using the representation $l = \sum_{i=1}^n \beta_i 3^{i-1}$, we can weigh load l by placing all the weights $w_i = 3^{i-1}$ for negative β_i 's along with the load on one cup of the scale and all the weights $w_i = 3^{i-1}$ for positive β_i 's on the opposite cup.

Now we'll prove that no set of n weights can cover a larger range of consecutive integral loads than $(3^n - 1)/2$. Each of the n weights can be either put on the left cup of the scale, or put on the right cup, or not to be used at all. Hence, there are $3^n - 1$ possible arrangements of the weights on the scale, with each of them having its mirror image (where all the weights are switched to the opposite pan of the scale). Eliminating this symmetry, leaves us with

just $(3^n - 1)/2$ arrangements, which can weight at most $(3^n - 1)/2$ different integral loads. Therefore, the largest range of consecutive integral loads they can cover cannot exceed $(3^n - 1)/2$.


7. a. Apply Prim's algorithm to the following graph:



Tree vertices	Priority queue of remaining vertices
a(-,-)	b(a,5) c(a,7) d(a,∞) e(a,2)
e(a,2)	b(e,3) c(e,4) d(e,5)
b(e,3)	c(e,4) d(e,5)
c(e,4)	d(c,4)
d(c,4)	

The minimum spanning tree found by the algorithm comprises the edges *ae*, *eb*, *ec*, and *cd*.

b. Apply Prim's algorithm to the following graph:



Tree vertices	Priority queue of fringe vertices
a(-,-)	b(a,3) c(a,5) d(a,4)
b(a,3)	c(a,5) d(a,4) e(b,3) f(b,6)
e(b,3)	c(a,5) d(e,1) f(e,2) i(e,4)
d(e,1)	c(d,2) f(e,2) i(e,4) h(d,5)
c(d,2)	f(e,2) i(e,4) h(d,5) g(c,4)
f(e,2)	i(e,4) h(d,5) g(c,4) j(f,5)
i(e,4)	h(d,5) g(c,4) j(i,3) l(i,5)
j(i,3)	h(d,5) g(c,4) l(i,5)
g(c,4)	h(g,3) l(i,5) k(g,6)
h(g,3)	l(i,5) k(g,6)
l(i,5)	k(g,6)
k(g,6)	

The minimum spanning tree found by the algorithm comprises the edges *ab*, *be*, *ed*, *dc*, *ef*, *ei*, *ij*, *cg*, *gh*, *il*, *gk*.

8. There is no need to check the graph's connectivity because Prim's algorithm can do it itself. If the algorithm reaches all the graph's vertices (via edges of finite lengths), the graph is connected, otherwise, it is not.
9. a. The simplest and most logical solution is to assign all the edge weights to 1.

- b. Applying a depth-first search (or breadth-first search) traversal to get a depth-first search tree (or a breadth-first search tree), is conceptually simpler and for sparse graphs represented by their adjacency lists faster.
10. The number of spanning trees for any weighted connected graph is a positive finite number. (At least one spanning tree exists, e.g., the one obtained by a depth-first search traversal of the graph. And the number of spanning trees must be finite because any such tree comprises a subset of edges of the finite set of edges of the given graph.) Hence, one can always find a spanning tree with the smallest total weight among the finite number of the candidates.


Let's prove now that the minimum spanning tree is unique if all the weights are distinct. We'll do this by contradiction, i.e., by assuming that there exists a graph $G = (V, E)$ with all distinct weights but with more than one minimum spanning tree. Let $e_1, \dots, e_{|V|-1}$ be the list of edges composing the minimum spanning tree T_P obtained by Prim's algorithm with some specific vertex as the algorithm's starting point and let T' be another minimum spanning tree. Let $e_i = (v, u)$ be the first edge in the list $e_1, \dots, e_{|V|-1}$ of the edges of T_P which is not in T' (if $T_P \neq T'$, such edge must exist) and let (v, u') be an edge of T' connecting v with a vertex not in the subtree T_{i-1} formed by $\{e_1, \dots, e_{i-1}\}$ (if $i = 1$, T_{i-1} consists of vertex v only). Similarly to the proof of Prim's algorithm's correctness, let us replace (v, u') by $e_i = (v, u)$ in T' . It will create another spanning tree, whose weight is smaller than the weight of T' because the weight of $e_i = (v, u)$ is smaller than the weight of (v, u') . (Since e_i was chosen by Prim's algorithm, its weight is the smallest among all the weights on the edges connecting the tree vertices of the subtree T_{i-1} and the vertices adjacent to it. And since all the weights are distinct, the weight of (v, u') must be strictly greater than the weight of $e_i = (v, u)$.) This contradicts the assumption that T' was a minimum spanning tree.

11. If a key's value in a min-heap was decreased, it may need to be pushed up (via swaps) along the chain of its ancestors until it is smaller than or equal to its parent or reaches the root. If a key's value in a min-heap was increased, it may need to be pushed down by swaps with the smaller of its current children until it is smaller than or equal to its children or reaches a leaf. Since the height of a min-heap with n nodes is equal to $\lfloor \log_2 n \rfloor$ (by the same reason the height of a max-heap is given by this formula—see Section 6.4), the operation's efficiency is in $O(\log n)$. (Note: The old value of the key in question need not be known, of course. Comparing the new value with that of the parent and, if the min-heap condition holds, with the smaller of the two children, will suffice.)


Exercises 9.2

1. Apply Kruskal's algorithm to find a minimum spanning tree of the following graphs.

a.



b.



2. Indicate whether the following statements are true or false:

- a. If e is a minimum-weight edge in a connected weighted graph, it must be among edges of at least one minimum spanning tree of the graph.
 - b. If e is a minimum-weight edge in a connected weighted graph, it must be among edges of each minimum spanning tree of the graph.
 - c. If edge weights of a connected weighted graph are all distinct, the graph must have exactly one minimum spanning tree.
 - d. If edge weights of a connected weighted graph are not all distinct, the graph must have more than one minimum spanning tree.
3. What changes, if any, need to be made in algorithm *Kruskal* to make it find a **minimum spanning forest** for an arbitrary graph? (A minimum spanning forest is a forest whose trees are minimum spanning trees of the graph's connected components.)


4. Will either Kruskal's or Prim's algorithm work correctly on graphs that have negative edge weights?
5. Design an algorithm for finding a *maximum spanning tree*—a spanning tree with the largest possible edge weight—of a weighted connected graph.
6. Rewrite the pseudocode of Kruskal's algorithm in terms of the operations of the disjoint subsets' ADT.
7. \triangleright Prove the correctness of Kruskal's algorithm.
8. Prove that the time efficiency of $find(x)$ is in $O(\log n)$ for the union-by-size version of quick union.
9. Find at least two Web sites with animations of Kruskal's and Prim's algorithms. Discuss their merits and demerits..
10. Design and conduct an experiment to empirically compare the efficiencies of Prim's and Kruskal's algorithms on random graphs of different sizes and densities.
11. \blacktriangleright *Steiner tree* Four villages are located at the vertices of a unit square in the Euclidean plane. You are asked to connect them by the shortest network of roads so that there is a path between every pair of the villages along those roads. Find such a network.

Hints to Exercises 9.2

1. Trace the algorithm for the given graphs the same way it is done for another input in the section.
2. Two of the four assertions are true, the other two are false.
3. Applying Kruskal's algorithm to a disconnected graph should help to answer the question.
4. The answer is the same for both algorithms. If you believe that the algorithms work correctly on graphs with negative weights, prove this assertion; if you believe this is not to be the case, give a counterexample for each algorithm.
5. Is the general trick of transforming maximization problems to their minimization counterparts (see Section 6.6) applicable here?
6. Substitute the three operations of the disjoint subsets' ADT—*makeset*(x), *find*(x), and *union*(x, y)—in the appropriate places of the pseudocode given in the section.
7. Follow the plan used in Section 9.1 to prove the correctness of Prim's algorithm.
8. The argument is very similar to the one made in the section for the union-by-size version of quick find.
9. You may want to take advantage of the list of desirable characteristics in algorithm visualizations, which is given in Section 2.7.
10. n/a
11. The question is not trivial because introducing extra points (called *Steiner points*) may make the total length of the network smaller than that of a minimum spanning tree of the square. Solving first the problem for three equidistant points might give you an indication how a solution to the problem in question could look like.

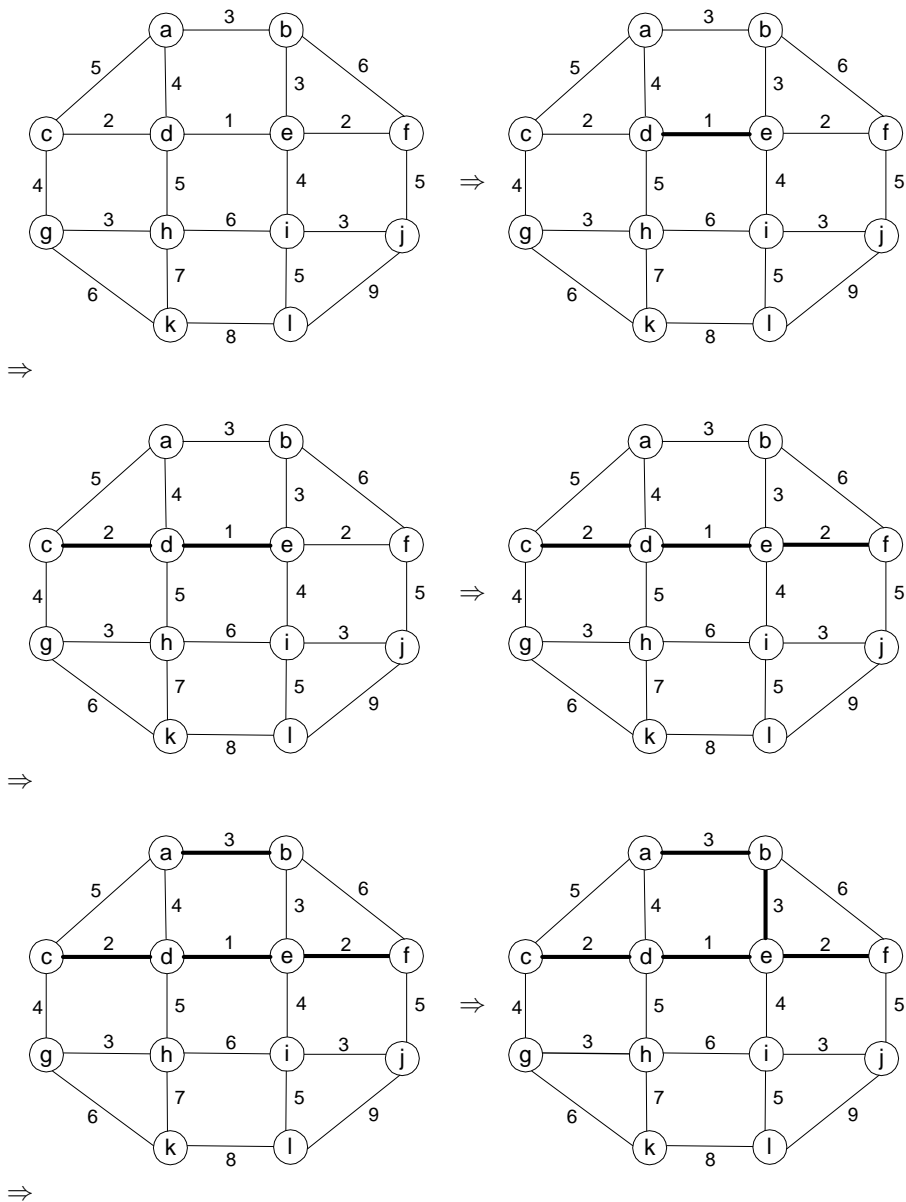
Solutions to Exercises 9.2


1. a.



Tree edges	Sorted list of edges (selected edges are shown in bold)	Illustration
	bc ₁ de ₂ bd ₃ cd ₄ ab ₅ ad ₆ ce ₆	
bc ₁	bc ₁ de ₂ bd ₃ cd ₄ ab ₅ ad ₆ ce ₆	
de ₂	bc ₁ de ₂ bd ₃ cd ₄ ab ₅ ad ₆ ce ₆	
bd ₃	bc ₁ de ₂ bd ₃ cd ₄ ab ₅ ad ₆ ce ₆	
ab ₅		

b.






2. a. True. (Otherwise, Kruskal's algorithm would be invalid.)

b. False. As a simple counterexample, consider a complete graph with three vertices and the same weight on its three edges

c. True (Problem 10 in Exercises 9.1).

- d. False (see, for example, the graph of Problem 1a).
3. Since the number of edges in a minimum spanning forest of a graph with $|V|$ vertices and $|C|$ connected components is equal to $|V| - |C|$ (this formula is a simple generalization of $|E| = |V| - 1$ for connected graphs), $Kruskal(G)$ will never get to $|V| - 1$ tree edges unless the graph is connected. A simple remedy is to replace the loop **while** $ecounter < |V| - 1$ with **while** $k < |E|$ to make the algorithm stop after exhausting the sorted list of its edges.
4. Both algorithms work correctly for graphs with negative edge weights. One way of showing this is to add to all the weights of a graph with negative weights some large positive number. This makes all the new weights positive, and one can “translate” the algorithms’ actions on the new graph to the corresponding actions on the old one. Alternatively, you can check that the proofs justifying the algorithms’ correctness do not depend on the edge weights being nonnegative.
5. Replace each weight $w(u, v)$ by $-w(u, v)$ and apply any minimum spanning tree algorithm that works on graphs with arbitrary weights (e.g., Prim’s or Kruskal’s algorithm) to the graph with the new weights.
6. **Algorithm** $Kruskal(G)$
 //Kruskal’s algorithm with explicit disjoint-subsets operations
 //Input: A weighted connected graph $G = \langle V, E \rangle$
 //Output: E_T , the set of edges composing a minimum spanning tree of G
 sort E in nondecreasing order of the edge weights $w(e_{i_1}) \leq \dots \leq w(e_{i_{|E|}})$
for each vertex $v \in V$ *make*(v)
 $E_T \leftarrow \emptyset$; $ecounter \leftarrow 0$ //initialize the set of tree edges and its size
 $k \leftarrow 0$ //the number of processed edges
while $ecounter < |V| - 1$
 $k \leftarrow k + 1$
 if $find(u) \neq find(v)$ // u, v are the endpoints of edge e_{i_k}
 $E_T \leftarrow E_T \cup \{e_{i_k}\}$; $ecounter \leftarrow ecounter + 1$
 union(u, v)
return E_T
7. Let us prove by induction that each of the forests F_i , $i = 0, \dots, |V| - 1$, of Kruskal’s algorithm is a part (i.e., a subgraph) of some minimum spanning tree. (This immediately implies, of course, that the last forest in the sequence, $F_{|V|-1}$, is a minimum spanning tree itself. Indeed, it contains all vertices of the graph, and it is connected because it is both acyclic and has $|V| - 1$ edges.) The basis of the induction is trivial, since F_0 is

made up of $|V|$ single-vertex trees and therefore must be a subgraph of any spanning tree of the graph. For the inductive step, let us assume that F_{i-1} is a subgraph of some minimum spanning tree T . We need to prove that F_i , generated from F_{i-1} by Kruskal's algorithm, is also a part of a minimum spanning tree. We prove this by contradiction by assuming that no minimum spanning tree of the graph can contain F_i . Let $e_i = (v, u)$ be the minimum weight edge added by Kruskal's algorithm to forest F_{i-1} to obtain forest F_i . (Note that vertices v and u must belong to different trees of F_{i-1} —otherwise, edge (v, u) would've created a cycle.) By our assumption, e_i cannot belong to T . Therefore, if we add e_i to T , a cycle must be formed (see the figure below). In addition to edge $e_i = (v, u)$, this cycle must contain another edge (v', u') connecting a vertex v' in the same tree of F_{i-1} as v to a vertex u' not in that tree. (It is possible that v' coincides with v or u' coincides with u but not both.) If we now delete the edge (v', u') from this cycle, we will obtain another spanning tree of the entire graph whose weight is less than or equal to the weight of T since the weight of e_i is less than or equal to the weight of (v', u') . Hence, this spanning tree is a minimum spanning tree, which contradicts the assumption that no minimum spanning tree contains F_i . This completes the correctness proof of Kruskal's algorithm.




8. In the *union-by-size* version of *quick-union*, each vertex starts at depth 0 of its own tree. The depth of a vertex increases by 1 when the tree it is in is attached to a tree with at least as many nodes during a union operation. Since the total number of nodes in the new tree containing the node is at least twice as much as in the old one, the number of such increases cannot exceed $\log_2 n$. Therefore the height of any tree (which is the largest depth of the tree's nodes) generated by a legitimate sequence of unions will not exceed $\log_2 n$. Hence, the efficiency of $find(x)$ is in $O(\log n)$ because $find(x)$ traverses the pointer chain from the x 's node to the tree's root.

9. n/a

10. n/a


11. The minimum Steiner tree that solves the problem is shown below. (The other solution can be obtained by rotating the figure 90° .)




A popular discussion of Steiner trees can be found in “Last Recreations: Hydras, Eggs, and Other Mathematical Mystifications” by Martin Gardner. In general, no polynomial time algorithm is known for finding a minimum Steiner tree; moreover, the problem is known to be *NP*-hard (see Section 11.3). For the state-of-the-art information, see, e.g., The Steiner Tree Page at <http://ganley.org/steiner/>.

Exercises 9.3

- Explain what adjustments if any need to be made in Dijkstra's algorithm and/or in an underlying graph to solve the following problems.
 - Solve the single-source shortest-paths problem for directed weighted graphs.
 - Find a shortest path between two given vertices of a weighted graph or digraph. (This variation is called the *single-pair shortest-path problem*.)
 - Find the shortest paths to a given vertex from each other vertex of a weighted graph or digraph. (This variation is called the *single-destination shortest-paths problem*.)
 - Solve the single-source shortest-path problem in a graph with nonnegative numbers assigned to its vertices (and the length of a path defined as the sum of the vertex numbers on the path).
- Solve the following instances of the single-source shortest-paths problem with vertex *a* as the source:
 -



b.



- Give a counterexample that shows that Dijkstra's algorithm may not work for a weighted connected graph with negative weights.

4. Let T be a tree constructed by Dijkstra's algorithm in the process of solving the single-source shortest-path problem for a weighted connected graph G .
 - a. True or false: T is a spanning tree of G ?
 - b. True or false: T is a minimum spanning tree of G ?
5. Write a pseudocode of a simpler version of Dijkstra's algorithm that finds only the distances (i.e., the lengths of shortest paths but not shortest paths themselves) from a given vertex to all other vertices of a graph represented by its weight matrix.
6. \triangleright Prove the correctness of Dijkstra's algorithm for graphs with positive weights.
7. Design a linear-time algorithm for solving the single-source shortest-paths problem for dags (directed acyclic graphs) represented by their adjacency lists.
8. Design an efficient algorithm for finding the length of a longest path in a dag. (This problem is important because it determines a lower bound on the total time needed for completing a project composed of precedence-constrained tasks.)
9. *Shortest-path modeling* Assume you have a model of a weighted connected graph made of balls (representing the vertices) connected by strings of appropriate lengths (representing the edges).
 - a. Describe how you can solve the single-pair shortest-path problem with this model.
 - b. Describe how you can solve the single-source shortest-paths problem with this model.
10. Revisit Problem 6 in Exercises 1.3 about determining the best route for a subway passenger to take from one designated station to another in a well-developed subway system like those in Washington, DC and London, UK. Write a program for this task.


Hints to Exercises 9.3

1. One of the questions requires no changes in either the algorithm or the graph; the others require simple adjustments.
2. Just trace the algorithm on the given graphs the same way it is done for an example in the section.
3. The nearest vertex does not have to be adjacent to the source in such a graph.
4. Only one of the assertions is correct. Find a small counterexample for the other.
5. Simplify the pseudocode given in the section by implementing the priority queue as an unordered array and ignoring the parental labeling of vertices.
6. Prove it by induction on the number of vertices included in the tree constructed by the algorithm.
7. Topologically sort the dag's vertices first.
8. Topologically sort the dag's vertices first.
9. Take advantage of the ways of thinking used in geometry and physics.
10. Before you embark on implementing a shortest-path algorithm, you have to decide what criterion determines the "best route". Of course, it would be highly desirable to have a program asking the user which of several possible criteria he or she wants applied.

Solutions to Exercises 9.3

1. a. It will suffice to take into account edge directions in processing adjacent vertices.
- b. Start the algorithm at one of the given vertices and stop it as soon as the other vertex is added to the tree.
- c. If the given graph is undirected, solve the single-source problem with the destination vertex as the source and reverse all the paths obtained in the solution. If the given graph is directed, reverse all its edges first, solve the single-source problem for the new digraph with the destination vertex as the source, and reverse the direction of all the paths obtained in the solution.
- d. Create a new graph by replacing every vertex v with two vertices v' and v'' connected by an edge whose weight is equal to the given weight of vertex v . All the edges entering and leaving v in the original graph will enter v' and leave v'' in the new graph, respectively. Assign the weight of 0 to each original edge. Applying Dijkstra's algorithm to the new graph will solve the problem.

2. a.




Tree vertices	Remaining vertices			
a(-,0)	b(-,∞)	c(-,∞)	d(a,7)	e(-,∞)
d(a,7)	b(d,7+2)	c(d,7+5)	e(-,∞)	
b(d,9)	c(d,12)	e(-,∞)		
c(d,12)	e(c,12+6)			
e(c,18)				

The shortest paths (identified by following nonnumeric labels backwards from a destination vertex to the source) and their lengths are:


- from a to d : $a - d$ of length 7
- from a to b : $a - d - b$ of length 9
- from a to c : $a - d - c$ of length 12
- from a to e : $a - d - c - e$ of length 18

b.



Tree vertices	Fringe vertices	Shortest paths
a(-,0)	b(a,3) c(a,5) d(a,4)	to b: $a - b$ of length 3
b(a,3)	c(a,5) d(a,4) e(b,3+3) f(b,3+6)	to d: $a - d$ of length 4
d(a,4)	c(a,5) e(d,4+1) f(a,9) h(d,4+5)	to c: $a - c$ of length 5
c(a,5)	e(d,5) f(a,9) h(d,9) g(c,5+4)	to e: $a - d - e$ of length 5
e(d,5)	f(e,5+2) h(d,9) g(c,9) i(e,5+4)	to f: $a - d - e - f$ of length 7
f(e,7)	h(d,9) g(c,9) i(e,9) j(f,7+5)	to h: $a - d - h$ of length 9
h(d,9)	g(c,9) i(e,9) j(f,12) k(h,9+7)	to g: $a - c - g$ of length 9
g(c,9)	i(e,9) j(f,12) k(g,9+6)	to i: $a - d - e - i$ of length 9
i(e,9)	j(f,12) k(g,15) l(i,9+5)	to j: $a - d - e - f$ of length 12
j(f,12)	k(g,15) l(i,14)	to l: $a - d - e - i - l$ of length 14
l(i,14)	k(g,15)	to k: $a - c - g - k$ of length 15
k(g,15)		


3. Consider, for example, the graph



As the shortest path from a to b , Dijkstra's algorithm yields $a - b$ of length 2, which is longer than $a - c - b$ of length 1.

4. a. True: On each iteration, we add to a previously constructed tree an edge connecting a vertex in the tree to a vertex that is not in the tree. So, the resulting structure must be a tree. And, after the last operation, it includes all the vertices of the graph. Hence, it's a spanning tree.

b. False. Here is a simple counterexample:



With vertex a as the source, Dijkstra's algorithm yields, as the shortest path tree, the tree composed of edges (a, b) and (a, c) . The graph's minimum spanning tree is composed of (a, b) and (b, c) .

5. **Algorithm** *SimpleDijkstra*($W[0..n-1, 0..n-1], s$)
 //Input: A matrix of nonnegative edge weights W and
 // integer s between 0 and $n-1$ indicating the source
 //Output: An array $D[0..n-1]$ of the shortest path lengths
 // from s to all vertices
for $i \leftarrow 0$ **to** $n-1$ **do**
 $D[i] \leftarrow \infty$; *treeflag*[i] \leftarrow **false**
 $D[s] \leftarrow 0$
for $i \leftarrow 0$ **to** $n-1$ **do**
 $dmin \leftarrow \infty$
 for $j \leftarrow 0$ **to** $n-1$ **do**
 if not *treeflag*[j] **and** $D[j] < dmin$
 $jmin \leftarrow j$; $dmin \leftarrow D[jmin]$
 treeflag[$jmin$] \leftarrow **true**
 for $j \leftarrow 0$ **to** $n-1$ **do**
 if not *treeflag*[j] **and** $dmin + W[jmin, j] < \infty$
 $D[j] \leftarrow dmin + W[jmin, j]$
return D

6. We will prove by induction on the number of vertices i in tree T_i constructed by Dijkstra's algorithm that this tree contains i closest vertices to source s (including the source itself), for each of which the tree path from s to v is a shortest path of length d_v . For $i = 1$, the assertion is obviously true for the trivial path from the source to itself. For the general step, assume that it is true for the algorithm's tree T_i with i vertices. Let v_{i+1} be the vertex added next to the tree by the algorithm. All the vertices on a shortest path from s to v_{i+1} preceding v_{i+1} must be in T_i because they are closer to s than v_{i+1} . (Otherwise, the first vertex on the path from s to v_{i+1} that is not in T_i would've been added to T_i instead of v_{i+1} .) Hence, the $(i+1)$ st closest vertex can be selected as the algorithm does: by minimizing the sum of d_v (the shortest distance from s to $v \in T_i$ by the assumption of the induction) and the length of the edge from v to an adjacent vertex not in the tree.

7. **Algorithm** *DagShortestPaths*(G, s)
 //Solves the single-source shortest paths problem for a dag
 //Input: A weighted dag $G = \langle V, E \rangle$ and its vertex s
 //Output: The length d_v of a shortest path from s to v and
 // its penultimate vertex p_v for every vertex v in V
 topologically sort the vertices of G
for every vertex v **do**
 $d_v \leftarrow \infty$; $p_v \leftarrow \text{null}$
 $d_s \leftarrow 0$
for every vertex v taken in topological order **do**
 for every vertex u adjacent to v **do**
 if $d_v + w(v, u) < d_u$
 $d_u \leftarrow d_v + w(v, u)$; $p_u \leftarrow v$

Topological sorting can be done in $\Theta(|V|+|E|)$ time (see Section 5.3). The distance initialization takes $\Theta(|V|)$ time. The innermost loop is executed for every edge of the dag. Hence, the total running time is in $\Theta(|V|+|E|)$.

8. **Algorithm** *DagLongestPath*(G)
 //Finds the length of a longest path in a dag
 //Input: A weighted dag $G = \langle V, E \rangle$
 //Output: The length of its longest path $dmax$
 topologically sort the vertices of G
for every vertex v **do**
 $d_v \leftarrow 0$ //the length of the longest path to v seen so far
for every vertex v taken in topological order **do**
 for every vertex u adjacent to v **do**
 if $d_v + w(v, u) > d_u$
 $d_u \leftarrow d_v + w(v, u)$
 $dmax \leftarrow 0$
for every vertex v **do**
 if $d_v > dmax$
 $dmax \leftarrow d_v$
return $dmax$

9. a. Take the two balls representing the two singled out vertices in two hands and stretch the model to get the shortest path in question as a straight line between the two ball-vertices.

b. Hold the ball representing the source in one hand and let the rest of the model hang down: The force of gravity will make the shortest path to each of the other balls be on a straight line down.

10. n/a

Exercises 9.4

1. a. Construct a Huffman code for the following data:

character	A	B	C	D	
probability	0.4	0.1	0.2	0.15	0.15

- b. Encode the text **ABACABAD** using the code of question a.

- c. Decode the text whose encoding is 100010111001010 in the code of question a.

2. For data transmission purposes, it is often desirable to have a code with a minimum variance of the codeword lengths (among codes of the same average length). Compute the average and variance of the codeword length in two Huffman codes that result from a different tie breaking during a Huffman code construction for the following data:

character	A	B	C	D	E
probability	0.1	0.1	0.2	0.2	0.4

3. Indicate whether each of the following properties are true for every Huffman code.
- The codewords of the two least frequent characters have the same length.
 - The codeword's length of a more frequent character is always smaller than or equal to the codeword's length of a less frequent one.
4. What is the maximal length of a codeword possible in a Huffman encoding of an alphabet of n characters?
5. a. Write a pseudocode for the Huffman tree construction algorithm.
- b. What is the time efficiency class of the algorithm for constructing a Huffman tree as a function of the alphabet's size?
6. Show that a Huffman tree can be constructed in linear time if the alphabet's characters are given in a sorted order of their frequencies.
7. Given a Huffman coding tree, which algorithm would you use to get the codewords for all the characters? What is its time-efficiency class as a function of the alphabet's size?
8. Explain how one can generate a Huffman code without an explicit generation of a Huffman coding tree.
9. a. Write a program that constructs a Huffman code for a given English text and encode it.

- b. Write a program for decoding an English text which has been encoded with a Huffman code.
 - c. Experiment with your encoding program to find a range of typical compression ratios for Huffman's encoding of English texts of, say, 1000 words.
 - d. Experiment with your encoding program to find out how sensitive the compression ratios are to using standard estimates of frequencies instead of actual frequencies of character occurrences in English texts.
10. *Card guessing* Design a strategy that minimizes the expected number of questions asked in the following game [Gar94], #52. You have a deck of cards that consists of one ace of spades, two deuces of spades, three threes, and on up to nine nines, making 45 cards in all. Someone draws a card from the shuffled deck, which you have to identify by asking questions answerable with yes or no.






Hints to Exercises 9.4

1. See the example given in the section.
2. After combining the two nodes with the lowest probabilities, resolve the tie arising on the next iteration in two different ways. For each of the two Huffman codes obtained, compute the mean and variance of the codeword length.
3. You may base your answers on the way Huffman's algorithm works or on the fact that Huffman codes are known to be optimal prefix codes.
4. The maximal length of a codeword relates to the height of Huffman's coding tree in an obvious fashion. Try to find a set of n specific frequencies for an alphabet of size n for which the tree has the shape yielding the longest codeword possible.
5. a. What is the most appropriate data structure for an algorithm whose principal operation is finding the two smallest elements in a given set, deleting them, and then adding a new item to the remaining ones?

b. Identify the principal operations of the algorithm, the number of times they are executed, and their efficiencies for the data structure used.
6. Maintain two queues: one for given frequencies, the other for weights of new trees.
7. It would be natural to use one of the standard traversal algorithms.
8. Generate the codewords right to left.
9. n/a
10. A similar example was discussed at the end of Section 9.4. Construct Huffman's tree and then come up with specific questions that would yield that tree. (You are allowed to ask questions such as: Is this card the ace, or a seven, or an eight?)

Solutions to Exercises 9.4

1. a.



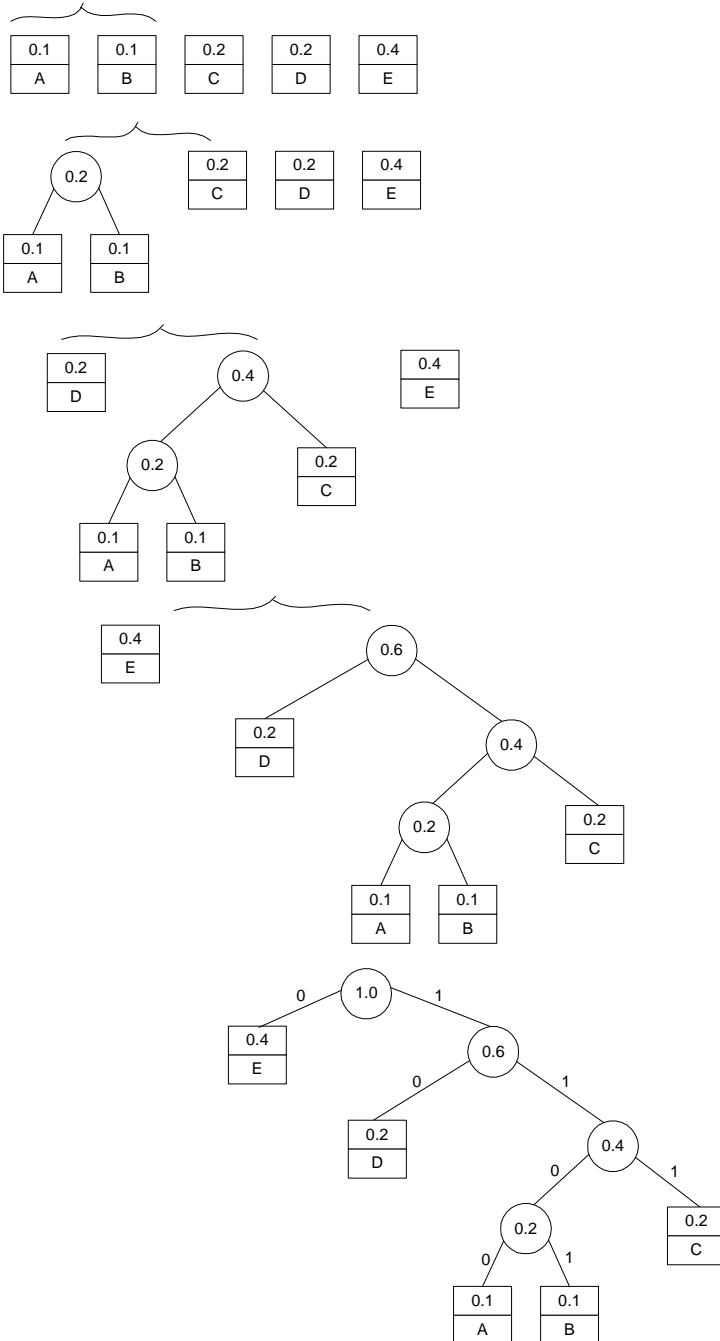
character	A	B	C	D	—
probability	0.4	0.1	0.2	0.15	0.15
codeword	0	100	111	101	110

b. The text **ABACABAD** will be encoded as 0100011101000101.

c. With the code of part a, 100010111001010 will be decoded as

$$\begin{array}{ccccccc} 100 & | & 0 & | & 101 & | & 110 & | & 0 & | & 101 & | & 0 \\ \textit{B} & \textit{A} & \textit{D} & \text{—} & \textit{A} & \textit{D} & \textit{A} & \end{array}$$

2. Here is one way:



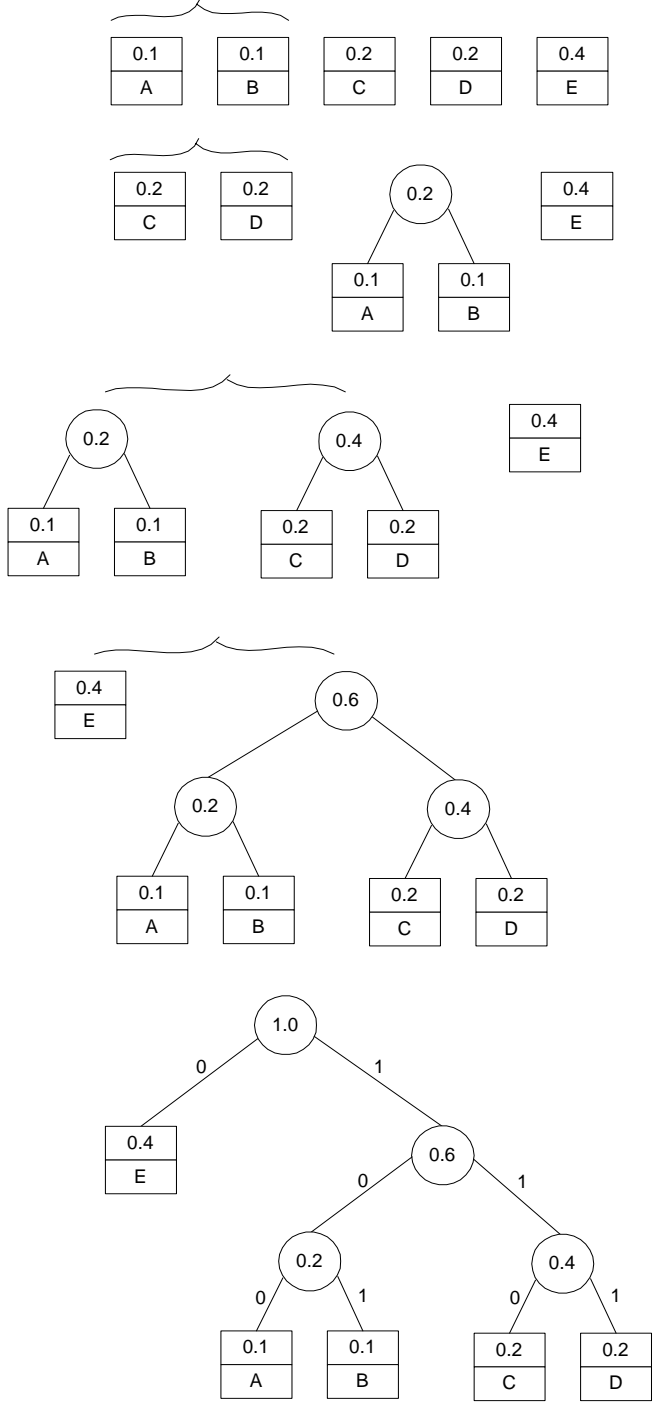
character	A	B	C	D	E
probability	0.1	0.1	0.2	0.2	0.4
codeword	1100	1101	111	10	0
length	4	4	3	2	1

Thus, the mean and variance of the codeword's length are, respectively,

$$\bar{l} = \sum_{i=1}^5 l_i p_i = 4 \cdot 0.1 + 4 \cdot 0.1 + 3 \cdot 0.2 + 2 \cdot 0.2 + 1 \cdot 0.4 = 2.6$$

$$Var = \sum_{i=1}^5 (l_i - \bar{l})^2 p_i = (4-2.6)^2 0.1 + (4-2.6)^2 0.1 + (3-2.6)^2 0.2 + (2-2.6)^2 0.2 + (1-2.6)^2 0.4 = 1.64$$

Here is another way:



character	A	B	C	D	E
probability	0.1	0.1	0.2	0.2	0.4
codeword	100	101	110	111	0
length	3	3	3	3	1

Thus, the mean and variance of the codeword's length are, respectively,

$$\bar{l} = \sum_{i=1}^5 l_i p_i = 2.2 \quad \text{and} \quad Var = \sum_{i=1}^5 (l_i - \bar{l})^2 p_i = 0.96.$$

3. a. Yes. This follows immediately from the way Huffman's algorithm operates: after each of its iterations, the two least frequent characters that are combined on the first iteration are always on the same level of their tree in the algorithm's forest. An easy formal proof of this obvious observation is by induction.

(Note that if there are more than two least frequent characters, the assertion may be false for some pair of them, e.g., $A(\frac{1}{3})$, $B(\frac{1}{3})$, $C(\frac{1}{3})$.)

b. Yes. Let's use the optimality of Huffman codes to prove this property by contradiction. Assume that there exists a Huffman code containing two characters c_i and c_j such that $p(c_i) > p(c_j)$ and $l(w(c_i)) > l(w(c_j))$, where $p(c_i)$ and $l(w(c_i))$ are the probability and codeword's length of c_i , respectively, and $p(c_j)$ and $l(w(c_j))$ are the probability and codeword's length of c_j , respectively. Let's create a new code by simply swapping the codewords of c_1 and c_2 and leaving the codewords for all the other characters the same. The new code will obviously remain prefix-free and its expected length $\sum_{k=1}^n l(w(c_k))p(c_k)$ will be smaller than that of the initial code. This contradicts the optimality of the initial Huffman code and, hence, proves the property in question.

4. The answer is $n - 1$. Since two leaves corresponding to the two least frequent characters must be on the same level of the tree, the tallest Huffman coding tree has to have the remaining leaves each on its own level. The height of such a tree is $n - 1$. An easy and natural way to get a Huffman tree of this shape is by assuming that $p_1 \leq p_2 < \dots < p_n$ and having the weight W_i of a tree created on the i th iteration of Huffman's algorithm, $i = 1, 2, \dots, n - 2$, be less than or equal to p_{i+2} . (Note that for such inputs, $W_i = \sum_{k=1}^{i+1} p_k$ for every $i = 1, 2, \dots, n - 1$.)

As a specific example, it's convenient to consider consecutive powers of 2:

$$p_1 = p_2 \quad \text{and} \quad p_i = 2^{i-n-1} \quad \text{for } i = 2, \dots, n.$$

(For, say, $n = 4$, we have $p_1 = p_2 = 1/8$, $p_3 = 1/4$ and $p_4 = 1/2$.)

Indeed, $p_i = 2^i / 2^{n+1}$ is an increasing sequence as a function of i . Further, $W_i = p_{i+2}$ for every $i = 1, 2, \dots, n - 2$, since

$$\begin{aligned} W_i &= \sum_{k=1}^{i+1} p_k = p_1 + \sum_{k=2}^{i+1} p_k = 2^2 / 2^{n+1} + \sum_{k=2}^{i+1} 2^k / 2^{n+1} = \frac{1}{2^{n+1}} (2^2 + \sum_{k=2}^{i+1} 2^k) \\ &= \frac{1}{2^{n+1}} (2^2 + (2^{i+2} - 4)) = \frac{2^{i+2}}{2^{n+1}} = p_{i+2}. \end{aligned}$$

5. a. The following pseudocode is based on maintaining a priority queue of trees, with the priorities equal the trees' weights.

Algorithm *Huffman*($W[0..n - 1]$)

//Constructs Huffman's tree

//Input: An array $W[0..n - 1]$ of weights

//Output: A Huffman tree with the given weights assigned to its leaves
initialize priority queue Q of size n with one-node trees and priorities equal to the elements of $W[0..n - 1]$

while Q has more than one element **do**

$T_l \leftarrow$ the minimum-weight tree in Q

 delete the minimum-weight tree in Q

$T_r \leftarrow$ the minimum-weight tree in Q

 delete the minimum-weight tree in Q

 create a new tree T with T_l and T_r as its left and right subtrees
 and the weight equal to the sum of T_l and T_r weights

 insert T into Q

return T

Note: See also Problem 6 for an alternative algorithm.

b. The algorithm requires the following operations: initializing a priority queue, deleting its smallest element $2(n - 1)$ times, computing the weight of a combined tree and inserting it into the priority queue $n - 1$ times. The overall running time will be dominated by the time spent on deletions, even taking into account that the size of the priority queue will be decreasing from n to 2. For the min-heap implementation, the time efficiency will be in $O(n \log n)$; for the array or linked list representations, it will be in $O(n^2)$. (Note: For the coding application of Huffman trees, the size of the underlying alphabet is typically not large; hence, a simpler data structure for the priority queue might well suffice.)

6. The critical insight here is that the weights of the trees generated by Huffman's algorithm for nonnegative weights (frequencies) form a nondecreasing sequence. As the hint to this problem suggests, we can then maintain


two queues: one for given frequencies in nondecreasing order, the other for weights of new trees. On each iteration, we do the following: find the two smallest elements among the first two (ordered) elements in the queues (the second queue is empty on the first iteration and can contain just one element thereafter); add their sum to the second queue; and then delete these two elements from their queues. The algorithm stops after $n - 1$ iterations (where n is the alphabet's size), each of which requiring a constant time.

7. Use one of the standard traversals of the binary tree and generate a bit string for each node of the tree as follows:. Starting with the empty bit string for the root, append 0 to the node's string when visiting the node's left subtree begins and append 1 to the node's string when visiting the node's right subtree begins. At a leaf, print out the current bit string as the leaf's codeword. Since Huffman's tree with n leaves has a total of $2n - 1$ nodes (see Sec. 4.4), the efficiency will be in $\Theta(n)$.
8. We can generate the codewords right to left by the following method that stems immediately from Huffman's algorithm: when two trees are combined, append 0 in front of the current bit strings for each leaf in the left subtree and append 1 in front of the current bit strings for each leaf in the right subtree. (The substrings associated with the initial one-node trees are assumed to be empty.)
9. n/a
10. See the next page

10. The probabilities of a selected card be of a particular type is given in the following table:

card	ace	deuce	three	four	five	six	seven	eight	nine
probability	$1/45$	$2/45$	$3/45$	$4/45$	$5/45$	$6/45$	$7/45$	$8/45$	$9/45$

Huffman's tree for this data looks as follows:



The first question this tree implies can be phrased as follows: "Is the selected card a four, a five, or a nine?" . (The other questions can be phrased in a similar fashion.)

The expected number of questions needed to identify a card is equal to the weighted path length from the root to the leaves in the tree:

$$\bar{l} = \sum_{i=1}^9 l_i p_i = \frac{5 \cdot 1}{45} + \frac{5 \cdot 2}{45} + \frac{4 \cdot 3}{45} + \frac{3 \cdot 5}{45} + \frac{3 \cdot 6}{45} + \frac{3 \cdot 7}{45} + \frac{3 \cdot 8}{45} + \frac{2 \cdot 9}{45} = \frac{135}{45} = 3.$$

This file contains the exercises, hints, and solutions for Chapter 10 of the book "Introduction to the Design and Analysis of Algorithms," 2nd edition, by A. Levitin. The problems that might be challenging for at least some students are marked by \triangleright ; those that might be difficult for a majority of students are marked by \blacktriangleright .

Exercises 10.1

1. Solve the following linear programming problems geometrically.

(a)

$$\begin{array}{ll}\text{maximize} & 3x + y \\ \text{subject to} & -x + y \leq 1 \\ & 2x + y \leq 4 \\ & x \geq 0, \ y \geq 0\end{array}$$

(b)

$$\begin{array}{ll}\text{maximize} & x + 2y \\ \text{subject to} & 4x \geq y \\ & y \leq 3 + x \\ & x \geq 0, \ y \geq 0\end{array}$$

2. Consider the linear programming problem

$$\begin{array}{ll}\text{minimize} & c_1x + c_2y \\ \text{subject to} & x + y \geq 4 \\ & x + 3y \geq 6 \\ & x \geq 0, \ y \geq 0\end{array}$$

where c_1 and c_2 are some real numbers not both equal to zero.

- (a) Give an example of the coefficient values c_1 and c_2 for which the problem has a unique optimal solution.
 - (b) Give an example of the coefficient values c_1 and c_2 for which the problem has infinitely many optimal solutions.
 - (c) Give an example of the coefficient values c_1 and c_2 for which the problem does not have an optimal solution.
3. Would the solution to problem (10.2) be different if its inequality constraints were strict, i.e., $x + y < 4$ and $x + 3y < 6$, respectively?
 4. Trace the simplex method on
 - (a) the problem of Exercise 1a.
 - (b) the problem of Exercise 1b.
 5. Trace the simplex method on the problem of Example 1 in Section 6.6

(a) \triangleright by hand.

(b) by using one of the implementations available on the Internet.

6. Determine how many iterations the simplex method needs to solve the problem

$$\begin{array}{ll}\text{maximize} & \sum_{j=1}^n x_j \\ \text{subject to} & 0 \leq x_j \leq b_j, \text{ where } b_j > 0 \text{ for } j = 1, 2, \dots, n.\end{array}$$

7. Can we apply the simplex method to solve the knapsack problem (see Example 2 in Section 6.6)? If you answer yes, indicate whether it is a good algorithm for the problem in question; if you answer no, explain why not.

8. \triangleright Prove that no linear programming problem can have exactly $k \geq 1$ optimal solutions unless $k = 1$.

9. If a linear programming problem

$$\begin{array}{ll}\text{maximize} & \sum_{j=1}^n c_j x_j \\ \text{subject to} & \sum_{j=1}^n a_{ij} x_j \leq b_i \text{ for } i = 1, 2, \dots, m \\ & x_1, x_2, \dots, x_n \geq 0\end{array}$$

is considered as **primal**, then its **dual** is defined as the linear programming problem

$$\begin{array}{ll}\text{minimize} & \sum_{i=1}^m b_i y_i \\ \text{subject to} & \sum_{i=1}^m a_{ij} y_i \geq c_j \text{ for } j = 1, 2, \dots, n \\ & y_1, y_2, \dots, y_m \geq 0.\end{array}$$

(a) Express the primal and dual problems in matrix notations.

(b) Find the dual of the following linear programming problem

$$\begin{array}{ll}\text{maximize} & x_1 + 4x_2 - x_3 \\ \text{subject to} & x_1 + x_2 + x_3 \leq 6 \\ & x_1 - x_2 - 2x_3 \leq 2 \\ & x_1, x_2, x_3 \geq 0.\end{array}$$

(c) Solve the primal and dual problems and compare the optimal values of their objective functions.

10. \triangleright *Parliament pacification* In a parliament, each parliamentarian has at most three enemies. Design an algorithm that divides the parliament into two chambers in such a way that no parliamentarian has more than one enemy in his or her chamber (after [Sav03], p.1, #4).

Hints to Exercises 10.1


1. Sketch the feasible region of the problem in question. Follow this up by either applying the Extreme-Point Theorem or by inspecting level lines, whichever is more appropriate. Both methods were illustrated in the text.
2. Sketch the feasible region of the problem. Then choose values of the parameters c_1 and c_2 to obtain a desired behavior of the objective function's level lines.
3. What is the principal difference between maximizing a linear function, say $f(x) = 2x$, on a closed vs. semi-open interval, e.g., $0 \leq x \leq 1$ vs. $0 \leq x < 1$?
4. Trace the simplex method on the instances given, as it was done for an example in the text.
5. When solving the problem by hand, you might want to start by getting rid of fractional coefficients in the problem's statement. Also note that the problem's specifics make it possible to replace its equality constraint by one inequality constraint. You were asked to solve this problem directly in Exercises 6.6 (see Problem 8).
6. The specifics of the problem make it possible to see the optimal solution at once. Sketching its feasible region for $n = 2$ or $n = 3$, though not necessary, may help to see both this solution and the number of iterations needed by the simplex method to solve it.
7. Consider separately two versions of the problem: continuous and 0-1 (see Example 2 in Section 6.6).
8. If $x' = (x'_1, x'_2, \dots, x'_n)$ and $x'' = (x''_1, x''_2, \dots, x''_n)$ are two distinct optimal solutions to the same linear programming problem, what can we say about any point of the line segment with the endpoints at x' and x'' ? Note that any such point x can be expressed as $x = tx' + (1 - t)x'' = (tx'_1 + (1 - t)x''_1, tx'_2 + (1 - t)x''_2, \dots, tx'_n + (1 - t)x''_n)$, where $0 \leq t \leq 1$.
9.
 - a. You will need to use the notion of a matrix transpose, defined as the matrix whose rows are the columns of the given matrix.
 - b. Apply the general definition to the specific problem given. Note the change from maximization to minimization, the change of the roles played by objective function's coefficients and constraints' right-hand sides, the transposition of the constraints and the reversal of their signs.
 - c. You may use either the simplex method or the geometric approach.
10. The problem can be solved by an iterative improvement algorithm that decreases the number of enemy pairs who are members of the same chamber.

Solutions to Exercises 10.1

1. a. The feasible region of the linear programming problem

$$\begin{array}{ll}\text{maximize} & 3x + y \\ \text{subject to} & -x + y \leq 1 \\ & 2x + y \leq 4 \\ & x \geq 0, y \geq 0\end{array}$$

is given in the following figure:



Since the feasible region of the problem is nonempty and bounded, an optimal solution exists and can be found at one of the extreme points of the feasible region. The extreme points and the corresponding values of the objective function are given in the following table:


Extreme point	Value of $3x + y$
$(0, 0)$	0
$(2, 0)$	6
$(1, 2)$	5
$(0, 1)$	1

Hence, the optimal solution is $x = 2$, $y = 0$, with the corresponding value of the objective function equal to 6.

b. The feasible region of the linear programming problem


$$\begin{array}{ll}\text{maximize} & x + 2y \\ \text{subject to} & 4x \geq y \\ & x + 3 \geq y \\ & x \geq 0, \ y \geq 0\end{array}$$

is given in the following figure:



The feasible region is unbounded. On considering a few level lines $x+2y = z$, it is easy to see that level lines can be moved in the north-east direction, which corresponds to increasing values of z , as far as we wish without losing common points with the feasible region. (Alternatively, we can consider a simple sequence of points, such as $(n,0)$, which are feasible for any nonnegative integer value of n and make the objective function $z = x + 2y$ as large as we wish as n goes to infinity.) Hence, the problem is unbounded and therefore does not have a finite optimal solution.

2. The feasible region of the problem with the constraints $x + y \geq 4$, $x + 3y \geq 6$, $x \geq 0$, $y \geq 0$ is given in Figure 10.3 of Section 10.1:



a. Minimization of $z = c_1x + c_2y$ on this feasible region, with $c_1, c_2 > 0$, pushes level lines $c_1x + c_2y = z$ in the south-west direction. Any family of such lines with a slope strictly between the slopes of the boundary lines $x + y = 4$ and $x + 3y = 6$ will hit the extreme point $(3, 1)$ as the only minimum solution to the problem. For example, as mentioned in Section 10.1, we can minimize $3x + 5y$ among infinitely many other possible answers.

b. For a linear programming problem to have infinitely many solutions, the optimal level line of its objective function must contain a line segment that is a part of the feasible region's boundary. Hence, there are four qualitatively distinct answers:

- minimize $1 \cdot x + 0 \cdot y$ (or, more generally, any objective function of the form $c \cdot x + 0 \cdot y$, where $c > 0$);
- minimize $x + y$ (or, more generally, any objective function of the form $cx + cy$, where $c > 0$);
- minimize $x + 3y$ (or, more generally, any objective function of the form $cx + 3cy$, where $c > 0$);
- minimize $0 \cdot x + 1 \cdot y$ (or, more generally, any objective function of the form $0 \cdot x + c \cdot y$, where $c > 0$).

3. The problem with the strict inequalities does not have an optimal solution: The objective function values of any sequence of feasible points approaching $x_1 = 3, x_2 = 1$ (the optimal solution to Problem (10.2)) will approach $z = 14$ as its limit, but this value will not be attained at any feasible point of the problem with the strict inequalities.

4. a. The standard form of the problem given is

$$\begin{aligned} &\text{maximize} && 3x + y \\ &\text{subject to} && -x + y + u = 1 \\ &&& 2x + y + v = 4 \\ &&& x, y, u, v \geq 0. \end{aligned}$$

Here are the tableaux generated by the simplex method in solving this problem:

	x	y	u	v	
u	-1	1	1	0	1
v	2	1	0	1	4
	-3	-1	0	0	0

\leftarrow \uparrow

	x	y	u	v	
u	0	$\frac{3}{2}$	1	$\frac{1}{2}$	3
x	1	$\frac{1}{2}$	0	$\frac{1}{2}$	2
	0	$\frac{1}{2}$	0	$\frac{3}{2}$	6

$\theta_v = \frac{4}{2}$

The optimal solution found is $x = 2, y = 0$, with the maximal value of the objective function equal to 6.

b. The standard form of the problem given is

$$\begin{array}{ll} \text{maximize} & x + 2y \\ \text{subject to} & -4x + y + u = 0 \\ & -x + y + v = 3 \\ & x, y, u, v \geq 0. \end{array}$$

Here are the tableaux generated by the simplex method in solving this problem:

		x	y	u	v		
←	u	-4	1	1	0		0
	v	-1	1	0	1		3
		-1	-2	0	0		0

↑

$\theta_u = \frac{0}{1}$
 $\theta_v = \frac{3}{1}$

		x	y	u	v		
←	y	-4	1	1	0		0
	v	3	0	-1	1		3
		-9	0	2	0		0

↑

$\theta_v = \frac{3}{3}$

		x	y	u	v		
←	y	0	1	$-\frac{1}{3}$	$\frac{4}{3}$		4
	x	1	0	$-\frac{1}{3}$	$\frac{1}{3}$		1
		0	0	-1	3		9

↑

Since there are no positive elements in the pivot column of the last tableau, the problem is unbounded.

5. To simplify the task of getting an initial basic feasible solution here, we can replace the equality constraint $x + y + z = 100$ by the inequality $x + y + z \leq 100$, because an optimal solution (x^*, y^*, z^*) to the problem with the latter constraint must satisfy the former one.. (Otherwise, we would've been able to increase the maximal value of the objective function by increasing the value of z^* .) then, after an optional replacement of the objective function and the constraints by the equivalent ones without fractional coefficients, the problem can be presented in the standard form as follows

$$\begin{aligned} &\text{maximize} && 10x + 7y + 3z \\ &\text{subject to} && x + y + z + u = 100 \\ &&& 3x - y + + v = 0 \\ &&& x + y - 4z + + w = 0 \\ &&& x, y, z, u, v, w \geq 0. \end{aligned}$$

Here are the tableaux generated by the simplex method in solving this problem:

	x	y	z	u	v	w	
u	1	1	1	1	0	0	100
v	3	-1	0	0	1	0	0
w	1	1	-4	0	0	1	0
	-10	-7	-3	0	0	0	0

$\theta_u = \frac{100}{1}$
 $\theta_v = \frac{0}{3}$
 $\theta_w = \frac{0}{1}$

	x	y	z	u	v	w	
u	0	$\frac{4}{3}$	1	1	$-\frac{1}{3}$	0	$\theta_u = \frac{100}{4/3}$
x	1	$-\frac{1}{3}$	0	0	$\frac{1}{3}$	0	
w	0	$\frac{4}{3}$	-4	0	$-\frac{1}{3}$	1	$\theta_w = \frac{0}{4/3}$
	0	$-\frac{31}{3}$	-3	0	$\frac{10}{3}$	0	

	x	y	z	u	v	w	
$\leftarrow u$	0	0	5	1	0	-1	100
x	1	0	-1	0	$\frac{1}{4}$	$\frac{1}{4}$	0
y	0	1	-3	0	$-\frac{1}{4}$	$\frac{3}{4}$	0
	0	0	-34	0	$\frac{3}{4}$	$\frac{31}{4}$	0

\uparrow

$\theta_u = \frac{100}{5}$

	x	y	z	u	v	w	
z	0	0	1	$\frac{1}{5}$	0	$-\frac{1}{5}$	20
x	1	0	0	$\frac{1}{5}$	$\frac{1}{4}$	$\frac{1}{20}$	20
y	0	1	0	$\frac{3}{5}$	$-\frac{1}{4}$	$\frac{3}{20}$	60
	0	0	0	$\frac{34}{5}$	$\frac{3}{4}$	$\frac{19}{20}$	680

The optimal solution found is $x = 20$, $y = 60$, $z = 20$, with the maximal value of the objective function equal to 680.

6. The optimal solution to the problem is $x_1 = b_1, \dots, x_n = b_n$. After introducing a slack variable s_i in the i th inequality $x_i = b_i$ to get to the standard form and starting with $x_1 = 0, \dots, x_n = 0$, $s_1 = b_1, \dots, s_n = b_n$, the simplex method will need n iterations to get to an optimal solution $x_1 = b_1, \dots, x_n = b_n$, $s_1 = 0, \dots, s_n = 0$. It follows from the fact that on each iteration the simplex method replaces only one basic variable. Here, on each of its n iterations, it will replace some slack variable s_i by the corresponding x_i .

7. The continuous version of the knapsack problem can be solved by the simplex method, because it is a special case of the general linear programming problem (see Example 2 in Section 6.6). However, it is hardly a good method for solving this problem because it can be solved more efficiently by a much simpler algorithm based on the greedy approach. You may want to design such an algorithm by yourself before looking it up in the book.

The 0-1 version of the knapsack problem cannot be solved by the simplex method because of the integrality (0-1) constraints imposed on the problem's variables.

8. The assertion follows immediately from the fact that if $x' = (x'_1, \dots, x'_n)$ and $x'' = (x''_1, \dots, x''_n)$ are two distinct optimal solutions to the same linear programming problem, then any of the infinite number of points of the line segment with endpoints at x' and x'' will be an optimal solution to this problem as well. Indeed, let x^t be such a point:

$$x^t = tx' + (1-t)x'' = (tx'_1 + (1-t)x''_1, \dots, tx'_n + (1-t)x''_n), \text{ where } 0 \leq t \leq 1.$$

First, x^t will satisfy all the constraints of the problem, whether they are linear inequalities or linear equations, because both x' and x'' do. Indeed, let the i th constraint be inequality $\sum_{j=1}^n a_{ij}x_j \leq b_i$. Then $\sum_{j=1}^n a_{ij}x'_j \leq b_i$ and $\sum_{j=1}^n a_{ij}x''_j \leq b_i$. Multiplying these inequalities by t and $1-t$, respectively, and adding the results, we obtain

$$t \sum_{j=1}^n a_{ij}x'_j + (1-t) \sum_{j=1}^n a_{ij}x''_j \leq tb_i + (1-t)b_i$$

or

$$\sum_{j=1}^n (ta_{ij}x'_j + (1-t)a_{ij}x''_j) = \sum_{j=1}^n a_{ij}(tx'_j + (1-t)x''_j) = \sum_{j=1}^n a_{ij}x^t_j \leq b_i,$$

i.e., x satisfies the inequality. The same argument holds for inequalities $\sum_{j=1}^n a_{ij}x_j \geq b_i$ and equations $\sum_{j=1}^n a_{ij}x_j = b_i$.

Second, $x^t = tx' + (1-t)x''$ will maximize the value of the objective function. Indeed, if the maximal value of the objective function is z^* , then

$$\sum_{j=1}^n c_jx'_j = z^* \quad \text{and} \quad \sum_{j=1}^n c_jx''_j = z^*.$$

Multiplying these equalities by t and $1-t$, respectively, and adding the results, we will obtain

$$t \sum_{j=1}^n c_jx'_j + (1-t) \sum_{j=1}^n c_jx''_j = tz^* + (1-t)z^*$$

or

$$\sum_{j=1}^n (tc_jx'_j + (1-t)c_jx''_j) = \sum_{j=1}^n c_j(tx'_j + (1-t)x''_j) = \sum_{j=1}^n c_jx^t_j = z^*.$$

I.e., we proved that x^t does maximize the objective function and hence will be an optimal solution to the problem in question for any $0 \leq t \leq 1$.

Note: What we actually proved is the fact that the set of optimal solutions to a linear programming problem is convex. And any nonempty convex set can contain either a single point or infinitely many points.

9. a. A linear programming problem

$$\begin{aligned} & \text{maximize} && \sum_{j=1}^n c_j x_j \\ & \text{subject to} && \sum_{j=1}^n a_{ij} x_j \leq b_i \quad \text{for } i = 1, 2, \dots, m \\ & && x_1, x_2, \dots, x_n \geq 0 \end{aligned}$$

can be compactly written using the matrix notations as follows:

$$\begin{aligned} & \text{maximize} && cx \\ & \text{subject to} && Ax \leq b \\ & && x \geq 0, \end{aligned}$$

where

$$c = [c_1 \quad \dots \quad c_n], \quad x = \begin{bmatrix} x_1 \\ \vdots \\ x_n \end{bmatrix}, \quad A = \begin{bmatrix} a_{11} & \dots & a_{1n} \\ \vdots & & \vdots \\ a_{m1} & \dots & a_{mn} \end{bmatrix}, \quad b = \begin{bmatrix} b_1 \\ \vdots \\ b_m \end{bmatrix},$$

$Ax \leq b$ holds if and only if each coordinate of the product Ax is less than or equal to the corresponding coordinate of vector b , and $x \geq 0$ is shorthand for the nonnegativity requirement for all the variables. The **dual** can be written as follows:

$$\begin{aligned} & \text{minimize} && b^T y \\ & \text{subject to} && A^T y \geq c^T \\ & && y \geq 0, \end{aligned}$$

where b^T is the transpose of b (i.e., $b^T = [b_1, \dots, b_m]$), c^T is the transpose of c (i.e., c^T is the column-vector made up of the coordinates of the row-vector c), A^T is the transpose of A (i.e., the n -by- m matrix whose j th row is the j th column of matrix A , $j = 1, 2, \dots, n$), and y is the vector-column of m new unknowns y_1, \dots, y_m .

b. The dual of the linear programming problem

$$\begin{aligned} & \text{maximize} && x_1 + 4x_2 - x_3 \\ & \text{subject to} && x_1 + x_2 + x_3 \leq 6 \\ & && x_1 - x_2 - 2x_3 \leq 2 \\ & && x_1, x_2, x_3 \geq 0 \end{aligned}$$

is

$$\begin{aligned} & \text{minimize} && 6y_1 + 2y_2 \\ & \text{subject to} && y_1 + y_2 \geq 1 \\ & && y_1 - y_2 \geq 4 \\ & && y_1 - 2y_2 \geq -1 \\ & && y_1, y_2 \geq 0. \end{aligned}$$

c. The standard form of the primal problem is

$$\begin{aligned} & \text{maximize} && x_1 + 4x_2 - x_3 \\ & \text{subject to} && x_1 + x_2 + x_3 + x_4 = 6 \\ & && x_1 - x_2 - 2x_3 + x_5 = 2 \\ & && x_1, x_2, x_3, x_4, x_5 \geq 0. \end{aligned}$$

The simplex method yields the following tableaux:

	x_1	x_2	x_3	x_4	x_5	
$\leftarrow x_4$	1	1	1	1	0	6
x_5	1	-1	-2	0	1	2
	-1	-4	1	0	0	0


$\theta_{x_4} = \frac{6}{1}$

\uparrow

	x_1	x_2	x_3	x_4	x_5	
x_2	1	1	1	1	0	6
x_5	2	0	-1	1	1	8
	3	0	5	4	0	24

The found optimal solution is $x_1 = 0, x_2 = 6, x_3 = 0$.

Since the dual problem has just two variables, it is easier to solve it geometrically. Its feasible region is presented in the following figure:



Although it is unbounded, the minimization problem in question does have a finite optimal solution $y_1 = 4$, $y_2 = 0$. Note that the optimal values of the objective functions in the primal and dual problems are equal to each other:

$$0 + 4 \cdot 6 - 0 = 6 \cdot 4 + 2 \cdot 0.$$


This is the principal assertion of the Duality Theorem, one of the most important facts in linear programming theory.

10. Partition the parliamentarians into two chambers in an arbitrary way. Let p be the total number of enemy pairs who are members of the same chamber. While there is a parliamentarian who has at least two enemies in his or her chamber, transform this parliamentarian into the other chamber. (Since the transformed parliamentarian will have no more than one enemy in his or her new chamber and the number of enemy pairs in his or her old chamber will decrease by two, the total number of enemy pairs will decrease at least by one. Hence, the algorithm will terminate after no more than p iterations, with the final partition of the parliamentarians in which every parliamentarian has no more than one enemy in his or her chamber.)


Exercises 10.2

1. Since maximum-flow algorithms require processing edges in both directions, it is convenient to modify the adjacency matrix representation of a network as follows: If there is a directed edge from vertex i to vertex j of capacity u_{ij} , then the element in the i th row and the j th column is set to u_{ij} , while the element in the j th row and the i th column is set to $-u_{ij}$; if there is no edge between vertices i and j , both these elements are set to zero. Outline a simple algorithm for identifying a source and a sink in a network presented by such a matrix and indicate its time efficiency.
2. Apply the shortest-augmenting path algorithm to find a maximum flow and a minimum cut in the following networks:

a.



b.



3. a. Does the maximum-flow problem always have a unique solution? Would your answer be different for networks with different capacities on all their edges?

b. Answer the same questions for the minimum-cut problem of finding a cut of the smallest capacity in a given network.
4. a. Explain how the maximum-flow problem for a network with several sources and sinks can be transformed to the same problem for a network with a single source and a single sink.

b. Some networks have capacity constraints on flow amounts that can flow through their intermediate vertices. Explain how the maximum-flow problem for such a network can be transformed to the maximum-flow problem for a network with edge capacity constraints only.
5. \triangleright Consider a network that is a rooted tree, with the root as its source, the leaves as its sinks, and all the edges directed along the paths from the root to the leaves. Design an efficient algorithm for finding of a maximum flow in such a network. What is the time efficiency of your algorithm?

6. \triangleright a. Prove equality (10.9).
 - \triangleright b. Prove that for any flow in a network and any cut in it, the value of the flow is equal to the flow across the cut (see equality (10.12)). Explain the relationship between this property and equality (10.9).
7. a. Express the maximum-flow problem for the network of Figure 10.4 as a linear programming problem.
 - b. Solve this linear programming problem by the simplex method.
8. As an alternative to the shortest-augmenting-path algorithm, Edmonds and Karp [Edm72] suggested the maximal-augmenting-path algorithm in which a flow is augmented along the path that increases the flow by the largest amount. Implement both these algorithms in the language of your choice and perform an empirical investigation of their relative efficiency.
9. Write a report on a more advanced maximum-flow algorithm such as (i) Dinitz's algorithm; (ii) Karzanov's algorithm; (iii) Malhotra-Kamar-Maheshwari algorithm; (iv) Goldberg-Tarjan algorithm.
10. *Dining problem* Several families go out to dinner together. To increase their social interaction, they would like to sit at tables so that no two members of the same family are at the same table. Show how to find a seating arrangement that meets this objective (or prove that no such arrangement exists) by using a maximum flow problem. Assume that the dinner contingent has p families and that the i th family has a_i members. Also assume that q tables are available and the j th table has a seating capacity of b_j [Ahu93].

Hints to Exercises 10.2

1. What properties of the adjacency matrix elements stem from the source and sink definitions, respectively?
2. See the algorithm and an example illustrating it in the text.
3. Of course, the value (capacity) of an optimal flow (cut) is the same for any optimal solution. The question is whether distinct flows (cuts) can yield the same optimal value.
4. a. Add extra vertices and edges to the network given.


b. If an intermediate vertex has a constraint on a flow amount that can flow through it, split the vertex into two.
5. Take advantage of the recursive structure of a rooted tree.
6. a. Sum the equations expressing the flow-conservation requirement.

b. Sum the equations defining the flow value and flow-conservation requirements for the vertices in set X inducing the cut.
7. a. Use template (10.11) given in the text.

b. Use either an add-on tool of your spreadsheet or some software available on the Internet.
10. Use edge capacities to impose the problem's constraints. Also, take advantage of the solution to Problem 4a.

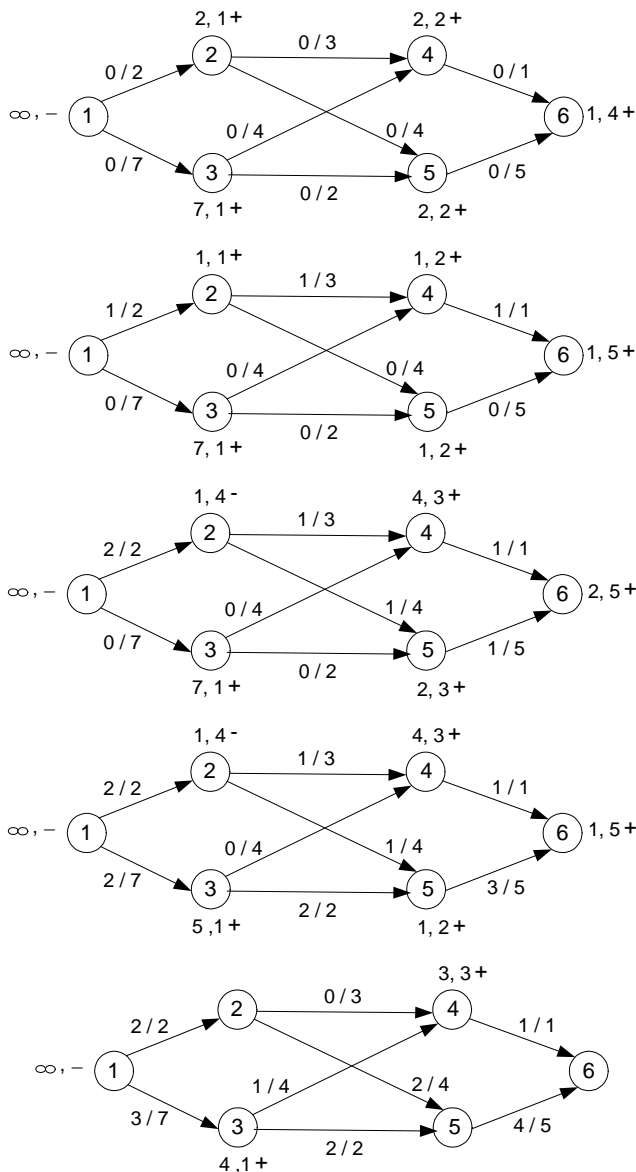
Solutions to Exercises 10.2

1. The definition of a source implies that a vertex is a source if and only if there are no negative elements in its row in the adjacency matrix. Similarly, the definition of a sink implies that a vertex is a sink if and only if there are no positive elements in its row of the adjacency matrix. Thus a simple scan of the adjacency matrix rows solves the problem in $O(n^2)$ time, where n is the number of vertices in the network.
2. a. Here is an application of the shortest-augmenting path algorithm to the network of Problem 2a:




The maximum flow is shown on the last diagram above. The minimum cut found is $\{(2, 5), (4, 6)\}$.

b. Here is an application of the shortest-augmenting path algorithm to the network of Problem 10.2b:




The maximum flow of value 5 is shown on the last diagram above. The minimum cut found is $\{(1, 2), (3, 5), (4, 6)\}$.

3. a. The maximum-flow problem may have more than one optimal solution. In fact, there may be infinitely many of them if we allow (as the definition does) non-integer edge flows. For example, for any $0 \leq t \leq 1$, the flow depicted in the diagram below is a maximum flow of value 1. Exactly two of them—for $t = 0$ and $t = 1$ —are integer flows.




The answer does not change for networks with distinct capacities: e.g., consider the previous example with the capacities of edges $(2,3)$, $(3,5)$, $(2,4)$, and $(4,5)$ changed to, say, 2, 3, 4, and 5, respectively.

- b. The answer for the number of distinct minimum cuts is analogous to that for maximum flows: there can be more than one of them in the same network (though, of course, their number must always be finite because the number of all edge subsets is finite to begin with). For example, the network




has four minimum cuts: $\{(1,2), (1,3)\}$, $\{(1,2), (3,4)\}$, $\{(1,3), (2,4)\}$, and $\{(2,4), (3,4)\}$. The answer does not change for networks with distinct edge capacities. For example, the network



has two minimum cuts: $\{(1,2), (1,3)\}$ and $\{(4,5)\}$.

4. a. Add two vertices to the network given to serve as the source and sink of the new network, respectively. Connect the new source to each of the original sources and each of the original sinks to the new sink with edges of some large capacity M . (It suffices to take M greater than or equal to the sum of the capacities of the edges leaving each source of the original network.)
- b. Replace each intermediate vertex i with an upper bound b_i on a flow amount that can flow through it with two vertices i' and i'' connected by an edge of capacity b_i as shown below:



Note that all the edges entering and leaving i in the original network should enter i' and leave i'' , respectively, in the new one.

5. The problem can be solved by calling $TreeFlow(T(r), \infty)$, where

Algorithm $TreeFlow(T(r), v)$

//Finds a maximum flow for tree $T(r)$ rooted at r ,

//whose value doesn't exceed v (available at the root),

//and returns its value

if r is a leaf $maxflowval \leftarrow v$

else

$maxflowval \leftarrow 0$

for every child c of r **do**

$x_{rc} \leftarrow TreeFlow(T(c), \min\{u_{rc}, v\})$

$v \leftarrow v - x_{rc}$

$maxflowval \leftarrow maxflowval + x_{rc}$

return $maxflowval$

The efficiency of the algorithm is clearly linear because a $\Theta(1)$ call is made for each of the nodes of the tree.

6. a. Adding the $i - 2$ equalities expressing the flow conservation requirements yields

$$\sum_{i=2}^{n-1} \sum_j x_{ji} = \sum_{i=2}^{n-1} \sum_j x_{ij}.$$

For any edge from an intermediate vertex p to an intermediate vertex q ($2 \leq p, q \leq n - 1$), the edge flow x_{pq} occurs once in both the left- and

right-hand sides of the last equality and hence will cancel out. The remaining terms yield:

$$\sum_{i=2}^{n-1} x_{1i} = \sum_{i=2}^{n-1} x_{in}.$$

Adding x_{1n} to both sides of the last equation, if there is an edge from source 1 to sink n , results in the desired equality:

$$\sum_i x_{1i} = \sum_i x_{in}.$$

b. Summing up the flow-value definition $v = \sum_j x_{1j}$ and the flow-conservation requirement $\sum_j x_{ji} = \sum_j x_{ij}$ for every $i \in X$ ($i > 1$), we obtain

$$v + \sum_{i: i \in X, i > 1} \sum_j x_{ji} = \sum_j x_{1j} + \sum_{i: i \in X, i > 1} \sum_j x_{ij},$$

or, since there are no edges entering the source,

$$v + \sum_{i \in X} \sum_j x_{ji} = \sum_{i \in X} \sum_j x_{ij}.$$

Moving the summation from the left-hand side to the right-hand side and splitting the sum into the sum over the vertices in X and the sum over the vertices in \bar{X} , we obtain:

$$\begin{aligned} v &= \sum_{i \in X} \sum_j x_{ij} - \sum_{i \in X} \sum_j x_{ji} \\ &= \sum_{i \in X} \sum_{j \in X} x_{ij} + \sum_{i \in X} \sum_{j \in \bar{X}} x_{ij} - \sum_{i \in X} \sum_{j \in X} x_{ji} - \sum_{i \in X} \sum_{j \in \bar{X}} x_{ji} \\ &= \sum_{i \in X} \sum_{j \in \bar{X}} x_{ij} - \sum_{i \in X} \sum_{j \in \bar{X}} x_{ji}, \quad \text{Q.E.D.} \end{aligned}$$

Note that equation (10.9) expresses this general property for two special cuts: $C_1(X_1, \bar{X}_1)$ induced by $X_1 = \{1\}$ and $C_2(X_2, \bar{X}_2)$ induced by $X_2 = V - \{n\}$.

7. a.

$$\text{maximize } v = x_{12} + x_{14}$$


$$\begin{aligned} \text{subject to } & x_{12} - x_{23} - x_{25} = 0 \\ & x_{23} + x_{43} - x_{36} = 0 \\ & x_{14} - x_{43} = 0 \\ & x_{25} - x_{56} = 0 \\ & 0 \leq x_{12} \leq 2 \\ & 0 \leq x_{14} \leq 3 \\ & 0 \leq x_{23} \leq 5 \\ & 0 \leq x_{25} \leq 3 \\ & 0 \leq x_{36} \leq 2 \\ & 0 \leq x_{43} \leq 1 \\ & 0 \leq x_{56} \leq 4. \end{aligned}$$

b. The optimal solution is $x_{12} = 2$, $x_{14} = 1$, $x_{23} = 1$, $x_{25} = 1$, $x_{36} = 2$, $x_{43} = 1$, $x_{56} = 1$.

8. n/a

9. n/a

10. Solve the maximum flow problem for the following network:




If the maximum flow value is equal to $\sum_{i=1}^p a_i$, then the problem has a solution indicated by the full edges of capacity 1 in the maximum flow; otherwise, the problem does not have a solution.


Exercises 10.3

1. For each matching shown below in bold, find an augmentation or explain why no augmentation exists.


a.



b.




2. Apply the maximum-matching algorithm to the following bipartite graph:



3. a. What is the largest and what is the smallest possible cardinality of a matching in a bipartite graph $G = \langle V, U, E \rangle$ with n vertices in each vertex set V and U and at least n edges?

b. What is the largest and what is the smallest number of distinct solutions the maximum-cardinality matching problem can have for a bipartite graph $G = \langle V, U, E \rangle$ with n vertices in each vertex set V and U and at least n edges?


4. a. **Hall's Marriage Theorem** asserts that a bipartite graph $G = \langle V, U, E \rangle$ has a matching that matches all vertices of the set V if and only if for each subset $S \subseteq V$, $|R(S)| \geq |S|$ where $R(S)$ is the set of all vertices adjacent to a vertex in S . Check this property for the following graph with (i) $V = \{1, 2, 3, 4\}$ and (ii) $V = \{5, 6, 7\}$.




b. You have to devise an algorithm that returns “yes” if there is a matching in a bipartite graph $G = \langle V, U, E \rangle$ that matches all vertices in V and returns “no” otherwise. Would you base your algorithm on checking the condition of Hall's Marriage Theorem?

5. Suppose there are five committees A , B , C , D , and E composed of six persons a , b , c , d , e , f as follows: committee A 's members are b and e ; committee B 's members are b , d and e ; committee C 's members are a , c , d , e , f ; committee D 's members are b , d , e ; committee E 's members are b and e . Is there a **system of distinct representatives**, i.e., is it possible to select a representative from each committee so that all the selected persons are distinct?
6. Show how the maximum-cardinality matching problem in a bipartite graph can be reduced to the maximum-flow problem discussed in Section 10.2.
7. Consider the following greedy algorithm for finding a maximum matching in a bipartite graph $G = \langle V, U, E \rangle$: Sort all the vertices in nondecreasing order of their degrees. Scan this sorted list to add to the current matching (initially empty) the edge from the list's free vertex to an adjacent free vertex of the lowest degree. If the list's vertex is matched or if there are no adjacent free vertices for it, the vertex is simply skipped. Does this algorithm always produce a maximum matching in a bipartite graph?
8. Design a linear algorithm for finding a maximum matching in a tree.
9. Implement the maximum-matching algorithm of this section in the language of your choice. Experiment with its performance on bipartite graphs with n vertices in each of the vertex sets and randomly generated edges (in both dense and sparse modes) to compare the observed running time with the algorithm's theoretical efficiency.
10. *Domino puzzle* A domino is a 2-by-1 tile that can be oriented either horizontally or vertically. A tiling of a given board composed of 1-by-1 squares is covering it with dominoes exactly without overlap. Which of the following boards can be tiled with dominoes? If you answer yes, indicate a tiling; if you answer no, explain why.


a.



b.



c.




Hints to Exercises 10.3

1. You may but do not have to use the algorithm described in the section.
2. See an application of this algorithm to another bipartite graph in the section.
3. The definition of a matching and its cardinality should lead you to the answers to these questions with no difficulty.
4. a. You do not have to check the inequality for each subset S of V if you can point out a subset for which the inequality does not hold. Otherwise, fill in a table for all the subsets S of the indicated set V with columns for S , $R(S)$, and $|R(S)| \geq |S|$.


b. Think time efficiency.
5. Reduce the problem to finding a maximum matching in a bipartite graph.
6. Transform a given bipartite graph into a network by making vertices of the former be intermediate vertices of the latter.
7. Since this greedy algorithm is arguably simpler than the augmenting path algorithm given in the section, should we expect a positive or negative answer? Of course, this point cannot be substituted for a more specific argument or a counterexample.
8. Start by presenting a tree given as a BFS tree.
9. For pointers regarding an efficient implementation of the algorithm, see [Pap82], Section 10.2.
10. Although it is possible to answer the questions without referring to the matching problem explicitly, you should ultimately coach your answers using the notions discussed in the section.

Solutions to Exercises 10.3


1. a. The matching given in the exercise is reproduced below:



Its possible augmentations are:




- b. No augmentation of the matching given in part b (reproduced below) is possible.



This conclusion can be arrived at either by applying the maximum-matching algorithm or by simply noting that only one of vertices 5 and 8 can be matched (and, hence, no more than three of vertices 5, 6, 7, and 8 can be matched in any matching).

2. Here is a trace of the maximum-matching algorithm applied to the bipartite graph in question:



3. a. The largest cardinality of a matching is n when all the vertices of a graph are matched (perfect matching). For example, if $V = \{v_1, v_2, \dots, v_n\}$, $U = \{u_1, u_2, \dots, u_n\}$ and $E = \{(v_1, u_1), (v_2, u_2), \dots, (v_n, u_n)\}$, then $M = E$

is a perfect matching of size n .

The smallest cardinality of a matching is 1. (It cannot be zero because the number of edges is assumed to be at least $n \geq 1$.) For example, in the bipartite graph with $V = \{v_1, v_2, \dots, v_n\}$, $U = \{u_1, u_2, \dots, u_n\}$, and $E = \{(v_1, u_1), (v_1, u_2), \dots, (v_1, u_n)\}$, the size of any matching is 1.

b. Consider $K_{n,n}$, the bipartite graph in which each of the n vertices in V is connected to each of the n vertices in U . To obtain a perfect matching, there are n possible mates for vertex v_1 , $n-1$ possible remaining mates for v_2 , and so on until there is just one possible remaining mate for v_n . Therefore the total number of distinct perfect matchings for $K_{n,n}$ is $n(n-1)\dots 1 = n!$.

The smallest number of distinct maximum matchings is 1. For example, if $V = \{v_1, v_2, \dots, v_n\}$, $U = \{u_1, u_2, \dots, u_n\}$, and $E = \{(v_1, u_1), (v_2, u_2), \dots, (v_n, u_n)\}$, $M = E$ is the only perfect (and hence maximum) matching.

4. a. (i) For $V = \{1, 2, 3, 4\}$, the inequality obviously fails for $S = V$ since $|R(S)| = |U| = 3$ while $|S| = 4$.

Hence, according to Hall's Marriage Theorem, there is no matching that matches all the vertices of the set $\{1, 2, 3, 4\}$.


- (ii) For subsets S of $V = \{5, 6, 7\}$, we have the following table:

S	$R(S)$	$ R(S) \geq S $
$\{5\}$	$\{1, 2, 3\}$	$3 \geq 1$
$\{6\}$	$\{1\}$	$1 \geq 1$
$\{7\}$	$\{2, 3, 4\}$	$3 \geq 1$
$\{5, 6\}$	$\{1, 2, 3\}$	$3 \geq 2$
$\{5, 7\}$	$\{1, 2, 3, 4\}$	$4 \geq 2$
$\{6, 7\}$	$\{1, 2, 3, 4\}$	$4 \geq 2$
$\{5, 6, 7\}$	$\{1, 2, 3, 4\}$	$4 \geq 3$

Hence, according to Hall's Marriage Theorem, there is a matching that matches all the vertices of the set $\{5, 6, 7\}$. (Obviously, such a matching must be maximal.)

b. Since Hall's theorem requires checking an inequality for each subset $S \subseteq V$, the worst-case time efficiency of an algorithm based on it would be in $\Omega(2^{|V|})$. A much better solution is to find a maximum matching M^* (e.g., by the section's algorithm) and return "yes" if $|M^*| = |V|$ and "no" otherwise.


5. It is convenient to model the situation by a bipartite graph $G = \langle V, U, E \rangle$ where V represents the committees, U represents the committee members, and $(v, u) \in E$ if and only if u belongs to committee v :



There exists no matching that would match all the vertices of the set V . One way to prove it is based on Hall's Marriage Theorem (see Problem 4) whose necessary condition is violated for set $S = \{A, B, D, E\}$ with $R(S) = \{b, d, e\}$. Another way is to select a matching such as $M = \{(B, b), (C, c), (D, d), (E, e)\}$ (shown in bold) and check that the maximum-matching algorithm fails to find an augmenting path for it.

6. Add one source vertex s and connect it to each of the vertices in the set V by directed edges leaving s . Add one sink vertex t and connect each of the vertices in the set U to t by a directed edge entering t . Direct all the edges of the original graph to point from V to U . Assign 1 as the capacity of every edge in the network. A solution to the maximum-flow problem for the network yields a maximum matching for the original bipartite graph: it consists of the full edges (edges with the unit flow on them) between vertices of the original graph.

7. The greedy algorithm does not always find a maximum matching. As a counterexample, consider the bipartite graph shown below:



Since all its vertices have the same degree of 2, we can order them in numerical order of their labels: 1, 2, 3, 4, 5, 6. Using the same rule to break ties for selecting mates for vertices on the list, the greedy algorithm yields the matching $M = \{(1, 4), (2, 5)\}$, which is smaller than, say, $M^* = \{(1, 4), (2, 6), (3, 5)\}$.

8. A crucial observation is that for any edge between a leaf and its parent there is a maximum matching containing this edge. Indeed, consider a leaf v and its parent p . Let M be a maximum matching in the tree. If (v, p) is in M , we have a maximum matching desired. If M does not include (v, p) , it must contain an edge (u, p) from some vertex $u \neq v$ to p .

because otherwise we could have added (v, p) to M to get a larger matching. But then simply replacing (u, p) by (v, p) in the matching M yields a maximum matching containing (v, p) . This operation is to be repeated recursively for the smaller forest obtained.

Based on this observation, Manber ([Man89], p. 431) suggested the following recursive algorithm. Take an arbitrary leaf v and match it to its parent p . Remove from the tree both v and p along with all the edges incident to p . Also remove all the former sibling leaves of v , if any. This operation is to be repeated recursively for the smaller forest obtained.


Thieling Chen [Thieling Chen, "Maximum Matching and Minimum Vertex Covers of Trees," *The Western Journal of Graduate Research*, vol. 10, no. 1, 2001, pp. 10-14] suggested to implement the same idea by processing vertices of the tree in the reverse BFS order (i.e., bottom up and right to left across each level) as in the following pseudocode:

Algorithm *Tree-Max-Matching*

```
//Constructs a maximum matching in a free tree
//Input: A tree  $T$ 
//Output: A maximum cardinality matching  $M$  in  $T$ 
initialize matching  $M$  to the empty set
starting at an arbitrary vertex, traverse  $T$  by BFS, numbering the visited
vertices sequentially from 0 to  $n - 1$  and saving pointers to a visited
vertex and its parent
for  $i \leftarrow n - 1$  downto 0 do
    if vertex numbered  $i$  and its parent are both not marked
        add the edge between them to  $M$ 
        mark the parent
return  $M$ 
```


The time efficiency of this algorithm is obviously in $\Theta(n)$, where n is the number of vertices in an input tree.

10. a. No domino tiling of the 7-by-7 board exists because the total number of the board's squares is odd. For a domino tiling of a board to exist, the total number of squares on the board must be a multiple of two squares of a single domino. This is necessary (but not sufficient) condition for tiling with dominoes.
- b. Note that any n -by-1 or 1-by- n board where n is even can be tiled with dominoes in an obvious way. Hence any board that is composed of such strips, where n 's need not be the same, can be tiled as well. Here is a tiling of the board given in part b:




c. No domino tiling of the board given in part c is possible. Think of the board as being an 8-by-8 chessboard (with two missing squares at the diagonally opposite corners) whose squares of the board colored alternately black and white. Since a single domino would cover exactly one black and one white square, no tiling of the board is possible because it has 32 squares of one color and 30 squares of the other color.


Note: Obviously, a domino tiling can be considered a matching of the board's adjacent squares. The board in question can be represented by a bipartite graph, with the sets V and U representing diagonally adjacent squares (i.e., colored the same color as on the chessboard). Such graphs for the boards given are shown below:



a



b



c

The domino tiling question then becomes the perfect matching problem for the graph representing the board. For the graph of part a no perfect matching exists because the total number of vertices in the sets V and U is odd. A perfect matching for part b was outlined above. Finally, no perfect matching exists for the graph of part c because the number of vertices in the sets V and U is not the same. Note that while we could answer the domino tiling question for the three boards given without introducing the boards' graphs, this might be necessary for a more complex configuration of the board.

Exercises 10.4

1. Consider an instance of the stable marriage problem given by the ranking matrix

	A	B	C
α	1, 3	2, 2	3, 1
β	3, 1	1, 3	2, 2
γ	2, 2	3, 1	1, 3

For each of its marriage matchings, indicate whether it is stable or not. For the unstable matchings, specify a blocking pair. For the stable matchings, indicate whether they are man-optimal, woman-optimal, or neither. (Assume that the Greek and Roman letters denote the men and women, respectively.)

2. Design a simple algorithm for checking whether a given marriage matching is stable and determine its time efficiency class.
3. Find a stable-marriage matching for the instance given in Problem 1 by applying the stable-marriage algorithm

- (a) in its men-proposing version.
- (b) in its women-proposing version.

4. Find a stable-marriage matching for the instance defined by the following ranking matrix:

	A	B	C	D
α	1, 3	2, 3	3, 2	4, 3
β	1, 4	4, 1	3, 4	2, 2
γ	2, 2	1, 4	3, 3	4, 1
δ	4, 1	2, 2	3, 1	1, 4

5. Determine the time-efficiency class of the stable-marriage algorithm
 - (a) in the worst case.
 - (b) in the best case.
6. Prove that a man-optimal stable marriage set is always unique. Is it also true for a woman-optimal stable marriage matching?
7. Prove that in the man-optimal stable matching, each woman has the worst partner that she can have in any stable marriage matching.
8. Implement the stable-marriage algorithm given in Section 10.4 so that its running time is in $O(n^2)$. Run an experiment to ascertain its average-case efficiency.
9. Write a report on the college admission problem (residents-hospitals assignment) that generalizes the stable marriage problem.

10. \triangleright Consider the ***problem of the roommates***, which is related to but more difficult than the stable marriage problem: “An even number of boys wish to divide up into pairs of roommates. A set of pairings is called stable if under it there are no two boys who are not roommates and who prefer each other to their actual roommates.” [Gal62] Give an instance of this problem that does *not* have a stable pairing.

Hints to Exercises 10.4

1. A marriage matching is obtained by selecting three matrix cells, one cell from each row and column. To determine the stability of a given marriage matching, check each of the remaining matrix cells for containing a blocking pair.
2. It suffices to consider each member of one sex (say, the men) as a potential member of a blocking pair.
3. An application of the men-proposing version to another instance is given in the section. For the women-proposing version, reverse the roles of the sexes.
4. You may use either the men-proposing or women-proposing version of the algorithm.
5. The time efficiency is clearly defined by the number of proposals made. You may but are not required to provide the exact number of proposals in the worst and best cases, respectively; an appropriate Θ class will suffice.
6. Prove it by contradiction.
7. Prove it by contradiction.
8. Choose data structures so that the innermost loop of the algorithm can run in constant time.
9. The principal references are [Gal62] and [Gus89].
10. Consider four boys, three of whom rate the fourth boy as the least desired roommate. Complete these rankings to obtain an instance with no stable pairing.

Solutions to Exercises 10.4

1. There are the total of $3! = 6$ one-one matchings of two disjoint 3-element sets:

	A	B	C
α	$\boxed{1,3}$	2, 2	3, 1
β	3, 1	$\boxed{1,3}$	2, 2
γ	2, 2	3, 1	$\boxed{1,3}$

$\{(\alpha, A), (\beta, B), (\gamma, C)\}$ is stable: no other cell can be blocking since each man has his best choice. This is obviously the man-optimal matching.

	A	B	C
α	$\boxed{1,3}$	2, 2	3, 1
β	3, 1	1, 3	$\boxed{2,2}$
γ	2, 2	$\boxed{3,1}$	1, 3

$\{(\alpha, A), (\beta, C), (\gamma, B)\}$ is unstable: (γ, A) is a blocking pair.

	A	B	C
α	1, 3	$\boxed{2,2}$	3, 1
β	$\boxed{3,1}$	1, 3	2, 2
γ	2, 2	3, 1	$\boxed{1,3}$

$\{(\alpha, B), (\beta, A), (\gamma, C)\}$ is unstable: (β, C) is a blocking pair.

	A	B	C
α	1, 3	$\boxed{2,2}$	3, 1
β	3, 1	1, 3	$\boxed{2,2}$
γ	$\boxed{2,2}$	3, 1	1, 3

$\{(\alpha, B), (\beta, C), (\gamma, A)\}$ is stable: all the other cells contain a 3 (the lowest rank) and hence cannot be a blocking pair. This is neither a man-optimal nor a woman-optimal matching since it's inferior to $\{(\alpha, A), (\beta, B), (\gamma, C)\}$ for the men and inferior to $\{(\alpha, C), (\beta, A), (\gamma, B)\}$ for the women.

	A	B	C
α	1, 3	2, 2	$\boxed{3,1}$
β	$\boxed{3,1}$	1, 3	2, 2
γ	2, 2	$\boxed{3,1}$	1, 3

$\{(\alpha, C), (\beta, A), (\gamma, B)\}$ is stable: no other cell can be blocking since each woman has her best choice. This is obviously the woman-optimal matching.

	A	B	C
α	1, 3	2, 2	3, 1
β	3, 1	1, 3	2, 2
γ	2, 2	3, 1	1, 3

$\{(\alpha, C), (\beta, B), (\gamma, A)\}$ is unstable: (α, B) is a blocking pair.

2. Stability-checking algorithm

Input: A marriage matching M of n (m, w) pairs along with rankings of the women by each man and rankings of the men by each woman

Output: “yes” if the input is stable and a blocking pair otherwise

for $m \leftarrow 1$ **to** n **do**

for each w such that m prefers w to his mate in M **do**

if w prefers m to her mate in M

return (m, w)

return “yes”

With appropriate data structures, it is not difficult to implement this algorithm to run in $O(n^2)$ time. For example, the mates of the men and the mates of the women in a current matching can be stored in two arrays of size n and all the preferences can be stored in the n -by- n ranking matrix containing two rankings in each cell.

3. a.

		<i>A</i>	<i>B</i>	<i>C</i>	
Free men:	α	1,3	2, 2	3, 1	α proposed to <i>A</i>
α, β, γ	β	3, 1	1, 3	2, 2	<i>A</i> accepted
	γ	2, 2	3, 1	1, 3	
		<i>A</i>	<i>B</i>	<i>C</i>	
Free men:	α	1,3	2, 2	3, 1	β proposed to <i>B</i>
β, γ	β	3, 1	1,3	2, 2	<i>B</i> accepted
	γ	2, 2	3, 1	1, 3	
		<i>A</i>	<i>B</i>	<i>C</i>	
Free men:	α	1,3	2, 2	3, 1	γ proposed to <i>C</i>
γ	β	3, 1	1,3	2, 2	<i>C</i> accepted
	γ	2, 2	3, 1	1,3	

The (man-optimal) stable marriage matching is $M = \{(\alpha, A), (\beta, B), (\gamma, C)\}$.

b.

		A	B	C	
Free women:	α	1, 3	2, 2	3, 1	A proposed to β
A, B, C	β	3, 1	1, 3	2, 2	β accepted
	γ	2, 2	3, 1	1, 3	

		A	B	C	
Free women:	α	1, 3	2, 2	3, 1	B proposed to γ
B, C	β	3, 1	1, 3	2, 2	γ accepted
	γ	2, 2	3, 1	1, 3	

		A	B	C	
Free women:	α	1, 3	2, 2	3, 1	C proposed to α
C	β	3, 1	1, 3	2, 2	α accepted
	γ	2, 2	3, 1	1, 3	

The (woman-optimal) stable marriage matching is $M = \{(\beta, A), (\gamma, B), (\alpha, C)\}$.

4. The solution is given on the next page.

iteration 1
Free men: $\alpha, \beta, \gamma, \delta$

	A	B	C	D
α	1,3	2,3	3,2	4,3
β	1,4	4,1	3,4	2,2
γ	2,2	1,4	3,3	4,1
δ	4,1	2,2	3,1	1,4

α proposed to A; A accepted

iteration 2
Free men: β, γ, δ

	A	B	C	D
α	1,3	2,3	3,2	4,3
β	1,4	4,1	3,4	2,2
γ	2,2	1,4	3,3	4,1
δ	4,1	2,2	3,1	1,4

β proposed to A; A rejected

iteration 3
Free men: β, γ, δ

	A	B	C	D
α	1,3	2,3	3,2	4,3
β	1,4	4,1	3,4	2,2
γ	2,2	1,4	3,3	4,1
δ	4,1	2,2	3,1	1,4

β proposed to D; D accepted

iteration 4
Free men: γ, δ

	A	B	C	D
α	1,3	2,3	3,2	4,3
β	1,4	4,1	3,4	2,2
γ	2,2	1,4	3,3	4,1
δ	4,1	2,2	3,1	1,4

γ proposed to B; B accepted

iteration 5
Free men: δ

	A	B	C	D
α	1,3	2,3	3,2	4,3
β	1,4	4,1	3,4	2,2
γ	2,2	1,4	3,3	4,1
δ	4,1	2,2	3,1	1,4

δ proposed to D; D rejected

iteration 6
Free men: δ

	A	B	C	D
α	1,3	2,3	3,2	4,3
β	1,4	4,1	3,4	2,2
γ	2,2	1,4	3,3	4,1
δ	4,1	2,2	3,1	1,4

δ proposed to B; B replaced γ with δ

iteration 7
Free men: γ

	A	B	C	D
α	1,3	2,3	3,2	4,3
β	1,4	4,1	3,4	2,2
γ	2,2	1,4	3,3	4,1
δ	4,1	2,2	3,1	1,4

γ proposed to A; A replaced α with γ

iteration 8
Free men: α

	A	B	C	D
α	1,3	2,3	3,2	4,3
β	1,4	4,1	3,4	2,2
γ	2,2	1,4	3,3	4,1
δ	4,1	2,2	3,1	1,4

α proposed to B; B rejected

iteration 9
Free men: α

	A	B	C	D
α	1,3	2,3	3,2	4,3
β	1,4	4,1	3,4	2,2
γ	2,2	1,4	3,3	4,1
δ	4,1	2,2	3,1	1,4

α proposed to C; C accepted

Free men: none

$$M = \{(\alpha, C), (\beta, D), (\gamma, A), (\delta, B)\}$$

5. a. The worst-case time efficiency of the algorithm is in $\Theta(n^2)$. On the one hand, the total number of the proposals, $P(n)$, cannot exceed n^2 , the total number of possible partners for n men, because a man does not propose to the same woman more than once. On the other hand, for the instance of size n where all the men and women have the identical preference list $1, 2, \dots, n$, $P(n) = \sum_{i=1}^n i = n(n+1)/2$. Thus, if $P_w(n)$ is the number of proposals made by the algorithm in the worst case,

$$n(n+1)/2 \leq P_w(n) \leq n^2,$$

i.e., $P_w(n) \in \Theta(n^2)$.

- b. The best-case time efficiency of the algorithm is in $\Theta(n)$: the algorithm makes the minimum of n proposals, one by each man, on the input that ranks first a different woman for each of the n men.

6. Assume that there are two distinct man-optimal solutions to an instance of the stable marriage problem. Then there must exist at least one man m matched to two different women w_1 and w_2 in these solutions. Since no ties are allowed in the rankings, m must prefer one of these two women to the other, say, w_1 to w_2 . But then the marriage matching in which m is matched with w_2 is not man-optimal in contradiction to the assumption.

Of course, a woman-optimal solution is unique too due to the complete symmetry of these notions.

7. Assume that on the contrary there exists a man-optimal stable matching M in which some woman w doesn't have the worst possible partner in a stable matching, i.e., w prefers her partner m in M to her partner \bar{m} in another stable matching \bar{M} . Since $(m, w) \notin \bar{M}$, m must prefer his partner \bar{w} in \bar{M} to w because otherwise (m, w) would be a blocking pair for \bar{M} . But this contradicts the assumption that M is a man-optimal stable matching in which every man, including m , has the best possible partner in a stable matching.

8. n/a

9. n/a

10. Consider an instance of the problem of the roommates with four boys α , β , γ , and δ and the following preference lists (* stands for any legitimate

rating):

boy	rank 1	rank 2	rank 3
α	β	γ	δ
β	γ	α	δ
γ	α	β	δ
δ	*	*	*

Any pairing would have to pair one of the boys α, β, γ with δ . But any such pairing would be unstable since whoever is paired with δ will want to move out and one of the other two boys, having him rated first, will prefer him to his current roommate. For example, if α is paired with δ then β and γ are paired too while γ prefers α to β (in addition to α preferring γ to δ).

Note: This example is from the seminal paper by D. Gale and L. S. Shapley "College Admissions and the Stability of Marriage", *American Mathematical Monthly*, vol. 69 (Jan. 1962), 9-15. For an in-depth discussion of the problem of the roommates see the monograph by D. Gusfield and R.W. Irwing *The Stable Marriage Problem: Structure and Algorithms*,. MIT Press, 1989.